

# Using Strimzi

# Table of Contents

1. Overview of Strimzi .....	1
1.1. Kafka capabilities .....	1
1.2. Kafka use cases .....	1
1.3. How Strimzi supports Kafka .....	2
1.4. Strimzi Operators .....	2
1.4.1. Cluster Operator .....	3
1.4.2. Topic Operator .....	4
1.4.3. User Operator .....	5
1.5. Strimzi custom resources .....	6
1.5.1. Strimzi custom resource example .....	6
1.6. Listener configuration .....	8
1.7. Document Conventions .....	9
2. Deployment configuration .....	11
2.1. Kafka cluster configuration .....	11
2.1.1. Configuring Kafka .....	12
2.1.2. Configuring the Entity Operator .....	17
2.1.3. Kafka and ZooKeeper storage types .....	21
2.1.4. Scaling clusters .....	30
2.1.5. Retrieving JMX metrics with JmxTrans .....	39
2.1.6. Maintenance time windows for rolling updates .....	41
2.1.7. Connecting to ZooKeeper from a terminal .....	43
2.1.8. Deleting Kafka nodes manually .....	44
2.1.9. Deleting ZooKeeper nodes manually .....	44
2.1.10. List of Kafka cluster resources .....	45
2.2. Kafka Connect/S2I cluster configuration .....	50
2.2.1. Configuring Kafka Connect .....	50
2.2.2. Kafka Connect configuration for multiple instances .....	54
2.2.3. Configuring Kafka Connect user authorization .....	55
2.2.4. List of Kafka Connect cluster resources .....	58
2.2.5. List of Kafka Connect (S2I) cluster resources .....	58
2.3. Kafka MirrorMaker cluster configuration .....	59
2.3.1. Configuring Kafka MirrorMaker .....	59
2.3.2. List of Kafka MirrorMaker cluster resources .....	63
2.4. Kafka MirrorMaker 2.0 cluster configuration .....	63
2.4.1. MirrorMaker 2.0 data replication .....	64
2.4.2. Cluster configuration .....	64
2.4.3. ACL rules synchronization .....	66
2.4.4. Synchronizing data between Kafka clusters using MirrorMaker 2.0 .....	67

2.5. Kafka Bridge cluster configuration .....	72
2.5.1. Configuring the Kafka Bridge .....	72
2.5.2. List of Kafka Bridge cluster resources .....	75
2.6. Customizing Kubernetes resources .....	75
2.6.1. Customizing the image pull policy .....	77
2.7. Configuring pod scheduling .....	77
2.7.1. Specifying affinity, tolerations, and topology spread constraints .....	78
2.7.2. Configuring pod anti-affinity in Kafka components .....	79
2.7.3. Configuring node affinity in Kafka components .....	80
2.7.4. Setting up dedicated nodes and scheduling pods on them .....	81
2.8. External logging .....	82
2.8.1. Creating a ConfigMap for logging .....	83
3. Configuring external listeners .....	85
3.1. Accessing Kafka using node ports .....	85
3.2. Accessing Kafka using loadbalancers .....	86
3.3. Accessing Kafka using ingress .....	88
3.4. Accessing Kafka using OpenShift routes .....	90
4. Managing secure access to Kafka .....	92
4.1. Security options for Kafka .....	92
4.1.1. Listener authentication .....	92
4.1.2. Kafka authorization .....	95
4.2. Security options for Kafka clients .....	96
4.2.1. Identifying a Kafka cluster for user handling .....	96
4.2.2. User authentication .....	97
4.2.3. User authorization .....	99
4.3. Securing access to Kafka brokers .....	100
4.3.1. Securing Kafka brokers .....	101
4.3.2. Securing user access to Kafka .....	103
4.3.3. Restricting access to Kafka listeners using network policies .....	104
4.4. Using OAuth 2.0 token-based authentication .....	105
4.4.1. OAuth 2.0 authentication mechanism .....	106
4.4.2. OAuth 2.0 Kafka broker configuration .....	106
4.4.3. Session re-authentication for Kafka brokers .....	110
4.4.4. OAuth 2.0 Kafka client configuration .....	111
4.4.5. OAuth 2.0 client authentication flow .....	111
4.4.6. Configuring OAuth 2.0 authentication .....	114
4.4.7. Authorization server examples .....	122
4.5. Using OAuth 2.0 token-based authorization .....	123
4.5.1. OAuth 2.0 authorization mechanism .....	123
4.5.2. Configuring OAuth 2.0 authorization support .....	123
4.5.3. Managing policies and permissions in Keycloak Authorization Services .....	126

4.5.4. Example authorization rules configuration using Authorization Services .....	132
5. Using Strimzi Operators .....	143
5.1. Using the Cluster Operator .....	143
5.1.1. Cluster Operator configuration .....	143
5.1.2. Provisioning Role-Based Access Control (RBAC) .....	147
5.2. Using the Topic Operator .....	160
5.2.1. Kafka topic resource .....	160
5.2.2. Configuring a Kafka topic .....	163
5.2.3. Configuring the Topic Operator with resource requests and limits .....	164
5.3. Using the User Operator .....	165
5.3.1. Configuring the User Operator with resource requests and limits .....	165
5.4. Monitoring operators using Prometheus metrics .....	166
6. Kafka Bridge .....	167
6.1. Kafka Bridge overview .....	167
6.1.1. Kafka Bridge interface .....	167
6.1.2. Supported clients for the Kafka Bridge .....	168
6.1.3. Securing the Kafka Bridge .....	168
6.1.4. Accessing the Kafka Bridge outside of Kubernetes .....	169
6.1.5. Requests to the Kafka Bridge .....	170
6.1.6. CORS .....	172
6.1.7. Kafka Bridge API resources .....	174
6.1.8. Kafka Bridge deployment .....	174
6.2. Kafka Bridge quickstart .....	174
6.2.1. Deploying the Kafka Bridge to your Kubernetes cluster .....	175
6.2.2. Exposing the Kafka Bridge service to your local machine .....	176
6.2.3. Producing messages to topics and partitions .....	177
6.2.4. Creating a Kafka Bridge consumer .....	179
6.2.5. Subscribing a Kafka Bridge consumer to topics .....	180
6.2.6. Retrieving the latest messages from a Kafka Bridge consumer .....	181
6.2.7. Committing offsets to the log .....	182
6.2.8. Seeking to offsets for a partition .....	183
6.2.9. Deleting a Kafka Bridge consumer .....	184
7. Cruise Control for cluster rebalancing .....	185
7.1. Why use Cruise Control? .....	185
7.2. Optimization goals overview .....	185
7.3. Optimization proposals overview .....	190
7.4. Rebalance performance tuning overview .....	192
7.5. Cruise Control configuration .....	194
7.6. Deploying Cruise Control .....	197
7.7. Generating optimization proposals .....	201
7.8. Approving an optimization proposal .....	203

7.9. Stopping a cluster rebalance .....	205
7.10. Fixing problems with a <b>KafkaRebalance</b> resource .....	205
8. Distributed tracing .....	207
8.1. Overview of OpenTracing and Jaeger .....	208
8.2. Setting up tracing for Kafka clients .....	208
8.2.1. Initializing a Jaeger tracer for Kafka clients .....	209
8.2.2. Environment variables for tracing .....	209
8.3. Instrumenting Kafka clients with tracers .....	211
8.3.1. Instrumenting producers and consumers for tracing .....	211
8.3.2. Instrumenting Kafka Streams applications for tracing .....	215
8.4. Setting up tracing for MirrorMaker, Kafka Connect, and the Kafka Bridge .....	216
8.4.1. Enabling tracing in MirrorMaker, Kafka Connect, and Kafka Bridge resources .....	216
9. Managing TLS certificates .....	220
9.1. Certificate Authorities .....	220
9.1.1. CA certificates .....	221
9.1.2. Installing your own CA certificates .....	221
9.2. Secrets .....	223
9.2.1. PKCS #12 storage .....	223
9.2.2. Cluster CA Secrets .....	223
9.2.3. Client CA Secrets .....	225
9.2.4. Disabling <b>ownerReference</b> in the CA Secrets .....	225
9.2.5. User Secrets .....	226
9.3. Certificate renewal and validity periods .....	226
9.3.1. Renewal process with generated CAs .....	227
9.3.2. Client applications .....	227
9.3.3. Renewing CA certificates manually .....	228
9.3.4. Renewing your own CA certificates .....	230
9.4. Replacing private keys .....	232
9.5. TLS connections .....	232
9.5.1. ZooKeeper communication .....	233
9.5.2. Kafka interbroker communication .....	233
9.5.3. Topic and User Operators .....	233
9.5.4. Cruise Control .....	233
9.5.5. Kafka Client connections .....	233
9.6. Configuring internal clients to trust the cluster CA .....	233
9.7. Configuring external clients to trust the cluster CA .....	235
9.8. Kafka listener certificates .....	236
9.8.1. Providing your own Kafka listener certificates .....	237
9.8.2. Alternative subjects in server certificates for Kafka listeners .....	238
10. Managing Strimzi .....	241
10.1. Working with custom resources .....	241

10.1.1. Performing <code>kubectl</code> operations on custom resources .....	241
10.1.2. Strimzi custom resource status information .....	243
10.1.3. Finding the status of a custom resource .....	246
10.2. Manually starting rolling updates of Kafka and ZooKeeper clusters .....	246
10.2.1. Prerequisites .....	247
10.2.2. Performing a rolling update using a StatefulSet annotation .....	247
10.2.3. Performing a rolling update using a Pod annotation .....	247
10.3. Discovering services using labels and annotations .....	248
10.3.1. Returning connection details on services .....	249
10.4. Recovering a cluster from persistent volumes .....	250
10.4.1. Recovery from namespace deletion .....	250
10.4.2. Recovery from loss of a Kubernetes cluster .....	251
10.4.3. Recovering a deleted cluster from persistent volumes .....	251
10.5. Tuning client configuration .....	256
10.5.1. Kafka producer configuration tuning .....	256
10.5.2. Kafka consumer configuration tuning .....	261
10.6. Uninstalling Strimzi .....	268
Appendix A: Frequently asked questions .....	269
A.1. Questions related to the Cluster Operator .....	269
A.1.1. Why do I need cluster administrator privileges to install Strimzi? .....	269
A.1.2. Why does the Cluster Operator need to create <code>ClusterRoleBindings</code> ? .....	269
A.1.3. Can standard Kubernetes users create Kafka custom resources? .....	270
A.1.4. What do the <i>failed to acquire lock</i> warnings in the log mean? .....	270
A.1.5. Why is hostname verification failing when connecting to NodePorts using TLS? .....	270
Appendix B: Custom Resource API Reference .....	272
B.1. Common configuration properties .....	272
B.1.1. <code>replicas</code> .....	272
B.1.2. <code>bootstrapServers</code> .....	272
B.1.3. <code>ssl</code> .....	272
B.1.4. <code>trustedCertificates</code> .....	273
B.1.5. <code>resources</code> .....	274
B.1.6. <code>image</code> .....	276
B.1.7. <code>livenessProbe</code> and <code>readinessProbe</code> healthchecks .....	278
B.1.8. <code>metricsConfig</code> .....	279
B.1.9. <code>jvmOptions</code> .....	281
B.1.10. Garbage collector logging .....	283
B.2. <code>Kafka</code> schema reference .....	283
B.3. <code>KafkaSpec</code> schema reference .....	283
B.4. <code>KafkaClusterSpec</code> schema reference .....	284
B.4.1. <code>listeners</code> .....	284
B.4.2. <code>config</code> .....	285

B.4.3. <code>brokerRackInitImage</code>	287
B.4.4. <code>jmxOptions</code>	288
B.4.5. <code>logging</code>	289
B.4.6. Schema reference	290
B.5. <code>EphemeralStorage</code> schema reference	293
B.6. <code>PersistentClaimStorage</code> schema reference	293
B.7. <code>PersistentClaimStorageOverride</code> schema reference	294
B.8. <code>JbodStorage</code> schema reference	294
B.9. <code>GenericKafkaListener</code> schema reference	294
B.9.1. <code>listeners</code>	295
B.9.2. <code>type</code>	296
B.9.3. <code>port</code>	299
B.9.4. <code>tls</code>	300
B.9.5. <code>authentication</code>	300
B.9.6. <code>networkPolicyPeers</code>	300
B.9.7. Schema reference	302
B.10. <code>KafkaListenerAuthenticationTls</code> schema reference	303
B.11. <code>KafkaListenerAuthenticationScramSha512</code> schema reference	303
B.12. <code>KafkaListenerAuthenticationOAuth</code> schema reference	304
B.13. <code>GenericSecretSource</code> schema reference	306
B.14. <code>CertSecretSource</code> schema reference	306
B.15. <code>GenericKafkaListenerConfiguration</code> schema reference	307
B.15.1. <code>brokerCertChainAndKey</code>	307
B.15.2. <code>externalTrafficPolicy</code>	307
B.15.3. <code>loadBalancerSourceRanges</code>	307
B.15.4. <code>class</code>	308
B.15.5. <code>preferredNodePortAddressType</code>	308
B.15.6. <code>useServiceDnsDomain</code>	309
B.15.7. Schema reference	309
B.16. <code>CertAndKeySecretSource</code> schema reference	311
B.17. <code>GenericKafkaListenerConfigurationBootstrap</code> schema reference	312
B.17.1. <code>alternativeNames</code>	312
B.17.2. <code>host</code>	312
B.17.3. <code>nodePort</code>	314
B.17.4. <code>loadBalancerIP</code>	314
B.17.5. <code>annotations</code>	315
B.17.6. Schema reference	316
B.18. <code>GenericKafkaListenerConfigurationBroker</code> schema reference	317
B.18.1. Schema reference	318
B.19. <code>KafkaListeners</code> schema reference	319
B.20. <code>KafkaListenerPlain</code> schema reference	319

B.21. <code>KafkaListenerTls</code> schema reference .....	320
B.22. <code>TlsListenerConfiguration</code> schema reference .....	321
B.23. <code>KafkaListenerExternalRoute</code> schema reference .....	321
B.24. <code>RouteListenerOverride</code> schema reference .....	322
B.25. <code>RouteListenerBootstrapOverride</code> schema reference .....	322
B.26. <code>RouteListenerBrokerOverride</code> schema reference .....	322
B.27. <code>KafkaListenerExternalConfiguration</code> schema reference .....	323
B.28. <code>KafkaListenerExternalLoadBalancer</code> schema reference .....	323
B.29. <code>LoadBalancerListenerOverride</code> schema reference .....	324
B.30. <code>LoadBalancerListenerBootstrapOverride</code> schema reference .....	324
B.31. <code>LoadBalancerListenerBrokerOverride</code> schema reference .....	324
B.32. <code>KafkaListenerExternalNodePort</code> schema reference .....	325
B.33. <code>NodePortListenerOverride</code> schema reference .....	326
B.34. <code>NodePortListenerBootstrapOverride</code> schema reference .....	326
B.35. <code>NodePortListenerBrokerOverride</code> schema reference .....	327
B.36. <code>NodePortListenerConfiguration</code> schema reference .....	327
B.37. <code>KafkaListenerExternalIngress</code> schema reference .....	328
B.38. <code>IngressListenerConfiguration</code> schema reference .....	328
B.39. <code>IngressListenerBootstrapConfiguration</code> schema reference .....	329
B.40. <code>IngressListenerBrokerConfiguration</code> schema reference .....	329
B.41. <code>KafkaAuthorizationSimple</code> schema reference .....	330
B.41.1. <code>superUsers</code> .....	330
B.41.2. Schema reference .....	330
B.42. <code>KafkaAuthorizationOpa</code> schema reference .....	331
B.42.1. <code>url</code> .....	331
B.42.2. <code>allowOnError</code> .....	331
B.42.3. <code>initialCacheCapacity</code> .....	331
B.42.4. <code>maximumCacheSize</code> .....	331
B.42.5. <code>expireAfterMs</code> .....	331
B.42.6. <code>superUsers</code> .....	331
B.42.7. Schema reference .....	332
B.43. <code>KafkaAuthorizationKeycloak</code> schema reference .....	333
B.44. <code>Rack</code> schema reference .....	334
B.44.1. Schema reference .....	336
B.45. <code>Probe</code> schema reference .....	336
B.46. <code>JvmOptions</code> schema reference .....	336
B.47. <code>SystemProperty</code> schema reference .....	337
B.48. <code>KafkaJmxOptions</code> schema reference .....	337
B.49. <code>KafkaJmxAuthenticationPassword</code> schema reference .....	337
B.50. <code>JmxPrometheusExporterMetrics</code> schema reference .....	338
B.51. <code>ExternalConfigurationMetrics</code> schema reference .....	338

B.52. <code>InlineLogging</code> schema reference .....	338
B.53. <code>ExternalLogging</code> schema reference .....	339
B.54. <code>TlsSidecar</code> schema reference .....	339
B.54.1. Schema reference.....	341
B.55. <code>KafkaClusterTemplate</code> schema reference .....	342
B.56. <code>StatefulSetTemplate</code> schema reference .....	343
B.57. <code>MetadataTemplate</code> schema reference .....	343
B.57.1. Schema reference.....	344
B.58. <code>PodTemplate</code> schema reference .....	344
B.58.1. <code>hostAliases</code> .....	345
B.58.2. Schema reference.....	345
B.59. <code>ResourceTemplate</code> schema reference .....	347
B.60. <code>ExternalServiceTemplate</code> schema reference .....	347
B.60.1. Schema reference.....	347
B.61. <code>PodDisruptionBudgetTemplate</code> schema reference.....	348
B.61.1. Schema reference.....	349
B.62. <code>ContainerTemplate</code> schema reference .....	349
B.62.1. Schema reference.....	350
B.63. <code>ContainerEnvVar</code> schema reference .....	350
B.64. <code>ZookeeperClusterSpec</code> schema reference .....	350
B.64.1. <code>config</code> .....	350
B.64.2. <code>logging</code> .....	352
B.64.3. Schema reference.....	353
B.65. <code>ZookeeperClusterTemplate</code> schema reference .....	355
B.66. <code>TopicOperatorSpec</code> schema reference .....	355
B.67. <code>EntityOperatorSpec</code> schema reference .....	356
B.68. <code>EntityTopicOperatorSpec</code> schema reference .....	357
B.68.1. <code>logging</code> .....	357
B.68.2. Schema reference.....	358
B.69. <code>EntityUserOperatorSpec</code> schema reference .....	359
B.69.1. <code>logging</code> .....	359
B.69.2. Schema reference.....	361
B.70. <code>EntityOperatorTemplate</code> schema reference .....	361
B.71. <code>CertificateAuthority</code> schema reference .....	362
B.72. <code>CruiseControlSpec</code> schema reference .....	363
B.73. <code>CruiseControlTemplate</code> schema reference .....	364
B.74. <code>BrokerCapacity</code> schema reference .....	365
B.75. <code>JmxTransSpec</code> schema reference .....	365
B.76. <code>JmxTransOutputDefinitionTemplate</code> schema reference .....	366
B.77. <code>JmxTransQueryTemplate</code> schema reference .....	366
B.78. <code>JmxTransTemplate</code> schema reference .....	367

B.79. <code>KafkaExporterSpec</code> schema reference .....	367
B.80. <code>KafkaExporterTemplate</code> schema reference .....	368
B.81. <code>KafkaStatus</code> schema reference .....	368
B.82. <code>Condition</code> schema reference .....	369
B.83. <code>ListenerStatus</code> schema reference .....	369
B.84. <code>ListenerAddress</code> schema reference .....	370
B.85. <code>KafkaConnect</code> schema reference .....	370
B.86. <code>KafkaConnectSpec</code> schema reference .....	370
B.86.1. <code>config</code> .....	370
B.86.2. <code>logging</code> .....	372
B.86.3. Schema reference .....	373
B.87. <code>KafkaConnectTls</code> schema reference .....	376
B.87.1. <code>trustedCertificates</code> .....	376
B.87.2. Schema reference .....	376
B.88. <code>KafkaClientAuthenticationTls</code> schema reference .....	376
B.88.1. <code>certificateAndKey</code> .....	376
B.88.2. Schema reference .....	377
B.89. <code>KafkaClientAuthenticationScramSha512</code> schema reference .....	377
B.89.1. <code>username</code> .....	377
B.89.2. <code>passwordSecret</code> .....	377
B.89.3. Schema reference .....	378
B.90. <code>PasswordSecretSource</code> schema reference .....	379
B.91. <code>KafkaClientAuthenticationPlain</code> schema reference .....	379
B.91.1. <code>username</code> .....	379
B.91.2. <code>passwordSecret</code> .....	379
B.91.3. Schema reference .....	380
B.92. <code>KafkaClientAuthenticationOAuth</code> schema reference .....	380
B.92.1. Schema reference .....	383
B.93. <code>JaegerTracing</code> schema reference .....	384
B.94. <code>KafkaConnectTemplate</code> schema reference .....	384
B.95. <code>DeploymentTemplate</code> schema reference .....	385
B.96. <code>ExternalConfiguration</code> schema reference .....	385
B.96.1. <code>env</code> .....	385
B.96.2. <code>volumes</code> .....	387
B.96.3. Schema reference .....	388
B.97. <code>ExternalConfigurationEnv</code> schema reference .....	388
B.98. <code>ExternalConfigurationEnvVarSource</code> schema reference .....	389
B.99. <code>ExternalConfigurationVolumeSource</code> schema reference .....	389
B.100. <code>Build</code> schema reference .....	389
B.100.1. <code>output</code> .....	389
B.100.2. <code>plugins</code> .....	391

B.100.3. Schema reference .....	394
B.101. <b>DockerOutput</b> schema reference .....	394
B.102. <b>ImageStreamOutput</b> schema reference .....	395
B.103. <b>Plugin</b> schema reference .....	395
B.104. <b>JarArtifact</b> schema reference .....	395
B.105. <b>TgzArtifact</b> schema reference .....	396
B.106. <b>KafkaConnectStatus</b> schema reference .....	396
B.107. <b>ConnectorPlugin</b> schema reference .....	397
B.108. <b>KafkaConnectS2I</b> schema reference .....	397
B.109. <b>KafkaConnectS2ISpec</b> schema reference .....	397
B.109.1. Schema reference .....	398
B.110. <b>KafkaConnectS2IStatus</b> schema reference .....	400
B.111. <b>KafkaTopic</b> schema reference .....	400
B.112. <b>KafkaTopicSpec</b> schema reference .....	401
B.113. <b>KafkaTopicStatus</b> schema reference .....	401
B.114. <b>KafkaUser</b> schema reference .....	401
B.115. <b>KafkaUserSpec</b> schema reference .....	402
B.116. <b>KafkaUserTlsClientAuthentication</b> schema reference .....	402
B.117. <b>KafkaUserScramSha512ClientAuthentication</b> schema reference .....	403
B.118. <b>KafkaUserAuthorizationSimple</b> schema reference .....	403
B.119. <b>AclRule</b> schema reference .....	403
B.119.1. <b>resource</b> .....	404
B.119.2. <b>type</b> .....	405
B.119.3. <b>operation</b> .....	405
B.119.4. <b>host</b> .....	405
B.119.5. Schema reference .....	405
B.120. <b>AclRuleTopicResource</b> schema reference .....	406
B.121. <b>AclRuleGroupResource</b> schema reference .....	406
B.122. <b>AclRuleClusterResource</b> schema reference .....	407
B.123. <b>AclRuleTransactionalIdResource</b> schema reference .....	407
B.124. <b>KafkaUserQuotas</b> schema reference .....	408
B.124.1. <b>quotas</b> .....	408
B.124.2. Schema reference .....	408
B.125. <b>KafkaUserTemplate</b> schema reference .....	409
B.125.1. Schema reference .....	409
B.126. <b>KafkaUserStatus</b> schema reference .....	410
B.127. <b>KafkaMirrorMaker</b> schema reference .....	410
B.128. <b>KafkaMirrorMakerSpec</b> schema reference .....	410
B.128.1. <b>whitelist</b> .....	410
B.128.2. <b>KafkaMirrorMakerConsumerSpec</b> and <b>KafkaMirrorMakerProducerSpec</b> .....	410
B.128.3. <b>logging</b> .....	411

B.128.4. Schema reference .....	411
B.129. <i>KafkaMirrorMakerConsumerSpec</i> schema reference .....	413
B.129.1. <i>numStreams</i> .....	413
B.129.2. <i>offsetCommitInterval</i> .....	413
B.129.3. <i>config</i> .....	414
B.129.4. <i>groupId</i> .....	415
B.129.5. Schema reference .....	415
B.130. <i>KafkaMirrorMakerTls</i> schema reference .....	416
B.130.1. <i>trustedCertificates</i> .....	416
B.130.2. Schema reference .....	416
B.131. <i>KafkaMirrorMakerProducerSpec</i> schema reference .....	416
B.131.1. <i>abortOnSendFailure</i> .....	416
B.131.2. <i>config</i> .....	416
B.131.3. Schema reference .....	417
B.132. <i>KafkaMirrorMakerTemplate</i> schema reference .....	418
B.133. <i>KafkaMirrorMakerStatus</i> schema reference .....	418
B.134. <i>KafkaBridge</i> schema reference .....	419
B.135. <i>KafkaBridgeSpec</i> schema reference .....	419
B.135.1. <i>logging</i> .....	419
B.135.2. Schema reference .....	421
B.136. <i>KafkaBridgeTls</i> schema reference .....	423
B.137. <i>KafkaBridgeHttpConfig</i> schema reference .....	423
B.137.1. <i>cors</i> .....	423
B.137.2. Schema reference .....	423
B.138. <i>KafkaBridgeHttpCors</i> schema reference .....	424
B.139. <i>KafkaBridgeConsumerSpec</i> schema reference .....	424
B.139.1. Schema reference .....	425
B.140. <i>KafkaBridgeProducerSpec</i> schema reference .....	425
B.140.1. Schema reference .....	426
B.141. <i>KafkaBridgeTemplate</i> schema reference .....	427
B.142. <i>KafkaBridgeStatus</i> schema reference .....	427
B.143. <i>KafkaConnector</i> schema reference .....	427
B.144. <i>KafkaConnectorSpec</i> schema reference .....	428
B.145. <i>KafkaConnectorStatus</i> schema reference .....	428
B.146. <i>KafkaMirrorMaker2</i> schema reference .....	428
B.147. <i>KafkaMirrorMaker2Spec</i> schema reference .....	429
B.148. <i>KafkaMirrorMaker2ClusterSpec</i> schema reference .....	430
B.148.1. <i>config</i> .....	430
B.148.2. Schema reference .....	431
B.149. <i>KafkaMirrorMaker2Tls</i> schema reference .....	431
B.150. <i>KafkaMirrorMaker2MirrorSpec</i> schema reference .....	431

B.151. <code>KafkaMirrorMaker2ConnectorSpec</code> schema reference .....	432
B.152. <code>KafkaMirrorMaker2Status</code> schema reference .....	433
B.153. <code>KafkaRebalance</code> schema reference .....	433
B.154. <code>KafkaRebalanceSpec</code> schema reference .....	433
B.155. <code>KafkaRebalanceStatus</code> schema reference .....	434

# Chapter 1. Overview of Strimzi

Strimzi simplifies the process of running Apache Kafka in a Kubernetes cluster.

This guide provides instructions for configuring Kafka components and using Strimzi Operators. Procedures relate to how you might want to modify your deployment and introduce additional features, such as Cruise Control or distributed tracing.

You can configure your deployment using [Strimzi custom resources](#). The [Custom resource API reference](#) describes the properties you can use in your configuration.

**NOTE**

Looking to get started with Strimzi? For step-by-step deployment instructions, see the [Deploying and Upgrading Strimzi](#) guide.

## 1.1. Kafka capabilities

The underlying data stream-processing capabilities and component architecture of Kafka can deliver:

- Microservices and other applications to share data with extremely high throughput and low latency
- Message ordering guarantees
- Message rewind/replay from data storage to reconstruct an application state
- Message compaction to remove old records when using a key-value log
- Horizontal scalability in a cluster configuration
- Replication of data to control fault tolerance
- Retention of high volumes of data for immediate access

## 1.2. Kafka use cases

Kafka's capabilities make it suitable for:

- Event-driven architectures
- Event sourcing to capture changes to the state of an application as a log of events
- Message brokering
- Website activity tracking
- Operational monitoring through metrics
- Log collection and aggregation
- Commit logs for distributed systems
- Stream processing so that applications can respond to data in real time

## 1.3. How Strimzi supports Kafka

Strimzi provides container images and Operators for running Kafka on Kubernetes. Strimzi Operators are fundamental to the running of Strimzi. The Operators provided with Strimzi are purpose-built with specialist operational knowledge to effectively manage Kafka.

Operators simplify the process of:

- Deploying and running Kafka clusters
- Deploying and running Kafka components
- Configuring access to Kafka
- Securing access to Kafka
- Upgrading Kafka
- Managing brokers
- Creating and managing topics
- Creating and managing users

## 1.4. Strimzi Operators

Strimzi supports Kafka using *Operators* to deploy and manage the components and dependencies of Kafka to Kubernetes.

Operators are a method of packaging, deploying, and managing a Kubernetes application. Strimzi Operators extend Kubernetes functionality, automating common and complex tasks related to a Kafka deployment. By implementing knowledge of Kafka operations in code, Kafka administration tasks are simplified and require less manual intervention.

### Operators

Strimzi provides Operators for managing a Kafka cluster running within a Kubernetes cluster.

#### Cluster Operator

Deploys and manages Apache Kafka clusters, Kafka Connect, Kafka MirrorMaker, Kafka Bridge, Kafka Exporter, and the Entity Operator

#### Entity Operator

Comprises the Topic Operator and User Operator


#### Topic Operator

Manages Kafka topics

#### User Operator

Manages Kafka users

The Cluster Operator can deploy the Topic Operator and User Operator as part of an **Entity Operator** configuration at the same time as a Kafka cluster.



### 1.4.1. Cluster Operator

Strimzi uses the Cluster Operator to deploy and manage clusters for:

- Kafka (including ZooKeeper, Entity Operator, Kafka Exporter, and Cruise Control)
- Kafka Connect
- Kafka MirrorMaker
- Kafka Bridge

Custom resources are used to deploy the clusters.

For example, to deploy a Kafka cluster:

- A **Kafka** resource with the cluster configuration is created within the Kubernetes cluster.
- The Cluster Operator deploys a corresponding Kafka cluster, based on what is declared in the **Kafka** resource.


The Cluster Operator can also deploy (through configuration of the **Kafka** resource):

- A Topic Operator to provide operator-style topic management through **KafkaTopic** custom resources

- A User Operator to provide operator-style user management through **KafkaUser** custom resources

The Topic Operator and User Operator function within the Entity Operator on deployment.


*Example architecture for the Cluster Operator*



## 1.4.2. Topic Operator

The Topic Operator provides a way of managing topics in a Kafka cluster through Kubernetes resources.

*Example architecture for the Topic Operator*



The role of the Topic Operator is to keep a set of **KafkaTopic** Kubernetes resources describing Kafka topics in-sync with corresponding Kafka topics.

Specifically, if a **KafkaTopic** is:

- Created, the Topic Operator creates the topic
- Deleted, the Topic Operator deletes the topic
- Changed, the Topic Operator updates the topic

Working in the other direction, if a topic is:

- Created within the Kafka cluster, the Operator creates a [KafkaTopic](#)
- Deleted from the Kafka cluster, the Operator deletes the [KafkaTopic](#)
- Changed in the Kafka cluster, the Operator updates the [KafkaTopic](#)

This allows you to declare a [KafkaTopic](#) as part of your application's deployment and the Topic Operator will take care of creating the topic for you. Your application just needs to deal with producing or consuming from the necessary topics.

If the topic is reconfigured or reassigned to different Kafka nodes, the [KafkaTopic](#) will always be up to date.

### 1.4.3. User Operator

The User Operator manages Kafka users for a Kafka cluster by watching for [KafkaUser](#) resources that describe Kafka users, and ensuring that they are configured properly in the Kafka cluster.

For example, if a [KafkaUser](#) is:

- Created, the User Operator creates the user it describes
- Deleted, the User Operator deletes the user it describes
- Changed, the User Operator updates the user it describes

Unlike the Topic Operator, the User Operator does not sync any changes from the Kafka cluster with the Kubernetes resources. Kafka topics can be created by applications directly in Kafka, but it is not expected that the users will be managed directly in the Kafka cluster in parallel with the User Operator.

The User Operator allows you to declare a [KafkaUser](#) resource as part of your application's deployment. You can specify the authentication and authorization mechanism for the user. You can also configure *user quotas* that control usage of Kafka resources to ensure, for example, that a user does not monopolize access to a broker.

When the user is created, the user credentials are created in a [Secret](#). Your application needs to use the user and its credentials for authentication and to produce or consume messages.

In addition to managing credentials for authentication, the User Operator also manages authorization rules by including a description of the user's access rights in the [KafkaUser](#) declaration.

## 1.5. Strimzi custom resources

A deployment of Kafka components to a Kubernetes cluster using Strimzi is highly configurable through the application of custom resources. Custom resources are created as instances of APIs added by Custom resource definitions (CRDs) to extend Kubernetes resources.

CRDs act as configuration instructions to describe the custom resources in a Kubernetes cluster, and are provided with Strimzi for each Kafka component used in a deployment, as well as users and topics. CRDs and custom resources are defined as YAML files. Example YAML files are provided with the Strimzi distribution.

CRDs also allow Strimzi resources to benefit from native Kubernetes features like CLI accessibility and configuration validation.

### *Additional resources*

- [Extend the Kubernetes API with CustomResourceDefinitions](#)

### 1.5.1. Strimzi custom resource example

CRDs require a one-time installation in a cluster to define the schemas used to instantiate and manage Strimzi-specific resources.

After a new custom resource type is added to your cluster by installing a CRD, you can create instances of the resource based on its specification.

Depending on the cluster setup, installation typically requires cluster admin privileges.

**NOTE** Access to manage custom resources is limited to Strimzi administrators. For more information, see [Designating Strimzi administrators](#) in the *Deploying and Upgrading Strimzi* guide.

A CRD defines a new **kind** of resource, such as **kind:Kafka**, within a Kubernetes cluster.

The Kubernetes API server allows custom resources to be created based on the **kind** and understands from the CRD how to validate and store the custom resource when it is added to the Kubernetes cluster.

**WARNING** When CRDs are deleted, custom resources of that type are also deleted. Additionally, the resources created by the custom resource, such as pods and statefulsets are also deleted.

Each Strimzi-specific custom resource conforms to the schema defined by the CRD for the resource's **kind**. The custom resources for Strimzi components have common configuration properties, which are defined under **spec**.

To understand the relationship between a CRD and a custom resource, let's look at a sample of the CRD for a Kafka topic.

## Kafka topic CRD

```
apiVersion: kafka.strimzi.io/v1beta1
kind: CustomResourceDefinition
metadata: ①
  name: kafkatopics.kafka.strimzi.io
  labels:
    app: strimzi
spec: ②
  group: kafka.strimzi.io
  versions:
    v1beta1
  scope: Namespaced
  names:
    # ...
    singular: kafkatopic
    plural: kafkatopics
    shortNames:
      - kt ③
  additionalPrinterColumns: ④
    # ...
  subresources:
    status: {} ⑤
  validation: ⑥
  openAPIV3Schema:
    properties:
      spec:
        type: object
        properties:
          partitions:
            type: integer
            minimum: 1
          replicas:
            type: integer
            minimum: 1
            maximum: 32767
    # ...
```

- ① The metadata for the topic CRD, its name and a label to identify the CRD.
- ② The specification for this CRD, including the group (domain) name, the plural name and the supported schema version, which are used in the URL to access the API of the topic. The other names are used to identify instance resources in the CLI. For example, `kubectl get kafkatopic my-topic` or `kubectl get kafkatopics`.
- ③ The shortname can be used in CLI commands. For example, `kubectl get kt` can be used as an abbreviation instead of `kubectl get kafkatopic`.
- ④ The information presented when using a `get` command on the custom resource.
- ⑤ The current status of the CRD as described in the [schema reference](#) for the resource.
- ⑥ openAPIV3Schema validation provides validation for the creation of topic custom resources. For

example, a topic requires at least one partition and one replica.

**NOTE**

You can identify the CRD YAML files supplied with the Strimzi installation files, because the file names contain an index number followed by ‘Crd’.

Here is a corresponding example of a **KafkaTopic** custom resource.

*Kafka topic custom resource*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaTopic ①
metadata:
  name: my-topic
  labels:
    strimzi.io/cluster: my-cluster ②
spec: ③
  partitions: 1
  replicas: 1
  config:
    retention.ms: 7200000
    segment.bytes: 1073741824
status:
  conditions: ④
    lastTransitionTime: "2019-08-20T11:37:00.706Z"
    status: "True"
    type: Ready
  observedGeneration: 1
/ ...
```

① The **kind** and **apiVersion** identify the CRD of which the custom resource is an instance.

② A label, applicable only to **KafkaTopic** and **KafkaUser** resources, that defines the name of the Kafka cluster (which is same as the name of the **Kafka** resource) to which a topic or user belongs.

③ The spec shows the number of partitions and replicas for the topic as well as the configuration parameters for the topic itself. In this example, the retention period for a message to remain in the topic and the segment file size for the log are specified.

④ Status conditions for the **KafkaTopic** resource. The **type** condition changed to **Ready** at the **lastTransitionTime**.

Custom resources can be applied to a cluster through the platform CLI. When the custom resource is created, it uses the same validation as the built-in resources of the Kubernetes API.

After a **KafkaTopic** custom resource is created, the Topic Operator is notified and corresponding Kafka topics are created in Strimzi.

## 1.6. Listener configuration

Listeners are used to connect to Kafka brokers.

Strimzi provides a generic [GenericKafkaListener](#) schema with properties to configure listeners through the [Kafka](#) resource.

The [GenericKafkaListener](#) provides a flexible approach to listener configuration.

You can specify properties to configure *internal* listeners for connecting within the Kubernetes cluster, or *external* listeners for connecting outside the Kubernetes cluster.

#### *Generic listener configuration*

Each listener is [defined as an array in the Kafka resource](#).

For more information on listener configuration, see the [GenericKafkaListener schema reference](#).

Generic listener configuration replaces the previous approach to listener configuration using the [KafkaListeners schema reference](#), which is [deprecated](#). However, you can [convert the old format into the new format](#) with backwards compatibility.

The [KafkaListeners](#) schema uses sub-properties for [plain](#), [tls](#) and [external](#) listeners, with fixed ports for each. Because of the limits inherent in the architecture of the schema, it is only possible to configure three listeners, with configuration options limited to the type of listener.

With the [GenericKafkaListener](#) schema, you can configure as many listeners as required, as long as their names and ports are unique.

You might want to configure multiple external listeners, for example, to handle access from networks that require different authentication mechanisms. Or you might need to join your Kubernetes network to an outside network. In which case, you can configure internal listeners (using the [useServiceDnsDomain](#) property) so that the Kubernetes service DNS domain (typically [.cluster.local](#)) is not used.

#### *Configuring listeners to secure access to Kafka brokers*

You can configure listeners for secure connection using authentication. For more information on securing access to Kafka brokers, see [Managing access to Kafka](#).

#### *Configuring external listeners for client access outside Kubernetes*

You can configure external listeners for client access outside a Kubernetes environment using a specified connection mechanism, such as a loadbalancer. For more information on the configuration options for connecting an external client, see [Configuring external listeners](#).

#### *Listener certificates*

You can provide your own server certificates, called *Kafka listener certificates*, for TLS listeners or external listeners which have TLS encryption enabled. For more information, see [Kafka listener certificates](#).

## 1.7. Document Conventions

### *Replaceables*

In this document, replaceable text is styled in [monospace](#), with italics, uppercase, and hyphens.

For example, in the following code, you will want to replace `MY-NAMESPACE` with the name of your namespace:

```
sed -i 's/namespace: .*namespace: MY-NAMESPACE/' install/cluster-operator/*RoleBinding*.yaml
```

# Chapter 2. Deployment configuration

This chapter describes how to configure different aspects of the supported deployments using custom resources:

- Kafka clusters
- Kafka Connect clusters
- Kafka Connect clusters with *Source2Image* support
- Kafka MirrorMaker
- Kafka Bridge
- Cruise Control

**NOTE** Labels applied to a custom resource are also applied to the Kubernetes resources comprising Kafka MirrorMaker. This provides a convenient mechanism for resources to be labeled as required.

The *Deploying and Upgrading Strimzi* guide describes how to [monitor your Strimzi deployment](#).

## 2.1. Kafka cluster configuration

This section describes how to configure a Kafka deployment in your Strimzi cluster. A Kafka cluster is deployed with a ZooKeeper cluster. The deployment can also include the Topic Operator and User Operator, which manage Kafka topics and users.

You configure Kafka using the [Kafka](#) resource. Configuration options are also available for ZooKeeper and the Entity Operator within the [Kafka](#) resource. The Entity Operator comprises the Topic Operator and User Operator.

The full schema of the [Kafka](#) resource is described in the [Kafka schema reference](#).

### *Listener configuration*

You configure listeners for connecting clients to Kafka brokers. For more information on configuring listeners for connecting brokers, see [Listener configuration](#).

### *Authorizing access to Kafka*

You can configure your Kafka cluster to allow or decline actions executed by users. For more information on securing access to Kafka brokers, see [Managing access to Kafka](#).

### *Managing TLS certificates*

When deploying Kafka, the Cluster Operator automatically sets up and renews TLS certificates to enable encryption and authentication within your cluster. If required, you can manually renew the cluster and client CA certificates before their renewal period ends. You can also replace the keys used by the cluster and client CA certificates. For more information, see [Renewing CA certificates manually](#) and [Replacing private keys](#).

### *Additional resources*

- For more information about Apache Kafka, see the [Apache Kafka website](#).

## 2.1.1. Configuring Kafka

Use the properties of the **Kafka** resource to configure your Kafka deployment.

As well as configuring Kafka, you can add configuration for ZooKeeper and the Strimzi Operators. Common configuration properties, such as logging and healthchecks, are configured independently for each component.

This procedure shows only some of the possible configuration options, but those that are particularly important include:

- Resource requests (CPU / Memory)
- JVM options for maximum and minimum memory allocation
- Listeners (and authentication of clients)
- Authentication
- Storage
- Rack awareness
- Metrics
- Cruise Control for cluster rebalancing

### Prerequisites

- An OpenShift cluster
- A running Cluster Operator

See the *Deploying and Upgrading Strimzi* guide for instructions on deploying a:

- [Cluster Operator](#)
- [Kafka cluster](#)

### Procedure

1. Edit the **spec** properties for the **Kafka** resource.

The properties you can configure are shown in this example configuration:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    replicas: 3 ①
    version: 0.21.1 ②
    logging: ③
    type: inline
```

```

loggers:
  kafka.root.logger.level: "INFO"
resources: ④
  requests:
    memory: 64Gi
    cpu: "8"
  limits:
    memory: 64Gi
    cpu: "12"
readinessProbe: ⑤
  initialDelaySeconds: 15
  timeoutSeconds: 5
livenessProbe:
  initialDelaySeconds: 15
  timeoutSeconds: 5
jvmOptions: ⑥
  -Xms: 8192m
  -Xmx: 8192m
image: my-org/my-image:latest ⑦
listeners: ⑧
  - name: plain ⑨
    port: 9092 ⑩
    type: internal ⑪
    tls: false ⑫
    configuration:
      useServiceDnsDomain: true ⑬
  - name: tls
    port: 9093
    type: internal
    tls: true
    authentication: ⑭
      type: tls
  - name: external ⑮
    port: 9094
    type: route
    tls: true
    configuration:
      brokerCertChainAndKey: ⑯
        secretName: my-secret
        certificate: my-certificate.crt
        key: my-key.key
authorization: ⑰
  type: simple
config: ⑱
  auto.create.topics.enable: "false"
  offsets.topic.replication.factor: 3
  transaction.state.log.replication.factor: 3
  transaction.state.log.min_isr: 2
  ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384" ⑲
  ssl.enabled.protocols: "TLSv1.2"
  ssl.protocol: "TLSv1.2"

```

```
storage: ②
  type: persistent-claim
  size: 10000Gi
rack:
  topologyKey: topology.kubernetes.io/zone
metricsConfig:
  type: jmxPrometheusExporter
  valueFrom:
    configMapKeyRef:
      name: my-config-map
      key: my-key
jmxOptions:
  authentication:
    type: "password"
# ...
zookeeper:
  replicas: 3
  logging:
    type: inline
    loggers:
      zookeeper.root.logger: "INFO"
resources:
  requests:
    memory: 8Gi
    cpu: "2"
  limits:
    memory: 8Gi
    cpu: "2"
jvmOptions:
  -Xms: 4096m
  -Xmx: 4096m
storage:
  type: persistent-claim
  size: 1000Gi
metricsConfig:
  # ...
entityOperator:
  tlsSidecar:
    resources:
      requests:
        cpu: 200m
        memory: 64Mi
      limits:
        cpu: 500m
        memory: 128Mi
topicOperator:
  watchedNamespace: my-topic-namespace
  reconciliationIntervalSeconds: 60
  logging:
    type: inline
    loggers:
```

```

rootLogger.level: "INFO"
resources:
  requests:
    memory: 512Mi
    cpu: "1"
  limits:
    memory: 512Mi
    cpu: "1"
userOperator:
  watchedNamespace: my-topic-namespace
  reconciliationIntervalSeconds: 60
  logging:
    type: inline
    loggers:
      rootLogger.level: INFO
resources:
  requests:
    memory: 512Mi
    cpu: "1"
  limits:
    memory: 512Mi
    cpu: "1"
kafkaExporter:
  # ...
cruiseControl:
  # ...
  tlsSidecar:
    # ...

```

- ① The number of replica nodes. If your cluster already has topics defined, you can [scale clusters](#).
- ② Kafka version, which can be changed to a supported version by following [the upgrade procedure](#).
- ③ Specified [Kafka loggers and log levels](#) added directly ([inline](#)) or indirectly ([external](#)) through a ConfigMap. A custom ConfigMap must be placed under the `log4j.properties` key. For the Kafka `kafka.root.logger.level` logger, you can set the log level to INFO, ERROR, WARN, TRACE, DEBUG, FATAL or OFF.
- ④ Requests for reservation of [supported resources](#), currently `cpu` and `memory`, and limits to specify the maximum resources that can be consumed.
- ⑤ [Healthchecks](#) to know when to restart a container (liveness) and when a container can accept traffic (readiness).
- ⑥ [JVM configuration options](#) to optimize performance for the Virtual Machine (VM) running Kafka.
- ⑦ ADVANCED OPTION: [Container image configuration](#), which is recommended only in special situations.
- ⑧ Listeners configure how clients connect to the Kafka cluster via bootstrap addresses. Listeners are [configured as \*internal\* or \*external\* listeners for connection from inside or](#)

[outside the Kubernetes cluster](#).

- ⑨ Name to identify the listener. Must be unique within the Kafka cluster.
- ⑩ Port number used by the listener inside Kafka. The port number has to be unique within a given Kafka cluster. Allowed port numbers are 9092 and higher with the exception of ports 9404 and 9999, which are already used for Prometheus and JMX. Depending on the listener type, the port number might not be the same as the port number that connects Kafka clients.
- ⑪ Listener type specified as `internal`, or for external listeners, as `route`, `loadbalancer`, `nodeport` or `ingress`.
- ⑫ Enables TLS encryption for each listener. Default is `false`. TLS encryption is not required for `route` listeners.
- ⑬ Defines whether the fully-qualified DNS names including the cluster service suffix (usually `.cluster.local`) are assigned.
- ⑭ Listener authentication mechanism specified as mutual TLS, SCRAM-SHA-512 or token-based OAuth 2.0.
- ⑮ External listener configuration specifies how the Kafka cluster is exposed outside Kubernetes, such as through a `route`, `loadbalancer` or `nodeport`.
- ⑯ Optional configuration for a Kafka listener certificate managed by an external Certificate Authority. The `brokerCertChainAndKey` specifies a Secret that contains a server certificate and a private key. You can configure Kafka listener certificates on any listener with enabled TLS encryption.
- ⑰ Authorization enables simple, OAUTH 2.0, or OPA authorization on the Kafka broker. Simple authorization uses the `AclAuthorizer` Kafka plugin.
- ⑱ The `config` specifies the broker configuration. Standard Apache Kafka configuration may be provided, restricted to those properties not managed directly by Strimzi.
- ⑲ SSL properties for listeners with TLS encryption enabled to enable a specific *cipher suite* or TLS version.
- ⑳ Storage is configured as `ephemeral`, `persistent-claim` or `jbod`.

Storage size for `persistent volumes` may be increased and additional `volumes` may be added to `JBOD storage`.

Persistent storage has `additional configuration options`, such as a storage `id` and `class` for dynamic volume provisioning.

`Rack awareness` is configured to spread replicas across different racks. A `topologykey` must match the label of a cluster node.

`Prometheus metrics` enabled. In this example, metrics are configured for the Prometheus JMX Exporter (the default metrics exporter).

Prometheus rules for exporting metrics to a Grafana dashboard through the Prometheus JMX Exporter, which are enabled by referencing a ConfigMap containing configuration for the Prometheus JMX exporter. You can enable metrics without further configuration using a reference to a ConfigMap containing an empty file under `metricsConfig.valueFrom.configMapKeyRef.key`.

`JMX options` to open JMX port 9999 to obtain JMX metrics, which is configured for password

protection in this example. You can open the port without protection using `jmxOptions: {}`.

ZooKeeper-specific configuration, which contains properties similar to the Kafka configuration.

**The number of ZooKeeper nodes.** ZooKeeper clusters or ensembles usually run with an odd number of nodes, typically three, five, or seven. The majority of nodes must be available in order to maintain an effective quorum. If the ZooKeeper cluster loses its quorum, it will stop responding to clients and the Kafka brokers will stop working. Having a stable and highly available ZooKeeper cluster is crucial for Strimzi.

Specified [ZooKeeper loggers and log levels](#).

Entity Operator configuration, which [specifies the configuration for the Topic Operator and User Operator](#).

Entity Operator [TLS sidecar configuration](#). Entity Operator uses the TLS sidecar for secure communication with ZooKeeper.

Specified [Topic Operator loggers and log levels](#). This example uses [inline](#) logging.

Specified [User Operator loggers and log levels](#).

Kafka Exporter configuration. [Kafka Exporter](#) is an optional component for extracting metrics data from Kafka brokers, in particular consumer lag data.

Optional configuration for Cruise Control, which is used to [rebalance the Kafka cluster](#).

Cruise Control [TLS sidecar configuration](#). Cruise Control uses the TLS sidecar for secure communication with ZooKeeper.

## 2. Create or update the resource:

```
kubectl apply -f KAFKA-CONFIG-FILE
```

### 2.1.2. Configuring the Entity Operator

The Entity Operator is responsible for managing Kafka-related entities in a running Kafka cluster.

The Entity Operator comprises the:

- [Topic Operator](#) to manage Kafka topics
- [User Operator](#) to manage Kafka users

Through [Kafka](#) resource configuration, the Cluster Operator can deploy the Entity Operator, including one or both operators, when deploying a Kafka cluster.

**NOTE**

When deployed, the Entity Operator contains the operators according to the deployment configuration.

The operators are automatically configured to manage the topics and users of the Kafka cluster.

## Entity Operator configuration properties

Use the `entityOperator` property in `Kafka.spec` to configure the Entity Operator.

The `entityOperator` property supports several sub-properties:

- `tlsSidecar`
- `topicOperator`
- `userOperator`
- `template`

The `tlsSidecar` property contains the configuration of the TLS sidecar container, which is used to communicate with ZooKeeper.

The `template` property contains the configuration of the Entity Operator pod, such as labels, annotations, affinity, and tolerations. For more information on configuring templates, see [Customizing Kubernetes resources](#).

The `topicOperator` property contains the configuration of the Topic Operator. When this option is missing, the Entity Operator is deployed without the Topic Operator.

The `userOperator` property contains the configuration of the User Operator. When this option is missing, the Entity Operator is deployed without the User Operator.

For more information on the properties used to configure the Entity Operator, see the [EntityUserOperatorSpec schema reference](#).

*Example of basic configuration enabling both operators*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  entityOperator:
    topicOperator: {}
    userOperator: {}
```

If an empty object (`{}`) is used for the `topicOperator` and `userOperator`, all properties use their default values.

When both `topicOperator` and `userOperator` properties are missing, the Entity Operator is not deployed.

## Topic Operator configuration properties

Topic Operator deployment can be configured using additional options inside the `topicOperator` object. The following properties are supported:

### `watchedNamespace`

The Kubernetes namespace in which the topic operator watches for `KafkaTopics`. Default is the namespace where the Kafka cluster is deployed.

### `reconciliationIntervalSeconds`

The interval between periodic reconciliations in seconds. Default `90`.

### `zookeeperSessionTimeoutSeconds`

The ZooKeeper session timeout in seconds. Default `20`.

### `topicMetadataMaxAttempts`

The number of attempts at getting topic metadata from Kafka. The time between each attempt is defined as an exponential back-off. Consider increasing this value when topic creation might take more time due to the number of partitions or replicas. Default `6`.

### `image`

The `image` property can be used to configure the container image which will be used. For more details about configuring custom container images, see [image](#).

### `resources`

The `resources` property configures the amount of resources allocated to the Topic Operator. For more details about resource request and limit configuration, see [resources](#).

### `logging`

The `logging` property configures the logging of the Topic Operator. For more details, see [logging](#).

### *Example Topic Operator configuration*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  entityOperator:
    # ...
  topicOperator:
    watchedNamespace: my-topic-namespace
    reconciliationIntervalSeconds: 60
    # ...
```

## User Operator configuration properties

User Operator deployment can be configured using additional options inside the `userOperator` object. The following properties are supported:

### `watchedNamespace`

The Kubernetes namespace in which the user operator watches for `KafkaUsers`. Default is the namespace where the Kafka cluster is deployed.

### `reconciliationIntervalSeconds`

The interval between periodic reconciliations in seconds. Default `120`.

### `zookeeperSessionTimeoutSeconds`

The ZooKeeper session timeout in seconds. Default `6`.

### `image`

The `image` property can be used to configure the container image which will be used. For more details about configuring custom container images, see [image](#).

### `resources`

The `resources` property configures the amount of resources allocated to the User Operator. For more details about resource request and limit configuration, see [resources](#).

### `logging`

The `logging` property configures the logging of the User Operator. For more details, see [logging](#).

### `secretPrefix`

The `secretPrefix` property adds a prefix to the name of all Secrets created from the KafkaUser resource. For example, `STRIMZI_SECRET_PREFIX=kafka-` would prefix all Secret names with `kafka-`. So a KafkaUser named `my-user` would create a Secret named `kafka-my-user`.

### *Example User Operator configuration*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  entityOperator:
    # ...
  userOperator:
    watchedNamespace: my-user-namespace
    reconciliationIntervalSeconds: 60
    # ...
```

### 2.1.3. Kafka and ZooKeeper storage types

As stateful applications, Kafka and ZooKeeper need to store data on disk. Strimzi supports three storage types for this data:

- Ephemeral
- Persistent
- JBOD storage

**NOTE** JBOD storage is supported only for Kafka, not for ZooKeeper.

When configuring a [Kafka](#) resource, you can specify the type of storage used by the Kafka broker and its corresponding ZooKeeper node. You configure the storage type using the [storage](#) property in the following resources:

- [Kafka.spec.kafka](#)
- [Kafka.spec.zookeeper](#)

The storage type is configured in the [type](#) field.

**WARNING** The storage type cannot be changed after a Kafka cluster is deployed.

#### *Additional resources*

- For more information about ephemeral storage, see [ephemeral storage schema reference](#).
- For more information about persistent storage, see [persistent storage schema reference](#).
- For more information about JBOD storage, see [JBOD schema reference](#).
- For more information about the schema for [Kafka](#), see [Kafka schema reference](#).

### Data storage considerations

An efficient data storage infrastructure is essential to the optimal performance of Strimzi.

Block storage is required. File storage, such as NFS, does not work with Kafka.

For your block storage, you can choose, for example:

- Cloud-based block storage solutions, such as [Amazon Elastic Block Store \(EBS\)](#)
- [Local persistent volumes](#)
- Storage Area Network (SAN) volumes accessed by a protocol such as *Fibre Channel* or *iSCSI*

**NOTE** Strimzi does not require Kubernetes raw block volumes.

### File systems

It is recommended that you configure your storage system to use the *XFS* file system. Strimzi is also compatible with the *ext4* file system, but this might require additional configuration for best

results.

### Apache Kafka and ZooKeeper storage

Use separate disks for Apache Kafka and ZooKeeper.

Three types of data storage are supported:

- Ephemeral (Recommended for development only)
- Persistent
- JBOD (Just a Bunch of Disks, suitable for Kafka only)

For more information, see [Kafka and ZooKeeper storage](#).

Solid-state drives (SSDs), though not essential, can improve the performance of Kafka in large clusters where data is sent to and received from multiple topics asynchronously. SSDs are particularly effective with ZooKeeper, which requires fast, low latency data access.

**NOTE**

You do not need to provision replicated storage because Kafka and ZooKeeper both have built-in data replication.

### Ephemeral storage

Ephemeral storage uses `emptyDir` volumes to store data. To use ephemeral storage, set the `type` field to `ephemeral`.

**IMPORTANT**

`emptyDir` volumes are not persistent and the data stored in them is lost when the pod is restarted. After the new pod is started, it must recover all data from the other nodes of the cluster. Ephemeral storage is not suitable for use with single-node ZooKeeper clusters or for Kafka topics with a replication factor of 1. This configuration will cause data loss.

#### *An example of Ephemeral storage*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    storage:
      type: ephemeral
      # ...
  zookeeper:
    # ...
    storage:
      type: ephemeral
      # ...
```

## Log directories

The ephemeral volume is used by the Kafka brokers as log directories mounted into the following path:

```
/var/lib/kafka/data/kafka-logIDX
```

Where `IDX` is the Kafka broker pod index. For example `/var/lib/kafka/data/kafka-log0`.

## Persistent storage

Persistent storage uses [Persistent Volume Claims](#) to provision persistent volumes for storing data. Persistent Volume Claims can be used to provision volumes of many different types, depending on the [Storage Class](#) which will provision the volume. The data types which can be used with persistent volume claims include many types of SAN storage as well as [Local persistent volumes](#).

To use persistent storage, the `type` has to be set to `persistent-claim`. Persistent storage supports additional configuration options:

### **id (optional)**

Storage identification number. This option is mandatory for storage volumes defined in a JBOD storage declaration. Default is `0`.

### **size (required)**

Defines the size of the persistent volume claim, for example, "1000Gi".

### **class (optional)**

The Kubernetes [Storage Class](#) to use for dynamic volume provisioning.

### **selector (optional)**

Allows selecting a specific persistent volume to use. It contains key:value pairs representing labels for selecting such a volume.

### **deleteClaim (optional)**

Boolean value which specifies if the Persistent Volume Claim has to be deleted when the cluster is undeployed. Default is `false`.

Increasing the size of persistent volumes in an existing Strimzi cluster is only supported in Kubernetes versions that support persistent volume resizing. The persistent volume to be resized must use a storage class that supports volume expansion. For other versions of Kubernetes and storage classes which do not support volume expansion, you must decide the necessary storage size before deploying the cluster. Decreasing the size of existing persistent volumes is not possible.

### **WARNING**

*Example fragment of persistent storage configuration with 1000Gi size*

```
# ...
storage:
  type: persistent-claim
  size: 1000Gi
# ...
```

The following example demonstrates the use of a storage class.

*Example fragment of persistent storage configuration with specific Storage Class*

```
# ...
storage:
  type: persistent-claim
  size: 1Gi
  class: my-storage-class
# ...
```

Finally, a **selector** can be used to select a specific labeled persistent volume to provide needed features such as an SSD.

*Example fragment of persistent storage configuration with selector*

```
# ...
storage:
  type: persistent-claim
  size: 1Gi
  selector:
    hdd-type: ssd
    deleteClaim: true
# ...
```

### Storage class overrides

You can specify a different storage class for one or more Kafka brokers or ZooKeeper nodes, instead of using the default storage class. This is useful if, for example, storage classes are restricted to different availability zones or data centers. You can use the **overrides** field for this purpose.

In this example, the default storage class is named **my-storage-class**:

## Example Strimzi cluster using storage class overrides

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  labels:
    app: my-cluster
    name: my-cluster
    namespace: myproject
spec:
  # ...
  kafka:
    replicas: 3
    storage:
      deleteClaim: true
      size: 100Gi
      type: persistent-claim
      class: my-storage-class
      overrides:
        - broker: 0
          class: my-storage-class-zone-1a
        - broker: 1
          class: my-storage-class-zone-1b
        - broker: 2
          class: my-storage-class-zone-1c
  # ...
  zookeeper:
    replicas: 3
    storage:
      deleteClaim: true
      size: 100Gi
      type: persistent-claim
      class: my-storage-class
      overrides:
        - broker: 0
          class: my-storage-class-zone-1a
        - broker: 1
          class: my-storage-class-zone-1b
        - broker: 2
          class: my-storage-class-zone-1c
  # ...
```

As a result of the configured `overrides` property, the volumes use the following storage classes:

- The persistent volumes of ZooKeeper node 0 will use `my-storage-class-zone-1a`.
- The persistent volumes of ZooKeeper node 1 will use `my-storage-class-zone-1b`.
- The persistent volumes of Zookeepr node 2 will use `my-storage-class-zone-1c`.
- The persistent volumes of Kafka broker 0 will use `my-storage-class-zone-1a`.

- The persistent volumes of Kafka broker 1 will use `my-storage-class-zone-1b`.
- The persistent volumes of Kafka broker 2 will use `my-storage-class-zone-1c`.

The `overrides` property is currently used only to override storage class configurations. Overriding other storage configuration fields is not currently supported. Other fields from the storage configuration are currently not supported.

### Persistent Volume Claim naming

When persistent storage is used, it creates Persistent Volume Claims with the following names:

#### `data-cluster-name-kafka-idx`

Persistent Volume Claim for the volume used for storing data for the Kafka broker pod `idx`.

#### `data-cluster-name-zookeeper-idx`

Persistent Volume Claim for the volume used for storing data for the ZooKeeper node pod `idx`.

### Log directories

The persistent volume is used by the Kafka brokers as log directories mounted into the following path:

```
/var/lib/kafka/data/kafka-logIDX
```

Where `IDX` is the Kafka broker pod index. For example `/var/lib/kafka/data/kafka-log0`.

### Resizing persistent volumes

You can provision increased storage capacity by increasing the size of the persistent volumes used by an existing Strimzi cluster. Resizing persistent volumes is supported in clusters that use either a single persistent volume or multiple persistent volumes in a JBOD storage configuration.

**NOTE**

You can increase but not decrease the size of persistent volumes. Decreasing the size of persistent volumes is not currently supported in Kubernetes.

#### Prerequisites

- A Kubernetes cluster with support for volume resizing.
- The Cluster Operator is running.
- A Kafka cluster using persistent volumes created using a storage class that supports volume expansion.

#### Procedure

1. In a `Kafka` resource, increase the size of the persistent volume allocated to the Kafka cluster, the ZooKeeper cluster, or both.
  - To increase the volume size allocated to the Kafka cluster, edit the `spec.kafka.storage` property.

- To increase the volume size allocated to the ZooKeeper cluster, edit the `spec.zookeeper.storage` property.

For example, to increase the volume size from `1000Gi` to `2000Gi`:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    storage:
      type: persistent-claim
      size: 2000Gi
      class: my-storage-class
    # ...
  zookeeper:
    # ...
```

## 2. Create or update the resource:

```
kubectl apply -f KAFKA-CONFIG-FILE
```

Kubernetes increases the capacity of the selected persistent volumes in response to a request from the Cluster Operator. When the resizing is complete, the Cluster Operator restarts all pods that use the resized persistent volumes. This happens automatically.

### *Additional resources*

For more information about resizing persistent volumes in Kubernetes, see [Resizing Persistent Volumes using Kubernetes](#).

## JBOD storage overview

You can configure Strimzi to use JBOD, a data storage configuration of multiple disks or volumes. JBOD is one approach to providing increased data storage for Kafka brokers. It can also improve performance.

A JBOD configuration is described by one or more volumes, each of which can be either [ephemeral](#) or [persistent](#). The rules and constraints for JBOD volume declarations are the same as those for ephemeral and persistent storage. For example, you cannot decrease the size of a persistent storage volume after it has been provisioned, or you cannot change the value of `sizeLimit` when `type=ephemeral`.

### JBOD configuration

To use JBOD with Strimzi, the storage `type` must be set to `jbd`. The `volumes` property allows you to describe the disks that make up your JBOD storage array or configuration. The following fragment

shows an example JBOD configuration:

```
# ...
storage:
  type: jbod
  volumes:
    - id: 0
      type: persistent-claim
      size: 100Gi
      deleteClaim: false
    - id: 1
      type: persistent-claim
      size: 100Gi
      deleteClaim: false
# ...
```

The ids cannot be changed once the JBOD volumes are created.

Users can add or remove volumes from the JBOD configuration.

#### JBOD and Persistent Volume Claims

When persistent storage is used to declare JBOD volumes, the naming scheme of the resulting Persistent Volume Claims is as follows:

##### **data-*id*-cluster-name-kafka-*idx***

Where *id* is the ID of the volume used for storing data for Kafka broker pod *idx*.

##### Log directories

The JBOD volumes will be used by the Kafka brokers as log directories mounted into the following path:

##### **/var/lib/kafka/data-*id*/kafka-log\_*idx*\_**

Where *id* is the ID of the volume used for storing data for Kafka broker pod *idx*. For example `/var/lib/kafka/data-0/kafka-log0`.

#### Adding volumes to JBOD storage

This procedure describes how to add volumes to a Kafka cluster configured to use JBOD storage. It cannot be applied to Kafka clusters configured to use any other storage type.

**NOTE** When adding a new volume under an *id* which was already used in the past and removed, you have to make sure that the previously used `PersistentVolumeClaims` have been deleted.

##### *Prerequisites*

- A Kubernetes cluster
- A running Cluster Operator

- A Kafka cluster with JBOD storage

#### *Procedure*

1. Edit the `spec.kafka.storage.volumes` property in the `Kafka` resource. Add the new volumes to the `volumes` array. For example, add the new volume with id `2`:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    storage:
      type: jbod
      volumes:
        - id: 0
          type: persistent-claim
          size: 100Gi
          deleteClaim: false
        - id: 1
          type: persistent-claim
          size: 100Gi
          deleteClaim: false
        - id: 2
          type: persistent-claim
          size: 100Gi
          deleteClaim: false
    # ...
  zookeeper:
    # ...
```

2. Create or update the resource:

```
kubectl apply -f KAFKA-CONFIG-FILE
```

3. Create new topics or reassign existing partitions to the new disks.

#### *Additional resources*

For more information about reassigning topics, see [Partition reassignment](#).

### **Removing volumes from JBOD storage**

This procedure describes how to remove volumes from Kafka cluster configured to use JBOD storage. It cannot be applied to Kafka clusters configured to use any other storage type. The JBOD storage always has to contain at least one volume.

**IMPORTANT**

To avoid data loss, you have to move all partitions before removing the volumes.

*Prerequisites*

- A Kubernetes cluster
- A running Cluster Operator
- A Kafka cluster with JBOD storage with two or more volumes

*Procedure*

1. Reassign all partitions from the disks which are you going to remove. Any data in partitions still assigned to the disks which are going to be removed might be lost.
2. Edit the `spec.kafka.storage.volumes` property in the `Kafka` resource. Remove one or more volumes from the `volumes` array. For example, remove the volumes with ids **1** and **2**:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    storage:
      type: jbod
      volumes:
        - id: 0
          type: persistent-claim
          size: 100Gi
          deleteClaim: false
    # ...
  zookeeper:
    # ...
```

3. Create or update the resource:

```
kubectl apply -f KAFKA-CONFIG-FILE
```

*Additional resources*

For more information about reassigning topics, see [Partition reassignment](#).

## 2.1.4. Scaling clusters

### Scaling Kafka clusters

## Adding brokers to a cluster

The primary way of increasing throughput for a topic is to increase the number of partitions for that topic. That works because the extra partitions allow the load of the topic to be shared between the different brokers in the cluster. However, in situations where every broker is constrained by a particular resource (typically I/O) using more partitions will not result in increased throughput. Instead, you need to add brokers to the cluster.

When you add an extra broker to the cluster, Kafka does not assign any partitions to it automatically. You must decide which partitions to move from the existing brokers to the new broker.

Once the partitions have been redistributed between all the brokers, the resource utilization of each broker should be reduced.

## Removing brokers from a cluster

Because Strimzi uses [StatefulSets](#) to manage broker pods, you cannot remove *any* pod from the cluster. You can only remove one or more of the highest numbered pods from the cluster. For example, in a cluster of 12 brokers the pods are named `cluster-name-kafka-0` up to `cluster-name-kafka-11`. If you decide to scale down by one broker, the `cluster-name-kafka-11` will be removed.

Before you remove a broker from a cluster, ensure that it is not assigned to any partitions. You should also decide which of the remaining brokers will be responsible for each of the partitions on the broker being decommissioned. Once the broker has no assigned partitions, you can scale the cluster down safely.

## Partition reassignment

The Topic Operator does not currently support reassigning replicas to different brokers, so it is necessary to connect directly to broker pods to reassign replicas to brokers.

Within a broker pod, the `kafka-reassign-partitions.sh` utility allows you to reassign partitions to different brokers.

It has three different modes:

### `--generate`

Takes a set of topics and brokers and generates a *reassignment JSON file* which will result in the partitions of those topics being assigned to those brokers. Because this operates on whole topics, it cannot be used when you only want to reassign some partitions of some topics.

### `--execute`

Takes a *reassignment JSON file* and applies it to the partitions and brokers in the cluster. Brokers that gain partitions as a result become followers of the partition leader. For a given partition, once the new broker has caught up and joined the ISR (in-sync replicas) the old broker will stop being a follower and will delete its replica.

### `--verify`

Using the same *reassignment JSON file* as the `--execute` step, `--verify` checks whether all the

partitions in the file have been moved to their intended brokers. If the reassignment is complete, `--verify` also removes any [throttles](#) that are in effect. Unless removed, throttles will continue to affect the cluster even after the reassignment has finished.

It is only possible to have one reassignment running in a cluster at any given time, and it is not possible to cancel a running reassignment. If you need to cancel a reassignment, wait for it to complete and then perform another reassignment to revert the effects of the first reassignment. The [kafka-reassign-partitions.sh](#) will print the reassignment JSON for this reversion as part of its output. Very large reassessments should be broken down into a number of smaller reassessments in case there is a need to stop in-progress reassignment.

### Reassignment JSON file

The *reassignment JSON file* has a specific structure:

```
{  
  "version": 1,  
  "partitions": [  
    <PartitionObjects>  
  ]  
}
```

Where *<PartitionObjects>* is a comma-separated list of objects like:

```
{  
  "topic": <TopicName>,  
  "partition": <Partition>,  
  "replicas": [ <AssignedBrokerIds> ]  
}
```

**NOTE**

Although Kafka also supports a "[log\\_dirs](#)" property this should not be used in Strimzi.

The following is an example reassignment JSON file that assigns partition [4](#) of topic [topic-a](#) to brokers [2](#), [4](#) and [7](#), and partition [2](#) of topic [topic-b](#) to brokers [1](#), [5](#) and [7](#):

```
{
  "version": 1,
  "partitions": [
    {
      "topic": "topic-a",
      "partition": 4,
      "replicas": [2,4,7]
    },
    {
      "topic": "topic-b",
      "partition": 2,
      "replicas": [1,5,7]
    }
  ]
}
```

Partitions not included in the JSON are not changed.

#### Reassigning partitions between JBOD volumes

When using JBOD storage in your Kafka cluster, you can choose to reassign the partitions between specific volumes and their log directories (each volume has a single log directory). To reassign a partition to a specific volume, add the `log_dirs` option to `<PartitionObjects>` in the reassignment JSON file.

```
{
  "topic": <TopicName>,
  "partition": <Partition>,
  "replicas": [ <AssignedBrokerIds> ],
  "log_dirs": [ <AssignedLogDirs> ]
}
```

The `log_dirs` object should contain the same number of log directories as the number of replicas specified in the `replicas` object. The value should be either an absolute path to the log directory, or the `any` keyword.

For example:

```
{
  "topic": "topic-a",
  "partition": 4,
  "replicas": [2,4,7],
  "log_dirs": [ "/var/lib/kafka/data-0/kafka-log2", "/var/lib/kafka/data-0/kafka-log4", "/var/lib/kafka/data-0/kafka-log7" ]
}
```

## Generating reassignment JSON files

This procedure describes how to generate a reassignment JSON file that reassigns all the partitions for a given set of topics using the `kafka-reassign-partitions.sh` tool.

### Prerequisites

- A running Cluster Operator
- A Kafka resource
- A set of topics to reassign the partitions of

### Procedure

1. Prepare a JSON file named `topics.json` that lists the topics to move. It must have the following structure:

```
{  
  "version": 1,  
  "topics": [  
    <TopicObjects>  
  ]  
}
```

where `<TopicObjects>` is a comma-separated list of objects like:

```
{  
  "topic": <TopicName>  
}
```

For example if you want to reassign all the partitions of `topic-a` and `topic-b`, you would need to prepare a `topics.json` file like this:

```
{  
  "version": 1,  
  "topics": [  
    { "topic": "topic-a"},  
    { "topic": "topic-b"}  
  ]  
}
```

2. Copy the `topics.json` file to one of the broker pods:

```
cat topics.json | kubectl exec -c kafka <BrokerPod> -i -- \  
/bin/bash -c \  
'cat > /tmp/topics.json'
```

3. Use the `kafka-reassign-partitions.sh` command to generate the reassignment JSON.

```
kubectl exec <BrokerPod> -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--topics-to-move-json-file /tmp/topics.json \
--broker-list <BrokerList> \
--generate
```

For example, to move all the partitions of **topic-a** and **topic-b** to brokers **4** and **7**

```
kubectl exec <BrokerPod> -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--topics-to-move-json-file /tmp/topics.json \
--broker-list 4,7 \
--generate
```

## Creating reassignment JSON files manually

You can manually create the reassignment JSON file if you want to move specific partitions.

### Reassignment throttles

Partition reassignment can be a slow process because it involves transferring large amounts of data between brokers. To avoid a detrimental impact on clients, you can throttle the reassignment process. This might cause the reassignment to take longer to complete.

- If the throttle is too low then the newly assigned brokers will not be able to keep up with records being published and the reassignment will never complete.
- If the throttle is too high then clients will be impacted.

For example, for producers, this could manifest as higher than normal latency waiting for acknowledgement. For consumers, this could manifest as a drop in throughput caused by higher latency between polls.

## Scaling up a Kafka cluster

This procedure describes how to increase the number of brokers in a Kafka cluster.

### Prerequisites

- An existing Kafka cluster.
- A *reassignment JSON file* named `reassignment.json` that describes how partitions should be reassigned to brokers in the enlarged cluster.

### Procedure

1. Add as many new brokers as you need by increasing the `Kafka.spec.kafka.replicas` configuration option.
2. Verify that the new broker pods have started.
3. Copy the `reassignment.json` file to the broker pod on which you will later execute the

commands:

```
cat reassignment.json | \
kubectl exec broker-pod -c kafka -i -- /bin/bash -c \
'cat > /tmp/reassignment.json'
```

For example:

```
cat reassignment.json | \
kubectl exec my-cluster-kafka-0 -c kafka -i -- /bin/bash -c \
'cat > /tmp/reassignment.json'
```

4. Execute the partition reassignment using the [kafka-reassign-partitions.sh](#) command line tool from the same broker pod.

```
kubectl exec broker-pod -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--execute
```

If you are going to throttle replication you can also pass the [--throttle](#) option with an inter-broker throttled rate in bytes per second. For example:

```
kubectl exec my-cluster-kafka-0 -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--throttle 5000000 \
--execute
```

This command will print out two reassignment JSON objects. The first records the current assignment for the partitions being moved. You should save this to a local file (not a file in the pod) in case you need to revert the reassignment later on. The second JSON object is the target reassignment you have passed in your reassignment JSON file.

5. If you need to change the throttle during reassignment you can use the same command line with a different throttled rate. For example:

```
kubectl exec my-cluster-kafka-0 -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--throttle 10000000 \
--execute
```

6. Periodically verify whether the reassignment has completed using the [kafka-reassign-partitions.sh](#) command line tool from any of the broker pods. This is the same command as the

previous step but with the `--verify` option instead of the `--execute` option.

```
kubectl exec broker-pod -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--verify
```

For example,

```
kubectl exec my-cluster-kafka-0 -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--verify
```

7. The reassignment has finished when the `--verify` command reports each of the partitions being moved as completed successfully. This final `--verify` will also have the effect of removing any reassignment throttles. You can now delete the revert file if you saved the JSON for reverting the assignment to their original brokers.

## Scaling down a Kafka cluster

This procedure describes how to decrease the number of brokers in a Kafka cluster.

### Prerequisites

- An existing Kafka cluster.
- A *reassignment JSON file* named `reassignment.json` describing how partitions should be reassigned to brokers in the cluster once the broker(s) in the highest numbered `Pod(s)` have been removed.

### Procedure

1. Copy the `reassignment.json` file to the broker pod on which you will later execute the commands:

```
cat reassignment.json | \
kubectl exec broker-pod -c kafka -i -- /bin/bash -c \
'cat > /tmp/reassignment.json'
```

For example:

```
cat reassignment.json | \
kubectl exec my-cluster-kafka-0 -c kafka -i -- /bin/bash -c \
'cat > /tmp/reassignment.json'
```

2. Execute the partition reassignment using the `kafka-reassign-partitions.sh` command line tool from the same broker pod.

```
kubectl exec broker-pod -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--execute
```

If you are going to throttle replication you can also pass the `--throttle` option with an inter-broker throttled rate in bytes per second. For example:

```
kubectl exec my-cluster-kafka-0 -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--throttle 5000000 \
--execute
```

This command will print out two reassignment JSON objects. The first records the current assignment for the partitions being moved. You should save this to a local file (not a file in the pod) in case you need to revert the reassignment later on. The second JSON object is the target reassignment you have passed in your reassignment JSON file.

3. If you need to change the throttle during reassignment you can use the same command line with a different throttled rate. For example:

```
kubectl exec my-cluster-kafka-0 -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--throttle 10000000 \
--execute
```

4. Periodically verify whether the reassignment has completed using the `kafka-reassign-partitions.sh` command line tool from any of the broker pods. This is the same command as the previous step but with the `--verify` option instead of the `--execute` option.

```
kubectl exec broker-pod -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--verify
```

For example,

```
kubectl exec my-cluster-kafka-0 -c kafka -it -- \
bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 \
--reassignment-json-file /tmp/reassignment.json \
--verify
```

- The reassignment has finished when the `--verify` command reports each of the partitions being moved as completed successfully. This final `--verify` will also have the effect of removing any reassignment throttles. You can now delete the revert file if you saved the JSON for reverting the assignment to their original brokers.
- Once all the partition reassessments have finished, the broker(s) being removed should not have responsibility for any of the partitions in the cluster. You can verify this by checking that the broker's data log directory does not contain any live partition logs. If the log directory on the broker contains a directory that does not match the extended regular expression `[a-zA-Z0-9.-]+\.+[a-zA-Z0-9]+-delete$` then the broker still has live partitions and it should not be stopped.

You can check this by executing the command:

```
kubectl exec my-cluster-kafka-0 -c kafka -it -- \
/bin/bash -c \
"ls -l /var/lib/kafka/kafka-log_<N>_ | grep -E '^d' | grep -vE '[a-zA-Z0-9.-]+\.+[a-zA-Z0-9]+-delete$'"
```

where *N* is the number of the `Pod(s)` being deleted.

If the above command prints any output then the broker still has live partitions. In this case, either the reassignment has not finished, or the reassignment JSON file was incorrect.

- Once you have confirmed that the broker has no live partitions you can edit the `Kafka.spec.kafka.replicas` of your `Kafka` resource, which will scale down the `StatefulSet`, deleting the highest numbered broker `Pod(s)`.

## 2.1.5. Retrieving JMX metrics with JmxTrans

`JmxTrans` is a tool for retrieving JMX metrics data from Java processes and pushing that data, in various formats, to remote sinks inside or outside the cluster. `JmxTrans` can communicate with a secure JMX port.

Strimzi supports using `JmxTrans` to read JMX metrics from Kafka brokers.

`JmxTrans` reads JMX metrics data from secure or insecure Kafka brokers and pushes the data to remote sinks in various data formats. For example, `JmxTrans` can obtain JMX metrics about the request rate of each Kafka broker's network and then push the data to a Logstash database outside the Kubernetes cluster.

### Configuring a `JmxTrans` deployment

#### Prerequisites

- A running Kubernetes cluster

You can configure a `JmxTrans` deployment by using the `Kafka.spec.jmxTrans` property. A `JmxTrans` deployment can read from a secure or insecure Kafka broker. To configure a `JmxTrans` deployment, define the following properties:

- `Kafka.spec.jmxTrans.outputDefinitions`
- `Kafka.spec.jmxTrans.kafkaQueries`

For more information on these properties, see the [JmxTransSpec schema reference](#).

**NOTE** To use JMXTrans, `jmxOptions` must be configured on the Kafka broker.

### Configuring JmxTrans output definitions

Output definitions specify where JMX metrics are pushed to, and in which data format. For information about supported data formats, see [Data formats](#). How many seconds JmxTrans agent waits for before pushing new data can be configured through the `flushDelay` property. The `host` and `port` properties define the target host address and target port the data is pushed to. The `name` property is a required property that is referenced by the `Kafka.spec.jmxTrans.kafkaQueries` property.

Here is an example configuration pushing JMX data in the Graphite format every 5 seconds to a Logstash database on `http://myLogstash:9999`, and another pushing to `standardOut` (standard output):

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  jmxTrans:
    outputDefinitions:
      - outputType: "com.googlecode.jmxtrans.model.output.GraphiteWriter"
        host: "http://myLogstash"
        port: 9999
        flushDelay: 5
        name: "logstash"
      - outputType: "com.googlecode.jmxtrans.model.output.StdOutWriter"
        name: "standardOut"
        # ...
      # ...
  zookeeper:
    # ...
```

### Configuring JmxTrans queries

JmxTrans queries specify what JMX metrics are read from the Kafka brokers. Currently JmxTrans queries can only be sent to the Kafka Brokers. Configure the `targetMBean` property to specify which target MBean on the Kafka broker is addressed. Configuring the `attributes` property specifies which MBean attribute is read as JMX metrics from the target MBean. JmxTrans supports wildcards to read from target MBeans, and filter by specifying the `typenames`. The `outputs` property defines where the metrics are pushed to by specifying the name of the output definitions.

The following JmxTrans deployment reads from all MBeans that match the pattern

`kafka.server:type=BrokerTopicMetrics,name=*` and have `name` in the target MBean's name. From those Mbeans, it obtains JMX metrics about the `Count` attribute and pushes the metrics to standard output as defined by `outputs`.

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  # ...
  jmxTrans:
    kafkaQueries:
      - targetMBean: "kafka.server:type=BrokerTopicMetrics,*"
        typeNames: ["name"]
        attributes: ["Count"]
        outputs: ["standardOut"]
  zookeeper:
    # ...
```

## Additional resources

For more information about JmxTrans, see the [JmxTrans github](#).

### 2.1.6. Maintenance time windows for rolling updates

Maintenance time windows allow you to schedule certain rolling updates of your Kafka and ZooKeeper clusters to start at a convenient time.

#### Maintenance time windows overview

In most cases, the Cluster Operator only updates your Kafka or ZooKeeper clusters in response to changes to the corresponding `Kafka` resource. This enables you to plan when to apply changes to a `Kafka` resource to minimize the impact on Kafka client applications.

However, some updates to your Kafka and ZooKeeper clusters can happen without any corresponding change to the `Kafka` resource. For example, the Cluster Operator will need to perform a rolling restart if a CA (Certificate Authority) certificate that it manages is close to expiry.

While a rolling restart of the pods should not affect *availability* of the service (assuming correct broker and topic configurations), it could affect *performance* of the Kafka client applications. Maintenance time windows allow you to schedule such spontaneous rolling updates of your Kafka and ZooKeeper clusters to start at a convenient time. If maintenance time windows are not configured for a cluster then it is possible that such spontaneous rolling updates will happen at an inconvenient time, such as during a predictable period of high load.

#### Maintenance time window definition

You configure maintenance time windows by entering an array of strings in the `Kafka.spec.maintenanceTimeWindows` property. Each string is a [cron expression](#) interpreted as being

in UTC (Coordinated Universal Time, which for practical purposes is the same as Greenwich Mean Time).

The following example configures a single maintenance time window that starts at midnight and ends at 01:59am (UTC), on Sundays, Mondays, Tuesdays, Wednesdays, and Thursdays:

```
# ...
maintenanceTimeWindows:
- "* * 0-1 ? * SUN,MON,TUE,WED,THU *"
# ...
```

In practice, maintenance windows should be set in conjunction with the `Kafka.spec.clusterCa.renewalDays` and `Kafka.spec.clientsCa.renewalDays` properties of the `Kafka` resource, to ensure that the necessary CA certificate renewal can be completed in the configured maintenance time windows.

**NOTE**

Strimzi does not schedule maintenance operations exactly according to the given windows. Instead, for each reconciliation, it checks whether a maintenance window is currently "open". This means that the start of maintenance operations within a given time window can be delayed by up to the Cluster Operator reconciliation interval. Maintenance time windows must therefore be at least this long.

#### *Additional resources*

- For more information about the Cluster Operator configuration, see [Cluster Operator configuration](#).

### Configuring a maintenance time window

You can configure a maintenance time window for rolling updates triggered by supported processes.

#### *Prerequisites*

- A Kubernetes cluster.
- The Cluster Operator is running.

#### *Procedure*

1. Add or edit the `maintenanceTimeWindows` property in the `Kafka` resource. For example to allow maintenance between 0800 and 1059 and between 1400 and 1559 you would set the `maintenanceTimeWindows` as shown below:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  maintenanceTimeWindows:
    - "* * 8-10 * * ?"
    - "* * 14-15 * * ?"
```

## 2. Create or update the resource:

```
kubectl apply -f KAFKA-CONFIG-FILE
```

### *Additional resources*

Performing rolling updates:

- [Performing a rolling update using a StatefulSet annotation](#)
- [Performing a rolling update using a Pod annotation](#)

### **2.1.7. Connecting to ZooKeeper from a terminal**

Most Kafka CLI tools can connect directly to Kafka, so under normal circumstances you should not need to connect to ZooKeeper. ZooKeeper services are secured with encryption and authentication and are not intended to be used by external applications that are not part of Strimzi.

However, if you want to use Kafka CLI tools that require a connection to ZooKeeper, you can use a terminal inside a ZooKeeper container and connect to `localhost:12181` as the ZooKeeper address.

#### *Prerequisites*

- A Kubernetes cluster is available.
- A Kafka cluster is running.
- The Cluster Operator is running.

#### *Procedure*

##### 1. Open the terminal using the Kubernetes console or run the `exec` command from your CLI.

For example:

```
kubectl exec -ti my-cluster-zookeeper-0 -- bin/kafka-topics.sh --list --zookeeper localhost:12181
```

Be sure to use `localhost:12181`.

You can now run Kafka commands to ZooKeeper.

## 2.1.8. Deleting Kafka nodes manually

This procedure describes how to delete an existing Kafka node by using a Kubernetes annotation. Deleting a Kafka node consists of deleting both the `Pod` on which the Kafka broker is running and the related `PersistentVolumeClaim` (if the cluster was deployed with persistent storage). After deletion, the `Pod` and its related `PersistentVolumeClaim` are recreated automatically.

**WARNING**

Deleting a `PersistentVolumeClaim` can cause permanent data loss. The following procedure should only be performed if you have encountered storage issues.

*Prerequisites*

See the *Deploying and Upgrading Strimzi* guide for instructions on running a:

- [Cluster Operator](#)
- [Kafka cluster](#)

*Procedure*

1. Find the name of the `Pod` that you want to delete.

For example, if the cluster is named *cluster-name*, the pods are named *cluster-name-kafka-index*, where *index* starts at zero and ends at the total number of replicas.

2. Annotate the `Pod` resource in Kubernetes.

Use `kubectl annotate`:

```
kubectl annotate pod cluster-name-kafka-index strimzi.io/delete-pod-and-pvc=true
```

3. Wait for the next reconciliation, when the annotated pod with the underlying persistent volume claim will be deleted and then recreated.

## 2.1.9. Deleting ZooKeeper nodes manually

This procedure describes how to delete an existing ZooKeeper node by using a Kubernetes annotation. Deleting a ZooKeeper node consists of deleting both the `Pod` on which ZooKeeper is running and the related `PersistentVolumeClaim` (if the cluster was deployed with persistent storage). After deletion, the `Pod` and its related `PersistentVolumeClaim` are recreated automatically.

**WARNING**

Deleting a `PersistentVolumeClaim` can cause permanent data loss. The following procedure should only be performed if you have encountered storage issues.

*Prerequisites*

See the *Deploying and Upgrading Strimzi* guide for instructions on running a:

- Cluster Operator
- Kafka cluster

#### Procedure

1. Find the name of the Pod that you want to delete.

For example, if the cluster is named *cluster-name*, the pods are named *cluster-name-zookeeper-index*, where *index* starts at zero and ends at the total number of replicas.

2. Annotate the Pod resource in Kubernetes.

Use `kubectl annotate`:

```
kubectl annotate pod cluster-name-zookeeper-index strimzi.io/delete-pod-and-
pvc=true
```

3. Wait for the next reconciliation, when the annotated pod with the underlying persistent volume claim will be deleted and then recreated.

### 2.1.10. List of Kafka cluster resources

The following resources are created by the Cluster Operator in the Kubernetes cluster:

#### Shared resources

##### `cluster-name-cluster-ca`

Secret with the Cluster CA private key used to encrypt the cluster communication.

##### `cluster-name-cluster-ca-cert`

Secret with the Cluster CA public key. This key can be used to verify the identity of the Kafka brokers.

##### `cluster-name-clients-ca`

Secret with the Clients CA private key used to sign user certificates

##### `cluster-name-clients-ca-cert`

Secret with the Clients CA public key. This key can be used to verify the identity of the Kafka users.

##### `cluster-name-cluster-operator-certs`

Secret with Cluster operators keys for communication with Kafka and ZooKeeper.

#### Zookeeper nodes

##### `cluster-name-zookeeper`

StatefulSet which is in charge of managing the ZooKeeper node pods.

##### `cluster-name-zookeeper-idx`

Pods created by the Zookeeper StatefulSet.

## **cluster-name-zookeeper-nodes**

Headless Service needed to have DNS resolve the ZooKeeper pods IP addresses directly.

## **cluster-name-zookeeper-client**

Service used by Kafka brokers to connect to ZooKeeper nodes as clients.

## **cluster-name-zookeeper-config**

ConfigMap that contains the ZooKeeper ancillary configuration, and is mounted as a volume by the ZooKeeper node pods.

## **cluster-name-zookeeper-nodes**

Secret with ZooKeeper node keys.

## **cluster-name-zookeeper**

Service account used by the Zookeeper nodes.

## **cluster-name-zookeeper**

Pod Disruption Budget configured for the ZooKeeper nodes.

## **cluster-name-network-policy-zookeeper**

Network policy managing access to the ZooKeeper services.

## **data-cluster-name-zookeeper-idx**

Persistent Volume Claim for the volume used for storing data for the ZooKeeper node pod `idx`. This resource will be created only if persistent storage is selected for provisioning persistent volumes to store data.

## *Kafka brokers*

## **cluster-name-kafka**

StatefulSet which is in charge of managing the Kafka broker pods.

## **cluster-name-kafka-idx**

Pods created by the Kafka StatefulSet.

## **cluster-name-kafka-brokers**

Service needed to have DNS resolve the Kafka broker pods IP addresses directly.

## **cluster-name-kafka-bootstrap**

Service can be used as bootstrap servers for Kafka clients connecting from within the Kubernetes cluster.

## **cluster-name-kafka-external-bootstrap**

Bootstrap service for clients connecting from outside the Kubernetes cluster. This resource is created only when an external listener is enabled. The old service name will be used for backwards compatibility when the listener name is `external` and port is `9094`.

## **cluster-name-kafka-pod-id**

Service used to route traffic from outside the Kubernetes cluster to individual pods. This

resource is created only when an external listener is enabled. The old service name will be used for backwards compatibility when the listener name is `external` and port is `9094`.

#### **cluster-name-kafka-external-bootstrap**

Bootstrap route for clients connecting from outside the Kubernetes cluster. This resource is created only when an external listener is enabled and set to type `route`. The old route name will be used for backwards compatibility when the listener name is `external` and port is `9094`.

#### **cluster-name-kafka-pod-id**

Route for traffic from outside the Kubernetes cluster to individual pods. This resource is created only when an external listener is enabled and set to type `route`. The old route name will be used for backwards compatibility when the listener name is `external` and port is `9094`.

#### **cluster-name-kafka-listener-name-bootstrap**

Bootstrap service for clients connecting from outside the Kubernetes cluster. This resource is created only when an external listener is enabled. The new service name will be used for all other external listeners.

#### **cluster-name-kafka-listener-name-pod-id**

Service used to route traffic from outside the Kubernetes cluster to individual pods. This resource is created only when an external listener is enabled. The new service name will be used for all other external listeners.

#### **cluster-name-kafka-listener-name-bootstrap**

Bootstrap route for clients connecting from outside the Kubernetes cluster. This resource is created only when an external listener is enabled and set to type `route`. The new route name will be used for all other external listeners.

#### **cluster-name-kafka-listener-name-pod-id**

Route for traffic from outside the Kubernetes cluster to individual pods. This resource is created only when an external listener is enabled and set to type `route`. The new route name will be used for all other external listeners.

#### **cluster-name-kafka-config**

ConfigMap which contains the Kafka ancillary configuration and is mounted as a volume by the Kafka broker pods.

#### **cluster-name-kafka-brokers**

Secret with Kafka broker keys.

#### **cluster-name-kafka**

Service account used by the Kafka brokers.

#### **cluster-name-kafka**

Pod Disruption Budget configured for the Kafka brokers.

#### **cluster-name-network-policy-kafka**

Network policy managing access to the Kafka services.

## **strimzi-namespace-name-cluster-name-kafka-init**

Cluster role binding used by the Kafka brokers.

## **cluster-name-jmx**

Secret with JMX username and password used to secure the Kafka broker port. This resource is created only when JMX is enabled in Kafka.

## **data-cluster-name-kafka-idx**

Persistent Volume Claim for the volume used for storing data for the Kafka broker pod `idx`. This resource is created only if persistent storage is selected for provisioning persistent volumes to store data.

## **data-id-cluster-name-kafka-idx**

Persistent Volume Claim for the volume `id` used for storing data for the Kafka broker pod `idx`. This resource is created only if persistent storage is selected for JBOD volumes when provisioning persistent volumes to store data.

### *Entity Operator*

These resources are only created if the Entity Operator is deployed using the Cluster Operator.

## **cluster-name-entity-operator**

Deployment with Topic and User Operators.

## **cluster-name-entity-operator-random-string**

Pod created by the Entity Operator deployment.

## **cluster-name-entity-topic-operator-config**

ConfigMap with ancillary configuration for Topic Operators.

## **cluster-name-entity-user-operator-config**

ConfigMap with ancillary configuration for User Operators.

## **cluster-name-entity-operator-certs**

Secret with Entity Operator keys for communication with Kafka and ZooKeeper.

## **cluster-name-entity-operator**

Service account used by the Entity Operator.

## **strimzi-cluster-name-entity-topic-operator**

Role binding used by the Entity Topic Operator.

## **strimzi-cluster-name-entity-user-operator**

Role binding used by the Entity User Operator.

### *Kafka Exporter*

These resources are only created if the Kafka Exporter is deployed using the Cluster Operator.

**cluster-name-kafka-exporter**

Deployment with Kafka Exporter.

**cluster-name-kafka-exporter-random-string**

Pod created by the Kafka Exporter deployment.

**cluster-name-kafka-exporter**

Service used to collect consumer lag metrics.

**cluster-name-kafka-exporter**

Service account used by the Kafka Exporter.

*Cruise Control*

These resources are only created if Cruise Control was deployed using the Cluster Operator.

**cluster-name-cruise-control**

Deployment with Cruise Control.

**cluster-name-cruise-control-random-string**

Pod created by the Cruise Control deployment.

**cluster-name-cruise-control-config**

ConfigMap that contains the Cruise Control ancillary configuration, and is mounted as a volume by the Cruise Control pods.

**cluster-name-cruise-control-certs**

Secret with Cruise Control keys for communication with Kafka and ZooKeeper.

**cluster-name-cruise-control**

Service used to communicate with Cruise Control.

**cluster-name-cruise-control**

Service account used by Cruise Control.

**cluster-name-network-policy-cruise-control**

Network policy managing access to the Cruise Control service.

*JMXTrans*

These resources are only created if JMXTrans is deployed using the Cluster Operator.

**cluster-name-jmxtrans**

Deployment with JMXTrans.

**cluster-name-jmxtrans-random-string**

Pod created by the JMXTrans deployment.

**cluster-name-jmxtrans-config**

ConfigMap that contains the JMXTrans ancillary configuration, and is mounted as a volume by

the JMXTrans pods.

#### **cluster-name-jmxtrans**

Service account used by JMXTrans.

## 2.2. Kafka Connect/S2I cluster configuration

This section describes how to configure a Kafka Connect or Kafka Connect with Source-to-Image (S2I) deployment in your Strimzi cluster.

Kafka Connect is an integration toolkit for streaming data between Kafka brokers and other systems using [Connector](#) plugins. Kafka Connect provides a framework for integrating Kafka with an external data source or target, such as a database, for import or export of data using connectors. Connectors are plugins that provide the connection configuration needed.

If you are using Kafka Connect, you configure either the [KafkaConnect](#) or the [KafkaConnectS2I](#) resource. Use the [KafkaConnectS2I](#) resource if you are using the [Source-to-Image \(S2I\)](#) framework to deploy Kafka Connect.

- The full schema of the [KafkaConnect](#) resource is described in [KafkaConnect schema reference](#).
- The full schema of the [KafkaConnectS2I](#) resource is described in [KafkaConnectS2I schema reference](#).

#### *Additional resources*

- [Creating and managing connectors](#)
- [Deploying a KafkaConnector resource to Kafka Connect](#)

### 2.2.1. Configuring Kafka Connect

Use Kafka Connect to set up external data connections to your Kafka cluster.

Use the properties of the [KafkaConnect](#) or [KafkaConnectS2I](#) resource to configure your Kafka Connect deployment. The example shown in this procedure is for the [KafkaConnect](#) resource, but the properties are the same for the [KafkaConnectS2I](#) resource.

#### *Kafka connector configuration*

[KafkaConnector](#) resources allow you to create and manage connector instances for Kafka Connect in a Kubernetes-native way.

In the configuration, you enable [KafkaConnectors](#) for a Kafka Connect cluster by adding the `strimzi.io/use-connector-resources` annotation. You can also specify external configuration for Kafka Connect connectors through the `externalConfiguration` property.

Connectors are created, reconfigured, and deleted using the Kafka Connect HTTP REST interface, or by using [KafkaConnectors](#). For more information on these methods, see [Creating and managing connectors](#) in the *Deploying and Upgrading Strimzi* guide.

The connector configuration is passed to Kafka Connect as part of an HTTP request and stored

within Kafka itself. ConfigMaps and Secrets are standard Kubernetes resources used for storing configurations and confidential data. You can use ConfigMaps and Secrets to configure certain elements of a connector. You can then reference the configuration values in HTTP REST commands (this keeps the configuration separate and more secure, if needed). This method applies especially to confidential data, such as usernames, passwords, or certificates.

#### Prerequisites

- A Kubernetes cluster
- A running Cluster Operator

See the *Deploying and Upgrading Strimzi* guide for instructions on running a:

- [Cluster Operator](#)
- [Kafka cluster](#)

#### Procedure

1. Edit the `spec` properties for the `KafkaConnect` or `KafkaConnectS2I` resource.

The properties you can configure are shown in this example configuration:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect ①
metadata:
  name: my-connect-cluster
  annotations:
    strimzi.io/use-connector-resources: "true" ②
spec:
  replicas: 3 ③
  authentication: ④
    type: tls
    certificateAndKey:
      certificate: source.crt
      key: source.key
      secretName: my-user-source
  bootstrapServers: my-cluster-kafka-bootstrap:9092 ⑤
  tls: ⑥
    trustedCertificates:
      - secretName: my-cluster-cluster-cert
        certificate: ca.crt
      - secretName: my-cluster-cluster-cert
        certificate: ca2.crt
  config: ⑦
    group.id: my-connect-cluster
    offset.storage.topic: my-connect-cluster-offsets
    config.storage.topic: my-connect-cluster-configs
    status.storage.topic: my-connect-cluster-status
    key.converter: org.apache.kafka.connect.json.JsonConverter
    value.converter: org.apache.kafka.connect.json.JsonConverter
    key.converter.schemas.enable: true
```

```

value.converter.schemas.enable: true
config.storage.replication.factor: 3
offset.storage.replication.factor: 3
status.storage.replication.factor: 3
externalConfiguration: ⑧
env:
  - name: AWS_ACCESS_KEY_ID
    valueFrom:
      secretKeyRef:
        name: aws-creds
        key: awsAccessKey
  - name: AWS_SECRET_ACCESS_KEY
    valueFrom:
      secretKeyRef:
        name: aws-creds
        key: awsSecretAccessKey
resources: ⑨
  requests:
    cpu: "1"
    memory: 2Gi
  limits:
    cpu: "2"
    memory: 2Gi
logging: ⑩
  type: inline
  loggers:
    log4j.rootLogger: "INFO"
readinessProbe: ⑪
  initialDelaySeconds: 15
  timeoutSeconds: 5
livenessProbe:
  initialDelaySeconds: 15
  timeoutSeconds: 5
metricsConfig: ⑫
  type: jmxPrometheusExporter
  valueFrom:
    configMapKeyRef:
      name: my-config-map
      key: my-key
jvmOptions: ⑬
  "-Xmx": "1g"
  "-Xms": "1g"
image: my-org/my-image:latest ⑭
rack:
  topologyKey: topology.kubernetes.io/zone ⑮
template: ⑯
pod:
  affinity:
    podAntiAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        - labelSelector:

```

```

matchExpressions:
  - key: application
    operator: In
    values:
      - postgresql
      - mongodb
    topologyKey: "kubernetes.io/hostname"
connectContainer: ⑯
env:
  - name: JAEGER_SERVICE_NAME
    value: my-jaeger-service
  - name: JAEGER_AGENT_HOST
    value: jaeger-agent-name
  - name: JAEGER_AGENT_PORT
    value: "6831"

```

- ① Use [KafkaConnect](#) or [KafkaConnectS2I](#), as required.
- ② Enables [KafkaConnectors](#) for the Kafka Connect cluster.
- ③ [The number of replica nodes](#).
- ④ Authentication for the Kafka Connect cluster, using the [TLS mechanism](#), as shown here, using [OAuth bearer tokens](#), or a SASL-based [SCRAM-SHA-512](#) or [PLAIN](#) mechanism. By default, Kafka Connect connects to Kafka brokers using a plain text connection.
- ⑤ [Bootstrap server](#) for connection to the Kafka Connect cluster.
- ⑥ [TLS encryption](#) with key names under which TLS certificates are stored in X.509 format for the cluster. If certificates are stored in the same secret, it can be listed multiple times.
- ⑦ [Kafka Connect configuration](#) of *workers* (not connectors). Standard Apache Kafka configuration may be provided, restricted to those properties not managed directly by Strimzi.
- ⑧ [External configuration for Kafka connectors](#) using environment variables, as shown here, or volumes.
- ⑨ Requests for reservation of [supported resources](#), currently [cpu](#) and [memory](#), and limits to specify the maximum resources that can be consumed.
- ⑩ Specified [Kafka Connect loggers and log levels](#) added directly ([inline](#)) or indirectly ([external](#)) through a ConfigMap. A custom ConfigMap must be placed under the [log4j.properties](#) or [log4j2.properties](#) key. For the Kafka Connect [log4j.rootLogger](#) logger, you can set the log level to INFO, ERROR, WARN, TRACE, DEBUG, FATAL or OFF.
- ⑪ [Healthchecks](#) to know when to restart a container (liveness) and when a container can accept traffic (readiness).
- ⑫ [Prometheus metrics](#), which are enabled by referencing a ConfigMap containing configuration for the Prometheus JMX exporter in this example. You can enable metrics without further configuration using a reference to a ConfigMap containing an empty file under [metricsConfig.valueFrom.configMapKeyRef.key](#).
- ⑬ [JVM configuration options](#) to optimize performance for the Virtual Machine (VM) running Kafka Connect.

- ⑯ ADVANCED OPTION: [Container image configuration](#), which is recommended only in special situations.
- ⑰ Rack awareness is configured to spread replicas across different racks. A [topologykey](#) must match the label of a cluster node.
- ⑱ Template customization. Here a pod is scheduled with anti-affinity, so the pod is not scheduled on nodes with the same hostname.
- ⑲ Environment variables are also [set for distributed tracing using Jaeger](#).

2. Create or update the resource:

```
kubectl apply -f KAFKA-CONNECT-CONFIG-FILE
```

3. If authorization is enabled for Kafka Connect, [configure Kafka Connect users to enable access to the Kafka Connect consumer group and topics](#).

### 2.2.2. Kafka Connect configuration for multiple instances

If you are running multiple instances of Kafka Connect, you have to change the default configuration of the following [config](#) properties:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect
spec:
  # ...
  config:
    group.id: connect-cluster ①
    offset.storage.topic: connect-cluster-offsets ②
    config.storage.topic: connect-cluster-configs ③
    status.storage.topic: connect-cluster-status ④
    # ...
# ...
```

- ① Kafka Connect cluster group that the instance belongs to.
- ② Kafka topic that stores connector offsets.
- ③ Kafka topic that stores connector and task status configurations.
- ④ Kafka topic that stores connector and task status updates.

**NOTE**

Values for the three topics must be the same for all Kafka Connect instances with the same [group.id](#).

Unless you change the default settings, each Kafka Connect instance connecting to the same Kafka cluster is deployed with the same values. What happens, in effect, is all instances are coupled to run in a cluster and use the same topics.

If multiple Kafka Connect clusters try to use the same topics, Kafka Connect will not work as expected and generate errors.

If you wish to run multiple Kafka Connect instances, change the values of these properties for each instance.

### 2.2.3. Configuring Kafka Connect user authorization

This procedure describes how to authorize user access to Kafka Connect.

When any type of authorization is being used in Kafka, a Kafka Connect user requires read/write access rights to the consumer group and the internal topics of Kafka Connect.

The properties for the consumer group and internal topics are automatically configured by Strimzi, or they can be specified explicitly in the `spec` of the [KafkaConnect](#) or [KafkaConnectS2I](#) resource.

*Example configuration properties in the [KafkaConnect](#) resource*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect
spec:
  # ...
  config:
    group.id: my-connect-cluster ①
    offset.storage.topic: my-connect-cluster-offsets ②
    config.storage.topic: my-connect-cluster-configs ③
    status.storage.topic: my-connect-cluster-status ④
    # ...
  # ...
```

① Kafka Connect cluster group that the instance belongs to.

② Kafka topic that stores connector offsets.

③ Kafka topic that stores connector and task status configurations.

④ Kafka topic that stores connector and task status updates.

This procedure shows how access is provided when `simple` authorization is being used.

Simple authorization uses ACL rules, handled by the Kafka [AclAuthorizer](#) plugin, to provide the right level of access. For more information on configuring a [KafkaUser](#) resource to use simple authorization, see the [AclRule schema reference](#).

**NOTE**

The default values for the consumer group and topics will differ when [running multiple instances](#).

#### Prerequisites

- A Kubernetes cluster

- A running Cluster Operator

#### *Procedure*

1. Edit the `authorization` property in the `KafkaUser` resource to provide access rights to the user.

In the following example, access rights are configured for the Kafka Connect topics and consumer group using `literal` name values:

Property	Name
<code>offset.storage.topic</code>	<code>connect-cluster-offsets</code>
<code>status.storage.topic</code>	<code>connect-cluster-status</code>
<code>config.storage.topic</code>	<code>connect-cluster-configs</code>
<code>group</code>	<code>connect-cluster</code>

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: my-user
  labels:
    strimzi.io/cluster: my-cluster
spec:
  # ...
  authorization:
    type: simple
    acls:
      # access to offset.storage.topic
      - resource:
          type: topic
          name: connect-cluster-offsets
          patternType: literal
          operation: Write
          host: "*"
      - resource:
          type: topic
          name: connect-cluster-offsets
          patternType: literal
          operation: Create
          host: "*"
      - resource:
          type: topic
          name: connect-cluster-offsets
          patternType: literal
          operation: Describe
          host: "*"
      - resource:
          type: topic
          name: connect-cluster-offsets
          patternType: literal
          operation: Read
```

```
host: "*"
# access to status.storage.topic
- resource:
    type: topic
    name: connect-cluster-status
    patternType: literal
    operation: Write
    host: "*"
- resource:
    type: topic
    name: connect-cluster-status
    patternType: literal
    operation: Create
    host: "*"
- resource:
    type: topic
    name: connect-cluster-status
    patternType: literal
    operation: Describe
    host: "*"
- resource:
    type: topic
    name: connect-cluster-status
    patternType: literal
    operation: Read
    host: "*"
# access to config.storage.topic
- resource:
    type: topic
    name: connect-cluster-configs
    patternType: literal
    operation: Write
    host: "*"
- resource:
    type: topic
    name: connect-cluster-configs
    patternType: literal
    operation: Create
    host: "*"
- resource:
    type: topic
    name: connect-cluster-configs
    patternType: literal
    operation: Describe
    host: "*"
- resource:
    type: topic
    name: connect-cluster-configs
    patternType: literal
    operation: Read
    host: "*"
```

```
# consumer group
- resource:
  type: group
  name: connect-cluster
  patternType: literal
  operation: Read
  host: "*"
```

2. Create or update the resource.

```
kubectl apply -f KAFKA-USER-CONFIG-FILE
```

## 2.2.4. List of Kafka Connect cluster resources

The following resources are created by the Cluster Operator in the Kubernetes cluster:

### ***connect-cluster-name-connect***

Deployment which is in charge to create the Kafka Connect worker node pods.

### ***connect-cluster-name-connect-api***

Service which exposes the REST interface for managing the Kafka Connect cluster.

### ***connect-cluster-name-config***

ConfigMap which contains the Kafka Connect ancillary configuration and is mounted as a volume by the Kafka broker pods.

### ***connect-cluster-name-connect***

Pod Disruption Budget configured for the Kafka Connect worker nodes.

## 2.2.5. List of Kafka Connect (S2I) cluster resources

The following resources are created by the Cluster Operator in the Kubernetes cluster:

### ***connect-cluster-name-connect-source***

ImageStream which is used as the base image for the newly-built Docker images.

### ***connect-cluster-name-connect***

BuildConfig which is responsible for building the new Kafka Connect Docker images.

### ***connect-cluster-name-connect***

ImageStream where the newly built Docker images will be pushed.

### ***connect-cluster-name-connect***

DeploymentConfig which is in charge of creating the Kafka Connect worker node pods.

### ***connect-cluster-name-connect-api***

Service which exposes the REST interface for managing the Kafka Connect cluster.

### **connect-cluster-name-config**

ConfigMap which contains the Kafka Connect ancillary configuration and is mounted as a volume by the Kafka broker pods.

### **connect-cluster-name-connect**

Pod Disruption Budget configured for the Kafka Connect worker nodes.

## **2.3. Kafka MirrorMaker cluster configuration**

This chapter describes how to configure a Kafka MirrorMaker deployment in your Strimzi cluster to replicate data between Kafka clusters.

You can use Strimzi with MirrorMaker or [MirrorMaker 2.0](#). MirrorMaker 2.0 is the latest version, and offers a more efficient way to mirror data between Kafka clusters.

If you are using MirrorMaker, you configure the [KafkaMirrorMaker](#) resource.

The following procedure shows how the resource is configured:

- [Configuring Kafka MirrorMaker](#)

The full schema of the [KafkaMirrorMaker](#) resource is described in the [KafkaMirrorMaker schema reference](#).

### **2.3.1. Configuring Kafka MirrorMaker**

Use the properties of the [KafkaMirrorMaker](#) resource to configure your Kafka MirrorMaker deployment.

You can configure access control for producers and consumers using TLS or SASL authentication. This procedure shows a configuration that uses TLS encryption and authentication on the consumer and producer side.

#### *Prerequisites*

- See the [Deploying and Upgrading Strimzi](#) guide for instructions on running a:
  - [Cluster Operator](#)
  - [Kafka cluster](#)
- Source and target Kafka clusters must be available

#### *Procedure*

1. Edit the [spec](#) properties for the [KafkaMirrorMaker](#) resource.

The properties you can configure are shown in this example configuration:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaMirrorMaker
metadata:
  name: my-mirror-maker
```

```

spec:
  replicas: 3 ①
  consumer:
    bootstrapServers: my-source-cluster-kafka-bootstrap:9092 ②
    groupId: "my-group" ③
    numStreams: 2 ④
    offsetCommitInterval: 120000 ⑤
    tls: ⑥
      trustedCertificates:
        - secretName: my-source-cluster-ca-cert
          certificate: ca.crt
    authentication: ⑦
      type: tls
      certificateAndKey:
        secretName: my-source-secret
        certificate: public.crt
        key: private.key
    config: ⑧
      max.poll.records: 100
      receive.buffer.bytes: 32768
      ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384" ⑨
      ssl.enabled.protocols: "TLSv1.2"
      ssl.protocol: "TLSv1.2"
      ssl.endpoint.identification.algorithm: HTTPS ⑩
  producer:
    bootstrapServers: my-target-cluster-kafka-bootstrap:9092
    abortOnSendFailure: false ⑪
    tls:
      trustedCertificates:
        - secretName: my-target-cluster-ca-cert
          certificate: ca.crt
    authentication:
      type: tls
      certificateAndKey:
        secretName: my-target-secret
        certificate: public.crt
        key: private.key
    config:
      compression.type: gzip
      batch.size: 8192
      ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384" ⑫
      ssl.enabled.protocols: "TLSv1.2"
      ssl.protocol: "TLSv1.2"
      ssl.endpoint.identification.algorithm: HTTPS ⑬
  whitelist: "my-topic|other-topic" ⑭
  resources: ⑮
    requests:
      cpu: "1"
      memory: 2Gi
    limits:
      cpu: "2"

```

```

    memory: 2Gi
  logging: ⑯
    type: inline
  loggers:
    mirrormaker.root.logger: "INFO"
  readinessProbe: ⑰
    initialDelaySeconds: 15
    timeoutSeconds: 5
  livenessProbe:
    initialDelaySeconds: 15
    timeoutSeconds: 5
  metricsConfig: ⑱
    type: jmxPrometheusExporter
  valueFrom:
    configMapKeyRef:
      name: my-config-map
      key: my-key
  jvmOptions: ⑲
    "-Xmx": "1g"
    "-Xms": "1g"
  image: my-org/my-image:latest ⑳
  template:
    pod:
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            - labelSelector:
                matchExpressions:
                  - key: application
                    operator: In
                    values:
                      - postgresql
                      - mongodb
              topologyKey: "kubernetes.io/hostname"
  connectContainer:
    env:
      - name: JAEGER_SERVICE_NAME
        value: my-jaeger-service
      - name: JAEGER_AGENT_HOST
        value: jaeger-agent-name
      - name: JAEGER_AGENT_PORT
        value: "6831"
  tracing:
    type: jaeger

```

① The number of replica nodes.

② Bootstrap servers for consumer and producer.

③ Group ID for the consumer.

④ The number of consumer streams.

- ⑤ The offset auto-commit interval in milliseconds.
- ⑥ TLS encryption with key names under which TLS certificates are stored in X.509 format for consumer or producer. If certificates are stored in the same secret, it can be listed multiple times.
- ⑦ Authentication for consumer or producer, using the [TLS mechanism](#), as shown here, using [OAuth bearer tokens](#), or a SASL-based [SCRAM-SHA-512](#) or [PLAIN](#) mechanism.
- ⑧ Kafka configuration options for [consumer](#) and [producer](#).
- ⑨ SSL properties for external listeners to run with a specific *cipher suite* for a TLS version.
- ⑩ Hostname verification is enabled by setting to [HTTPS](#). An empty string disables the verification.
- ⑪ If the `abortOnSendFailure` property is set to `true`, Kafka MirrorMaker will exit and the container will restart following a send failure for a message.
- ⑫ SSL properties for external listeners to run with a specific *cipher suite* for a TLS version.
- ⑬ Hostname verification is enabled by setting to [HTTPS](#). An empty string disables the verification.
- ⑭ A [whitelist](#) of topics mirrored from source to target Kafka cluster.
- ⑮ Requests for reservation of [supported resources](#), currently `cpu` and `memory`, and limits to specify the maximum resources that can be consumed.
- ⑯ Specified [loggers and log levels](#) added directly ([inline](#)) or indirectly ([external](#)) through a ConfigMap. A custom ConfigMap must be placed under the `log4j.properties` or `log4j2.properties` key. MirrorMaker has a single logger called `mirrormaker.root.logger`. You can set the log level to INFO, ERROR, WARN, TRACE, DEBUG, FATAL or OFF.
- ⑰ [Healthchecks](#) to know when to restart a container (liveness) and when a container can accept traffic (readiness).
- ⑱ [Prometheus metrics](#), which are enabled by referencing a ConfigMap containing configuration for the Prometheus JMX exporter in this example. You can enable metrics without further configuration using a reference to a ConfigMap containing an empty file under `metricsConfig.valueFrom.configMapKeyRef.key`.
- ⑲ [JVM configuration options](#) to optimize performance for the Virtual Machine (VM) running Kafka MirrorMaker.
- ⑳ ADVANCED OPTION: [Container image configuration](#), which is recommended only in special situations.

[Template customization](#). Here a pod is scheduled with anti-affinity, so the pod is not scheduled on nodes with the same hostname.

Environment variables are also [set for distributed tracing using Jaeger](#).

[Distributed tracing is enabled for Jaeger](#).

**WARNING**

With the `abortOnSendFailure` property set to `false`, the producer attempts to send the next message in a topic. The original message might be lost, as there is no attempt to resend a failed message.

## 2. Create or update the resource:

```
kubectl apply -f <your-file>
```

### 2.3.2. List of Kafka MirrorMaker cluster resources

The following resources are created by the Cluster Operator in the Kubernetes cluster:

#### **<mirror-maker-name>-mirror-maker**

Deployment which is responsible for creating the Kafka MirrorMaker pods.

#### **<mirror-maker-name>-config**

ConfigMap which contains ancillary configuration for the Kafka MirrorMaker, and is mounted as a volume by the Kafka broker pods.

#### **<mirror-maker-name>-mirror-maker**

Pod Disruption Budget configured for the Kafka MirrorMaker worker nodes.

## 2.4. Kafka MirrorMaker 2.0 cluster configuration

This section describes how to configure a Kafka MirrorMaker 2.0 deployment in your Strimzi cluster.

MirrorMaker 2.0 is used to replicate data between two or more active Kafka clusters, within or across data centers.

Data replication across clusters supports scenarios that require:

- Recovery of data in the event of a system failure
- Aggregation of data for analysis
- Restriction of data access to a specific cluster
- Provision of data at a specific location to improve latency

If you are using MirrorMaker 2.0, you configure the [KafkaMirrorMaker2](#) resource.

MirrorMaker 2.0 introduces an entirely new way of replicating data between clusters.

As a result, the resource configuration differs from the previous version of MirrorMaker. If you choose to use MirrorMaker 2.0, there is currently no legacy support, so any resources must be manually converted into the new format.

How MirrorMaker 2.0 replicates data is described here:

- [MirrorMaker 2.0 data replication](#)

The following procedure shows how the resource is configured for MirrorMaker 2.0:

- Synchronizing data between Kafka clusters

The full schema of the [KafkaMirrorMaker2](#) resource is described in the [KafkaMirrorMaker2 schema reference](#).

### 2.4.1. MirrorMaker 2.0 data replication

MirrorMaker 2.0 consumes messages from a source Kafka cluster and writes them to a target Kafka cluster.


MirrorMaker 2.0 uses:

- Source cluster configuration to consume data from the source cluster
- Target cluster configuration to output data to the target cluster

MirrorMaker 2.0 is based on the Kafka Connect framework, *connectors* managing the transfer of data between clusters. A MirrorMaker 2.0 [MirrorSourceConnector](#) replicates topics from a source cluster to a target cluster.

The process of *mirroring* data from one cluster to another cluster is asynchronous. The recommended pattern is for messages to be produced locally alongside the source Kafka cluster, then consumed remotely close to the target Kafka cluster.

MirrorMaker 2.0 can be used with more than one source cluster.



AMQ\_73\_0220

Figure 1. Replication across two clusters

### 2.4.2. Cluster configuration

You can use MirrorMaker 2.0 in *active/passive* or *active/active* cluster configurations.

- In an *active/active* configuration, both clusters are active and provide the same data simultaneously, which is useful if you want to make the same data available locally in different geographical locations.
- In an *active/passive* configuration, the data from an active cluster is replicated in a passive cluster, which remains on standby, for example, for data recovery in the event of system failure.

The expectation is that producers and consumers connect to active clusters only.

A MirrorMaker 2.0 cluster is required at each target destination.

### Bidirectional replication (active/active)

The MirrorMaker 2.0 architecture supports bidirectional replication in an *active/active* cluster configuration.

Each cluster replicates the data of the other cluster using the concept of *source* and *remote* topics. As the same topics are stored in each cluster, remote topics are automatically renamed by MirrorMaker 2.0 to represent the source cluster. The name of the originating cluster is prepended to the name of the topic.




Figure 2. Topic renaming

By flagging the originating cluster, topics are not replicated back to that cluster.

The concept of replication through *remote* topics is useful when configuring an architecture that requires data aggregation. Consumers can subscribe to source and remote topics within the same cluster, without the need for a separate aggregation cluster.

### Unidirectional replication (active/passive)

The MirrorMaker 2.0 architecture supports unidirectional replication in an *active/passive* cluster configuration.

You can use an *active/passive* cluster configuration to make backups or migrate data to another cluster. In this situation, you might not want automatic renaming of remote topics.

You can override automatic renaming by adding `IdentityReplicationPolicy` to the source connector

configuration of the [KafkaMirrorMaker2](#) resource. With this configuration applied, topics retain their original names.

## Topic configuration synchronization

Topic configuration is automatically synchronized between source and target clusters. By synchronizing configuration properties, the need for rebalancing is reduced.

## Data integrity

MirrorMaker 2.0 monitors source topics and propagates any configuration changes to remote topics, checking for and creating missing partitions. Only MirrorMaker 2.0 can write to remote topics.

## Offset tracking

MirrorMaker 2.0 tracks offsets for consumer groups using *internal topics*.

- The *offset sync* topic maps the source and target offsets for replicated topic partitions from record metadata
- The *checkpoint* topic maps the last committed offset in the source and target cluster for replicated topic partitions in each consumer group

Offsets for the *checkpoint* topic are tracked at predetermined intervals through configuration. Both topics enable replication to be fully restored from the correct offset position on failover.

MirrorMaker 2.0 uses its [MirrorCheckpointConnector](#) to emit *checkpoints* for offset tracking.

## Connectivity checks

A *heartbeat* internal topic checks connectivity between clusters.

The *heartbeat* topic is replicated from the source cluster.

Target clusters use the topic to check:

- The connector managing connectivity between clusters is running
- The source cluster is available

MirrorMaker 2.0 uses its [MirrorHeartbeatConnector](#) to emit *heartbeats* that perform these checks.

### 2.4.3. ACL rules synchronization

ACL access to remote topics is possible if you are **not** using the User Operator.

If [AclAuthorizer](#) is being used, without the User Operator, ACL rules that manage access to brokers also apply to remote topics. Users that can read a source topic can read its remote equivalent.

**NOTE**    OAuth 2.0 authorization does not support access to remote topics in this way.

## 2.4.4. Synchronizing data between Kafka clusters using MirrorMaker 2.0

Use MirrorMaker 2.0 to synchronize data between Kafka clusters through configuration.

The configuration must specify:

- Each Kafka cluster
- Connection information for each cluster, including TLS authentication
- The replication flow and direction
  - Cluster to cluster
  - Topic to topic

Use the properties of the [KafkaMirrorMaker2](#) resource to configure your Kafka MirrorMaker 2.0 deployment.

**NOTE** The previous version of MirrorMaker continues to be supported. If you wish to use the resources configured for the previous version, they must be updated to the format supported by MirrorMaker 2.0.

MirrorMaker 2.0 provides default configuration values for properties such as replication factors. A minimal configuration, with defaults left unchanged, would be something like this example:

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaMirrorMaker2
metadata:
  name: my-mirror-maker2
spec:
  version: 2.7.0
  connectCluster: "my-cluster-target"
  clusters:
    - alias: "my-cluster-source"
      bootstrapServers: my-cluster-source-kafka-bootstrap:9092
    - alias: "my-cluster-target"
      bootstrapServers: my-cluster-target-kafka-bootstrap:9092
  mirrors:
    - sourceCluster: "my-cluster-source"
      targetCluster: "my-cluster-target"
      sourceConnector: {}
```

You can configure access control for source and target clusters using TLS or SASL authentication. This procedure shows a configuration that uses TLS encryption and authentication for the source and target cluster.

### Prerequisites

- See the *Deploying and Upgrading Strimzi* guide for instructions on running a:
  - [Cluster Operator](#)

- Kafka cluster
- Source and target Kafka clusters must be available

*Procedure*

1. Edit the `spec` properties for the `KafkaMirrorMaker2` resource.

The properties you can configure are shown in this example configuration:

```

apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaMirrorMaker2
metadata:
  name: my-mirror-maker2
spec:
  version: 2.7.0 ①
  replicas: 3 ②
  connectCluster: "my-cluster-target" ③
  clusters:
    - alias: "my-cluster-source" ⑤
      authentication: ⑥
        certificateAndKey:
          certificate: source.crt
          key: source.key
          secretName: my-user-source
        type: tls
      bootstrapServers: my-cluster-source-kafka-bootstrap:9092 ⑦
      tls: ⑧
        trustedCertificates:
          - certificate: ca.crt
            secretName: my-cluster-source-cluster-ca-cert
    - alias: "my-cluster-target" ⑨
      authentication: ⑩
        certificateAndKey:
          certificate: target.crt
          key: target.key
          secretName: my-user-target
        type: tls
      bootstrapServers: my-cluster-target-kafka-bootstrap:9092 ⑪
      config: ⑫
        config.storage.replication.factor: 1
        offset.storage.replication.factor: 1
        status.storage.replication.factor: 1
        ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384" ⑬
        ssl.enabled.protocols: "TLSv1.2"
        ssl.protocol: "TLSv1.2"
        ssl.endpoint.identification.algorithm: HTTPS ⑭
      tls: ⑮
        trustedCertificates:
          - certificate: ca.crt
            secretName: my-cluster-target-cluster-ca-cert
  mirrors: ⑯

```

```

- sourceCluster: "my-cluster-source" ⑯
  targetCluster: "my-cluster-target" ⑰
  sourceConnector: ⑲
    config:
      replication.factor: 1 ⑳
      offset-syncs.topic.replication.factor: 1
      sync.topic.acls.enabled: "false"
      replication.policy.separator: ""
      replication.policy.class:
        "io.strimzi.kafka.connect.mirror.IdentityReplicationPolicy"
  heartbeatConnector:
    config:
      heartbeats.topic.replication.factor: 1
  checkpointConnector:
    config:
      checkpoints.topic.replication.factor: 1
  topicsPattern: ".*"
  groupsPattern: "group1|group2|group3"
resources:
  requests:
    cpu: "1"
    memory: 2Gi
  limits:
    cpu: "2"
    memory: 2Gi
logging:
  type: inline
  loggers:
    connect.root.logger.level: "INFO"
readinessProbe:
  initialDelaySeconds: 15
  timeoutSeconds: 5
livenessProbe:
  initialDelaySeconds: 15
  timeoutSeconds: 5
jvmOptions:
  "-Xmx": "1g"
  "-Xms": "1g"
image: my-org/my-image:latest
template:
  pod:
    affinity:
      podAntiAffinity:
        requiredDuringSchedulingIgnoredDuringExecution:
          - labelSelector:
              matchExpressions:
                - key: application
                  operator: In
                  values:
                    - postgresql
                    - mongodb

```

```

        topologyKey: "kubernetes.io/hostname"
connectContainer:
  env:
    - name: JAEGER_SERVICE_NAME
      value: my-jaeger-service
    - name: JAEGER_AGENT_HOST
      value: jaeger-agent-name
    - name: JAEGER_AGENT_PORT
      value: "6831"
tracing:
  type: jaeger
externalConfiguration:
  env:
    - name: AWS_ACCESS_KEY_ID
      valueFrom:
        secretKeyRef:
          name: aws-creds
          key: awsAccessKey
    - name: AWS_SECRET_ACCESS_KEY
      valueFrom:
        secretKeyRef:
          name: aws-creds
          key: awsSecretAccessKey

```

- ① The Kafka Connect and Mirror Maker 2.0 [version](#), which will always be the same.
- ② [The number of replica nodes](#).
- ③ [Kafka cluster alias](#) for Kafka Connect, which must specify the **target** Kafka cluster. The Kafka cluster is used by Kafka Connect for its internal topics.
- ④ [Specification](#) for the Kafka clusters being synchronized.
- ⑤ [Cluster alias](#) for the source Kafka cluster.
- ⑥ Authentication for the source cluster, using the [TLS mechanism](#), as shown here, using [OAuth bearer tokens](#), or a SASL-based [SCRAM-SHA-512](#) or [PLAIN](#) mechanism.
- ⑦ [Bootstrap server](#) for connection to the source Kafka cluster.
- ⑧ [TLS encryption](#) with key names under which TLS certificates are stored in X.509 format for the source Kafka cluster. If certificates are stored in the same secret, it can be listed multiple times.
- ⑨ [Cluster alias](#) for the target Kafka cluster.
- ⑩ Authentication for the target Kafka cluster is configured in the same way as for the source Kafka cluster.
- ⑪ [Bootstrap server](#) for connection to the target Kafka cluster.
- ⑫ [Kafka Connect configuration](#). Standard Apache Kafka configuration may be provided, restricted to those properties not managed directly by Strimzi.
- ⑬ [SSL properties](#) for external listeners to run with a specific *cipher suite* for a TLS version.
- ⑭ [Hostname verification is enabled](#) by setting to [HTTPS](#). An empty string disables the

verification.

⑯ TLS encryption for the target Kafka cluster is configured in the same way as for the source Kafka cluster.

⑯ [MirrorMaker 2.0 connectors](#).

⑯ [Cluster alias](#) for the source cluster used by the MirrorMaker 2.0 connectors.

⑯ [Cluster alias](#) for the target cluster used by the MirrorMaker 2.0 connectors.

⑯ [Configuration for the MirrorSourceConnector](#) that creates remote topics. The `config` overrides the default configuration options.

⑯ Replication factor for mirrored topics created at the target cluster.

Replication factor for the `MirrorSourceConnector offset-syncs` internal topic that maps the offsets of the source and target clusters.

When `ACL rules synchronization` is enabled, ACLs are applied to synchronized topics. The default is `true`.

Defines the separator used for the renaming of remote topics.

Adds a policy that overrides the automatic renaming of remote topics. Instead of prepending the name with the name of the source cluster, the topic retains its original name. This optional setting is useful for active/passive backups and data migration.

[Configuration for the MirrorHeartbeatConnector](#) that performs connectivity checks. The `config` overrides the default configuration options.

Replication factor for the heartbeat topic created at the target cluster.

[Configuration for the MirrorCheckpointConnector](#) that tracks offsets. The `config` overrides the default configuration options.

Replication factor for the checkpoints topic created at the target cluster.

Topic replication from the source cluster [defined as regular expression patterns](#). Here we request all topics.

Consumer group replication from the source cluster [defined as regular expression patterns](#). Here we request three consumer groups by name. You can use comma-separated lists.

Requests for reservation of [supported resources](#), currently `cpu` and `memory`, and limits to specify the maximum resources that can be consumed.

Specified [Kafka Connect loggers and log levels](#) added directly (`inline`) or indirectly (`external`) through a ConfigMap. A custom ConfigMap must be placed under the `log4j.properties` or `log4j2.properties` key. For the Kafka Connect `log4j.rootLogger` logger, you can set the log level to INFO, ERROR, WARN, TRACE, DEBUG, FATAL or OFF.

[Healthchecks](#) to know when to restart a container (liveness) and when a container can accept traffic (readiness).

[JVM configuration options](#) to optimize performance for the Virtual Machine (VM) running Kafka MirrorMaker.

ADVANCED OPTION: [Container image configuration](#), which is recommended only in special situations.

[Template customization](#). Here a pod is scheduled with anti-affinity, so the pod is not scheduled on nodes with the same hostname.

Environment variables are also [set for distributed tracing using Jaeger](#).

[Distributed tracing is enabled for Jaeger](#).

[External configuration](#) for a Kubernetes Secret mounted to Kafka MirrorMaker as an environment variable.

## 2. Create or update the resource:

```
kubectl apply -f <your-file>
```

## 2.5. Kafka Bridge cluster configuration

This section describes how to configure a Kafka Bridge deployment in your Strimzi cluster.

Kafka Bridge provides an API for integrating HTTP-based clients with a Kafka cluster.

If you are using the Kafka Bridge, you configure the [KafkaBridge](#) resource.

The full schema of the [KafkaBridge](#) resource is described in [KafkaBridge schema reference](#).

### 2.5.1. Configuring the Kafka Bridge

Use the Kafka Bridge to make HTTP-based requests to the Kafka cluster.

Use the properties of the [KafkaBridge](#) resource to configure your Kafka Bridge deployment.

In order to prevent issues arising when client consumer requests are processed by different Kafka Bridge instances, address-based routing must be employed to ensure that requests are routed to the right Kafka Bridge instance. Additionally, each independent Kafka Bridge instance must have a replica. A Kafka Bridge instance has its own state which is not shared with another instances.

#### Prerequisites

- A Kubernetes cluster
- A running Cluster Operator

See the [Deploying and Upgrading Strimzi](#) guide for instructions on running a:

- [Cluster Operator](#)
- [Kafka cluster](#)

#### Procedure

1. Edit the `spec` properties for the [KafkaBridge](#) resource.

The properties you can configure are shown in this example configuration:

```

apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaBridge
metadata:
  name: my-bridge
spec:
  replicas: 3 ①
  bootstrapServers: my-cluster-kafka-bootstrap:9092 ②
  tls: ③
    trustedCertificates:
      - secretName: my-cluster-cluster-cert
        certificate: ca.crt
      - secretName: my-cluster-cluster-cert
        certificate: ca2.crt
  authentication: ④
    type: tls
    certificateAndKey:
      secretName: my-secret
      certificate: public.crt
      key: private.key
  http: ⑤
    port: 8080
    cors: ⑥
      allowedOrigins: "https://strimzi.io"
      allowedMethods: "GET,POST,PUT,DELETE,OPTIONS,PATCH"
  consumer: ⑦
    config:
      auto.offset.reset: earliest
  producer: ⑧
    config:
      delivery.timeout.ms: 300000
  resources: ⑨
    requests:
      cpu: "1"
      memory: 2Gi
    limits:
      cpu: "2"
      memory: 2Gi
  logging: ⑩
    type: inline
    loggers:
      logger.bridge.level: "INFO"
      # enabling DEBUG just for send operation
      logger.send.name: "http.openapi.operation.send"
      logger.send.level: "DEBUG"
  jvmOptions: ⑪
    "-Xmx": "1g"
    "-Xms": "1g"
  readinessProbe: ⑫
    initialDelaySeconds: 15
    timeoutSeconds: 5
  livenessProbe:

```

```

initialDelaySeconds: 15
timeoutSeconds: 5
image: my-org/my-image:latest ⑬
template: ⑭
pod:
  affinity:
    podAntiAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        - labelSelector:
            matchExpressions:
              - key: application
                operator: In
                values:
                  - postgresql
                  - mongodb
      topologyKey: "kubernetes.io/hostname"
bridgeContainer: ⑮
env:
  - name: JAEGER_SERVICE_NAME
    value: my-jaeger-service
  - name: JAEGER_AGENT_HOST
    value: jaeger-agent-name
  - name: JAEGER_AGENT_PORT
    value: "6831"

```

- ① The number of replica nodes.
- ② Bootstrap server for connection to the target Kafka cluster.
- ③ TLS encryption with key names under which TLS certificates are stored in X.509 format for the source Kafka cluster. If certificates are stored in the same secret, it can be listed multiple times.
- ④ Authentication for the Kafka Bridge cluster, using the [TLS mechanism](#), as shown here, using [OAuth bearer tokens](#), or a SASL-based [SCRAM-SHA-512](#) or [PLAIN](#) mechanism. By default, the Kafka Bridge connects to Kafka brokers without authentication.
- ⑤ HTTP access to Kafka brokers.
- ⑥ CORS access specifying selected resources and access methods. Additional HTTP headers in requests describe the origins that are permitted access to the Kafka cluster.
- ⑦ Consumer configuration options.
- ⑧ Producer configuration options.
- ⑨ Requests for reservation of supported resources, currently `cpu` and `memory`, and limits to specify the maximum resources that can be consumed.
- ⑩ Specified Kafka Bridge loggers and log levels added directly ([inline](#)) or indirectly ([external](#)) through a ConfigMap. A custom ConfigMap must be placed under the `log4j.properties` or `log4j2.properties` key. For the Kafka Bridge loggers, you can set the log level to INFO, ERROR, WARN, TRACE, DEBUG, FATAL or OFF.
- ⑪ JVM configuration options to optimize performance for the Virtual Machine (VM) running

the Kafka Bridge.

- ⑫ [Healthchecks](#) to know when to restart a container (liveness) and when a container can accept traffic (readiness).
- ⑬ **ADVANCED OPTION:** [Container image configuration](#), which is recommended only in special situations.
- ⑭ [Template customization](#). Here a pod is scheduled with anti-affinity, so the pod is not scheduled on nodes with the same hostname.
- ⑮ Environment variables are also [set for distributed tracing using Jaeger](#).

2. Create or update the resource:

```
kubectl apply -f KAFKA-BRIDGE-CONFIG-FILE
```

### 2.5.2. List of Kafka Bridge cluster resources

The following resources are created by the Cluster Operator in the Kubernetes cluster:

#### ***bridge-cluster-name-bridge***

Deployment which is in charge to create the Kafka Bridge worker node pods.

#### ***bridge-cluster-name-bridge-service***

Service which exposes the REST interface of the Kafka Bridge cluster.

#### ***bridge-cluster-name-bridge-config***

ConfigMap which contains the Kafka Bridge ancillary configuration and is mounted as a volume by the Kafka broker pods.

#### ***bridge-cluster-name-bridge***

Pod Disruption Budget configured for the Kafka Bridge worker nodes.

## 2.6. Customizing Kubernetes resources

Strimzi creates several Kubernetes resources, such as [Deployments](#), [StatefulSets](#), [Pods](#), and [Services](#), which are managed by Strimzi operators. Only the operator that is responsible for managing a particular Kubernetes resource can change that resource. If you try to manually change an operator-managed Kubernetes resource, the operator will revert your changes back.

However, changing an operator-managed Kubernetes resource can be useful if you want to perform certain tasks, such as:

- Adding custom labels or annotations that control how [Pods](#) are treated by Istio or other services
- Managing how [Loadbalancer](#)-type Services are created by the cluster

You can make such changes using the [template](#) property in the Strimzi custom resources. The [template](#) property is supported in the following resources. The API reference provides more details

about the customizable fields.

#### **Kafka.spec.kafka**

See [KafkaClusterTemplate schema reference](#)

#### **Kafka.spec.zookeeper**

See [ZookeeperClusterTemplate schema reference](#)

#### **Kafka.spec.entityOperator**

See [EntityOperatorTemplate schema reference](#)

#### **Kafka.spec.kafkaExporter**

See [KafkaExporterTemplate schema reference](#)

#### **Kafka.spec.cruiseControl**

See [CruiseControlTemplate schema reference](#)

#### **Kafka.spec.jmxTrans**

See [JmxTransTemplate schema reference](#)

#### **KafkaConnect.spec**

See [KafkaConnectTemplate schema reference](#)

#### **KafkaConnectS2I.spec**

See [KafkaConnectTemplate schema reference](#)

#### **KafkaMirrorMaker.spec**

See [KafkaMirrorMakerTemplate schema reference](#)

#### **KafkaMirrorMaker2.spec**

See [KafkaConnectTemplate schema reference](#)

#### **KafkaBridge.spec**

See [KafkaBridgeTemplate schema reference](#)

#### **KafkaUser.spec**

See [KafkaUserTemplate schema reference](#)

In the following example, the `template` property is used to modify the labels in a Kafka broker's `StatefulSet`:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
  labels:
    app: my-cluster
spec:
  kafka:
    # ...
    template:
      statefulset:
        metadata:
          labels:
            mylabel: myvalue
    # ...
```

### 2.6.1. Customizing the image pull policy

Strimzi allows you to customize the image pull policy for containers in all pods deployed by the Cluster Operator. The image pull policy is configured using the environment variable `STRIMZI_IMAGE_PULL_POLICY` in the Cluster Operator deployment. The `STRIMZI_IMAGE_PULL_POLICY` environment variable can be set to three different values:

#### Always

Container images are pulled from the registry every time the pod is started or restarted.

#### IfNotPresent

Container images are pulled from the registry only when they were not pulled before.

#### Never

Container images are never pulled from the registry.

The image pull policy can be currently customized only for all Kafka, Kafka Connect, and Kafka MirrorMaker clusters at once. Changing the policy will result in a rolling update of all your Kafka, Kafka Connect, and Kafka MirrorMaker clusters.

#### *Additional resources*

- For more information about Cluster Operator configuration, see [Using the Cluster Operator](#).
- For more information about Image Pull Policies, see [Disruptions](#).

## 2.7. Configuring pod scheduling

When two applications are scheduled to the same Kubernetes node, both applications might use the same resources like disk I/O and impact performance. That can lead to performance degradation. Scheduling Kafka pods in a way that avoids sharing nodes with other critical workloads, using the right nodes or dedicated a set of nodes only for Kafka are the best ways how to avoid such

problems.

### 2.7.1. Specifying affinity, tolerations, and topology spread constraints

Use affinity, tolerations and topology spread constraints to schedule the pods of kafka resources onto nodes. Affinity, tolerations and topology spread constraints are configured using the `affinity`, `tolerations`, and `topologySpreadConstraint` properties in following resources:

- `Kafka.spec.kafka.template.pod`
- `Kafka.spec.zookeeper.template.pod`
- `Kafka.spec.entityOperator.template.pod`
- `KafkaConnect.spec.template.pod`
- `KafkaConnectS2I.spec.template.pod`
- `KafkaBridge.spec.template.pod`
- `KafkaMirrorMaker.spec.template.pod`
- `KafkaMirrorMaker2.spec.template.pod`

The format of the `affinity`, `tolerations`, and `topologySpreadConstraint` properties follows the Kubernetes specification. The affinity configuration can include different types of affinity:

- Pod affinity and anti-affinity
- Node affinity

**NOTE** On Kubernetes 1.16 and 1.17, the support for `topologySpreadConstraint` is disabled by default. In order to use `topologySpreadConstraint`, you have to enable the `EvenPodsSpread` feature gate in Kubernetes API server and scheduler.

#### Additional resources

- [Kubernetes node and pod affinity documentation](#)
- [Kubernetes taints and tolerations](#)
- [Kubernetes Topology Spread Constraints](#)

#### Use pod anti-affinity to avoid critical applications sharing nodes

Use pod anti-affinity to ensure that critical applications are never scheduled on the same disk. When running a Kafka cluster, it is recommended to use pod anti-affinity to ensure that the Kafka brokers do not share nodes with other workloads, such as databases.

#### Use node affinity to schedule workloads onto specific nodes

The Kubernetes cluster usually consists of many different types of worker nodes. Some are optimized for CPU heavy workloads, some for memory, while others might be optimized for storage (fast local SSDs) or network. Using different nodes helps to optimize both costs and performance. To achieve the best possible performance, it is important to allow scheduling of Strimzi components to use the right nodes.

Kubernetes uses node affinity to schedule workloads onto specific nodes. Node affinity allows you to create a scheduling constraint for the node on which the pod will be scheduled. The constraint is specified as a label selector. You can specify the label using either the built-in node label like `beta.kubernetes.io/instance-type` or custom labels to select the right node.

### Use node affinity and tolerations for dedicated nodes

Use taints to create dedicated nodes, then schedule Kafka pods on the dedicated nodes by configuring node affinity and tolerations.

Cluster administrators can mark selected Kubernetes nodes as tainted. Nodes with taints are excluded from regular scheduling and normal pods will not be scheduled to run on them. Only services which can tolerate the taint set on the node can be scheduled on it. The only other services running on such nodes will be system services such as log collectors or software defined networks.

Running Kafka and its components on dedicated nodes can have many advantages. There will be no other applications running on the same nodes which could cause disturbance or consume the resources needed for Kafka. That can lead to improved performance and stability.

### 2.7.2. Configuring pod anti-affinity in Kafka components

#### *Prerequisites*

- A Kubernetes cluster
- A running Cluster Operator

#### *Procedure*

1. Edit the `affinity` property in the resource specifying the cluster deployment. Use labels to specify the pods which should not be scheduled on the same nodes. The `topologyKey` should be set to `kubernetes.io/hostname` to specify that the selected pods should not be scheduled on nodes with the same hostname. For example:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    template:
      pod:
        affinity:
          podAntiAffinity:
            requiredDuringSchedulingIgnoredDuringExecution:
              - labelSelector:
                  matchExpressions:
                    - key: application
                      operator: In
                      values:
                        - postgresql
                        - mongodb
                topologyKey: "kubernetes.io/hostname"
    # ...
  zookeeper:
    # ...
```

## 2. Create or update the resource.

This can be done using `kubectl apply`:

```
kubectl apply -f KAFKA-CONFIG-FILE
```

### 2.7.3. Configuring node affinity in Kafka components

#### Prerequisites

- A Kubernetes cluster
- A running Cluster Operator

#### Procedure

##### 1. Label the nodes where Strimzi components should be scheduled.

This can be done using `kubectl label`:

```
kubectl label node NAME-OF-NODE node-type=fast-network
```

Alternatively, some of the existing labels might be reused.

##### 2. Edit the `affinity` property in the resource specifying the cluster deployment. For example:

```

apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    template:
      pod:
        affinity:
          nodeAffinity:
            requiredDuringSchedulingIgnoredDuringExecution:
              nodeSelectorTerms:
                - matchExpressions:
                    - key: node-type
                      operator: In
                      values:
                        - fast-network
    # ...
  zookeeper:
    # ...

```

3. Create or update the resource.

This can be done using `kubectl apply`:

```
kubectl apply -f KAFKA-CONFIG-FILE
```

#### 2.7.4. Setting up dedicated nodes and scheduling pods on them

##### *Prerequisites*

- A Kubernetes cluster
- A running Cluster Operator

##### *Procedure*

1. Select the nodes which should be used as dedicated.
2. Make sure there are no workloads scheduled on these nodes.
3. Set the taints on the selected nodes:

This can be done using `kubectl taint`:

```
kubectl taint node NAME-OF-NODE dedicated=Kafka:NoSchedule
```

4. Additionally, add a label to the selected nodes as well.

This can be done using `kubectl label`:

```
kubectl label node NAME-OF-NODE dedicated=Kafka
```

5. Edit the **affinity** and **tolerations** properties in the resource specifying the cluster deployment.

For example:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    template:
      pod:
        tolerations:
          - key: "dedicated"
            operator: "Equal"
            value: "Kafka"
            effect: "NoSchedule"
        affinity:
          nodeAffinity:
            requiredDuringSchedulingIgnoredDuringExecution:
              nodeSelectorTerms:
                - matchExpressions:
                    - key: dedicated
                      operator: In
                      values:
                        - Kafka
      # ...
  zookeeper:
    # ...
```

6. Create or update the resource.

This can be done using **kubectl apply**:

```
kubectl apply -f KAFKA-CONFIG-FILE
```

## 2.8. External logging

When setting the logging levels for a resource, you can specify them *inline* directly in the **spec.logging** property of the resource YAML:

```
spec:  
  # ...  
  logging:  
    type: inline  
    loggers:  
      kafka.root.logger.level: "INFO"
```

Or you can specify *external* logging:

```
spec:  
  # ...  
  logging:  
    type: external  
    name: customConfigMap
```

With external logging, logging properties are defined in a ConfigMap. The name of the ConfigMap is referenced in the `spec.logging.name` property.

The advantages of using a ConfigMap are that the logging properties are maintained in one place and are accessible to more than one resource.

### 2.8.1. Creating a ConfigMap for logging

To use a ConfigMap to define logging properties, you create the ConfigMap and then reference it as part of the logging definition in the `spec` of a resource.

The ConfigMap must contain the appropriate logging configuration.

- `log4j.properties` for Kafka components, ZooKeeper, and the Kafka Bridge
- `log4j2.properties` for the Topic Operator and User Operator

The configuration must be placed under these properties.

Here we demonstrate how a ConfigMap defines a root logger for a Kafka resource.

#### Procedure

1. Create the ConfigMap.

You can create the ConfigMap as a YAML file or from a properties file using `kubectl` at the command line.

ConfigMap example with a root logger definition for Kafka:

```
kind: ConfigMap
apiVersion: kafka.strimzi.io/v1beta1
metadata:
  name: logging-configmap
data:
  log4j.properties:
    kafka.root.logger.level="INFO"
```

From the command line, using a properties file:

```
kubectl create configmap logging-configmap --from-file=log4j.properties
```

The properties file defines the logging configuration:

```
# Define the logger
kafka.root.logger.level="INFO"
# ...
```

2. Define *external* logging in the `spec` of the resource, setting the `logging.name` to the name of the ConfigMap.

```
spec:
  # ...
  logging:
    type: external
    name: logging-configmap
```

3. Create or update the resource.

```
kubectl apply -f kafka.yaml
```

# Chapter 3. Configuring external listeners

Use an external listener to expose your Strimzi Kafka cluster to a client outside a Kubernetes environment.

Specify the connection `type` to expose Kafka in the external listener configuration.

- `nodeport` uses `NodePort` type `Services`
- `loadbalancer` uses `Loadbalancer` type `Services`
- `ingress` uses Kubernetes `Ingress` and the `NGINX Ingress Controller for Kubernetes`
- `route` uses OpenShift `Routes` and the HAProxy router

For more information on listener configuration, see [GenericKafkaListener schema reference](#).

**NOTE** `route` is only supported on OpenShift

*Additional resources*

- [Accessing Apache Kafka in Strimzi](#)

## 3.1. Accessing Kafka using node ports

This procedure describes how to access a Strimzi Kafka cluster from an external client using node ports.

To connect to a broker, you need a hostname and port number for the Kafka *bootstrap address*, as well as the certificate used for authentication.

*Prerequisites*

- A Kubernetes cluster
- A running Cluster Operator

*Procedure*

1. Configure a `Kafka` resource with an external listener set to the `nodeport` type.

For example:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    listeners:
      - name: external
        port: 9094
        type: nodeport
        tls: true
        authentication:
          type: tls
        # ...
      # ...
  zookeeper:
    # ...
```

2. Create or update the resource.

```
kubectl apply -f KAFKA-CONFIG-FILE
```

**NodePort** type services are created for each Kafka broker, as well as an external *bootstrap service*. The bootstrap service routes external traffic to the Kafka brokers. Node addresses used for connection are propagated to the **status** of the Kafka custom resource.

The cluster CA certificate to verify the identity of the kafka brokers is also created with the same name as the **Kafka** resource.

3. Retrieve the bootstrap address you can use to access the Kafka cluster from the status of the **Kafka** resource.

```
kubectl get kafka KAFKA-CLUSTER-NAME
-o=jsonpath='{.status.listeners[?(@.type=="external")].bootstrapServers}{"\n"}'
```

4. If TLS encryption is enabled, extract the public certificate of the broker certification authority.

```
kubectl get secret KAFKA-CLUSTER-NAME-cluster-ca-cert -o jsonpath='{.data.ca\.crt}'
| base64 -d > ca.crt
```

Use the extracted certificate in your Kafka client to configure TLS connection. If you enabled any authentication, you will also need to configure SASL or TLS authentication.

## 3.2. Accessing Kafka using loadbalancers

This procedure describes how to access a Strimzi Kafka cluster from an external client using

loadbalancers.

To connect to a broker, you need the address of the *bootstrap loadbalancer*, as well as the certificate used for TLS encryption.

#### Prerequisites

- A Kubernetes cluster
- A running Cluster Operator

#### Procedure

1. Configure a **Kafka** resource with an external listener set to the **loadbalancer** type.

For example:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    listeners:
      - name: external
        port: 9094
        type: loadbalancer
        tls: true
        # ...
        # ...
    zookeeper:
      # ...
```

2. Create or update the resource.

```
kubectl apply -f KAFKA-CONFIG-FILE
```

**loadbalancer** type services and loadbalancers are created for each Kafka broker, as well as an external *bootstrap service*. The bootstrap service routes external traffic to all Kafka brokers. DNS names and IP addresses used for connection are propagated to the **status** of each service.

The cluster CA certificate to verify the identity of the kafka brokers is also created with the same name as the **Kafka** resource.

3. Retrieve the address of the bootstrap service you can use to access the Kafka cluster from the status of the **Kafka** resource.

```
kubectl get kafka KAFKA-CLUSTER-NAME
-o=jsonpath='{.status.listeners[?(@.type=="external")].bootstrapServers}{"\n"}'
```

4. If TLS encryption is enabled, extract the public certificate of the broker certification authority.

```
kubectl get secret KAFKA-CLUSTER-NAME-cluster-ca-cert -o jsonpath='{.data.ca\.crt}' | base64 -d > ca.crt
```

Use the extracted certificate in your Kafka client to configure TLS connection. If you enabled any authentication, you will also need to configure SASL or TLS authentication.

### 3.3. Accessing Kafka using ingress

This procedure shows how to access a Strimzi Kafka cluster from an external client outside of Kubernetes using Nginx Ingress.

To connect to a broker, you need a hostname (advertised address) for the Ingress *bootstrap address*, as well as the certificate used for authentication.

For access using Ingress, the port is always 443.

#### *TLS passthrough*

Kafka uses a binary protocol over TCP, but the [NGINX Ingress Controller for Kubernetes](#) is designed to work with the HTTP protocol. To be able to pass the Kafka connections through the Ingress, Strimzi uses the TLS passthrough feature of the [NGINX Ingress Controller for Kubernetes](#). Ensure TLS passthrough is enabled in your [NGINX Ingress Controller for Kubernetes](#) deployment.

Because it is using the TLS passthrough functionality, TLS encryption cannot be disabled when exposing Kafka using [Ingress](#).

For more information about enabling TLS passthrough, see [TLS passthrough documentation](#).

#### *Prerequisites*

- Kubernetes cluster
- Deployed [NGINX Ingress Controller for Kubernetes](#) with TLS passthrough enabled
- A running Cluster Operator

#### *Procedure*

1. Configure a [Kafka](#) resource with an external listener set to the [ingress](#) type.

Specify the Ingress hosts for the bootstrap service and Kafka brokers.

For example:

```

apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    listeners:
      - name: external
        port: 9094
        type: ingress
        tls: true
        authentication:
          type: tls
    configuration: ①
    bootstrap:
      host: bootstrap.myingress.com
    brokers:
      - broker: 0
        host: broker-0.myingress.com
      - broker: 1
        host: broker-1.myingress.com
      - broker: 2
        host: broker-2.myingress.com
    # ...
  zookeeper:
    # ...

```

① Ingress hosts for the bootstrap service and Kafka brokers.

2. Create or update the resource.

```
kubectl apply -f KAFKA-CONFIG-FILE
```

**ClusterIP** type services are created for each Kafka broker, as well as an additional *bootstrap service*. These services are used by the Ingress controller to route traffic to the Kafka brokers. An **Ingress** resource is also created for each service to expose them using the Ingress controller. The Ingress hosts are propagated to the **status** of each service.

The cluster CA certificate to verify the identity of the kafka brokers is also created with the same name as the **Kafka** resource.

Use the address for the bootstrap host you specified in the **configuration** and port 443 (*BOOTSTRAP-HOST:443*) in your Kafka client as the *bootstrap address* to connect to the Kafka cluster.

3. Extract the public certificate of the broker certificate authority.

```
kubectl get secret KAFKA-CLUSTER-NAME-cluster-ca-cert -o jsonpath='{.data.ca\.crt}' | base64 -d > ca.crt
```

Use the extracted certificate in your Kafka client to configure the TLS connection. If you enabled any authentication, you will also need to configure SASL or TLS authentication.

## 3.4. Accessing Kafka using OpenShift routes

This procedure describes how to access a Strimzi Kafka cluster from an external client outside of OpenShift using routes.

To connect to a broker, you need a hostname for the route *bootstrap address*, as well as the certificate used for TLS encryption.

For access using routes, the port is always 443.

### Prerequisites

- An OpenShift cluster
- A running Cluster Operator

### Procedure

1. Configure a [Kafka](#) resource with an external listener set to the [route](#) type.

For example:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  labels:
    app: my-cluster
    name: my-cluster
    namespace: myproject
spec:
  kafka:
    # ...
    listeners:
      - name: listener1
        port: 9094
        type: route
        tls: true
        # ...
      # ...
  zookeeper:
    # ...
```

**WARNING**

An OpenShift Route address comprises the name of the Kafka cluster, the name of the listener, and the name of the namespace it is created in. For example, `my-cluster-kafka-listener1-bootstrap-myproject` (*CLUSTER-NAME*-*kafka-LISTENER-NAME*-*bootstrap-NAMESPACE*). Be careful that the whole length of the address does not exceed a maximum limit of 63 characters.

2. Create or update the resource.

```
kubectl apply -f KAFKA-CONFIG-FILE
```

`ClusterIP` type services are created for each Kafka broker, as well as an external *bootstrap service*. The services route the traffic from the OpenShift Routes to the Kafka brokers. An OpenShift `Route` resource is also created for each service to expose them using the HAProxy load balancer. DNS addresses used for connection are propagated to the `status` of each service.

The cluster CA certificate to verify the identity of the kafka brokers is also created with the same name as the `Kafka` resource.

3. Retrieve the address of the bootstrap service you can use to access the Kafka cluster from the status of the `Kafka` resource.

```
kubectl get kafka KAFKA-CLUSTER-NAME  
-o=jsonpath='{{.status.listeners[?(@.type=="external")].bootstrapServers}}{"\n"}'
```

4. Extract the public certificate of the broker certification authority.

```
kubectl get secret KAFKA-CLUSTER-NAME-cluster-ca-cert -o jsonpath='{{.data.ca\.crt}}'  
| base64 -d > ca.crt
```

Use the extracted certificate in your Kafka client to configure TLS connection. If you enabled any authentication, you will also need to configure SASL or TLS authentication.

# Chapter 4. Managing secure access to Kafka

You can secure your Kafka cluster by managing the access each client has to the Kafka brokers.

A secure connection between Kafka brokers and clients can encompass:

- Encryption for data exchange
- Authentication to prove identity
- Authorization to allow or decline actions executed by users

This chapter explains how to set up secure connections between Kafka brokers and clients, with sections describing:

- Security options for Kafka clusters and clients
- How to secure Kafka brokers
- How to use an authorization server for OAuth 2.0 token-based authentication and authorization

## 4.1. Security options for Kafka

Use the [Kafka](#) resource to configure the mechanisms used for Kafka authentication and authorization.

### 4.1.1. Listener authentication

For clients inside the Kubernetes cluster, you can create [plain](#) (without encryption) or [tls internal](#) listeners.


For clients outside the Kubernetes cluster, you create *external* listeners and specify a connection mechanism, which can be [nodeport](#), [loadbalancer](#), [ingress](#) or [route](#) (on OpenShift).

For more information on the configuration options for connecting an external client, see [Configuring external listeners](#).

Supported authentication options:

1. Mutual TLS authentication (only on the listeners with TLS enabled encryption)
2. SCRAM-SHA-512 authentication
3. [OAuth 2.0 token based authentication](#)

The authentication option you choose depends on how you wish to authenticate client access to Kafka brokers.



117\_AMQ\_0920

*Figure 3. Kafka listener authentication options*

The listener `authentication` property is used to specify an authentication mechanism specific to that listener.

If no `authentication` property is specified then the listener does not authenticate clients which connect through that listener. The listener will accept all connections without authentication.

Authentication must be configured when using the User Operator to manage `KafkaUsers`.

The following example shows:

- A `plain` listener configured for SCRAM-SHA-512 authentication
- A `tls` listener with mutual TLS authentication
- An `external` listener with mutual TLS authentication

Each listener is configured with a unique name and port within a Kafka cluster.

**NOTE**

Listeners cannot be configured to use the ports set aside for interbroker communication (9091) and metrics (9404).

## An example showing listener authentication configuration

```
# ...
listeners:
- name: plain
  port: 9092
  type: internal
  tls: true
  authentication:
    type: scram-sha-512
- name: tls
  port: 9093
  type: internal
  tls: true
  authentication:
    type: tls
- name: external
  port: 9094
  type: loadbalancer
  tls: true
  authentication:
    type: tls
# ...
```

## Mutual TLS authentication

Mutual TLS authentication is always used for the communication between Kafka brokers and ZooKeeper pods.

Strimzi can configure Kafka to use TLS (Transport Layer Security) to provide encrypted communication between Kafka brokers and clients either with or without mutual authentication. For mutual, or two-way, authentication, both the server and the client present certificates. When you configure mutual authentication, the broker authenticates the client (client authentication) and the client authenticates the broker (server authentication).

**NOTE** TLS authentication is more commonly one-way, with one party authenticating the identity of another. For example, when HTTPS is used between a web browser and a web server, the browser obtains proof of the identity of the web server.

## SCRAM-SHA-512 authentication

SCRAM (Salted Challenge Response Authentication Mechanism) is an authentication protocol that can establish mutual authentication using passwords. Strimzi can configure Kafka to use SASL (Simple Authentication and Security Layer) SCRAM-SHA-512 to provide authentication on both unencrypted and encrypted client connections.

When SCRAM-SHA-512 authentication is used with a TLS client connection, the TLS protocol provides the encryption, but is not used for authentication.

The following properties of SCRAM make it safe to use SCRAM-SHA-512 even on unencrypted connections:

- The passwords are not sent in the clear over the communication channel. Instead the client and the server are each challenged by the other to offer proof that they know the password of the authenticating user.
- The server and client each generate a new challenge for each authentication exchange. This means that the exchange is resilient against replay attacks.

When a `KafkaUser.spec.authentication.type` is configured with `scram-sha-512` the User Operator will generate a random 12-character password consisting of upper and lowercase ASCII letters and numbers.

## Network policies

Strimzi automatically creates a `NetworkPolicy` resource for every listener that is enabled on a Kafka broker. By default, a `NetworkPolicy` grants access to a listener to all applications and namespaces.

If you want to restrict access to a listener at the network level to only selected applications or namespaces, use the `networkPolicyPeers` property.

Use network policies as part of the listener authentication configuration. Each listener can have a different `networkPolicyPeers` configuration.

For more information, refer to the [Listener network policies](#) section and the [NetworkPolicyPeer API reference](#).

**NOTE**

Your configuration of Kubernetes must support ingress `NetworkPolicies` in order to use network policies in Strimzi.

## Additional listener configuration options

You can use the properties of the `GenericKafkaListenerConfiguration` schema to add further configuration to listeners.

### 4.1.2. Kafka authorization

You can configure authorization for Kafka brokers using the `authorization` property in the `Kafka.spec.kafka` resource. If the `authorization` property is missing, no authorization is enabled and clients have no restrictions. When enabled, authorization is applied to all enabled listeners. The authorization method is defined in the `type` field.

Supported authorization options:

- [Simple authorization](#)
- [OAuth 2.0 authorization](#) (if you are using OAuth 2.0 token based authentication)
- [Open Policy Agent \(OPA\) authorization](#)




Figure 4. Kafka cluster authorization options

## Super users

Super users can access all resources in your Kafka cluster regardless of any access restrictions, and are supported by all authorization mechanisms.

To designate super users for a Kafka cluster, add a list of user principals to the `superUsers` property. If a user uses TLS client authentication, their username is the common name from their certificate subject prefixed with `CN=`.

*An example configuration with super users*

```

authorization:
  type: simple
  superUsers:
    - CN=client_1
    - user_2
    - CN=client_3
  
```

## 4.2. Security options for Kafka clients

Use the `KafkaUser` resource to configure the authentication mechanism, authorization mechanism, and access rights for Kafka clients. In terms of configuring security, clients are represented as users.

You can authenticate and authorize user access to Kafka brokers. Authentication permits access, and authorization constrains the access to permissible actions.

You can also create *super users* that have unconstrained access to Kafka brokers.

The authentication and authorization mechanisms must match the [specification for the listener used to access the Kafka brokers](#).

### 4.2.1. Identifying a Kafka cluster for user handling

A `KafkaUser` resource includes a label that defines the appropriate name of the Kafka cluster (derived from the name of the `Kafka` resource) to which it belongs.

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: my-user
  labels:
    strimzi.io/cluster: my-cluster
```

The label is used by the User Operator to identify the `KafkaUser` resource and create a new user, and also in subsequent handling of the user.

If the label does not match the Kafka cluster, the User Operator cannot identify the `KafkaUser` and the user is not created.

If the status of the `KafkaUser` resource remains empty, check your label.

#### 4.2.2. User authentication

User authentication is configured using the `authentication` property in `KafkaUser.spec`. The authentication mechanism enabled for the user is specified using the `type` field.

Supported authentication mechanisms:

- TLS client authentication
- SCRAM-SHA-512 authentication

When no authentication mechanism is specified, the User Operator does not create the user or its credentials.

##### *Additional resources*

- [When to use mutual TLS authentication or SCRAM-SHA Authentication authentication for clients](#)

#### TLS Client Authentication

To use TLS client authentication, you set the `type` field to `tls`.

An example `KafkaUser` with TLS client authentication enabled

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: my-user
  labels:
    strimzi.io/cluster: my-cluster
spec:
  authentication:
    type: tls
    # ...
```

When the user is created by the User Operator, it creates a new Secret with the same name as the [KafkaUser](#) resource. The Secret contains a private and public key for TLS client authentication. The public key is contained in a user certificate, which is signed by the client Certificate Authority (CA).

All keys are in X.509 format.

Secrets provide private keys and certificates in PEM and PKCS #12 formats.

For more information on securing Kafka communication with Secrets, see [Managing TLS certificates](#).

*An example Secret with user credentials*

```
apiVersion: v1
kind: Secret
metadata:
  name: my-user
  labels:
    strimzi.io/kind: KafkaUser
    strimzi.io/cluster: my-cluster
type: Opaque
data:
  ca.crt: # Public key of the client CA
  user.crt: # User certificate that contains the public key of the user
  user.key: # Private key of the user
  user.p12: # PKCS #12 archive file for storing certificates and keys
  user.password: # Password for protecting the PKCS #12 archive file
```

## SCRAM-SHA-512 Authentication

To use the SCRAM-SHA-512 authentication mechanism, you set the `type` field to [scram-sha-512](#).

*An example KafkaUser with SCRAM-SHA-512 authentication enabled*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: my-user
  labels:
    strimzi.io/cluster: my-cluster
spec:
  authentication:
    type: scram-sha-512
    # ...
```

When the user is created by the User Operator, it creates a new secret with the same name as the [KafkaUser](#) resource. The secret contains the generated password in the `password` key, which is encoded with base64. In order to use the password, it must be decoded.

An example `Secret` with user credentials

```
apiVersion: v1
kind: Secret
metadata:
  name: my-user
  labels:
    strimzi.io/kind: KafkaUser
    strimzi.io/cluster: my-cluster
type: Opaque
data:
  password: Z2VuZXJhdGVkcGFzc3dvcmQ= ①
  sasl.jaas.config:
b3JnLmFwYWNoZS5rYWZrYS5jb21tb24uc2VjdXJpdHkuc2NyYW0uU2NyYW1Mb2dpbk1vZHVsZSBYZXF1aXJlZC
B1c2VybmFtZT0ibXktDXnlciIgcGFzc3dvcmQ9ImdlbmVyYXR1ZHhc3N3b3JkIjsK ②
```

① The generated password, base64 encoded.

② The JAAS configuration string for SASL SCRAM-SHA-512 authentication, base64 encoded.

Decoding the generated password:

```
echo "Z2VuZXJhdGVkcGFzc3dvcmQ=" | base64 --decode
```

### 4.2.3. User authorization

User authorization is configured using the `authorization` property in `KafkaUser.spec`. The authorization type enabled for a user is specified using the `type` field.

To use simple authorization, you set the `type` property to `simple` in `KafkaUser.spec.authorization`. Simple authorization uses the default Kafka authorization plugin, `AclAuthorizer`.

Alternatively, you can use [OPA authorization](#), or if you are already using OAuth 2.0 token based authentication, you can also use [OAuth 2.0 authorization](#).

If no authorization is specified, the User Operator does not provision any access rights for the user. Whether such a `KafkaUser` can still access resources depends on the authorizer being used. For example, for the `AclAuthorizer` this is determined by its `allow.everyone.if.no.acl.found` configuration.

#### ACL rules

`AclAuthorizer` uses ACL rules to manage access to Kafka brokers.

ACL rules grant access rights to the user, which you specify in the `acls` property.

For more information about the `AclRule` object, see the [AclRule schema reference](#).

## Super user access to Kafka brokers

If a user is added to a list of super users in a Kafka broker configuration, the user is allowed unlimited access to the cluster regardless of any authorization constraints defined in ACLs in [KafkaUser](#).

For more information on configuring super user access to brokers, see [Kafka authorization](#).

## User quotas

You can configure the `spec` for the [KafkaUser](#) resource to enforce quotas so that a user does not exceed access to Kafka brokers based on a byte threshold or a time limit of CPU utilization.

*An example KafkaUser with user quotas*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: my-user
  labels:
    strimzi.io/cluster: my-cluster
spec:
  # ...
  quotas:
    producerByteRate: 1048576 ①
    consumerByteRate: 2097152 ②
    requestPercentage: 55 ③
```

① Byte-per-second quota on the amount of data the user can push to a Kafka broker

② Byte-per-second quota on the amount of data the user can fetch from a Kafka broker

③ CPU utilization limit as a percentage of time for a client group

For more information on these properties, see the [KafkaUserQuotas schema reference](#).

## 4.3. Securing access to Kafka brokers

To establish secure access to Kafka brokers, you configure and apply:

- A [Kafka](#) resource to:
  - Create listeners with a specified authentication type
  - Configure authorization for the whole Kafka cluster
- A [KafkaUser](#) resource to access the Kafka brokers securely through the listeners

Configure the [Kafka](#) resource to set up:

- Listener authentication
- Network policies that restrict access to Kafka listeners
- Kafka authorization

- Super users for unconstrained access to brokers

Authentication is configured independently for each listener. Authorization is always configured for the whole Kafka cluster.

The Cluster Operator creates the listeners and sets up the cluster and client certificate authority (CA) certificates to enable authentication within the Kafka cluster.

You can replace the certificates generated by the Cluster Operator by [installing your own certificates](#). You can also [configure your listener to use a Kafka listener certificate managed by an external Certificate Authority](#). Certificates are available in PKCS #12 format (.p12) and PEM (.crt) formats.

Use [KafkaUser](#) to enable the authentication and authorization mechanisms that a specific client uses to access Kafka.

Configure the [KafkaUser](#) resource to set up:

- Authentication to match the enabled listener authentication
- Authorization to match the enabled Kafka authorization
- Quotas to control the use of resources by clients

The User Operator creates the user representing the client and the security credentials used for client authentication, based on the chosen authentication type.

#### *Additional resources*

For more information about the schema for:

- [Kafka](#), see the [Kafka schema reference](#).
- [KafkaUser](#), see the [KafkaUser schema reference](#).

### 4.3.1. Securing Kafka brokers

This procedure shows the steps involved in securing Kafka brokers when running Strimzi.

The security implemented for Kafka brokers must be compatible with the security implemented for the clients requiring access.

- `Kafka.spec.kafka.listeners[*].authentication` matches `KafkaUser.spec.authentication`
- `Kafka.spec.kafka.authorization` matches `KafkaUser.spec.authorization`

The steps show the configuration for simple authorization and a listener using TLS authentication. For more information on listener configuration, see [GenericKafkaListener schema reference](#).

Alternatively, you can use SCRAM-SHA or OAuth 2.0 for [listener authentication](#), and OAuth 2.0 or OPA for [Kafka authorization](#).

#### *Procedure*

1. Configure the [Kafka](#) resource.

- Configure the `authorization` property for authorization.
- Configure the `listeners` property to create a listener with authentication.

For example:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    authorization: ①
      type: simple
      superUsers: ②
        - CN=client_1
        - user_2
        - CN=client_3
    listeners:
      - name: tls
        port: 9093
        type: internal
        tls: true
        authentication:
          type: tls ③
    # ...
  zookeeper:
    # ...
```

① Authorization enables simple authorization on the Kafka broker using the [AclAuthorizer Kafka plugin](#).

② List of user principals with unlimited access to Kafka. *CN* is the common name from the client certificate when TLS authentication is used.

③ Listener authentication mechanisms may be configured for each listener, and specified as mutual TLS, SCRAM-SHA-512 or token-based OAuth 2.0.

If you are configuring an external listener, the configuration is dependent on the chosen connection mechanism.

## 2. Create or update the Kafka resource.

```
kubectl apply -f KAFKA-CONFIG-FILE
```

The Kafka cluster is configured with a Kafka broker listener using TLS authentication.

A service is created for each Kafka broker pod.

A service is created to serve as the *bootstrap address* for connection to the Kafka cluster.

The cluster CA certificate to verify the identity of the kafka brokers is also created with the same

name as the **Kafka** resource.

### 4.3.2. Securing user access to Kafka

Use the properties of the **KafkaUser** resource to configure a Kafka user.

You can use `kubectl apply` to create or modify users, and `kubectl delete` to delete existing users.

For example:

- `kubectl apply -f USER-CONFIG-FILE`
- `kubectl delete KafkaUser USER-NAME`

When you configure the **KafkaUser** authentication and authorization mechanisms, ensure they match the equivalent **Kafka** configuration:

- `KafkaUser.spec.authentication` matches `Kafka.spec.kafka.listeners[*].authentication`
- `KafkaUser.spec.authorization` matches `Kafka.spec.kafka.authorization`

This procedure shows how a user is created with TLS authentication. You can also create a user with SCRAM-SHA authentication.

The authentication required depends on the [type of authentication configured for the Kafka broker listener](#).

**NOTE** Authentication between Kafka users and Kafka brokers depends on the authentication settings for each. For example, it is not possible to authenticate a user with TLS if it is not also enabled in the Kafka configuration.

#### Prerequisites

- A running Kafka cluster [configured with a Kafka broker listener using TLS authentication and encryption](#).
- A running User Operator (typically [deployed with the Entity Operator](#)).

The authentication type in **KafkaUser** should match the authentication configured in **Kafka** brokers.

#### Procedure

1. Configure the **KafkaUser** resource.

For example:

```

apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: my-user
  labels:
    strimzi.io/cluster: my-cluster
spec:
  authentication: ①
    type: tls
  authorization:
    type: simple ②
  acls:
    - resource:
        type: topic
        name: my-topic
        patternType: literal
        operation: Read
    - resource:
        type: topic
        name: my-topic
        patternType: literal
        operation: Describe
    - resource:
        type: group
        name: my-group
        patternType: literal
        operation: Read

```

① User authentication mechanism, defined as mutual `tls` or `scram-sha-512`.

② Simple authorization, which requires an accompanying list of ACL rules.

## 2. Create or update the `KafkaUser` resource.

```
kubectl apply -f USER-CONFIG-FILE
```

The user is created, as well as a Secret with the same name as the `KafkaUser` resource. The Secret contains a private and public key for TLS client authentication.

For information on configuring a Kafka client with properties for secure connection to Kafka brokers, see [Setting up access for clients outside of Kubernetes](#) in the *Deploying Strimzi* guide.

### 4.3.3. Restricting access to Kafka listeners using network policies

You can restrict access to a listener to only selected applications by using the `networkPolicyPeers` property.

#### *Prerequisites*

- A Kubernetes cluster with support for Ingress NetworkPolicies.

- The Cluster Operator is running.

#### Procedure

1. Open the **Kafka** resource.
2. In the `networkPolicyPeers` property, define the application pods or namespaces that will be allowed to access the Kafka cluster.

For example, to configure a `tls` listener to allow connections only from application pods with the label `app` set to `kafka-client`:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    listeners:
      - name: tls
        port: 9093
        type: internal
        tls: true
        authentication:
          type: tls
      networkPolicyPeers:
        - podSelector:
            matchLabels:
              app: kafka-client
    # ...
  zookeeper:
    # ...
```

3. Create or update the resource.

Use `kubectl apply`:

```
kubectl apply -f your-file
```

#### Additional resources

- For more information about the schema, see the [NetworkPolicyPeer API reference](#) and the [GenericKafkaListener schema reference](#).

## 4.4. Using OAuth 2.0 token-based authentication

Strimzi supports the use of OAuth 2.0 authentication using the `SASL OAUTHBEARER` mechanism.

OAuth 2.0 enables standardized token-based authentication and authorization between applications, using a central authorization server to issue tokens that grant limited access to resources.

You can configure OAuth 2.0 authentication, then [OAuth 2.0 authorization](#).

OAuth 2.0 authentication can also be used in conjunction with [simple](#) or OPA-based [Kafka authorization](#).

Using OAuth 2.0 token-based authentication, application clients can access resources on application servers (called *resource servers*) without exposing account credentials.

The application client passes an access token as a means of authenticating, which application servers can also use to determine the level of access to grant. The authorization server handles the granting of access and inquiries about access.

In the context of Strimzi:

- Kafka brokers act as OAuth 2.0 resource servers
- Kafka clients act as OAuth 2.0 application clients

Kafka clients authenticate to Kafka brokers. The brokers and clients communicate with the OAuth 2.0 authorization server, as necessary, to obtain or validate access tokens.

For a deployment of Strimzi, OAuth 2.0 integration provides:

- Server-side OAuth 2.0 support for Kafka brokers
- Client-side OAuth 2.0 support for Kafka MirrorMaker, Kafka Connect and the Kafka Bridge

*Additional resources*

- [OAuth 2.0 site](#)

#### 4.4.1. OAuth 2.0 authentication mechanism

The Kafka *SASL OAUTHBearer* mechanism is used to establish authenticated sessions with a Kafka broker.

A Kafka client initiates a session with the Kafka broker using the *SASL OAUTHBearer* mechanism for credentials exchange, where credentials take the form of an access token.

Kafka brokers and clients need to be configured to use OAuth 2.0.

#### 4.4.2. OAuth 2.0 Kafka broker configuration

Kafka broker configuration for OAuth 2.0 involves:

- Creating the OAuth 2.0 client in the authorization server
- Configuring OAuth 2.0 authentication in the Kafka custom resource

**NOTE**

In relation to the authorization server, Kafka brokers and Kafka clients are both regarded as OAuth 2.0 clients.

## OAuth 2.0 client configuration on an authorization server

To configure a Kafka broker to validate the token received during session initiation, the recommended approach is to create an OAuth 2.0 *client* definition in an authorization server, configured as *confidential*, with the following client credentials enabled:

- Client ID of **kafka** (for example)
- Client ID and Secret as the authentication mechanism

**NOTE**

You only need to use a client ID and secret when using a non-public introspection endpoint of the authorization server. The credentials are not typically required when using public authorization server endpoints, as with fast local JWT token validation.

## OAuth 2.0 authentication configuration in the Kafka cluster

To use OAuth 2.0 authentication in the Kafka cluster, you specify, for example, a TLS listener configuration for your Kafka cluster custom resource with the authentication method **oauth**:

*Assigning the authentication method type for OAuth 2.0*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    listeners:
      - name: tls
        port: 9093
        type: internal
        tls: true
        authentication:
          type: oauth
#...
```

You can configure **plain**, **tls** and **external** listeners, but it is recommended not to use **plain** listeners or **external** listeners with disabled TLS encryption with OAuth 2.0 as this creates a vulnerability to network eavesdropping and unauthorized access through token theft.

You configure an **external** listener with **type: oauth** for a secure transport layer to communicate with the client.

```
# ...
listeners:
  - name: external
    port: 9094
    type: loadbalancer
    tls: true
    authentication:
      type: oauth
    #...
```

The `tls` property is *false* by default, so it must be enabled.

When you have defined the type of authentication as OAuth 2.0, you add configuration based on the type of validation, either as [fast local JWT validation](#) or [token validation using an introspection endpoint](#).

The procedure to configure OAuth 2.0 for listeners, with descriptions and examples, is described in [Configuring OAuth 2.0 support for Kafka brokers](#).

### Fast local JWT token validation configuration

Fast local JWT token validation checks a JWT token signature locally.

The local check ensures that a token:

- Conforms to type by containing a (`typ`) claim value of `Bearer` for an access token
- Is valid (not expired)
- Has an issuer that matches a `validIssuerURI`

You specify a `validIssuerURI` attribute when you configure the listener, so that any tokens not issued by the authorization server are rejected.

The authorization server does not need to be contacted during fast local JWT token validation. You activate fast local JWT token validation by specifying a `jwksEndpointUri` attribute, the endpoint exposed by the OAuth 2.0 authorization server. The endpoint contains the public keys used to validate signed JWT tokens, which are sent as credentials by Kafka clients.

**NOTE**

All communication with the authorization server should be performed using TLS encryption.

You can configure a certificate truststore as a Kubernetes Secret in your Strimzi project namespace, and use a `tlsTrustedCertificates` attribute to point to the Kubernetes Secret containing the truststore file.

You might want to configure a `userNameClaim` to properly extract a username from the JWT token. If you want to use Kafka ACL authorization, you need to identify the user by their username during authentication. (The `sub` claim in JWT tokens is typically a unique ID, not a username.)

## *Example configuration for fast local JWT token validation*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    #...
    listeners:
      - name: tls
        port: 9093
        type: internal
        tls: true
        authentication:
          type: oauth
          validIssuerUri: <https://<auth-server-address>/auth/realmstls>
          jwksEndpointUri: <https://<auth-server-
address>/auth/realmstls/protocol/openid-connect/certs>
          userNameClaim: preferred_username
          maxSecondsWithoutReauthentication: 3600
        tlsTrustedCertificates:
          - secretName: oauth-server-cert
            certificate: ca.crt
    #...
```

## **OAuth 2.0 introspection endpoint configuration**

Token validation using an OAuth 2.0 introspection endpoint treats a received access token as opaque. The Kafka broker sends an access token to the introspection endpoint, which responds with the token information necessary for validation. Importantly, it returns up-to-date information if the specific access token is valid, and also information about when the token expires.

To configure OAuth 2.0 introspection-based validation, you specify an `introspectionEndpointUri` attribute rather than the `jwksEndpointUri` attribute specified for fast local JWT token validation. Depending on the authorization server, you typically have to specify a `clientId` and `clientSecret`, because the introspection endpoint is usually protected.

## Example configuration for an introspection endpoint

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    listeners:
      - name: tls
        port: 9093
        type: internal
        tls: true
        authentication:
          type: oauth
          clientId: kafka-broker
          clientSecret:
            secretName: my-cluster-oauth
            key: clientSecret
          validIssuerUri: <https://<auth-server-address>/auth/realms/tls>
          introspectionEndpointUri: <https://<auth-server-
address>/auth/realms/tls/protocol/openid-connect/token/introspect>
          userNameClaim: preferred_username
          maxSecondsWithoutReauthentication: 3600
          tlsTrustedCertificates:
            - secretName: oauth-server-cert
              certificate: ca.crt
```

### 4.4.3. Session re-authentication for Kafka brokers

The Kafka *SASL OAUTHBEARER* mechanism, which is used for OAuth 2.0 authentication in Strimzi, supports a Kafka feature called the *re-authentication* mechanism.

When the re-authentication mechanism is enabled in the configuration of an `oauth` type listener, the broker's authenticated session expires when the access token expires. The client must then re-authenticate to the existing session by sending a new, valid access token to the broker, without dropping the connection.

If token validation is successful, a new client session is started using the existing connection. If the client fails to re-authenticate, the broker will close the connection if further attempts are made to send or receive messages. Java clients that use Kafka client library 2.2 or later automatically re-authenticate if the re-authentication mechanism is enabled on the broker.

You enable session re-authentication in the `Kafka` resource. Set the `maxSecondsWithoutReauthentication` property for a TLS listener with `type: oauth` authentication. Session re-authentication is supported for both types of token validation (fast local JWT and introspection endpoint). For an example configuration, see [Configuring OAuth 2.0 support for Kafka brokers](#).

For more information about the re-authentication mechanism, which was added in Kafka version 2.2, see [KIP-368](#).

## *Additional resources*

- [OAuth 2.0 Kafka broker configuration](#)
- [Configuring OAuth 2.0 support for Kafka brokers](#)
- [KafkaListenerAuthenticationOAuth schema reference](#)

### **4.4.4. OAuth 2.0 Kafka client configuration**

A Kafka client is configured with either:

- The credentials required to obtain a valid access token from an authorization server (client ID and Secret)
- A valid long-lived access token or refresh token, obtained using tools provided by an authorization server

The only information ever sent to the Kafka broker is an access token. The credentials used to authenticate with the authorization server to obtain the access token are never sent to the broker.

When a client obtains an access token, no further communication with the authorization server is needed.

The simplest mechanism is authentication with a client ID and Secret. Using a long-lived access token, or a long-lived refresh token, adds more complexity because there is an additional dependency on authorization server tools.

**NOTE**

If you are using long-lived access tokens, you may need to configure the client in the authorization server to increase the maximum lifetime of the token.

If the Kafka client is not configured with an access token directly, the client exchanges credentials for an access token during Kafka session initiation by contacting the authorization server. The Kafka client exchanges either:

- Client ID and Secret
- Client ID, refresh token, and (optionally) a Secret

### **4.4.5. OAuth 2.0 client authentication flow**

In this section, we explain and visualize the communication flow between Kafka client, Kafka broker, and authorization server during Kafka session initiation. The flow depends on the client and server configuration.

When a Kafka client sends an access token as credentials to a Kafka broker, the token needs to be validated.

Depending on the authorization server used, and the configuration options available, you may prefer to use:

- Fast local token validation based on JWT signature checking and local token introspection, without contacting the authorization server

- An OAuth 2.0 introspection endpoint provided by the authorization server

Using fast local token validation requires the authorization server to provide a JWKS endpoint with public certificates that are used to validate signatures on the tokens.

Another option is to use an OAuth 2.0 introspection endpoint on the authorization server. Each time a new Kafka broker connection is established, the broker passes the access token received from the client to the authorization server, and checks the response to confirm whether or not the token is valid.

Kafka client credentials can also be configured for:

- Direct local access using a previously generated long-lived access token
- Contact with the authorization server for a new access token to be issued

**NOTE**


An authorization server might only allow the use of opaque access tokens, which means that local token validation is not possible.

## Example client authentication flows

Here you can see the communication flows, for different configurations of Kafka clients and brokers, during Kafka session authentication.

- Client using client ID and secret, with broker delegating validation to authorization server
- Client using client ID and secret, with broker performing fast local token validation
- Client using long-lived access token, with broker delegating validation to authorization server
- Client using long-lived access token, with broker performing fast local validation

### *Client using client ID and secret, with broker delegating validation to authorization server*




AMQ\_46\_1019

1. Kafka client requests access token from authorization server, using client ID and secret, and optionally a refresh token.
2. Authorization server generates a new access token.
3. Kafka client authenticates with the Kafka broker using the *SASL OAUTHBEARER* mechanism to pass the access token.
4. Kafka broker validates the access token by calling a token introspection endpoint on authorization server, using its own client ID and secret.

- Kafka client session is established if the token is valid.


*Client using client ID and secret, with broker performing fast local token validation*



AMQ\_46\_1019

- Kafka client authenticates with authorization server from the token endpoint, using a client ID and secret, and optionally a refresh token.
- Authorization server generates a new access token.
- Kafka client authenticates with the Kafka broker using the *SASL OAUTHBEARER* mechanism to pass the access token.
- Kafka broker validates the access token locally using a JWT token signature check, and local token introspection.

*Client using long-lived access token, with broker delegating validation to authorization server*



AMQ\_46\_1019

- Kafka client authenticates with the Kafka broker using the *SASL OAUTHBEARER* mechanism to pass the long-lived access token.
- Kafka broker validates the access token by calling a token introspection endpoint on authorization server, using its own client ID and secret.
- Kafka client session is established if the token is valid.

*Client using long-lived access token, with broker performing fast local validation*



AMQ\_46\_1019

1. Kafka client authenticates with the Kafka broker using the *SASL OAUTHBearer* mechanism to pass the long-lived access token.
2. Kafka broker validates the access token locally using JWT token signature check, and local token introspection.

**WARNING**

Fast local JWT token signature validation is suitable only for short-lived tokens as there is no check with the authorization server if a token has been revoked. Token expiration is written into the token, but revocation can happen at any time, so cannot be accounted for without contacting the authorization server. Any issued token would be considered valid until it expires.

#### 4.4.6. Configuring OAuth 2.0 authentication

OAuth 2.0 is used for interaction between Kafka clients and Strimzi components.

In order to use OAuth 2.0 for Strimzi, you must:

1. [Configure an OAuth 2.0 authorization server for the Strimzi cluster and Kafka clients](#)
2. [Deploy or update the Kafka cluster with Kafka broker listeners configured to use OAuth 2.0](#)
3. [Update your Java-based Kafka clients to use OAuth 2.0](#)
4. [Update Kafka component clients to use OAuth 2.0](#)

##### Configuring an OAuth 2.0 authorization server

This procedure describes in general what you need to do to configure an authorization server for integration with Strimzi.

These instructions are not product specific.

The steps are dependent on the chosen authorization server. Consult the product documentation for the authorization server for information on how to set up OAuth 2.0 access.

**NOTE**

If you already have an authorization server deployed, you can skip the deployment step and use your current deployment.

##### Procedure

1. Deploy the authorization server to your cluster.

2. Access the CLI or admin console for the authorization server to configure OAuth 2.0 for Strimzi.

Now prepare the authorization server to work with Strimzi.

3. Configure a **kafka-broker** client.

4. Configure clients for each Kafka client component of your application.

#### *What to do next*

After deploying and configuring the authorization server, [configure the Kafka brokers to use OAuth 2.0](#).

### **Configuring OAuth 2.0 support for Kafka brokers**

This procedure describes how to configure Kafka brokers so that the broker listeners are enabled to use OAuth 2.0 authentication using an authorization server.

We advise use of OAuth 2.0 over an encrypted interface through configuration of TLS listeners. Plain listeners are not recommended.

If the authorization server is using certificates signed by the trusted CA and matching the OAuth 2.0 server hostname, TLS connection works using the default settings. Otherwise, you may need to configure the truststore with proper certificates or disable the certificate hostname validation.

When configuring the Kafka broker you have two options for the mechanism used to validate the access token during OAuth 2.0 authentication of the newly connected Kafka client:

- [Configuring fast local JWT token validation](#)
- [Configuring token validation using an introspection endpoint](#)

#### *Before you start*

For more information on the configuration of OAuth 2.0 authentication for Kafka broker listeners, see:

- [KafkaListenerAuthenticationOAuth schema reference](#)
- [Managing access to Kafka](#)

#### *Prerequisites*

- Strimzi and Kafka are running
- An OAuth 2.0 authorization server is deployed

#### *Procedure*

1. Update the Kafka broker configuration ([Kafka.spec.kafka](#)) of your **Kafka** resource in an editor.

```
kubectl edit kafka my-cluster
```

2. Configure the Kafka broker **listeners** configuration.

The configuration for each type of listener does not have to be the same, as they are

independent.

The examples here show the configuration options as configured for external listeners.

*Example 1: Configuring fast local JWT token validation*

```
#...
- name: external
  port: 9094
  type: loadbalancer
  tls: true
  authentication:
    type: oauth ①
    validIssuerUri: <https://<auth-server-address>/auth/realms/external> ②
    jwksEndpointUri: <https://<auth-server-
address>/auth/realms/external/protocol/openid-connect/certs> ③
    userNameClaim: preferred_username ④
    maxSecondsWithoutReauthentication: 3600 ⑤
    tlsTrustedCertificates: ⑥
    - secretName: oauth-server-cert
      certificate: ca.crt
    disableTlsHostnameVerification: true ⑦
    jwksExpirySeconds: 360 ⑧
    jwksRefreshSeconds: 300 ⑨
    jwksMinRefreshPauseSeconds: 1 ⑩
    enableECDSA: "true" ⑪
```

① Listener type set to `oauth`.

② URI of the token issuer used for authentication.

③ URI of the JWKS certificate endpoint used for local JWT validation.

④ The token claim (or key) that contains the actual user name in the token. The user name is the *principal* used to identify the user. The `userNameClaim` value will depend on the authentication flow and the authorization server used.

⑤ (Optional) Activates the Kafka re-authentication mechanism that enforces session expiry to the same length of time as the access token. If the specified value is less than the time left for the access token to expire, then the client will have to re-authenticate before the actual token expiry. By default, the session does not expire when the access token expires, and the client does not attempt re-authentication.

⑥ (Optional) Trusted certificates for TLS connection to the authorization server.

⑦ (Optional) Disable TLS hostname verification. Default is `false`.

⑧ The duration the JWKS certificates are considered valid before they expire. Default is `360` seconds. If you specify a longer time, consider the risk of allowing access to revoked certificates.

⑨ The period between refreshes of JWKS certificates. The interval must be at least 60 seconds shorter than the expiry interval. Default is `300` seconds.

⑩ The minimum pause in seconds between consecutive attempts to refresh JWKS public keys.

When an unknown signing key is encountered, the JWKS keys refresh is scheduled outside the regular periodic schedule with at least the specified pause since the last refresh attempt. The refreshing of keys follows the rule of exponential backoff, retrying on unsuccessful refreshes with ever increasing pause, until it reaches `jwksRefreshSeconds`. The default value is 1.

⑪ (Optional) If ECDSA is used for signing JWT tokens on authorization server, then this needs to be enabled. It installs additional crypto providers using BouncyCastle crypto library. Default is `false`.

*Example 2: Configuring token validation using an introspection endpoint*

```
- name: external
  port: 9094
  type: loadbalancer
  tls: true
  authentication:
    type: oauth
    validIssuerUri: <https://<auth-server-address>/auth/realms/external>
    introspectionEndpointUri: <https://<auth-server-
address>/auth/realms/external/protocol/openid-connect/token/introspect> ①
    clientId: kafka-broker ②
    clientSecret: ③
    secretName: my-cluster-oauth
    key: clientSecret
    userNameClaim: preferred_username ④
    maxSecondsWithoutReauthentication: 3600 ⑤
```

① URI of the token introspection endpoint.

② Client ID to identify the client.

③ Client Secret and client ID is used for authentication.

④ The token claim (or key) that contains the actual user name in the token. The user name is the *principal* used to identify the user. The `userNameClaim` value will depend on the authorization server used.

⑤ (Optional) Activates the Kafka re-authentication mechanism that enforces session expiry to the same length of time as the access token. If the specified value is less than the time left for the access token to expire, then the client will have to re-authenticate before the actual token expiry. By default, the session does not expire when the access token expires, and the client does not attempt re-authentication.

Depending on how you apply OAuth 2.0 authentication, and the type of authorization server, there are additional (optional) configuration settings you can use:

```

# ...
authentication:
  type: oauth
  # ...
  checkIssuer: false ①
  fallbackUserNameClaim: client_id ②
  fallbackUserNamePrefix: client-account- ③
  validTokenType: bearer ④
  userInfoEndpointUri: https://OAUTH-SERVER-
ADDRESS/auth/realms/external/protocol/openid-connect/userinfo ⑤

```

- ① If your authorization server does not provide an `iss` claim, it is not possible to perform an issuer check. In this situation, set `checkIssuer` to `false` and do not specify a `validIssuerUri`. Default is `true`.
- ② An authorization server may not provide a single attribute to identify both regular users and clients. When a client authenticates in its own name, the server might provide a *client ID*. When a user authenticates using a username and password, to obtain a refresh token or an access token, the server might provide a *username* attribute in addition to a client ID. Use this fallback option to specify the username claim (attribute) to use if a primary user ID attribute is not available.
- ③ In situations where `fallbackUserNameClaim` is applicable, it may also be necessary to prevent name collisions between the values of the `username` claim, and those of the fallback `username` claim. Consider a situation where a client called `producer` exists, but also a regular user called `producer` exists. In order to differentiate between the two, you can use this property to add a prefix to the user ID of the client.
- ④ (Only applicable when using `introspectionEndpointUri`) Depending on the authorization server you are using, the introspection endpoint may or may not return the *token type* attribute, or it may contain different values. You can specify a valid token type value that the response from the introspection endpoint has to contain.
- ⑤ (Only applicable when using `introspectionEndpointUri`) The authorization server may be configured or implemented in such a way to not provide any identifiable information in an Introspection Endpoint response. In order to obtain the user ID, you can configure the URI of the `userinfo` endpoint as a fallback. The `userNameClaim`, `fallbackUserNameClaim`, and `fallbackUserNamePrefix` settings are applied to the response of `userinfo` endpoint.

3. Save and exit the editor, then wait for rolling updates to complete.
4. Check the update in the logs or by watching the pod state transitions:

```

kubectl logs -f ${POD_NAME} -c ${CONTAINER_NAME}
kubectl get pod -w

```

The rolling update configures the brokers to use OAuth 2.0 authentication.

#### What to do next

- [Configure your Kafka clients to use OAuth 2.0](#)

## Configuring Kafka Java clients to use OAuth 2.0

This procedure describes how to configure Kafka producer and consumer APIs to use OAuth 2.0 for interaction with Kafka brokers.

Add a client callback plugin to your `pom.xml` file, and configure the system properties.

### Prerequisites

- Strimzi and Kafka are running
- An OAuth 2.0 authorization server is deployed and configured for OAuth access to Kafka brokers
- Kafka brokers are configured for OAuth 2.0

### Procedure

1. Add the client library with OAuth 2.0 support to the `pom.xml` file for the Kafka client:

```
<dependency>
  <groupId>io.strimzi</groupId>
  <artifactId>kafka-oauth-client</artifactId>
  <version>0.6.1</version>
</dependency>
```

2. Configure the system properties for the callback:

For example:

```
System.setProperty(ClientConfig.OAUTH_TOKEN_ENDPOINT_URI, "https://<auth-server-address>/auth/realms/master/protocol/openid-connect/token"); ①
System.setProperty(ClientConfig.OAUTH_CLIENT_ID, "<client-name>"); ②
System.setProperty(ClientConfig.OAUTH_CLIENT_SECRET, "<client-secret>"); ③
```

① URI of the authorization server token endpoint.

② Client ID, which is the name used when creating the *client* in the authorization server.

③ Client secret created when creating the *client* in the authorization server.

3. Enable the *SASL OAUTHBEARER* mechanism on a TLS encrypted connection in the Kafka client configuration:

For example:

```
props.put("sasl.jaas.config",
  "org.apache.kafka.common.security.oauthbearer.OAuthBearerLoginModule required;");
props.put("security.protocol", "SASL_SSL"); ①
props.put("sasl.mechanism", "OAUTHBEARER");
props.put("sasl.login.callback.handler.class",
  "io.strimzi.kafka.oauth.client.JaasClientOAuthLoginCallbackHandler");
```

① Here we use `SASL_SSL` for use over TLS connections. Use `SASL_PLAINTEXT` over unencrypted connections.

#### 4. Verify that the Kafka client can access the Kafka brokers.

##### *What to do next*

- [Configure Kafka components to use OAuth 2.0](#)

### Configuring OAuth 2.0 for Kafka components

This procedure describes how to configure Kafka components to use OAuth 2.0 authentication using an authorization server.

You can configure authentication for:

- Kafka Connect
- Kafka MirrorMaker
- Kafka Bridge

In this scenario, the Kafka component and the authorization server are running in the same cluster.

##### *Before you start*

For more information on the configuration of OAuth 2.0 authentication for Kafka components, see:

- [KafkaClientAuthenticationOAuth schema reference](#)

##### *Prerequisites*

- Strimzi and Kafka are running
- An OAuth 2.0 authorization server is deployed and configured for OAuth access to Kafka brokers
- Kafka brokers are configured for OAuth 2.0

##### *Procedure*

#### 1. Create a client secret and mount it to the component as an environment variable.

For example, here we are creating a client `Secret` for the Kafka Bridge:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Secret
metadata:
  name: my-bridge-oauth
type: Opaque
data:
  clientSecret: MGQ10TRmMzYtZTl1ZS00MDY2LWI50GEtMTM5MzM2NjdLZjQw ①
```

① The `clientSecret` key must be in base64 format.

#### 2. Create or edit the resource for the Kafka component so that OAuth 2.0 authentication is

configured for the authentication property.

For OAuth 2.0 authentication, you can use:

- Client ID and secret
- Client ID and refresh token
- Access token
- TLS

[KafkaClientAuthenticationOAuth schema reference](#) provides examples of each.

For example, here OAuth 2.0 is assigned to the Kafka Bridge client using a client ID and secret, and TLS:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaBridge
metadata:
  name: my-bridge
spec:
  # ...
  authentication:
    type: oauth ①
    tokenEndpointUri: https://<auth-server-
address>/auth/realms/master/protocol/openid-connect/token ②
    clientId: kafka-bridge
    clientSecret:
      secretName: my-bridge-oauth
      key: clientSecret
    tlsTrustedCertificates: ③
    - secretName: oauth-server-cert
      certificate: tls.crt
```

① Authentication type set to **oauth**.

② URI of the token endpoint for authentication.

③ Trusted certificates for TLS connection to the authorization server.

Depending on how you apply OAuth 2.0 authentication, and the type of authorization server, there are additional configuration options you can use:

```
# ...
spec:
  # ...
  authentication:
    # ...
    disableTlsHostnameVerification: true ①
    checkAccessTokenType: false ②
    accessTokenIsJwt: false ③
    scope: any ④
```

① (Optional) Disable TLS hostname verification. Default is `false`.

② If the authorization server does not return a `typ` (type) claim inside the JWT token, you can apply `checkAccessTokenType: false` to skip the token type check. Default is `true`.

③ If you are using opaque tokens, you can apply `accessTokenIsJwt: false` so that access tokens are not treated as JWT tokens.

④ (Optional) The `scope` for requesting the token from the token endpoint. An authorization server may require a client to specify the scope. In this case it is `any`.

3. Apply the changes to the deployment of your Kafka resource.

```
kubectl apply -f your-file
```

4. Check the update in the logs or by watching the pod state transitions:

```
kubectl logs -f ${POD_NAME} -c ${CONTAINER_NAME}
kubectl get pod -w
```

The rolling updates configure the component for interaction with Kafka brokers using OAuth 2.0 authentication.

#### 4.4.7. Authorization server examples

When choosing an authorization server, consider the features that best support configuration of your chosen authentication flow.

For the purposes of testing OAuth 2.0 with Strimzi, Keycloak and ORY Hydra were implemented as the OAuth 2.0 authorization server.

For more information, see:

- [Kafka authentication using OAuth 2.0](#)
- [Using Keycloak as the OAuth 2.0 authorization server](#)
- [Using Hydra as the OAuth 2.0 authorization server](#)

## 4.5. Using OAuth 2.0 token-based authorization

If you are using OAuth 2.0 with Keycloak for token-based authentication, you can also use Keycloak to configure authorization rules to constrain client access to Kafka brokers. Authentication establishes the identity of a user. Authorization decides the level of access for that user.

Strimzi supports the use of OAuth 2.0 token-based authorization through Keycloak [Authorization Services](#), which allows you to manage security policies and permissions centrally.

Security policies and permissions defined in Keycloak are used to grant access to resources on Kafka brokers. Users and clients are matched against policies that permit access to perform specific actions on Kafka brokers.

Kafka allows all users full access to brokers by default, and also provides the [AclAuthorizer](#) plugin to configure authorization based on Access Control Lists (ACLs).

ZooKeeper stores ACL rules that grant or deny access to resources based on *username*. However, OAuth 2.0 token-based authorization with Keycloak offers far greater flexibility on how you wish to implement access control to Kafka brokers. In addition, you can configure your Kafka brokers to use OAuth 2.0 authorization and ACLs.

### *Additional resources*

- [Using OAuth 2.0 token-based authentication](#)
- [Kafka Authorization](#)
- [Keycloak documentation](#)

### 4.5.1. OAuth 2.0 authorization mechanism

OAuth 2.0 authorization in Strimzi uses Keycloak server Authorization Services REST endpoints to extend token-based authentication with Keycloak by applying defined security policies on a particular user, and providing a list of permissions granted on different resources for that user. Policies use roles and groups to match permissions to users. OAuth 2.0 authorization enforces permissions locally based on the received list of grants for the user from Keycloak Authorization Services.

#### **Kafka broker custom authorizer**

A Keycloak *authorizer* ([KeycloakRBACAuthorizer](#)) is provided with Strimzi. To be able to use the Keycloak REST endpoints for Authorization Services provided by Keycloak, you configure a custom authorizer on the Kafka broker.

The authorizer fetches a list of granted permissions from the authorization server as needed, and enforces authorization locally on the Kafka Broker, making rapid authorization decisions for each client request.

### 4.5.2. Configuring OAuth 2.0 authorization support

This procedure describes how to configure Kafka brokers to use OAuth 2.0 authorization using

## Keycloak Authorization Services.

### *Before you begin*

Consider the access you require or want to limit for certain users. You can use a combination of Keycloak *groups*, *roles*, *clients*, and *users* to configure access in Keycloak.

Typically, groups are used to match users based on organizational departments or geographical locations. And roles are used to match users based on their function.

With Keycloak, you can store users and groups in LDAP, whereas clients and roles cannot be stored this way. Storage and access to user data may be a factor in how you choose to configure authorization policies.

#### **NOTE**

[Super users](#) always have unconstrained access to a Kafka broker regardless of the authorization implemented on the Kafka broker.

### *Prerequisites*

- Strimzi must be configured to use OAuth 2.0 with Keycloak for [token-based authentication](#). You use the same Keycloak server endpoint when you set up authorization.
- OAuth 2.0 authentication must be configured with the `maxSecondsWithoutReauthentication` option to enable re-authentication.

### *Procedure*

1. Access the Keycloak Admin Console or use the Keycloak Admin CLI to enable Authorization Services for the Kafka broker client you created when setting up OAuth 2.0 authentication.
2. Use Authorization Services to define resources, authorization scopes, policies, and permissions for the client.
3. Bind the permissions to users and clients by assigning them roles and groups.
4. Configure the Kafka brokers to use Keycloak authorization by updating the Kafka broker configuration ([Kafka.spec.kafka](#)) of your [Kafka](#) resource in an editor.

```
kubectl edit kafka my-cluster
```

5. Configure the Kafka broker `kafka` configuration to use `keycloak` authorization, and to be able to access the authorization server and Authorization Services.

For example:

```

apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka
  # ...
  authorization:
    type: keycloak ①
    tokenEndpointUri: <https://<auth-server-
address>/auth/realms/external/protocol/openid-connect/token> ②
    clientId: kafka ③
    delegateToKafkaAcls: false ④
    disableTlsHostnameVerification: false ⑤
    superUsers: ⑥
      - CN=fred
      - sam
      - CN=edward
    tlsTrustedCertificates: ⑦
      - secretName: oauth-server-cert
        certificate: ca.crt
    grantsRefreshPeriodSeconds: 60 ⑧
    grantsRefreshPoolSize: 5 ⑨
  #...

```

① Type **keycloak** enables Keycloak authorization.

② URI of the Keycloak token endpoint. For production, always use HTTPS.

③ The client ID of the OAuth 2.0 client definition in Keycloak that has Authorization Services enabled. Typically, **kafka** is used as the ID.

④ (Optional) Delegate authorization to Kafka **AclAuthorizer** if access is denied by Keycloak Authorization Services policies. Default is **false**.

⑤ (Optional) Disable TLS hostname verification. Default is **false**.

⑥ (Optional) Designated **super users**.

⑦ (Optional) Trusted certificates for TLS connection to the authorization server.

⑧ (Optional) The time between two consecutive grants refresh runs. That is the maximum time for active sessions to detect any permissions changes for the user on Keycloak. The default value is 60.

⑨ (Optional) The number of threads to use to refresh (in parallel) the grants for the active sessions. The default value is 5.

6. Save and exit the editor, then wait for rolling updates to complete.

7. Check the update in the logs or by watching the pod state transitions:

```
kubectl logs -f ${POD_NAME} -c kafka  
kubectl get pod -w
```

The rolling update configures the brokers to use OAuth 2.0 authorization.

8. Verify the configured permissions by accessing Kafka brokers as clients or users with specific roles, making sure they have the necessary access, or do not have the access they are not supposed to have.

#### 4.5.3. Managing policies and permissions in Keycloak Authorization Services

##### Kafka authorization model for resources

The [Kafka authorization model](#) defines resource types, and the permissions available for each type. When an action is performed by a Kafka client on a broker, the broker uses a configured authorizer to check permissions, depending on the action performed and the resource type.

Kafka has five resource types for controlling access: [Topic](#), [Group](#), [Cluster](#), [TransactionalId](#), [DelegationToken](#).

Each resource type has different permissions:

Topic:

- [Create](#)
- [Write](#)
- [Read](#)
- [Delete](#)
- [Describe](#)
- [DescribeConfigs](#)
- [Alter](#)
- [AlterConfigs](#)

Group:

- [Read](#)
- [Describe](#)
- [Delete](#)

Cluster:

- [Create](#)
- [Describe](#)
- [Alter](#)

- `DescribeConfigs`
- `AlterConfigs`
- `IdempotentWrite`
- `ClusterAction`

TransactionalId:

- `Describe`
- `Write`

DelegationToken:

- `Describe`

## **Keycloak Authorization Services model for managing permissions**

Keycloak Authorization Services use four concepts to define and grant permissions: *resources*, *authorization scopes*, *policies*, and *permissions*.

### **Resources**

Resources are a set of resource definitions that are used to match permitted actions. For example, a resource can be an individual topic, or it can be a set of all topics with names that start with the same prefix. The resource definition has a set of available authorization scopes associated with it, which represent a set of all actions available on the particular resource. Often, only a subset of these actions is actually permitted.

### **Authorization scopes**

Authorization scopes is simply a set of all available actions on all the different resource types. When defining a new resource, scopes are added from the set of all scopes.

### **Policies**

Policies are rules that use criteria to match a list of accounts. Policies can match service accounts based on client id or roles, or user accounts based on username, groups, or roles.

### **Permissions**

Permissions grant a subset of authorization scopes on a specific resource definition to a set of users.

## **Mapping Keycloak Authorization Services to the Kafka authorization model**

Use Keycloak Authorization Services rules on the OAuth client that represents the Kafka Broker to grant Kafka permissions to users or service accounts. Typically, the OAuth client has `kafka` as its client id.

The OAuth 2.0 client definition must have the *Authorization Enabled* option activated.

All permissions exist within the scope of this OAuth 2.0 client, which means that if you have different Kafka clusters configured with different OAuth 2.0 client IDs they would each have a

separate set of permissions even though they are part of the same realm.

When the Kafka client use the *SASL OAUTHBearer* mechanism, the Keycloak authorizer ([KeycloakRBACAuthorizer](#)) retrieves the list of grants for the current session from the Keycloak server using the access token of the current session. This list of grants is the result of evaluating the Keycloak Authorization Services policies and permissions.

#### *Introducing authorization scopes*

Typically, an initial configuration involves uploading the authorization scopes to create a list of all the possible actions that can be performed on all the Kafka resource types. This step is performed only once, before defining any permissions. Alternatively, you can add authorization scopes manually.

The authorization scopes should contain all the possible Kafka permissions regardless of the resource type:

- [Create](#)
- [Write](#)
- [Read](#)
- [Delete](#)
- [Describe](#)
- [Alter](#)
- [DescribeConfig](#)
- [AlterConfig](#)
- [ClusterAction](#)
- [IdempotentWrite](#)

#### *Defining resource patterns for permission checks*

The resources use pattern names for pattern matching against the targeted resources when performing permission checks.

The general pattern is as follows: [RESOURCE\\_TYPE:NAME\\_PATTERN](#)

The resource types mirror the Kafka authorization model. The pattern allows for the two matching options: exact matching (when the pattern does not end with `*`), and prefix matching (when the pattern ends with `*`).

#### *Example patterns for resources*

```
Topic:my-topic
Topic:orders-
Group:orders-
Cluster:*
```

In addition, the general pattern can be prefixed by [kafka-cluster:CLUSTER\\_NAME](#) followed by a comma, where the cluster name refers to the [metadata.name](#) in the Kafka custom resource.

## *Example patterns for resources with cluster prefix*

```
kafka-cluster:my-cluster,Topic:*
kafka-cluster:*,Group:b_*
```

When the `kafka-cluster` prefix is not present it is assumed to be `kafka-cluster:*`.

When defining a resource, you can associate a list of possible authorization scopes relevant to the resource. Set whatever actions make sense for the targeted resource type.

While you may add any authorization scope to any resource, only the scopes supported by the resource type are considered for access control.

### *Policies*

Policies are used to target permissions to one or more accounts. Targeting can refer to:

- \* Specific user or service accounts
- \* Realm roles or client roles
- \* User groups
- \* JS rule to match a client IP address

A policy is given a unique name, and can be reused to target multiple permissions to multiple resources.

### *Defining permissions based on scopes, resources and policies*

Use fine-grained permissions to pull together the policies, resources, and authorization scopes that grant access to users.

The name of the each permission should clearly define what permissions it grants to which users.

For more information on how to configure permissions through Keycloak Authorization Services, see [the authorization example](#).

## **Example permissions required for operations on Kafka**

The following examples demonstrate the permissions required for performing common operations on Kafka.

### *Creating a topic*

To create a topic the `Create` permission is required for the specific topic or for `Cluster:kafka-cluster`.

```
bin/kafka-topics.sh --create --topic my-topic \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command
-config=/tmp/config.properties
```

### *Listing the topic*

If user has `Describe` permission on the topic, the topic will be listed, otherwise not.

```
bin/kafka-topics.sh --list \  
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command  
-config=/tmp/config.properties
```

#### *Displaying the topic details*

To display the topic details the **Describe**, and **DescribeConfigs** permissions are required on the topic.

```
bin/kafka-topics.sh --describe --topic my-topic \  
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command  
-config=/tmp/config.properties
```

#### *Producing to the topic*

To produce to the topic the user needs **Describe** and **Write** permissions on the topic. If topic has not yet been created, and autolocation is enabled, the permissions to create the topic are required.

```
bin/kafka-console-producer.sh --topic my-topic \  
--broker-list my-cluster-kafka-bootstrap:9092  
--producer.config=/tmp/config.properties
```

#### *Consuming from the topic*

To consume from the topic the user needs **Describe** and **Read** permissions on the topic. Consuming from the topic normally relies on storing the consumer offsets in a consumer group. That requires additional **Describe** and **Read** permissions on the consumer group.

Two **resources** are needed for matching. For example:

```
Topic:my-topic  
Group:my-group-*
```

```
bin/kafka-console-consumer.sh --topic my-topic --group my-group-1 --from-beginning \  
--bootstrap-server my-cluster-kafka-bootstrap:9092 --consumer.config  
/tmp/config.properties
```

#### *Producing to the topic using an idempotent producer*

Besides needing the permissions for ordinary producing to the topic, an additional **IdempotentWrite** permission is required on the **Cluster** resource.

Two **resources** are needed for matching. For example:

```
Topic:my-topic  
Cluster:kafka-cluster
```

### *Listing consumer groups*

When listing consumer groups, only the groups on which the user has **Describe** permissions are returned. Alternatively, if the user has **Describe** permission on the **Cluster:kafka-cluster**, all the consumer groups are returned.

```
bin/kafka-consumer-groups.sh --list \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command
-config=/tmp/config.properties
```

### *Displaying the consumer group details*

To display the consumer group details the **Describe** permission is required on the group, and on the topic associated with the group.

```
bin/kafka-consumer-groups.sh --describe --group my-group-1 \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command
-config=/tmp/config.properties
```

### *Changing the topic configuration*

To change the topic configuration the **Describe** and **Alter** permissions are required on the topic.

```
bin/kafka-topics.sh --alter --topic my-topic --partitions 2 \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command
-config=/tmp/config.properties
```

### *Displaying the Kafka broker configuration*

To be able to use **kafka-configs.sh** to get the broker configuration, **DescribeConfigs** permission is required on the **Cluster:kafka-cluster**.

```
bin/kafka-configs.sh --entity-type brokers --entity-name 0 --describe --all \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command
-config=/tmp/config.properties
```

### *Changing the Kafka broker configuration*

To change the Kafka broker configuration the 'DescribeConfigs', and **AlterConfigs** permissions are required on **Cluster:kafka-cluster**.

```
bin/kafka-configs --entity-type brokers --entity-name 0 --alter --add-config
log.cleaner.threads=2 \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command
-config=/tmp/config.properties
```

### *Deleting a topic*

To delete the topic the **Describe**, and **Delete** permissions are required on the topic.

```
bin/kafka-topics.sh --delete --topic my-topic \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command-config=/tmp/config.properties
```

#### Selecting a leader partition

To run leader selection for topic partitions you need **Alter** permission on the **Cluster:kafka-cluster**.

```
bin/kafka-leader-election.sh --topic my-topic --partition 0 --election-type PREFERRED \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --admin.config=/tmp/config.properties
```

#### Reassigning partitions

To generate partition reassignment file the **Describe** on the topics involved.

```
bin/kafka-reassign-partitions.sh --topics-to-move-json-file /tmp/topics-to-move.json \
--broker-list "0,1" --generate \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command-config=/tmp/config.properties > /tmp/partition-reassignment.json
```

To execute the partition reassignment the **Describe**, and **Alter** permissions are required on **Cluster:kafka-cluster**, and **Describe** permissions are required on the topics involved.

```
bin/kafka-reassign-partitions.sh --reassignment-json-file /tmp/partition-reassignment.json --execute \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command-config=/tmp/config.properties
```

To verify partition reassignment the **Describe**, and **AlterConfigs** permissions are required on **Cluster:kafka-cluster**, and on each of the topics involved.

```
bin/kafka-reassign-partitions.sh --reassignment-json-file /tmp/partition-reassignment.json --verify \
--bootstrap-server my-cluster-kafka-bootstrap:9092 --command-config=/tmp/config.properties
```

### 4.5.4. Example authorization rules configuration using Authorization Services

This is an end-to-end example of using Keycloak Authorization Services to configure authorization rules for use with **keycloak** authorization in Strimzi. The example starts by deploying the Keycloak server with pre-configured realms, requiring no additional configuration. We deploy the Kafka cluster configured to use one of the pre-configured realms. We then connect to the Keycloak Admin

Console and demonstrate the sections of the user interface where authorizations rules are configured. We use Kafka CLI client tools with different personal and service accounts to demonstrate limited access based on the permissions granted to different accounts.

**CAUTION**

This example is focused on explaining how to use Keycloak Authorization Services. To simplify the example, the communication between components is not properly secured with TLS, and the components themselves are not configured with security in mind.

## Token-based authorization with Keycloak Authorization Services

Once the Kafka Broker has obtained an access token by using `oauth` authentication, it is possible to use centrally managed authorization rules to enforce access restrictions onto Kafka Clients. For this, Strimzi Kafka Operator comes with `keycloak` authorization which uses Keycloak Authorization Services to centrally manage permissions.

When using `keycloak` authorization, a custom authorizer is configured on the Kafka broker that uses Authorization Services REST endpoints available on Keycloak, which provide a list of granted permissions on resources for authenticated users. The list of grants (permissions) is fetched as the first action after an authenticated session is established by the Kafka client, and then regularly refreshed in the background. Grants are cached and enforced locally on the Kafka broker for each user session to provide fast authorization decisions. Because they are refreshed, any changes to the grants on the Keycloak server are detected and enforced.

### Starting the pods

For this example, we use the kubernetes deployment scripts available at <https://github.com/strimzi/strimzi-kafka-oauth/tree/0.6.1/examples/kubernetes>:

```
export ROOT=https://raw.githubusercontent.com/strimzi/strimzi-kafka-oauth/0.6.1
```

#### Deploy the Postgres database for Keycloak

```
kubectl apply -f $ROOT/examples/kubernetes/postgres-pvc.yaml  
kubectl apply -f $ROOT/examples/kubernetes/postgres.yaml
```

#### Deploy the Keycloak server

```
kubectl apply -f $ROOT/examples/kubernetes/keycloak-realms-configmap.yaml  
kubectl apply -f $ROOT/examples/kubernetes/keycloak-postgres.yaml
```

If your default namespace is not `myproject` (for example, `default`), use the following to deploy Keycloak:

```
curl -s $ROOT/examples/kubernetes/keycloak-postgres.yaml | sed -e  
"s#myproject#default#" | kubectl apply -f -
```

## Deploy the minimal Kafka cluster

Here we assume that the Strimzi Cluster Operator has already been installed on the Kubernetes cluster.

```
kubectl apply -f $ROOT/examples/kubernetes/kafka-oauth-single-authz.yaml
```

## Using the Keycloak Admin Console to configure authorization

We login to the Keycloak Admin Console by creating a tunnel to the keycloak pod:

```
kubectl port-forward keycloak 8080
```

Then we use a browser to connect to <http://localhost:8080/auth/admin> using `admin` as username and password.

The default view usually displays the **Master** realm. For this example we are interested in the `kafka-authz` realm, which you can select in the upper left corner.

Initially, the **Realm Settings** section is selected, but you can navigate to **Groups, Roles, Clients** and **Users**.

Under **Groups**, you can view groups to mark users as having some permissions. Groups are sets of users with a name assigned. Typically, they are used to compartmentalize users into geographical, organizational or departmental units, and so on.

In Keycloak, groups can be stored in an LDAP identity provider. That makes it possible to make a user a member of a group through a custom LDAP server admin user interface, for example, to grant them some permissions on Kafka resources.

Under **Users**, you can view all defined users. For this example, `alice` and `bob` are defined. `alice` is a member of the `ClusterManager Group`, and `bob` is a member of `ClusterManager-my-cluster Group`. In Keycloak, users can be stored in an LDAP identity provider.

Under **Roles**, you can view the realm roles to mark users or clients as having some permissions. Roles are a concept analogous to groups. They are usually used to *tag* users with organizational roles and have the requisite permissions. Roles cannot be stored in an LDAP identity provider. If LDAP is a requirement, you can use groups instead, and add Keycloak roles to the groups so that when users are assigned a group, they also get a corresponding role.

Under **Clients**, you can view the additional client configurations. For this example, `kafka`, `kafka-cli`, `team-a-client`, `team-b-client` are configured. The client with client id `kafka` is used by Kafka Brokers to perform the necessary OAuth 2.0 communication for access token validation, and to authenticate to other Kafka Broker instances using OAuth 2.0 client authentication. This client also contains the Authorization Services resource definitions, policies and authorization scopes used to perform authorization on the Kafka Brokers.

The client with client id `kafka-cli` is a public client that can be used by the Kafka command line tools when authenticating with username and password to obtain an access token or a refresh

token.

Clients `team-a-client`, and `team-b-client` are confidential clients representing services with partial access to certain Kafka topics.

The authorization configuration is defined in the `kafka` client from the **Authorization** tab, which becomes visible when **Authorization Enabled** is switched on from the **Settings** tab.

## Defining Authorization Services for access control

Keycloak Authorization Services use authorization scopes, policies and permissions to define and apply access control to resources, as explained in [Keycloak Authorization Services model for managing permissions](#).

From **Authorization / Permissions** you can see the granted permissions that use resources and policies defined from other **Resources** and **Policies** tabs. For example, the `kafka` client has the following permissions:

```
Dev Team A can write to topics that start with x_ on any cluster
Dev Team B can read from topics that start with x_ on any cluster
Dev Team B can update consumer group offsets that start with x_ on any cluster
ClusterManager of my-cluster Group has full access to cluster config on my-cluster
ClusterManager of my-cluster Group has full access to consumer groups on my-cluster
ClusterManager of my-cluster Group has full access to topics on my-cluster
```

`Dev Team A can write to topics that start with x_ on any cluster` combines a resource called `Topic:x_*`, scopes `Describe` and `Write`, and `Dev Team A` policy. The `Dev Team A` policy matches all users that have a realm role called `Dev Team A`.

`Dev Team B can read from topics that start with x_ on any cluster` combines `Topic:x_*`, and `Group:x_*` resources, scopes `Describe` and `Read`, and `Dev Team B` policy. The `Dev Team A` policy matches all users that have a realm role called `Dev Team B`. Matching users and clients have the ability to read from topics, and update the consumed offsets for topics and consumer groups that have names starting with `x_`.

## Targeting permissions using group or role policies

In Keycloak, confidential clients with service accounts enabled can authenticate to the server in their own name using a client ID and a secret. This is convenient for microservices which typically act in their own name, and not as agents of a particular user (like a web site would, for example). Service accounts can have roles assigned like regular users. They cannot, however, have groups assigned. As a consequence, if you want to target permissions to microservices using service accounts, you cannot use group policies, and should instead use role policies. Conversely, if you want to limit certain permissions only to regular user accounts where authentication with username and password is required, you can achieve that as a side effect of using the group policies, rather than the role policies. That is what is used for permissions that start with `ClusterManager`. Performing cluster management is usually done interactively - in person - using CLI tools. It makes sense to require the user to log in, before using the resulting access token to authenticate to the Kafka Broker. In this case, the access token represents the specific user, rather

than the client application.

## Authorization in action using CLI clients

Make sure that authorization rules have been properly imported.

From **Clients** › **kafka** › **Authorization** › **Settings** make sure that **Decision Strategy** is set to **Affirmative**, and NOT to **Unanimous**. Navigate in Keycloak to check that the expected resources, authorization claims, policies and permissions are defined.

With the configuration in place, we can check access to Kafka by using a producer and consumer to create topics using different user and service accounts.

First, a new interactive pod container is run using a Strimzi Kafka image to connect to a running Kafka broker.

```
kubectl run -ti --rm --restart=Never --image=quay.io/strimzi/kafka:latest-kafka-2.7.0
kafka-cli -- /bin/sh
```

**NOTE**

If `kubectl` times out waiting on the image download, subsequent attempts may result in an `AlreadyExists` error.

You can attach to the existing pod by running:

```
kubectl attach -ti kafka-cli
```

To produce messages as client `team-a-client`, we prepare a Kafka client configuration file with authentication parameters:

```
cat > /tmp/team-a-client.properties << EOF
security.protocol=SASL_PLAINTEXT
sasl.mechanism=OAUTHBEARER
sasl.jaas.config=org.apache.kafka.common.security.oauthbearer.OAuthBearerLoginModule
required \
    oauth.client.id="team-a-client" \
    oauth.client.secret="team-a-client-secret" \
    oauth.token.endpoint.uri="http://keycloak:8080/auth/realms/kafka-
authz/protocol/openid-connect/token" ;
sasl.login.callback.handler.class=io.strimzi.kafka.oauth.client.JaasClientOAuthLoginCa
llbackHandler
EOF
```

The roles assigned to a client, such as the `Dev Team A` realm role assigned to the `team-a-client` service account, are presented in Keycloak on the **Service Account Roles** tab from **Clients**.

We can use this configuration from the Kafka CLI to produce and consume messages, and perform other administration tasks.

## *Producing messages with authorized access*

The **team-a-client** configuration is used to produce messages to topic **my-topic**:

```
bin/kafka-console-producer.sh --broker-list my-cluster-kafka-bootstrap:9092 --topic  
my-topic \  
--producer.config=/tmp/team-a-client.properties  
First message
```

A **Not authorized to access topics: [my-topic]** error is returned when trying to push the first message.

**team-a-client** has a **Dev Team A** role that gives it permission to perform any supported actions on topics that start with **a\_**, but can only write to topics that start with **x\_**. The topic named **my-topic** matches neither of those rules.

The **team-a-client** configuration is then used to produce messages to topic **a\_messages**:

```
bin/kafka-console-producer.sh --broker-list my-cluster-kafka-bootstrap:9092 --topic  
a_messages \  
--producer.config /tmp/team-a-client.properties  
First message  
Second message
```

The messages are pushed out successfully, and in the Kafka container log there is DEBUG level output saying **Authorization GRANTED**.

Use CTRL-C to exit the CLI application.

You can see the Kafka container log by running:

```
kubectl logs my-cluster-kafka-0 -f
```

## *Consuming messages with authorized access*

The **team-a-client** configuration is used to consume messages from topic **a\_messages**:

```
bin/kafka-console-consumer.sh --bootstrap-server my-cluster-kafka-bootstrap:9092  
--topic a_messages \  
--from-beginning --consumer.config /tmp/team-a-client.properties
```

An error is returned as the **Dev Team A** role for **team-a-client** only has access to consumer groups that have names starting with **a\_**. The **team-a-client** configuration is then used to consume messages when specifying a custom consumer group with a name that starts with **a\_**:

```
bin/kafka-console-consumer.sh --bootstrap-server my-cluster-kafka-bootstrap:9092  
--topic a_messages \  
--from-beginning --consumer.config /tmp/team-a-client.properties --group  
a_consumer_group_1
```

This time the consumer receives all the messages from the `a_messages` topic.

#### *Administering Kafka with authorized access*

The `team-a-client` configuration is used in administrative operations.

Listing topics returns the `a_messages` topic:

```
bin/kafka-topics.sh --bootstrap-server my-cluster-kafka-bootstrap:9092 --command  
-config /tmp/team-a-client.properties --list
```

Listing consumer groups returns the `a_consumer_group_1` consumer group:

```
bin/kafka-consumer-groups.sh --bootstrap-server my-cluster-kafka-bootstrap:9092  
--command-config /tmp/team-a-client.properties --list
```

Fetching the default cluster configuration fails cluster authorization, because the operation requires cluster level permissions that `team-a-client` does not have:

```
bin/kafka-configs.sh --bootstrap-server my-cluster-kafka-bootstrap:9092 --command  
-config /tmp/team-a-client.properties \  
--entity-type brokers --describe --entity-default
```

#### *Using clients with different permissions*

As with `team-a-client`, we prepare a Kafka client configuration file with authentication parameters for `team-b-client`:

```
cat > /tmp/team-b-client.properties << EOF  
security.protocol=SASL_PLAINTEXT  
sasl.mechanism=OAUTHBEARER  
sasl.jaas.config=org.apache.kafka.common.security.oauthbearer.OAuthBearerLoginModule  
required \  
    oauth.client.id="team-b-client" \  
    oauth.client.secret="team-b-client-secret" \  
    oauth.token.endpoint.uri="http://keycloak:8080/auth/realms/kafka-  
authz/protocol/openid-connect/token" ;  
    sasl.login.callback.handler.class=io.strimzi.kafka.oauth.client.JaasClientOAuthLoginCa  
llbackHandler  
EOF
```

The `team-b-client` client configuration includes a `Dev Team B` realm role and permissions that start with `Dev Team B` ... These match the users and service accounts that have the `Dev Team B` realm role assigned to them. The `Dev Team B` users have full access to topics beginning with `b_` on the Kafka cluster `my-cluster`, the name of the designated cluster, and read access on topics that start with `x_`.

The `team-b-client` configuration is used to produce messages to topic `a_messages`:

```
bin/kafka-console-producer.sh --broker-list my-cluster-kafka-bootstrap:9092 --topic  
a_messages \  
--producer.config /tmp/team-b-client.properties  
Message 1
```

A `Not authorized to access topics: [a_messages]` error is returned when trying to push the first message, as expected, so we switch to topic `b_messages`:

```
bin/kafka-console-producer.sh --broker-list my-cluster-kafka-bootstrap:9092 --topic  
b_messages \  
--producer.config /tmp/team-b-client.properties  
Message 1  
Message 2  
Message 3
```

Producing messages to topic `b_messages` is authorized and successful.

We switch again, but this time to a topic that `team-b-client` can only read from, topic `x_messages`:

```
bin/kafka-console-producer.sh --broker-list my-cluster-kafka-bootstrap:9092 --topic  
x_messages \  
--producer.config /tmp/team-b-client.properties  
Message 1
```

A `Not authorized to access topics: [x_messages]` error is returned, as expected, so we switch to `team-a-client`:

```
bin/kafka-console-producer.sh --broker-list my-cluster-kafka-bootstrap:9092 --topic  
x_messages \  
--producer.config /tmp/team-a-client.properties  
Message 1
```

A `Not authorized to access topics: [x_messages]` error is returned again. Though `team-a-client` can write to the `x_messages` topic, it does not have a permission to create a topic if it does not yet exist.

Before `team-a-client` can write to the `x_messages` topic, a `admin power user` must create it with the correct configuration, such as the number of partitions and replicas.

## Managing Kafka with an authorized admin

Admin user **bob** is created with full access to manage everything on the Kafka cluster **my-cluster**.

Helper scripts are used to authenticate to the **keycloak** instance.

The following scripts are downloaded to **/tmp** dir and made executable:

```
curl https://raw.githubusercontent.com/stimzi/stimzi-kafka-oauth/0.6.1/examples/docker/kafka-oauth-stimzi/kafka/oauth.sh -s > /tmp/oauth.sh  
chmod +x /tmp/oauth.sh
```

```
curl https://raw.githubusercontent.com/stimzi/stimzi-kafka-oauth/0.6.1/examples/docker/kafka-oauth-stimzi/kafka/jwt.sh -s > /tmp/jwt.sh  
chmod +x /tmp/jwt.sh
```

User **bob** authenticates to the Keycloak server with his username and password to get a refresh token:

```
export TOKEN_ENDPOINT=http://keycloak:8080/auth/realms/kafka-authz/protocol/openid-connect/token  
REFRESH_TOKEN=$(./tmp/oauth.sh -q bob)
```

When prompted for a password, 'bob-password' is used.

The refresh token in this case is an offline token which is a long-lived refresh token that does not expire:

```
/tmp/jwt.sh $REFRESH_TOKEN
```

A configuration file is created for **bob**:

```
cat > /tmp/bob.properties << EOF  
security.protocol=SASL_PLAINTEXT  
sasl.mechanism=OAUTHBEARER  
sasl.jaas.config=org.apache.kafka.common.security.oauthbearer.OAuthBearerLoginModule  
required \  
oauth.refresh.token="$REFRESH_TOKEN" \  
oauth.client.id="kafka-cli" \  
oauth.token.endpoint.uri="http://keycloak:8080/auth/realms/kafka-authz/protocol/openid-connect/token" ;  
sasl.login.callback.handler.class=io.stimzi.kafka.oauth.client.JaasClientOAuthLoginCallbackHandler  
EOF
```

The **kafka-cli** public client is used for the **oauth.client.id** in the **sasl.jaas.config**. Since that is a public client it does not require a Secret. We can use this because we authenticate with a token

directly. In this case, the refresh token requests an access token behind the scenes, which is then sent to the Kafka broker for authentication. The refresh token has already been authenticated.

User **bob** has permission to create the **x\_messages** topic:

```
bin/kafka-topics.sh --bootstrap-server my-cluster-kafka-bootstrap:9092 --command
--config /tmp/bob.properties \
--topic x_messages --create --replication-factor 1 --partitions 1
```

User **bob** can list the topic, but **team-a-client** and **team-b-client** cannot:

```
bin/kafka-topics.sh --bootstrap-server my-cluster-kafka-bootstrap:9092 --command
--config /tmp/bob.properties --list
bin/kafka-topics.sh --bootstrap-server my-cluster-kafka-bootstrap:9092 --command
--config /tmp/team-a-client.properties --list
bin/kafka-topics.sh --bootstrap-server my-cluster-kafka-bootstrap:9092 --command
--config /tmp/team-b-client.properties --list
```

The **Dev Team A**, and **Dev Team B** roles both have **Describe** permission on topics that start with **x\_**, but they cannot see the other team's topics as they do not have **Describe** permissions on them.

The **team-a-client** can now successfully produce to the **x\_messages** topic:

```
bin/kafka-console-producer.sh --broker-list my-cluster-kafka-bootstrap:9092 --topic
x_messages \
--producer.config /tmp/team-a-client.properties
Message 1
Message 2
Message 3
```

As expected, **team-b-client** still cannot produce to the **x\_messages** topic, and the following operation returns an error:

```
bin/kafka-console-producer.sh --broker-list my-cluster-kafka-bootstrap:9092 --topic
x_messages \
--producer.config /tmp/team-b-client.properties
Message 4
Message 5
```

However, due to its Keycloak settings **team-b-client** can consume messages from the **x\_messages** topic:

```
bin/kafka-console-consumer.sh --bootstrap-server my-cluster-kafka-bootstrap:9092  
--topic x_messages \  
--from-beginning --consumer.config /tmp/team-b-client.properties --group  
x_consumer_group_b
```

Conversely, even though `team-a-client` can write to topic `x_messages`, the following read request returns a `Not authorized to access group: x_consumer_group_a` error:

```
bin/kafka-console-consumer.sh --bootstrap-server my-cluster-kafka-bootstrap:9092  
--topic x_messages \  
--from-beginning --consumer.config /tmp/team-a-client.properties --group  
x_consumer_group_a
```

A consumer group that begins with `a_` is used in the next read request:

```
bin/kafka-console-consumer.sh --bootstrap-server my-cluster-kafka-bootstrap:9092  
--topic x_messages \  
--from-beginning --consumer.config /tmp/team-a-client.properties --group  
a_consumer_group_a
```

An error is still returned, but this time it is `Not authorized to access topics: [x_messages]`.

`Dev Team A` has no `Read` access on topics that start with '`x_`'.

User `bob` can read from or write to any topic:

```
bin/kafka-console-consumer.sh --bootstrap-server my-cluster-kafka-bootstrap:9092  
--topic x_messages \  
--from-beginning --consumer.config /tmp/bob.properties
```

# Chapter 5. Using Strimzi Operators

Use the Strimzi operators to manage your Kafka cluster, and Kafka topics and users.

## 5.1. Using the Cluster Operator

The Cluster Operator is used to deploy a Kafka cluster and other Kafka components.

The Cluster Operator is deployed using YAML installation files.

For information on deploying the Cluster Operator, see [Deploying the Cluster Operator](#) in the *Deploying and Upgrading Strimzi* guide.

For information on the deployment options available for Kafka, see [Kafka Cluster configuration](#).

**NOTE** On OpenShift, a Kafka Connect deployment can incorporate a Source2Image feature to provide a convenient way to add additional connectors.

### 5.1.1. Cluster Operator configuration

The Cluster Operator can be configured through the following supported environment variables and through the logging configuration.

For more information of `image` configuration, see, [image](#).

#### STRIMZI\_NAMESPACE

A comma-separated list of namespaces that the operator should operate in. When not set, set to empty string, or to `*` the Cluster Operator will operate in all namespaces. The Cluster Operator deployment might use the [Kubernetes Downward API](#) to set this automatically to the namespace the Cluster Operator is deployed in. See the example below:

```
env:  
  - name: STRIMZI_NAMESPACE  
    valueFrom:  
      fieldRef:  
        fieldPath: metadata.namespace
```

#### STRIMZI\_FULL\_RECONCILIATION\_INTERVAL\_MS

Optional, default is 120000 ms. The interval between periodic reconciliations, in milliseconds.

#### STRIMZI\_OPERATION\_TIMEOUT\_MS

Optional, default 300000 ms. The timeout for internal operations, in milliseconds. This value should be increased when using Strimzi on clusters where regular Kubernetes operations take longer than usual (because of slow downloading of Docker images, for example).

#### STRIMZI\_OPERATOR\_NAMESPACE

The name of the namespace where the Strimzi Cluster Operator is running. Do not configure this

variable manually. Use the Kubernetes Downward API.

```
env:  
  - name: STRIMZI_OPERATOR_NAMESPACE  
    valueFrom:  
      fieldRef:  
        fieldPath: metadata.namespace
```

## STRIMZI\_OPERATOR\_NAMESPACE\_LABELS

Optional. The labels of the namespace where the Strimzi Cluster Operator is running. Namespace labels are used to configure the namespace selector in network policies to allow the Strimzi Cluster Operator to only have access to the operands from the namespace with these labels. When not set, the namespace selector in network policies is configured to allow access to the Strimzi Cluster Operator from any namespace in the Kubernetes cluster.

## STRIMZI\_KAFKA\_IMAGES

Required. This provides a mapping from Kafka version to the corresponding Docker image containing a Kafka broker of that version. The required syntax is whitespace or comma separated `<version>=<image>` pairs. For example `2.6.0=quay.io/strimzi/kafka:0.21.1-kafka-2.6.0`, `2.7.0=quay.io/strimzi/kafka:0.21.1-kafka-2.7.0`. This is used when a `Kafka.spec.kafka.version` property is specified but not the `Kafka.spec.kafka.image` in the `Kafka` resource.

## STRIMZI\_DEFAULT\_KAFKA\_INIT\_IMAGE

Optional, default `quay.io/strimzi/operator:0.21.1`. The image name to use as default for the init container started before the broker for initial configuration work (that is, rack support), if no image is specified as the `kafka-init-image` in the `Kafka` resource.

## STRIMZI\_KAFKA\_CONNECT\_IMAGES

Required. This provides a mapping from the Kafka version to the corresponding Docker image containing a Kafka connect of that version. The required syntax is whitespace or comma separated `<version>=<image>` pairs. For example `2.6.0=quay.io/strimzi/kafka:0.21.1-kafka-2.6.0`, `2.7.0=quay.io/strimzi/kafka:0.21.1-kafka-2.7.0`. This is used when a `KafkaConnect.spec.version` property is specified but not the `KafkaConnect.spec.image`.

## STRIMZI\_KAFKA\_CONNECT\_S2I\_IMAGES

Required. This provides a mapping from the Kafka version to the corresponding Docker image containing a Kafka connect of that version. The required syntax is whitespace or comma separated `<version>=<image>` pairs. For example `2.6.0=quay.io/strimzi/kafka:0.21.1-kafka-2.6.0`, `2.7.0=quay.io/strimzi/kafka:0.21.1-kafka-2.7.0`. This is used when a `KafkaConnectS2I.spec.version` property is specified but not the `KafkaConnectS2I.spec.image`.

## STRIMZI\_KAFKA\_MIRROR MAKER IMAGES

Required. This provides a mapping from the Kafka version to the corresponding Docker image containing a Kafka mirror maker of that version. The required syntax is whitespace or comma separated `<version>=<image>` pairs. For example `2.6.0=quay.io/strimzi/kafka:0.21.1-kafka-2.6.0`, `2.7.0=quay.io/strimzi/kafka:0.21.1-kafka-2.7.0`. This is used when a

`KafkaMirrorMaker.spec.version` property is specified but not the `KafkaMirrorMaker.spec.image`.

#### **STRIMZI\_DEFAULT\_TOPIC\_OPERATOR\_IMAGE**

Optional, default `quay.io/strimzi/operator:0.21.1`. The image name to use as the default when deploying the topic operator, if no image is specified as the `Kafka.spec.entityOperator.topicOperator.image` in `Kafka` resource.

#### **STRIMZI\_DEFAULT\_USER\_OPERATOR\_IMAGE**

Optional, default `quay.io/strimzi/operator:0.21.1`. The image name to use as the default when deploying the user operator, if no image is specified as the `Kafka.spec.entityOperator.userOperator.image` in the `Kafka` resource.

#### **STRIMZI\_DEFAULT\_TLS\_SIDECAR\_ENTITY\_OPERATOR\_IMAGE**

Optional, default `quay.io/strimzi/kafka:0.21.1-kafka-2.7.0`. The image name to use as the default when deploying the sidecar container which provides TLS support for the Entity Operator, if no image is specified as the `Kafka.spec.entityOperator.tlsSidecar.image` in the `Kafka` resource.

#### **STRIMZI\_IMAGE\_PULL\_POLICY**

Optional. The `ImagePullPolicy` which will be applied to containers in all pods managed by Strimzi Cluster Operator. The valid values are `Always`, `IfNotPresent`, and `Never`. If not specified, the Kubernetes defaults will be used. Changing the policy will result in a rolling update of all your Kafka, Kafka Connect, and Kafka MirrorMaker clusters.

#### **STRIMZI\_IMAGE\_PULL\_SECRETS**

Optional. A comma-separated list of `Secret` names. The secrets referenced here contain the credentials to the container registries where the container images are pulled from. The secrets are used in the `imagePullSecrets` field for all `Pods` created by the Cluster Operator. Changing this list results in a rolling update of all your Kafka, Kafka Connect, and Kafka MirrorMaker clusters.

#### **STRIMZI\_KUBERNETES\_VERSION**

Optional. Overrides the Kubernetes version information detected from the API server. See the example below:

```
env:  
  - name: STRIMZI_KUBERNETES_VERSION  
    value: |  
      major=1  
      minor=16  
      gitVersion=v1.16.2  
      gitCommit=c97fe5036ef3df2967d086711e6c0c405941e14b  
      gitTreeState=clean  
      buildDate=2019-10-15T19:09:08Z  
      goVersion=go1.12.10  
      compiler=gc  
      platform=linux/amd64
```

## KUBERNETES\_SERVICE\_DNS\_DOMAIN

Optional. Overrides the default Kubernetes DNS domain name suffix.

By default, services assigned in the Kubernetes cluster have a DNS domain name that uses the default suffix `cluster.local`.

For example, for broker `kafka-0`:

```
<cluster-name>-kafka-0.<cluster-name>-kafka-brokers.<namespace>.svc.cluster.local
```

The DNS domain name is added to the Kafka broker certificates used for hostname verification.

If you are using a different DNS domain name suffix in your cluster, change the `KUBERNETES_SERVICE_DNS_DOMAIN` environment variable from the default to the one you are using in order to establish a connection with the Kafka brokers.

## STRIMZI\_CONNECT\_BUILD\_TIMEOUT\_MS

Optional, default 300000 ms. The timeout for building new Kafka Connect images with additional connectots, in milliseconds. This value should be increased when using Strimzi to build container images containing many connectors or using a slow container registry.

## Configuration by ConfigMap

The Cluster Operator's logging is configured by the `strimzi-cluster-operator ConfigMap`.

A `ConfigMap` containing logging configuration is created when installing the Cluster Operator. This `ConfigMap` is described in the file `install/cluster-operator/050-ConfigMap-strimzi-cluster-operator.yaml`. You configure Cluster Operator logging by changing the data field `log4j2.properties` in this `ConfigMap`.

To update the logging configuration, you can edit the `050-ConfigMap-strimzi-cluster-operator.yaml` file and then run the following command:

```
kubectl apply -f install/cluster-operator/050-ConfigMap-strimzi-cluster-operator.yaml
```

Alternatively, edit the `ConfigMap` directly:

```
kubectl edit cm strimzi-cluster-operator
```

To change the frequency of the reload interval, set a time in seconds in the `monitorInterval` option in the created `ConfigMap`.

If the `ConfigMap` is missing when the Cluster Operator is deployed, the default logging values are used.

If the `ConfigMap` is accidentally deleted after the Cluster Operator is deployed, the most recently loaded logging configuration is used. Create a new `ConfigMap` to load a new logging configuration.

**NOTE** Do not remove the monitorInterval option from the ConfigMap.

## Periodic reconciliation

Although the Cluster Operator reacts to all notifications about the desired cluster resources received from the Kubernetes cluster, if the operator is not running, or if a notification is not received for any reason, the desired resources will get out of sync with the state of the running Kubernetes cluster.

In order to handle failovers properly, a periodic reconciliation process is executed by the Cluster Operator so that it can compare the state of the desired resources with the current cluster deployments in order to have a consistent state across all of them. You can set the time interval for the periodic reconciliations using the `STRIMZI_FULL_RECONCILIATION_INTERVAL_MS` variable.

### 5.1.2. Provisioning Role-Based Access Control (RBAC)

For the Cluster Operator to function it needs permission within the Kubernetes cluster to interact with resources such as `Kafka`, `KafkaConnect`, and so on, as well as the managed resources, such as `ConfigMaps`, `Pods`, `Deployments`, `StatefulSets` and `Services`. Such permission is described in terms of Kubernetes role-based access control (RBAC) resources:

- `ServiceAccount`,
- `Role` and `ClusterRole`,
- `RoleBinding` and `ClusterRoleBinding`.

In addition to running under its own `ServiceAccount` with a `ClusterRoleBinding`, the Cluster Operator manages some RBAC resources for the components that need access to Kubernetes resources.

Kubernetes also includes privilege escalation protections that prevent components operating under one `ServiceAccount` from granting other `ServiceAccounts` privileges that the granting `ServiceAccount` does not have. Because the Cluster Operator must be able to create the `ClusterRoleBindings`, and `RoleBindings` needed by resources it manages, the Cluster Operator must also have those same privileges.

## Delegated privileges

When the Cluster Operator deploys resources for a desired `Kafka` resource it also creates `ServiceAccounts`, `RoleBindings`, and `ClusterRoleBindings`, as follows:

- The Kafka broker pods use a `ServiceAccount` called `cluster-name-kafka`
  - When the rack feature is used, the `strimzi-cluster-name-kafka-init` `ClusterRoleBinding` is used to grant this `ServiceAccount` access to the nodes within the cluster via a `ClusterRole` called `strimzi-kafka-broker`
  - When the rack feature is not used no binding is created
- The ZooKeeper pods use a `ServiceAccount` called `cluster-name-zookeeper`
- The Entity Operator pod uses a `ServiceAccount` called `cluster-name-entity-operator`

- The Topic Operator produces Kubernetes events with status information, so the `ServiceAccount` is bound to a `ClusterRole` called `strimzi-entity-operator` which grants this access via the `strimzi-entity-operator RoleBinding`
- The pods for `KafkaConnect` and `KafkaConnectS2I` resources use a `ServiceAccount` called `cluster-name-cluster-connect`
- The pods for `KafkaMirrorMaker` use a `ServiceAccount` called `cluster-name-mirror-maker`
- The pods for `KafkaMirrorMaker2` use a `ServiceAccount` called `cluster-name-mirrormaker2`
- The pods for `KafkaBridge` use a `ServiceAccount` called `cluster-name-bridge`

## `ServiceAccount`

The Cluster Operator is best run using a `ServiceAccount`:

*Example `ServiceAccount` for the Cluster Operator*

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: strimzi-cluster-operator
  labels:
    app: strimzi
```

The `Deployment` of the operator then needs to specify this in its `spec.template.spec.serviceAccountName`:

*Partial example of `Deployment` for the Cluster Operator*

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: strimzi-cluster-operator
  labels:
    app: strimzi
spec:
  replicas: 1
  selector:
    matchLabels:
      name: strimzi-cluster-operator
      strimzi.io/kind: cluster-operator
  template:
    # ...
```

Note line 12, where the `strimzi-cluster-operator ServiceAccount` is specified as the `serviceAccountName`.

## `ClusterRoles`

The Cluster Operator needs to operate using `ClusterRoles` that gives access to the necessary resources. Depending on the Kubernetes cluster setup, a cluster administrator might be needed to

create the [ClusterRoles](#).

**NOTE**

Cluster administrator rights are only needed for the creation of the [ClusterRoles](#).  
The Cluster Operator will not run under the cluster admin account.

The [ClusterRoles](#) follow the *principle of least privilege* and contain only those privileges needed by the Cluster Operator to operate Kafka, Kafka Connect, and ZooKeeper clusters. The first set of assigned privileges allow the Cluster Operator to manage Kubernetes resources such as [StatefulSets](#), [Deployments](#), [Pods](#), and [ConfigMaps](#).

Cluster Operator uses ClusterRoles to grant permission at the namespace-scoped resources level and cluster-scoped resources level:

[ClusterRole](#) with namespaced resources for the Cluster Operator

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: strimzi-cluster-operator-namespaced
  labels:
    app: strimzi
rules:
- apiGroups:
  - "rbac.authorization.k8s.io"
  resources:
    # The cluster operator needs to access and manage rolebindings to grant Strimzi
    components cluster permissions
    - rolebindings
  verbs:
    - get
    - list
    - watch
    - create
    - delete
    - patch
    - update
- apiGroups:
  - "rbac.authorization.k8s.io"
  resources:
    # The cluster operator needs to access and manage roles to grant the entity
    operator permissions
    - roles
  verbs:
    - get
    - list
    - watch
    - create
    - delete
    - patch
    - update
- apiGroups:
```

```

- ""

resources:
  # The cluster operator needs to access and delete pods, this is to allow it to
  monitor pod health and coordinate rolling updates
  - pods
  # The cluster operator needs to access and manage service accounts to grant
  Strimzi components cluster permissions
  - serviceaccounts
  # The cluster operator needs to access and manage config maps for Strimzi
  components configuration
  - configmaps
  # The cluster operator needs to access and manage services and endpoints to
  expose Strimzi components to network traffic
  - services
  - endpoints
  # The cluster operator needs to access and manage secrets to handle credentials
  - secrets
  # The cluster operator needs to access and manage persistent volume claims to
  bind them to Strimzi components for persistent data
  - persistentvolumeclaims
verbs:
  - get
  - list
  - watch
  - create
  - delete
  - patch
  - update
- apiGroups:
  - "kafka.strimzi.io"

resources:
  # The cluster operator runs the KafkaAssemblyOperator, which needs to access and
  manage Kafka resources
  - kafkaes
  - kafkaes/status
  # The cluster operator runs the KafkaConnectAssemblyOperator, which needs to
  access and manage KafkaConnect resources
  - kafkaconnects
  - kafkaconnects/status
  # The cluster operator runs the KafkaConnectS2IAssemblyOperator, which needs to
  access and manage KafkaConnectS2I resources
  - kafkaconnects2is
  - kafkaconnects2is/status
  # The cluster operator runs the KafkaConnectorAssemblyOperator, which needs to
  access and manage KafkaConnector resources
  - kafkaconnectors
  - kafkaconnectors/status
  # The cluster operator runs the KafkaMirrorMakerAssemblyOperator, which needs to
  access and manage KafkaMirrorMaker resources
  - kafkamirrormakers
  - kafkamirrormakers/status

```

```

# The cluster operator runs the KafkaBridgeAssemblyOperator, which needs to
access and manage BridgeMaker resources
- kafkabridges
- kafkabridges/status

# The cluster operator runs the KafkaMirrorMaker2AssemblyOperator, which needs
to access and manage KafkaMirrorMaker2 resources
- kafkamirrormaker2s
- kafkamirrormaker2s/status

# The cluster operator runs the KafkaRebalanceAssemblyOperator, which needs to
access and manage KafkaRebalance resources
- kafkarebalances
- kafkarebalances/status

verbs:
- get
- list
- watch
- create
- delete
- patch
- update

- apiGroups:
  # The cluster operator needs the extensions api as the operator supports
  Kubernetes version 1.11+
  # apps/v1 was introduced in Kubernetes 1.14
  - "extensions"

resources:
  # The cluster operator needs to access and manage deployments to run deployment
  based Strimzi components
  - deployments
  - deployments/scale
  # The cluster operator needs to access replica sets to manage Strimzi components
  # and to determine error states
  - replicaset
  # The cluster operator needs to access and manage replication controllers to
  # manage replicaset
  - replicationcontrollers
  # The cluster operator needs to access and manage network policies to lock down
  communication between Strimzi components
  - networkpolicies
  # The cluster operator needs to access and manage ingresses which allow external
  access to the services in a cluster
  - ingresses

verbs:
- get
- list
- watch
- create
- delete
- patch
- update

- apiGroups:

```

```

- "apps"
resources:
  # The cluster operator needs to access and manage deployments to run deployment
  based Strimzi components
  - deployments
  - deployments/scale
  - deployments/status
  # The cluster operator needs to access and manage stateful sets to run stateful
  sets based Strimzi components
  - statefulsets
  # The cluster operator needs to access replica-sets to manage Strimzi components
  and to determine error states
  - replicaset
verbs:
  - get
  - list
  - watch
  - create
  - delete
  - patch
  - update
- apiGroups:
  - ""

resources:
  # The cluster operator needs to be able to create events and delegate
  permissions to do so
  - events
verbs:
  - create
- apiGroups:
  # OpenShift S2I requirements
  - apps.openshift.io
resources:
  - deploymentconfigs
  - deploymentconfigs/scale
  - deploymentconfigs/status
  - deploymentconfigs/finalizers
verbs:
  - get
  - list
  - watch
  - create
  - delete
  - patch
  - update
- apiGroups:
  # OpenShift S2I requirements
  - build.openshift.io
resources:
  - buildconfigs
  - buildconfigs/instantiate

```

```

    - builds
verbs:
    - get
    - list
    - watch
    - create
    - delete
    - patch
    - update
- apiGroups:
    # OpenShift S2I requirements
    - image.openshift.io
resources:
    - imagestreams
    - imagestreams/status
verbs:
    - get
    - list
    - watch
    - create
    - delete
    - patch
    - update
- apiGroups:
    - networking.k8s.io
resources:
    # The cluster operator needs to access and manage network policies to lock down
    communication between Strimzi components
    - networkpolicies
verbs:
    - get
    - list
    - watch
    - create
    - delete
    - patch
    - update
- apiGroups:
    - route.openshift.io
resources:
    # The cluster operator needs to access and manage routes to expose Strimzi
    components for external access
    - routes
    - routes/custom-host
verbs:
    - get
    - list
    - watch
    - create
    - delete
    - patch

```

```
- update
- apiGroups:
  - policy
resources:
  # The cluster operator needs to access and manage pod disruption budgets this
  # limits the number of concurrent disruptions
  # that a Strimzi component experiences, allowing for higher availability
  - poddisruptionbudgets
verbs:
- get
- list
- watch
- create
- delete
- patch
- update
```

The second includes the permissions needed for cluster-scoped resources.

## **ClusterRole** with cluster-scoped resources for the Cluster Operator

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: strimzi-cluster-operator-global
  labels:
    app: strimzi
rules:
- apiGroups:
  - "rbac.authorization.k8s.io"
  resources:
    # The cluster operator needs to create and manage cluster role bindings in the
    # case of an install where a user
    # has specified they want their cluster role bindings generated
    - clusterrolebindings
  verbs:
    - get
    - list
    - watch
    - create
    - delete
    - patch
    - update
- apiGroups:
  - storage.k8s.io
  resources:
    # The cluster operator requires "get" permissions to view storage class details
    # This is because only a persistent volume of a supported storage class type can
    be resized
    - storageclasses
  verbs:
    - get
- apiGroups:
  - ""
  resources:
    # The cluster operator requires "list" permissions to view all nodes in a
    cluster
    # The listing is used to determine the node addresses when NodePort access is
    configured
    # These addresses are then exposed in the custom resource states
    - nodes
  verbs:
    - list
```

The **strimzi-kafka-broker ClusterRole** represents the access needed by the init container in Kafka pods that is used for the rack feature. As described in the [Delegated privileges](#) section, this role is also needed by the Cluster Operator in order to be able to delegate this access.

[ClusterRole](#) for the Cluster Operator allowing it to delegate access to Kubernetes nodes to the Kafka broker pods

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: strimzi-kafka-broker
  labels:
    app: strimzi
rules:
  - apiGroups:
    - ""
      resources:
        # The Kafka Brokers require "get" permissions to view the node they are on
        # This information is used to generate a Rack ID that is used for High
        Availability configurations
        - nodes
      verbs:
        - get
```

The [strimzi-topic-operator ClusterRole](#) represents the access needed by the Topic Operator. As described in the [Delegated privileges](#) section, this role is also needed by the Cluster Operator in order to be able to delegate this access.

**ClusterRole** for the Cluster Operator allowing it to delegate access to events to the Topic Operator

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: strimzi-entity-operator
  labels:
    app: strimzi
rules:
  - apiGroups:
      - "kafka.strimzi.io"
    resources:
      # The entity operator runs the KafkaTopic assembly operator, which needs to
      # access and manage KafkaTopic resources
      - kafkatopics
      - kafkatopics/status
      # The entity operator runs the KafkaUser assembly operator, which needs to
      # access and manage KafkaUser resources
      - kafkausers
      - kafkausers/status
    verbs:
      - get
      - list
      - watch
      - create
      - patch
      - update
      - delete
  - apiGroups:
      - ""
    resources:
      - events
    verbs:
      # The entity operator needs to be able to create events
      - create
  - apiGroups:
      - ""
    resources:
      # The entity operator user-operator needs to access and manage secrets to store
      # generated credentials
      - secrets
    verbs:
      - get
      - list
      - watch
      - create
      - delete
      - patch
      - update
```

The `strimzi-kafka-client ClusterRole` represents the access needed by the components based on Kafka clients which use the client rack-awareness. As described in the [Delegated privileges](#) section, this role is also needed by the Cluster Operator in order to be able to delegate this access.

*ClusterRole for the Cluster Operator allowing it to delegate access to Kubernetes nodes to the Kafka client based pods*

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: strimzi-kafka-client
  labels:
    app: strimzi
rules:
  - apiGroups:
    - ""
      resources:
        # The Kafka clients (Connect, Mirror Maker, etc.) require "get" permissions to
        # view the node they are on
        # This information is used to generate a Rack ID (client.rack option) that is
        # used for consuming from the closest
        # replicas when enabled
        - nodes
    verbs:
      - get
```

## ClusterRoleBindings

The operator needs `ClusterRoleBindings` and `RoleBindings` which associates its `ClusterRole` with its `ServiceAccount`: `ClusterRoleBindings` are needed for `ClusterRoles` containing cluster-scoped resources.

*Example ClusterRoleBinding for the Cluster Operator*

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: strimzi-cluster-operator
  labels:
    app: strimzi
subjects:
  - kind: ServiceAccount
    name: strimzi-cluster-operator
    namespace: myproject
roleRef:
  kind: ClusterRole
  name: strimzi-cluster-operator-global
  apiGroup: rbac.authorization.k8s.io
```

`ClusterRoleBindings` are also needed for the `ClusterRoles` needed for delegation:

*Example ClusterRoleBinding for the Cluster Operator for the Kafka broker rack-awareness*

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: strimzi-cluster-operator-kafka-broker-delegation
  labels:
    app: strimzi
  # The Kafka broker cluster role must be bound to the cluster operator service account
  # so that it can delegate the cluster role to the Kafka brokers.
  # This must be done to avoid escalating privileges which would be blocked by
  # Kubernetes.
subjects:
  - kind: ServiceAccount
    name: strimzi-cluster-operator
    namespace: myproject
roleRef:
  kind: ClusterRole
  name: strimzi-kafka-broker
  apiGroup: rbac.authorization.k8s.io
```

and

*Example ClusterRoleBinding for the Cluster Operator for the Kafka client rack-awareness*

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: strimzi-cluster-operator-kafka-client-delegation
  labels:
    app: strimzi
  # The Kafka clients cluster role must be bound to the cluster operator service account
  # so that it can delegate the
  # cluster role to the Kafka clients using it for consuming from closest replica.
  # This must be done to avoid escalating privileges which would be blocked by
  # Kubernetes.
subjects:
  - kind: ServiceAccount
    name: strimzi-cluster-operator
    namespace: myproject
roleRef:
  kind: ClusterRole
  name: strimzi-kafka-client
  apiGroup: rbac.authorization.k8s.io
```

ClusterRoles containing only namespaced resources are bound using RoleBindings only.

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: strimzi-cluster-operator
  labels:
    app: strimzi
subjects:
  - kind: ServiceAccount
    name: strimzi-cluster-operator
    namespace: myproject
roleRef:
  kind: ClusterRole
  name: strimzi-cluster-operator-namespaced
  apiGroup: rbac.authorization.k8s.io
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: strimzi-cluster-operator-entity-operator-delegation
  labels:
    app: strimzi
# The Entity Operator cluster role must be bound to the cluster operator service
# account so that it can delegate the cluster role to the Entity Operator.
# This must be done to avoid escalating privileges which would be blocked by
# Kubernetes.
subjects:
  - kind: ServiceAccount
    name: strimzi-cluster-operator
    namespace: myproject
roleRef:
  kind: ClusterRole
  name: strimzi-entity-operator
  apiGroup: rbac.authorization.k8s.io
```

## 5.2. Using the Topic Operator

When you create, modify or delete a topic using the [KafkaTopic](#) resource, the Topic Operator ensures those changes are reflected in the Kafka cluster.

The [Deploying and Upgrading Strimzi](#) guide provides instructions to deploy the Topic Operator:

- [Using the Cluster Operator \(recommended\)](#)
- [Standalone to operate with Kafka clusters not managed by Strimzi](#)

### 5.2.1. Kafka topic resource

The [KafkaTopic](#) resource is used to configure topics, including the number of partitions and replicas.

The full schema for [KafkaTopic](#) is described in [KafkaTopic schema reference](#).

## Identifying a Kafka cluster for topic handling

A [KafkaTopic](#) resource includes a label that defines the appropriate name of the Kafka cluster (derived from the name of the [Kafka](#) resource) to which it belongs.

For example:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaTopic
metadata:
  name: topic-name-1
  labels:
    strimzi.io/cluster: my-cluster
```

The label is used by the Topic Operator to identify the [KafkaTopic](#) resource and create a new topic, and also in subsequent handling of the topic.

If the label does not match the Kafka cluster, the Topic Operator cannot identify the [KafkaTopic](#) and the topic is not created.

## Handling changes to topics

A fundamental problem that the Topic Operator has to solve is that there is no single source of truth: Both the [KafkaTopic](#) resource and the Kafka topic can be modified independently of the operator. Complicating this, the Topic Operator might not always be able to observe changes at each end in real time (for example, the operator might be down).

To resolve this, the operator maintains its own private copy of the information about each topic. When a change happens either in the Kafka cluster, or in Kubernetes, it looks at both the state of the other system and at its private copy in order to determine what needs to change to keep everything in sync. The same thing happens whenever the operator starts, and periodically while it is running.

For example, suppose the Topic Operator is not running, and a [KafkaTopic](#) `my-topic` gets created. When the operator starts it will lack a private copy of "my-topic", so it can infer that the [KafkaTopic](#) has been created since it was last running. The operator will create the topic corresponding to `my-topic`, and also store a private copy of the metadata for `my-topic`.

The private copy allows the operator to cope with scenarios where the topic configuration gets changed both in Kafka and in Kubernetes, so long as the changes are not incompatible (for example, both changing the same topic config key, but to different values). In the case of incompatible changes, the Kafka configuration wins, and the [KafkaTopic](#) will be updated to reflect that.

The private copy is held in the same ZooKeeper ensemble used by Kafka itself. This mitigates availability concerns, because if ZooKeeper is not running then Kafka itself cannot run, so the operator will be no less available than it would even if it was stateless.

## Kafka topic usage recommendations

When working with topics, be consistent. Always operate on either `KafkaTopic` resources or topics directly in Kubernetes. Avoid routinely switching between both methods for a given topic.

Use topic names that reflect the nature of the topic, and remember that names cannot be changed later.

If creating a topic in Kafka, use a name that is a valid Kubernetes resource name, otherwise the Topic Operator will need to create the corresponding `KafkaTopic` with a name that conforms to the Kubernetes rules.

**NOTE**

Recommendations for identifiers and names in Kubernetes are outlined in [Identifiers and Names in Kubernetes](#) community article.

## Kafka topic naming conventions

Kafka and Kubernetes impose their own validation rules for the naming of topics in Kafka and `KafkaTopic.metadata.name` respectively. There are valid names for each which are invalid in the other.

Using the `spec.topicName` property, it is possible to create a valid topic in Kafka with a name that would be invalid for the Kafka topic in Kubernetes.

The `spec.topicName` property inherits Kafka naming validation rules:

- The name must not be longer than 249 characters.
- Valid characters for Kafka topics are ASCII alphanumerics, `.`, `_`, and `-`.
- The name cannot be `.` or `..`, though `.` can be used in a name, such as `exampleTopic.` or `.exampleTopic`.

`spec.topicName` must not be changed.

For example:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaTopic
metadata:
  name: topic-name-1
spec:
  topicName: topicName-1 ①
  # ...
```

① Upper case is invalid in Kubernetes.

cannot be changed to:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaTopic
metadata:
  name: topic-name-1
spec:
  topicName: name-2
  # ...
```

Some Kafka client applications, such as Kafka Streams, can create topics in Kafka programmatically. If those topics have names that are invalid Kubernetes resource names, the Topic Operator gives them a valid `metadata.name` based on the Kafka name. Invalid characters are replaced and a hash is appended to the name. For example:

**NOTE**

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaTopic
metadata:
  name: mytopic--c55e57fe2546a33f9e603caf57165db4072e827e
spec:
  topicName: myTopic
  # ...
```

### 5.2.2. Configuring a Kafka topic

Use the properties of the `KafkaTopic` resource to configure a Kafka topic.

You can use `kubectl apply` to create or modify topics, and `kubectl delete` to delete existing topics.

For example:

- `kubectl apply -f <topic-config-file>`
- `kubectl delete KafkaTopic <topic-name>`

This procedure shows how to create a topic with 10 partitions and 2 replicas.

*Before you start*

It is important that you consider the following before making your changes:

- Kafka does *not* support making the following changes through the `KafkaTopic` resource:
  - Changing topic names using `spec.topicName`
  - Decreasing partition size using `spec.partitions`
- You cannot use `spec.replicas` to change the number of replicas that were initially specified.
- Increasing `spec.partitions` for topics with keys will change how records are partitioned, which can be particularly problematic when the topic uses *semantic partitioning*.

## Prerequisites

- A running Kafka cluster [configured with a Kafka broker listener using TLS authentication and encryption](#).
- A running Topic Operator (typically [deployed with the Entity Operator](#)).
- For deleting a topic, `delete.topic.enable=true` (default) in the `spec.kafka.config` of the [Kafka](#) resource.

## Procedure

1. Prepare a file containing the [KafkaTopic](#) to be created.

*An example KafkaTopic*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaTopic
metadata:
  name: orders
  labels:
    strimzi.io/cluster: my-cluster
spec:
  partitions: 10
  replicas: 2
```

**TIP**

When modifying a topic, you can get the current version of the resource using `kubectl get kafkatopic orders -o yaml`.

2. Create the [KafkaTopic](#) resource in Kubernetes.

```
kubectl apply -f TOPIC-CONFIG-FILE
```

## 5.2.3. Configuring the Topic Operator with resource requests and limits

You can allocate resources, such as CPU and memory, to the Topic Operator and set a limit on the amount of resources it can consume.

## Prerequisites

- The Cluster Operator is running.

## Procedure

1. Update the Kafka cluster configuration in an editor, as required:

```
kubectl edit kafka MY-CLUSTER
```

2. In the `spec.entityOperator.topicOperator.resources` property in the [Kafka](#) resource, set the resource requests and limits for the Topic Operator.

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  # Kafka and ZooKeeper sections...
  entityOperator:
    topicOperator:
      resources:
        requests:
          cpu: "1"
          memory: 500Mi
        limits:
          cpu: "1"
          memory: 500Mi
```

3. Apply the new configuration to create or update the resource.

```
kubectl apply -f KAFKA-CONFIG-FILE
```

## 5.3. Using the User Operator

When you create, modify or delete a user using the [KafkaUser](#) resource, the User Operator ensures those changes are reflected in the Kafka cluster.

The [Deploying and Upgrading Strimzi](#) guide provides instructions to deploy the User Operator:

- [Using the Cluster Operator \(recommended\)](#)
- [Standalone to operate with Kafka clusters not managed by Strimzi](#)

For more information about the schema, see [KafkaUser schema reference](#).

*Authenticating and authorizing access to Kafka*

Use [KafkaUser](#) to enable the authentication and authorization mechanisms that a specific client uses to access Kafka.

For more information on using [KafkaUser](#) to manage users and secure access to Kafka brokers, see [Securing access to Kafka brokers](#).

### 5.3.1. Configuring the User Operator with resource requests and limits

You can allocate resources, such as CPU and memory, to the User Operator and set a limit on the amount of resources it can consume.

*Prerequisites*

- The Cluster Operator is running.

*Procedure*

1. Update the Kafka cluster configuration in an editor, as required:

```
kubectl edit kafka MY-CLUSTER
```

2. In the `spec.entityOperator.userOperator.resources` property in the `Kafka` resource, set the resource requests and limits for the User Operator.

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  # Kafka and ZooKeeper sections...
  entityOperator:
    userOperator:
      resources:
        requests:
          cpu: "1"
          memory: 500Mi
        limits:
          cpu: "1"
          memory: 500Mi
```

Save the file and exit the editor. The Cluster Operator applies the changes automatically.

## 5.4. Monitoring operators using Prometheus metrics

Strimzi operators expose Prometheus metrics. The metrics are automatically enabled and contain information about:

- Number of reconciliations
- Number of Custom Resources the operator is processing
- Duration of reconciliations
- JVM metrics from the operators

Additionally, we provide an example Grafana dashboard.

For more information about Prometheus, see the [Introducing Metrics to Kafka](#) in the *Deploying and Upgrading Strimzi* guide.

# Chapter 6. Kafka Bridge

This chapter provides an overview of the Strimzi Kafka Bridge and helps you get started using its REST API to interact with Strimzi.

- To try out the Kafka Bridge in your local environment, see the [Kafka Bridge quickstart](#) later in this chapter.
- For detailed configuration steps, see [Kafka Bridge cluster configuration](#).
- To view the API documentation, see the [Kafka Bridge API reference](#).

## 6.1. Kafka Bridge overview

You can use the Strimzi Kafka Bridge as an interface to make specific types of HTTP requests to the Kafka cluster.

### 6.1.1. Kafka Bridge interface

The Kafka Bridge provides a RESTful interface that allows HTTP-based clients to interact with a Kafka cluster. It offers the advantages of a web API connection to Strimzi, without the need for client applications to interpret the Kafka protocol.

The API has two main resources — `consumers` and `topics` — that are exposed and made accessible through endpoints to interact with consumers and producers in your Kafka cluster. The resources relate only to the Kafka Bridge, not the consumers and producers connected directly to Kafka.

#### HTTP requests

The Kafka Bridge supports HTTP requests to a Kafka cluster, with methods to:

- Send messages to a topic.
- Retrieve messages from topics.
- Retrieve a list of partitions for a topic.
- Create and delete consumers.
- Subscribe consumers to topics, so that they start receiving messages from those topics.
- Retrieve a list of topics that a consumer is subscribed to.
- Unsubscribe consumers from topics.
- Assign partitions to consumers.
- Commit a list of consumer offsets.
- Seek on a partition, so that a consumer starts receiving messages from the first or last offset position, or a given offset position.

The methods provide JSON responses and HTTP response code error handling. Messages can be sent in JSON or binary formats.

Clients can produce and consume messages without the requirement to use the native Kafka protocol.

#### Additional resources

- To view the API documentation, including example requests and responses, see the [Kafka Bridge API reference](#).

### 6.1.2. Supported clients for the Kafka Bridge

You can use the Kafka Bridge to integrate both *internal* and *external* HTTP client applications with your Kafka cluster.


#### Internal clients

Internal clients are container-based HTTP clients running in *the same* Kubernetes cluster as the Kafka Bridge itself. Internal clients can access the Kafka Bridge on the host and port defined in the [KafkaBridge](#) custom resource.

#### External clients

External clients are HTTP clients running *outside* the Kubernetes cluster in which the Kafka Bridge is deployed and running. External clients can access the Kafka Bridge through an OpenShift Route, a loadbalancer service, or using an Ingress.

#### HTTP internal and external client integration



39\_0819

### 6.1.3. Securing the Kafka Bridge

Strimzi does not currently provide any encryption, authentication, or authorization for the Kafka

Bridge. This means that requests sent from external clients to the Kafka Bridge are:

- Not encrypted, and must use HTTP rather than HTTPS
- Sent without authentication

However, you can secure the Kafka Bridge using other methods, such as:

- Kubernetes Network Policies that define which pods can access the Kafka Bridge.
- Reverse proxies with authentication or authorization, for example, OAuth2 proxies.
- API Gateways.
- Ingress or OpenShift Routes with TLS termination.

The Kafka Bridge supports TLS encryption and TLS and SASL authentication when connecting to the Kafka Brokers. Within your Kubernetes cluster, you can configure:

- TLS or SASL-based authentication between the Kafka Bridge and your Kafka cluster
- A TLS-encrypted connection between the Kafka Bridge and your Kafka cluster.

For more information, see [Configuring the Kafka Bridge](#).

You can use ACLs in Kafka brokers to restrict the topics that can be consumed and produced using the Kafka Bridge.

#### 6.1.4. Accessing the Kafka Bridge outside of Kubernetes

After deployment, the Strimzi Kafka Bridge can only be accessed by applications running in the same Kubernetes cluster. These applications use the `kafka-bridge-name-bridge-service` Service to access the API.

If you want to make the Kafka Bridge accessible to applications running outside of the Kubernetes cluster, you can expose it manually by using one of the following features:

- Services of types LoadBalancer or NodePort
- Ingress resources
- OpenShift Routes

If you decide to create Services, use the following labels in the `selector` to configure the pods to which the service will route the traffic:

```
# ...
selector:
  strimzi.io/cluster: kafka-bridge-name ①
  strimzi.io/kind: KafkaBridge
#...
```

① Name of the Kafka Bridge custom resource in your Kubernetes cluster.

## 6.1.5. Requests to the Kafka Bridge

Specify data formats and HTTP headers to ensure valid requests are submitted to the Kafka Bridge.

### Content Type headers

API request and response bodies are always encoded as JSON.

- When performing consumer operations, **POST** requests must provide the following **Content-Type** header if there is a non-empty body:

```
Content-Type: application/vnd.kafka.v2+json
```

- When performing producer operations, **POST** requests must provide **Content-Type** headers specifying the *embedded data format* of the messages produced. This can be either **json** or **binary**.

Embedded data format	Content-Type header
JSON	<b>Content-Type: application/vnd.kafka.json.v2+json</b>
Binary	<b>Content-Type: application/vnd.kafka.binary.v2+json</b>

The embedded data format is set per consumer, as described in the next section.

The **Content-Type** must *not* be set if the **POST** request has an empty body. An empty body can be used to create a consumer with the default values.

### Embedded data format

The embedded data format is the format of the Kafka messages that are transmitted, over HTTP, from a producer to a consumer using the Kafka Bridge. Two embedded data formats are supported: JSON and binary.

When creating a consumer using the **/consumers/groupid** endpoint, the **POST** request body must specify an embedded data format of either JSON or binary. This is specified in the **format** field, for example:

```
{
  "name": "my-consumer",
  "format": "binary", ①
  ...
}
```

① A binary embedded data format.

The embedded data format specified when creating a consumer must match the data format of the Kafka messages it will consume.

If you choose to specify a binary embedded data format, subsequent producer requests must

provide the binary data in the request body as Base64-encoded strings. For example, when sending messages using the `/topics/topicname` endpoint, `records.value` must be encoded in Base64:

```
{  
  "records": [  
    {  
      "key": "my-key",  
      "value": "ZWR3YXJkdGhldGhyZWVsZWdnZWRjYXQ="  
    },  
  ]  
}
```

Producer requests must also provide a `Content-Type` header that corresponds to the embedded data format, for example, `Content-Type: application/vnd.kafka.binary.v2+json`.

## Message format

When sending messages using the `/topics` endpoint, you enter the message payload in the request body, in the `records` parameter.

The `records` parameter can contain any of these optional fields:

- Message `headers`
- Message `key`
- Message `value`
- Destination `partition`

*Example POST request to /topics*

```
curl -X POST \  
  http://localhost:8080/topics/my-topic \  
  -H 'content-type: application/vnd.kafka.json.v2+json' \  
  -d '{  
    "records": [  
      {  
        "key": "my-key",  
        "value": "sales-lead-0001"  
        "partition": 2  
        "headers": [  
          {  
            "key": "key1",  
            "value": "QXBhY2h1IEthZmthIGlzIHRoZSBib21iIQ==" ①  
          }  
        ]  
      },  
    ]  
  }'
```

- ① The header value in binary format and encoded as Base64.

## Accept headers

After creating a consumer, all subsequent GET requests must provide an **Accept** header in the following format:

```
Accept: application/vnd.kafka.EMBEDDED-DATA-FORMAT.v2+json
```

The **EMBEDDED-DATA-FORMAT** is either **json** or **binary**.

For example, when retrieving records for a subscribed consumer using an embedded data format of JSON, include this Accept header:

```
Accept: application/vnd.kafka.json.v2+json
```

### 6.1.6. CORS

Cross-Origin Resource Sharing (CORS) allows you to specify allowed methods and originating URLs for accessing the Kafka cluster in your [Kafka Bridge HTTP configuration](#).

*Example CORS configuration for Kafka Bridge*

```
# ...
cors:
  allowedOrigins: "https://strimzi.io"
  allowedMethods: "GET,POST,PUT,DELETE,OPTIONS,PATCH"
# ...
```

CORS allows for *simple* and *preflighted* requests between origin sources on different domains.

Simple requests are suitable for standard requests using **GET**, **HEAD**, **POST** methods.

A preflighted request sends a *HTTP OPTIONS* request as an initial check that the actual request is safe to send. On confirmation, the actual request is sent. Preflight requests are suitable for methods that require greater safeguards, such as **PUT** and **DELETE**, and use non-standard headers.

All requests require an **Origin** value in their header, which is the source of the HTTP request.

## Simple request

For example, this simple request header specifies the origin as <https://strimzi.io>.

```
Origin: https://strimzi.io
```

The header information is added to the request.

```
curl -v -X GET HTTP-ADDRESS/bridge-consumer/records \  
-H 'Origin: https://strimzi.io'\ \  
-H 'content-type: application/vnd.kafka.v2+json'
```

In the response from the Kafka Bridge, an **Access-Control-Allow-Origin** header is returned.

```
HTTP/1.1 200 OK  
Access-Control-Allow-Origin: * ①
```

① Returning an asterisk (\*) shows the resource can be accessed by any domain.

## Preflighted request

An initial preflight request is sent to Kafka Bridge using an **OPTIONS** method. The *HTTP OPTIONS* request sends header information to check that Kafka Bridge will allow the actual request.

Here the preflight request checks that a **POST** request is valid from <https://strimzi.io>.

```
OPTIONS /my-group/instances/my-user/subscription HTTP/1.1  
Origin: https://strimzi.io  
Access-Control-Request-Method: POST ①  
Access-Control-Request-Headers: Content-Type ②
```

① Kafka Bridge is alerted that the actual request is a **POST** request.

② The actual request will be sent with a **Content-Type** header.

**OPTIONS** is added to the header information of the preflight request.

```
curl -v -X OPTIONS -H 'Origin: https://strimzi.io' \  
-H 'Access-Control-Request-Method: POST' \  
-H 'content-type: application/vnd.kafka.v2+json'
```

Kafka Bridge responds to the initial request to confirm that the request will be accepted. The response header returns allowed origins, methods and headers.

```
HTTP/1.1 200 OK  
Access-Control-Allow-Origin: https://strimzi.io  
Access-Control-Allow-Methods: GET,POST,PUT,DELETE,OPTIONS,PATCH  
Access-Control-Allow-Headers: content-type
```

If the origin or method is rejected, an error message is returned.

The actual request does not require **Access-Control-Request-Method** header, as it was confirmed in the preflight request, but it does require the origin header.

```
curl -v -X POST HTTP-ADDRESS/topics/bridge-topic \
-H 'Origin: https://strimzi.io' \
-H 'content-type: application/vnd.kafka.v2+json'
```

The response shows the originating URL is allowed.

```
HTTP/1.1 200 OK
Access-Control-Allow-Origin: https://strimzi.io
```

*Additional resources*

[Fetch](#) CORS specification

### 6.1.7. Kafka Bridge API resources

For the full list of REST API endpoints and descriptions, including example requests and responses, see the [Kafka Bridge API reference](#).

### 6.1.8. Kafka Bridge deployment

You deploy the Kafka Bridge into your Kubernetes cluster by using the Cluster Operator.

After the Kafka Bridge is deployed, the Cluster Operator creates Kafka Bridge objects in your Kubernetes cluster. Objects include the *deployment*, *service*, and *pod*, each named after the name given in the custom resource for the Kafka Bridge.

*Additional resources*

- For deployment instructions, see [Deploying Kafka Bridge to your Kubernetes cluster](#) in the [Deploying and Upgrading Strimzi](#) guide.
- For detailed information on configuring the Kafka Bridge, see [Kafka Bridge cluster configuration](#)
- For information on configuring the host and port for the `KafkaBridge` resource, see [Configuring the Kafka Bridge](#).
- For information on integrating external clients, see [Accessing the Kafka Bridge outside of Kubernetes](#).

## 6.2. Kafka Bridge quickstart

Use this quickstart to try out the Strimzi Kafka Bridge in your local development environment. You will learn how to:

- Deploy the Kafka Bridge to your Kubernetes cluster
- Expose the Kafka Bridge service to your local machine by using port-forwarding
- Produce messages to topics and partitions in your Kafka cluster
- Create a Kafka Bridge consumer

- Perform basic consumer operations, such as subscribing the consumer to topics and retrieving the messages that you produced

In this quickstart, HTTP requests are formatted as curl commands that you can copy and paste to your terminal. Access to a Kubernetes cluster is required; to run and manage a local Kubernetes cluster, use a tool such as Minikube, CodeReady Containers, or MiniShift.

Ensure you have the prerequisites and then follow the tasks in the order provided in this chapter.

#### *About data formats*

In this quickstart, you will produce and consume messages in JSON format, not binary. For more information on the data formats and HTTP headers used in the example requests, see [Requests to the Kafka Bridge](#).

#### *Prerequisites for the quickstart*

- Cluster administrator access to a local or remote Kubernetes cluster.
- Strimzi is installed.
- A running Kafka cluster, deployed by the Cluster Operator, in a Kubernetes namespace.
- The Entity Operator is deployed and running as part of the Kafka cluster.

### 6.2.1. Deploying the Kafka Bridge to your Kubernetes cluster

Strimzi includes a YAML example that specifies the configuration of the Strimzi Kafka Bridge. Make some minimal changes to this file and then deploy an instance of the Kafka Bridge to your Kubernetes cluster.

#### *Procedure*

1. Edit the `examples/bridge/kafka-bridge.yaml` file.

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaBridge
metadata:
  name: quickstart ①
spec:
  replicas: 1
  bootstrapServers: <cluster-name>-kafka-bootstrap:9092 ②
  http:
    port: 8080
```

① When the Kafka Bridge is deployed, `-bridge` is appended to the name of the deployment and other related resources. In this example, the Kafka Bridge deployment is named `quickstart-bridge` and the accompanying Kafka Bridge service is named `quickstart-bridge-service`.

② In the `bootstrapServers` property, enter the name of the Kafka cluster as the `<cluster-name>`.

2. Deploy the Kafka Bridge to your Kubernetes cluster:

```
kubectl apply -f examples/bridge/kafka-bridge.yaml
```

A `quickstart-bridge` deployment, service, and other related resources are created in your Kubernetes cluster.

- Verify that the Kafka Bridge was successfully deployed:

```
kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
quickstart-bridge	1/1	1	1	34m
my-cluster-connect	1/1	1	1	24h
my-cluster-entity-operator	1/1	1	1	24h
#...				

#### What to do next

After deploying the Kafka Bridge to your Kubernetes cluster, [expose the Kafka Bridge service to your local machine](#).

#### Additional resources

- For more detailed information about configuring the Kafka Bridge, see [Kafka Bridge cluster configuration](#).

### 6.2.2. Exposing the Kafka Bridge service to your local machine

Next, use port forwarding to expose the Strimzi Kafka Bridge service to your local machine on <http://localhost:8080>.

**NOTE** Port forwarding is only suitable for development and testing purposes.

#### Procedure

- List the names of the pods in your Kubernetes cluster:

```
kubectl get pods -o name  
  
pod/kafka-consumer  
# ...  
pod/quickstart-bridge-589d78784d-9jcnr  
pod/strimzi-cluster-operator-76bcf9bc76-8dnfm
```

- Connect to the `quickstart-bridge` pod on port `8080`:

```
kubectl port-forward pod/quickstart-bridge-589d78784d-9jcnr 8080:8080 &
```

**NOTE**

If port 8080 on your local machine is already in use, use an alternative HTTP port, such as **8008**.

API requests are now forwarded from port 8080 on your local machine to port 8080 in the Kafka Bridge pod.

### 6.2.3. Producing messages to topics and partitions

Next, produce messages to topics in JSON format by using the [topics](#) endpoint. You can specify destination partitions for messages in the request body, as shown here. The [partitions](#) endpoint provides an alternative method for specifying a single destination partition for all messages as a path parameter.

*Procedure*

1. In a text editor, create a YAML definition for a Kafka topic with three partitions.

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaTopic
metadata:
  name: bridge-quickstart-topic
  labels:
    strimzi.io/cluster: <kafka-cluster-name> ①
spec:
  partitions: 3 ②
  replicas: 1
  config:
    retention.ms: 7200000
    segment.bytes: 1073741824
```

① The name of the Kafka cluster in which the Kafka Bridge is deployed.

② The number of partitions for the topic.

2. Save the file to the `examples/topic` directory as `bridge-quickstart-topic.yaml`.
3. Create the topic in your Kubernetes cluster:

```
kubectl apply -f examples/topic/bridge-quickstart-topic.yaml
```

4. Using the Kafka Bridge, produce three messages to the topic you created:

```

curl -X POST \
  http://localhost:8080/topics/bridge-quickstart-topic \
  -H 'content-type: application/vnd.kafka.json.v2+json' \
  -d '{
    "records": [
      {
        "key": "my-key",
        "value": "sales-lead-0001"
      },
      {
        "value": "sales-lead-0002",
        "partition": 2
      },
      {
        "value": "sales-lead-0003"
      }
    ]
  }'

```

- **sales-lead-0001** is sent to a partition based on the hash of the key.
  - **sales-lead-0002** is sent directly to partition 2.
  - **sales-lead-0003** is sent to a partition in the **bridge-quickstart-topic** topic using a round-robin method.
5. If the request is successful, the Kafka Bridge returns an **offsets** array, along with a **200** code and a **content-type** header of **application/vnd.kafka.v2+json**. For each message, the **offsets** array describes:
- The partition that the message was sent to
  - The current message offset of the partition

### Example response

```
#...
{
  "offsets": [
    {
      "partition": 0,
      "offset": 0
    },
    {
      "partition": 2,
      "offset": 0
    },
    {
      "partition": 0,
      "offset": 1
    }
  ]
}
```

### What to do next

After producing messages to topics and partitions, [create a Kafka Bridge consumer](#).

### Additional resources

- [POST /topics/{topicname}](#) in the API reference documentation.
- [POST /topics/{topicname}/partitions/{partitionid}](#) in the API reference documentation.

## 6.2.4. Creating a Kafka Bridge consumer

Before you can perform any consumer operations in the Kafka cluster, you must first create a consumer by using the `consumers` endpoint. The consumer is referred to as a *Kafka Bridge consumer*.

### Procedure

1. Create a Kafka Bridge consumer in a new consumer group named `bridge-quickstart-consumer-group`:

```
curl -X POST http://localhost:8080/consumers/bridge-quickstart-consumer-group \
-H 'content-type: application/vnd.kafka.v2+json' \
-d '{
  "name": "bridge-quickstart-consumer",
  "auto.offset.reset": "earliest",
  "format": "json",
  "enable.auto.commit": false,
  "fetch.min.bytes": 512,
  "consumer.request.timeout.ms": 30000
}'
```

- The consumer is named `bridge-quickstart-consumer` and the embedded data format is set as `json`.
- Some basic configuration settings are defined.
- The consumer will not commit offsets to the log automatically because the `enable.auto.commit` setting is `false`. You will commit the offsets manually later in this quickstart.

If the request is successful, the Kafka Bridge returns the consumer ID (`instance_id`) and base URL (`base_uri`) in the response body, along with a `200` code.

*Example response*

```
#...
{
  "instance_id": "bridge-quickstart-consumer",
  "base_uri": "http://<bridge-name>-bridge-service:8080/consumers/bridge-
  quickstart-consumer-group/instances/bridge-quickstart-consumer"
}
```

## 2. Copy the base URL (`base_uri`) to use in the other consumer operations in this quickstart.

*What to do next*

Now that you have created a Kafka Bridge consumer, you can [subscribe it to topics](#).

*Additional resources*

- [POST /consumers/{groupid}](#) in the API reference documentation.

### 6.2.5. Subscribing a Kafka Bridge consumer to topics

After you have created a Kafka Bridge consumer, subscribe it to one or more topics by using the [subscription](#) endpoint. Once subscribed, the consumer starts receiving all messages that are produced to the topic.

*Procedure*

- Subscribe the consumer to the `bridge-quickstart-topic` topic that you created earlier, in [Producing messages to topics and partitions](#):

```
curl -X POST http://localhost:8080/consumers/bridge-quickstart-consumer-
group/instances/bridge-quickstart-consumer/subscription \
-H 'content-type: application/vnd.kafka.v2+json' \
-d '{
  "topics": [
    "bridge-quickstart-topic"
  ]
}'
```

The `topics` array can contain a single topic (as shown here) or multiple topics. If you want to

subscribe the consumer to multiple topics that match a regular expression, you can use the `topic_pattern` string instead of the `topics` array.

If the request is successful, the Kafka Bridge returns a [204](#) (No Content) code only.

#### *What to do next*

After subscribing a Kafka Bridge consumer to topics, you can [retrieve messages from the consumer](#).

#### *Additional resources*

- [POST /consumers/{groupid}/instances/{name}/subscription](#) in the API reference documentation.

### **6.2.6. Retrieving the latest messages from a Kafka Bridge consumer**

Next, retrieve the latest messages from the Kafka Bridge consumer by requesting data from the [records](#) endpoint. In production, HTTP clients can call this endpoint repeatedly (in a loop).

#### *Procedure*

1. Produce additional messages to the Kafka Bridge consumer, as described in [Producing messages to topics and partitions](#).
2. Submit a `GET` request to the `records` endpoint:

```
curl -X GET http://localhost:8080/consumers/bridge-quickstart-consumer-group/instances/bridge-quickstart-consumer/records \
-H 'accept: application/vnd.kafka.json.v2+json'
```

After creating and subscribing to a Kafka Bridge consumer, a first GET request will return an empty response because the poll operation starts a rebalancing process to assign partitions.

3. Repeat step two to retrieve messages from the Kafka Bridge consumer.

The Kafka Bridge returns an array of messages—describing the topic name, key, value, partition, and offset—in the response body, along with a [200](#) code. Messages are retrieved from the latest offset by default.

```
HTTP/1.1 200 OK
content-type: application/vnd.kafka.json.v2+json
#...
[
  {
    "topic":"bridge-quickstart-topic",
    "key":"my-key",
    "value":"sales-lead-0001",
    "partition":0,
    "offset":0
  },
  {
    "topic":"bridge-quickstart-topic",
    "key":null,
    "value":"sales-lead-0003",
    "partition":0,
    "offset":1
  },
#...
```

**NOTE** If an empty response is returned, produce more records to the consumer as described in [Producing messages to topics and partitions](#), and then try retrieving messages again.

#### What to do next

After retrieving messages from a Kafka Bridge consumer, try [committing offsets to the log](#).

#### Additional resources

- [GET /consumers/{groupid}/instances/{name}/records](#) in the API reference documentation.

### 6.2.7. Committing offsets to the log

Next, use the [offsets](#) endpoint to manually commit offsets to the log for all messages received by the Kafka Bridge consumer. This is required because the Kafka Bridge consumer that you created earlier, in [Creating a Kafka Bridge consumer](#), was configured with the `enable.auto.commit` setting as `false`.

#### Procedure

- Commit offsets to the log for the `bridge-quickstart-consumer`:

```
curl -X POST http://localhost:8080/consumers/bridge-quickstart-consumer-
group/instances/bridge-quickstart-consumer/offsets
```

Because no request body is submitted, offsets are committed for all the records that have been received by the consumer. Alternatively, the request body can contain an array ([OffsetCommitSeekList](#)) that specifies the topics and partitions that you want to commit offsets

for.

If the request is successful, the Kafka Bridge returns a [204](#) code only.

#### *What to do next*

After committing offsets to the log, try out the endpoints for [seeking to offsets](#).

#### *Additional resources*

- [POST /consumers/{groupid}/instances/{name}/offsets](#) in the API reference documentation.

### 6.2.8. Seeking to offsets for a partition

Next, use the [positions](#) endpoints to configure the Kafka Bridge consumer to retrieve messages for a partition from a specific offset, and then from the latest offset. This is referred to in Apache Kafka as a seek operation.

#### *Procedure*

1. Seek to a specific offset for partition 0 of the [quickstart-bridge-topic](#) topic:

```
curl -X POST http://localhost:8080/consumers/bridge-quickstart-consumer-group/instances/bridge-quickstart-consumer/positions \
-H 'content-type: application/vnd.kafka.v2+json' \
-d '{
    "offsets": [
        {
            "topic": "bridge-quickstart-topic",
            "partition": 0,
            "offset": 2
        }
    ]
}'
```

If the request is successful, the Kafka Bridge returns a [204](#) code only.

2. Submit a [GET](#) request to the [records](#) endpoint:

```
curl -X GET http://localhost:8080/consumers/bridge-quickstart-consumer-group/instances/bridge-quickstart-consumer/records \
-H 'accept: application/vnd.kafka.json.v2+json'
```

The Kafka Bridge returns messages from the offset that you seeked to.

3. Restore the default message retrieval behavior by seeking to the last offset for the same partition. This time, use the [positions/end](#) endpoint.

```
curl -X POST http://localhost:8080/consumers/bridge-quickstart-consumer-group/instances/bridge-quickstart-consumer/positions/end \
-H 'content-type: application/vnd.kafka.v2+json' \
-d '{
  "partitions": [
    {
      "topic": "bridge-quickstart-topic",
      "partition": 0
    }
  ]
}'
```

If the request is successful, the Kafka Bridge returns another [204](#) code.

**NOTE**

You can also use the [positions/beginning](#) endpoint to seek to the first offset for one or more partitions.

*What to do next*

In this quickstart, you have used the Strimzi Kafka Bridge to perform several common operations on a Kafka cluster. You can now [delete the Kafka Bridge consumer](#) that you created earlier.

*Additional resources*

- [POST /consumers/{groupid}/instances/{name}/positions](#) in the API reference documentation.
- [POST /consumers/{groupid}/instances/{name}/positions/beginning](#) in the API reference documentation.
- [POST /consumers/{groupid}/instances/{name}/positions/end](#) in the API reference documentation.

## 6.2.9. Deleting a Kafka Bridge consumer

Finally, delete the Kafka Bridge consumer that you used throughout this quickstart.

*Procedure*

- Delete the Kafka Bridge consumer by sending a [DELETE](#) request to the [instances](#) endpoint.

```
curl -X DELETE http://localhost:8080/consumers/bridge-quickstart-consumer-group/instances/bridge-quickstart-consumer
```

If the request is successful, the Kafka Bridge returns a [204](#) code only.

*Additional resources*

- [DELETE /consumers/{groupid}/instances/{name}](#) in the API reference documentation.

# Chapter 7. Cruise Control for cluster rebalancing

You can deploy [Cruise Control](#) to your Strimzi cluster and use it to *rebalance* the Kafka cluster.

Cruise Control is an open source system for automating Kafka operations, such as monitoring cluster workload, rebalancing a cluster based on predefined constraints, and detecting and fixing anomalies. It consists of four main components—the Load Monitor, the Analyzer, the Anomaly Detector, and the Executor—and a REST API for client interactions. Strimzi utilizes the REST API to support the following Cruise Control features:

- Generating *optimization proposals* from multiple *optimization goals*.
- Rebalancing a Kafka cluster based on an optimization proposal.

Other Cruise Control features are not currently supported, including anomaly detection, notifications, write-your-own goals, and changing the topic replication factor.

Example YAML files for Cruise Control are provided in [examples/cruise-control/](#).

## 7.1. Why use Cruise Control?

Cruise Control reduces the time and effort involved in running an efficient and balanced Kafka cluster.

A typical cluster can become unevenly loaded over time. Partitions that handle large amounts of message traffic might be unevenly distributed across the available brokers. To rebalance the cluster, administrators must monitor the load on brokers and manually reassign busy partitions to brokers with spare capacity.

Cruise Control automates the cluster rebalancing process. It constructs a *workload model* of resource utilization for the cluster—based on CPU, disk, and network load—and generates optimization proposals (that you can approve or reject) for more balanced partition assignments. A set of configurable optimization goals is used to calculate these proposals.

When you approve an optimization proposal, Cruise Control applies it to your Kafka cluster. When the cluster rebalancing operation is complete, the broker pods are used more effectively and the Kafka cluster is more evenly balanced.

*Additional resources*

- [Cruise Control Wiki](#)

## 7.2. Optimization goals overview

To rebalance a Kafka cluster, Cruise Control uses optimization goals to generate [optimization proposals](#), which you can approve or reject.

Optimization goals are constraints on workload redistribution and resource utilization across a

Kafka cluster. Strimzi supports most of the optimization goals developed in the Cruise Control project. The supported goals, in the default descending order of priority, are as follows:

1. Rack-awareness
  2. Replica capacity
  3. *Capacity*: Disk capacity, Network inbound capacity, Network outbound capacity, CPU capacity
  4. Replica distribution
  5. Potential network output
  6. *Resource distribution*: Disk utilization distribution, Network inbound utilization distribution, Network outbound utilization distribution, CPU utilization distribution
- NOTE** The resource distribution goals are controlled using [capacity limits](#) on broker resources.
7. Leader bytes-in rate distribution
  8. Topic replica distribution
  9. Leader replica distribution
  10. Preferred leader election

For more information on each optimization goal, see [Goals](#) in the Cruise Control Wiki.

**NOTE** Intra-broker disk goals, "Write your own" goals, and Kafka assigner goals are not yet supported.

## Goals configuration in Strimzi custom resources

You configure optimization goals in [Kafka](#) and [KafkaRebalance](#) custom resources. Cruise Control has configurations for [hard](#) optimization goals that must be satisfied, as well as [master](#), [default](#), and [user-provided](#) optimization goals. Optimization goals for resource distribution (disk, network inbound, network outbound, and CPU) are subject to [capacity limits](#) on broker resources.

The following sections describe each goal configuration in more detail.

### Hard goals and soft goals

Hard goals are goals that *must* be satisfied in optimization proposals. Goals that are not configured as hard goals are known as *soft goals*. You can think of soft goals as *best effort* goals: they do *not* need to be satisfied in optimization proposals, but are included in optimization calculations. An optimization proposal that violates one or more soft goals, but satisfies all hard goals, is valid.

Cruise Control will calculate optimization proposals that satisfy all the hard goals and as many soft goals as possible (in their priority order). An optimization proposal that does *not* satisfy all the hard goals is rejected by Cruise Control and not sent to the user for approval.

**NOTE** For example, you might have a soft goal to distribute a topic's replicas evenly across the cluster (the topic replica distribution goal). Cruise Control will ignore this goal if doing so enables all the configured hard goals to be met.

In Cruise Control, the following [master optimization goals](#) are preset as hard goals:

```
RackAwareGoal; ReplicaCapacityGoal; DiskCapacityGoal; NetworkInboundCapacityGoal;  
NetworkOutboundCapacityGoal; CpuCapacityGoal
```

You configure hard goals in the Cruise Control deployment configuration, by editing the [hard.goals](#) property in [Kafka.spec.cruiseControl.config](#).

- To inherit the preset hard goals from Cruise Control, do not specify the [hard.goals](#) property in [Kafka.spec.cruiseControl.config](#)
- To change the preset hard goals, specify the desired goals in the [hard.goals](#) property, using their fully-qualified domain names.

*Example Kafka configuration for hard optimization goals*

```
apiVersion: kafka.strimzi.io/v1beta1  
kind: Kafka  
metadata:  
  name: my-cluster  
spec:  
  kafka:  
    # ...  
  zookeeper:  
    # ...  
  entityOperator:  
    topicOperator: {}  
    userOperator: {}  
  cruiseControl:  
    brokerCapacity:  
      inboundNetwork: 10000KB/s  
      outboundNetwork: 10000KB/s  
    config:  
      hard.goals: >  
        com.linkedin.kafka.cruisecontrol.analyzer.goals.NetworkInboundCapacityGoal,  
        com.linkedin.kafka.cruisecontrol.analyzer.goals.NetworkOutboundCapacityGoal  
      # ...
```

Increasing the number of configured hard goals will reduce the likelihood of Cruise Control generating valid optimization proposals.

If [skipHardGoalCheck: true](#) is specified in the [KafkaRebalance](#) custom resource, Cruise Control does *not* check that the list of user-provided optimization goals (in [KafkaRebalance.spec.goals](#)) contains *all* the configured hard goals ([hard.goals](#)). Therefore, if some, but not all, of the user-provided optimization goals are in the [hard.goals](#) list, Cruise Control will still treat them as hard goals even if

`skipHardGoalCheck: true` is specified.

## Master optimization goals

The *master optimization goals* are available to all users. Goals that are not listed in the master optimization goals are not available for use in Cruise Control operations.

Unless you change the Cruise Control [deployment configuration](#), Strimzi will inherit the following master optimization goals from Cruise Control, in descending priority order:

```
RackAwareGoal; ReplicaCapacityGoal; DiskCapacityGoal; NetworkInboundCapacityGoal;  
NetworkOutboundCapacityGoal; CpuCapacityGoal; ReplicaDistributionGoal;  
PotentialNwOutGoal; DiskUsageDistributionGoal; NetworkInboundUsageDistributionGoal;  
NetworkOutboundUsageDistributionGoal; CpuUsageDistributionGoal;  
TopicReplicaDistributionGoal; LeaderReplicaDistributionGoal;  
LeaderBytesInDistributionGoal; PreferredLeaderElectionGoal
```

Six of these goals are preset as [hard goals](#).

To reduce complexity, we recommend that you use the inherited master optimization goals, unless you need to *completely* exclude one or more goals from use in [KafkaRebalance](#) resources. The priority order of the master optimization goals can be modified, if desired, in the configuration for [default optimization goals](#).

You configure master optimization goals, if necessary, in the Cruise Control deployment configuration: `Kafka.spec.cruiseControl.config.goals`

- To accept the inherited master optimization goals, do not specify the `goals` property in `Kafka.spec.cruiseControl.config`.
- If you need to modify the inherited master optimization goals, specify a list of goals, in descending priority order, in the `goals` configuration option.

**NOTE**

If you change the inherited master optimization goals, you must ensure that the hard goals, if configured in the `hard.goals` property in `Kafka.spec.cruiseControl.config`, are a subset of the master optimization goals that you configured. Otherwise, errors will occur when generating optimization proposals.

## Default optimization goals

Cruise Control uses the *default optimization goals* to generate the *cached optimization proposal*. For more information about the cached optimization proposal, see [Optimization proposals overview](#).

You can override the default optimization goals by setting [user-provided optimization goals](#) in a [KafkaRebalance](#) custom resource.

Unless you specify `default.goals` in the Cruise Control [deployment configuration](#), the master optimization goals are used as the default optimization goals. In this case, the cached optimization proposal is generated using the master optimization goals.

- To use the master optimization goals as the default goals, do not specify the `default.goals` property in `Kafka.spec.cruiseControl.config`.
- To modify the default optimization goals, edit the `default.goals` property in `Kafka.spec.cruiseControl.config`. You must use a subset of the master optimization goals.

*Example Kafka configuration for default optimization goals*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  entityOperator:
    topicOperator: {}
    userOperator: {}
  cruiseControl:
    brokerCapacity:
      inboundNetwork: 10000KB/s
      outboundNetwork: 10000KB/s
    config:
      default.goals: >
        com.linkedin.kafka.cruisecontrol.analyzer.goals.RackAwareGoal,
        com.linkedin.kafka.cruisecontrol.analyzer.goals.ReplicaCapacityGoal,
        com.linkedin.kafka.cruisecontrol.analyzer.goals.DiskCapacityGoal
    # ...
```

If no default optimization goals are specified, the cached proposal is generated using the master optimization goals.

## User-provided optimization goals

*User-provided optimization goals* narrow down the configured default goals for a particular optimization proposal. You can set them, as required, in `spec.goals` in a `KafkaRebalance` custom resource:

`KafkaRebalance.spec.goals`

User-provided optimization goals can generate optimization proposals for different scenarios. For example, you might want to optimize leader replica distribution across the Kafka cluster without considering disk capacity or disk utilization. So, you create a `KafkaRebalance` custom resource containing a single user-provided goal for leader replica distribution.

User-provided optimization goals must:

- Include all configured [hard goals](#), or an error occurs
- Be a subset of the master optimization goals

To ignore the configured hard goals when generating an optimization proposal, add the `skipHardGoalCheck: true` property to the [KafkaRebalance](#) custom resource. See [Generating optimization proposals](#).

#### *Additional resources*

- [Cruise Control configuration](#)
- [Configurations](#) in the Cruise Control Wiki.

## 7.3. Optimization proposals overview

An *optimization proposal* is a summary of proposed changes that would produce a more balanced Kafka cluster, with partition workloads distributed more evenly among the brokers. Each optimization proposal is based on the set of [optimization goals](#) that was used to generate it, subject to any configured [capacity limits on broker resources](#).

An optimization proposal is contained in the `Status.Optimization.Result` property of a [KafkaRebalance](#) custom resource. The information provided is a summary of the full optimization proposal. Use the summary to decide whether to:

- Approve the optimization proposal. This instructs Cruise Control to apply the proposal to the Kafka cluster and start a cluster rebalance operation.
- Reject the optimization proposal. You can change the optimization goals and then generate another proposal.

All optimization proposals are *dry runs*: you cannot approve a cluster rebalance without first generating an optimization proposal. There is no limit to the number of optimization proposals that can be generated.

### Cached optimization proposal

Cruise Control maintains a *cached optimization proposal* based on the configured default optimization goals. Generated from the workload model, the cached optimization proposal is updated every 15 minutes to reflect the current state of the Kafka cluster. If you generate an optimization proposal using the default optimization goals, Cruise Control returns the most recent cached proposal.

To change the cached optimization proposal refresh interval, edit the `proposal.expiration.ms` setting in the Cruise Control deployment configuration. Consider a shorter interval for fast changing clusters, although this increases the load on the Cruise Control server.

### Contents of optimization proposals

The following table describes the properties contained in an optimization proposal:

*Table 1. Properties contained in an optimization proposal*

JSON property	Description
<code>numIntraBrokerReplicaMovements</code>	<p>The total number of partition replicas that will be transferred between the disks of the cluster's brokers.</p> <p><b>Performance impact during rebalance operation:</b> Relatively high, but lower than <code>numReplicaMovements</code>.</p>
<code>excludedBrokersForLeadership</code>	Not yet supported. An empty list is returned.
<code>numReplicaMovements</code>	<p>The number of partition replicas that will be moved between separate brokers.</p> <p><b>Performance impact during rebalance operation:</b> Relatively high.</p>
<code>onDemandBalancednessScoreBefore</code> , <code>onDemandBalancednessScoreAfter</code>	<p>A measurement of the overall <i>balancedness</i> of a Kafka Cluster, before and after the optimization proposal was generated.</p> <p>The score is calculated by subtracting the sum of the <code>BalancednessScore</code> of each violated soft goal from 100. Cruise Control assigns a <code>BalancednessScore</code> to every optimization goal based on several factors, including priority—the goal's position in the list of <code>default.goals</code> or user-provided goals.</p> <p>The <code>Before</code> score is based on the current configuration of the Kafka cluster. The <code>After</code> score is based on the generated optimization proposal.</p>
<code>intraBrokerDataToMoveMB</code>	<p>The sum of the size of each partition replica that will be moved between disks on the same broker (see also <code>numIntraBrokerReplicaMovements</code>).</p> <p><b>Performance impact during rebalance operation:</b> Variable. The larger the number, the longer the cluster rebalance will take to complete. Moving a large amount of data between disks on the same broker has less impact than between separate brokers (see <code>dataToMoveMB</code>).</p>
<code>recentWindows</code>	The number of metrics windows upon which the optimization proposal is based.
<code>dataToMoveMB</code>	<p>The sum of the size of each partition replica that will be moved to a separate broker (see also <code>numReplicaMovements</code>).</p> <p><b>Performance impact during rebalance operation:</b> Variable. The larger the number, the longer the cluster rebalance will take to complete.</p>
<code>monitoredPartitionsPercentage</code>	The percentage of partitions in the Kafka cluster covered by the optimization proposal. Affected by the number of <code>excludedTopics</code> .

JSON property	Description
<code>excludedTopics</code>	If you specified a regular expression in the <code>spec.excludedTopicsRegex</code> property in the <code>KafkaRebalance</code> resource, all topic names matching that expression are listed here. These topics are excluded from the calculation of partition replica/leader movements in the optimization proposal.
<code>numLeaderMovements</code>	The number of partitions whose leaders will be switched to different replicas. This involves a change to ZooKeeper configuration.  <b>Performance impact during rebalance operation:</b> Relatively low.
<code>excludedBrokersForReplicaMove</code>	Not yet supported. An empty list is returned.

#### *Additional resources*

- [Optimization goals overview](#)
- [Generating optimization proposals](#)
- [Approving an optimization proposal](#)

## 7.4. Rebalance performance tuning overview

You can adjust several performance tuning options for cluster rebalances. These options control how partition replica and leadership movements in a rebalance are executed, as well as the bandwidth that is allocated to a rebalance operation.

### Partition reassignment commands

[Optimization proposals](#) are comprised of separate partition reassignment commands. When you [approve](#) a proposal, the Cruise Control server applies these commands to the Kafka cluster.

A partition reassignment command consists of either of the following types of operations:

- Partition movement: Involves transferring the partition replica and its data to a new location. Partition movements can take one of two forms:
  - Inter-broker movement: The partition replica is moved to a log directory on a different broker.
  - Intra-broker movement: The partition replica is moved to a different log directory on the same broker.
- Leadership movement: This involves switching the leader of the partition's replicas.

Cruise Control issues partition reassignment commands to the Kafka cluster in batches. The performance of the cluster during the rebalance is affected by the number of each type of movement contained in each batch.

## Replica movement strategies

Cluster rebalance performance is also influenced by the *replica movement strategy* that is applied to the batches of partition reassignment commands. By default, Cruise Control uses the [BaseReplicaMovementStrategy](#), which simply applies the commands in the order they were generated. However, if there are some very large partition reassessments early in the proposal, this strategy can slow down the application of the other reassessments.

Cruise Control provides three alternative replica movement strategies that can be applied to optimization proposals:

- [PrioritizeSmallReplicaMovementStrategy](#): Order reassessments in order of ascending size.
- [PrioritizeLargeReplicaMovementStrategy](#): Order reassessments in order of descending size.
- [PostponeUrpReplicaMovementStrategy](#): Prioritize reassessments for replicas of partitions which have no out-of-sync replicas.

These strategies can be configured as a sequence. The first strategy attempts to compare two partition reassessments using its internal logic. If the reassessments are equivalent, then it passes them to the next strategy in the sequence to decide the order, and so on.

## Rebalance tuning options

Cruise Control provides several configuration options for tuning the rebalance parameters discussed above. You can set these tuning options at either the [Cruise Control server](#) or [optimization proposal](#) levels:

- The Cruise Control server setting can be set in the Kafka custom resource under [Kafka.spec.cruiseControl.config](#).
- The individual rebalance performance configurations can be set under [KafkaRebalance.spec](#).

The relevant configurations are summarized below:

Server and KafkaRebalance Configuration	Description	Default Value
<code>num.concurrent.partition.movements.per.broker</code>	The maximum number of inter-broker partition movements in each partition reassignment batch	5
<code>concurrentPartitionMovementsPerBroker</code>		
<code>num.concurrent.intra.broker.partition.movements</code>	The maximum number of intra-broker partition movements in each partition reassignment batch	2
<code>concurrentIntraBrokerPartitionMovements</code>		
<code>num.concurrent.leader.movements</code>	The maximum number of partition leadership changes in each partition reassignment batch	1000
<code>concurrentLeaderMovements</code>		

<b>Server and KafkaRebalance Configuration</b>	<b>Description</b>	<b>Default Value</b>
<code>default.replication.throttle</code>	The bandwidth (in bytes per second) to be assigned to the reassigning of partitions	No Limit
<code>replicationThrottle</code>		
<code>default.replica.movement.strategies</code>	The list of strategies (in priority order) used to determine the order in which partition reassignment commands are executed for generated proposals.	<code>BaseReplicaMovementStrategy</code>
<code>replicaMovementStrategies</code>	For the server setting, use a comma separated string with the fully qualified names of the strategy class (add <code>com.linkedin.kafka.cruisecontrol.executor.strategy.</code> to the start of each class name). For the <code>KafkaRebalance</code> resource setting use a YAML array of strategy class names.	

Changing the default settings affects the length of time that the rebalance takes to complete, as well as the load placed on the Kafka cluster during the rebalance. Using lower values reduces the load but increases the amount of time taken, and vice versa.

#### *Additional resources*

- [CruiseControlSpec schema reference](#).
- [KafkaRebalanceSpec schema reference](#).

## 7.5. Cruise Control configuration

The `config` property in `Kafka.spec.cruiseControl` contains configuration options as keys with values as one of the following JSON types:

- String
- Number
- Boolean

You can specify and configure all the options listed in the "Configurations" section of the [Cruise Control documentation](#), apart from those managed directly by Strimzi. Specifically, you **cannot** modify configuration options with keys equal to or starting with one of the keys mentioned [here](#).

If restricted options are specified, they are ignored and a warning message is printed to the Cluster Operator log file. All the supported options are passed to Cruise Control.

## An example Cruise Control configuration

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  # ...
  cruiseControl:
    # ...
    config:
      default.goals: >
        com.linkedin.kafka.cruisecontrol.analyzer.goals.RackAwareGoal,
        com.linkedin.kafka.cruisecontrol.analyzer.goals.ReplicaCapacityGoal
      cpu.balance.threshold: 1.1
      metadata.max.age.ms: 300000
      send.buffer.bytes: 131072
  # ...
```

## Cross-Origin Resource Sharing configuration

Cross-Origin Resource Sharing (CORS) allows you to specify allowed methods and originating URLs for accessing REST APIs.

By default, CORS is disabled for the Cruise Control REST API. When enabled, only [GET](#) requests for read-only access to the Kafka cluster state are allowed. This means that external applications, which are running in different origins than the Strimzi components, cannot make [POST](#) requests to the Cruise Control API. However, those applications can make [GET](#) requests to access read-only information about the Kafka cluster, such as the current cluster load or the most recent optimization proposal.

### Enabling CORS for Cruise Control

You enable and configure CORS in [Kafka.spec.cruiseControl.config](#).

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  # ...
  cruiseControl:
    # ...
    config:
      webserver.http.cors.enabled: true
      webserver.http.cors.origin: "*"
      webserver.http.cors.exposeheaders: "User-Task-ID,Content-Type"
  # ...
```

For more information, see [REST APIs in the Cruise Control Wiki](#).

## Capacity configuration

Cruise Control uses *capacity limits* to determine if optimization goals for resource distribution are being broken. There are four goals of this type:

- `DiskUsageDistributionGoal` - Disk utilization distribution
- `CpuUsageDistributionGoal` - CPU utilization distribution
- `NetworkInboundUsageDistributionGoal` - Network inbound utilization distribution
- `NetworkOutboundUsageDistributionGoal` - Network outbound utilization distribution

You specify capacity limits for Kafka broker resources in the `brokerCapacity` property in `Kafka.spec.cruiseControl`. They are enabled by default and you can change their default values. Capacity limits can be set for the following broker resources, using the standard Kubernetes byte units (K, M, G and T) or their bibyte (power of two) equivalents (Ki, Mi, Gi and Ti):

- `disk` - Disk storage per broker (Default: 100000Mi)
- `cpuUtilization` - CPU utilization as a percentage (Default: 100)
- `inboundNetwork` - Inbound network throughput in byte units per second (Default: 10000KiB/s)
- `outboundNetwork` - Outbound network throughput in byte units per second (Default: 10000KiB/s)

Because Strimzi Kafka brokers are homogeneous, Cruise Control applies the same capacity limits to every broker it is monitoring.

*An example Cruise Control brokerCapacity configuration using bibyte units*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  # ...
  cruiseControl:
    # ...
    brokerCapacity:
      disk: 100Gi
      cpuUtilization: 100
      inboundNetwork: 10000KiB/s
      outboundNetwork: 10000KiB/s
    # ...
```

*Additional resources*

For more information, refer to the [BrokerCapacity schema reference](#).

## Logging configuration

Cruise Control has its own configurable logger:

- `cruisecontrol.root.logger`

Cruise Control uses the Apache `log4j` logger implementation.

Use the `logging` property to configure loggers and logger levels.

You can set the log levels by specifying the logger and level directly (inline) or use a custom (external) ConfigMap. If a ConfigMap is used, you set `logging.name` property to the name of the ConfigMap containing the external logging configuration. Inside the ConfigMap, the logging configuration is described using `log4j.properties`.

Here we see examples of `inline` and `external` logging.

#### *Inline logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
# ...
spec:
  cruiseControl:
    # ...
    logging:
      type: inline
      loggers:
        cruisecontrol.root.logger: "INFO"
    # ...
```

#### *External logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
# ...
spec:
  cruiseControl:
    # ...
    logging:
      type: external
      name: customConfigMap
    # ...
```

## 7.6. Deploying Cruise Control

To deploy Cruise Control to your Strimzi cluster, define the configuration using the `cruiseControl` property in the `Kafka` resource, and then create or update the resource.

Deploy one instance of Cruise Control per Kafka cluster.

#### *Prerequisites*

- A Kubernetes cluster
- A running Cluster Operator

*Procedure*

1. Edit the `Kafka` resource and add the `cruiseControl` property.

The properties you can configure are shown in this example configuration:

```

apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  # ...
  cruiseControl:
    brokerCapacity: ①
      inboundNetwork: 10000KB/s
      outboundNetwork: 10000KB/s
    # ...
  config: ②
    default.goals: >
      com.linkedin.kafka.cruisecontrol.analyzer.goals.RackAwareGoal,
      com.linkedin.kafka.cruisecontrol.analyzer.goals.ReplicaCapacityGoal
    # ...
    cpu.balance.threshold: 1.1
    metadata.max.age.ms: 300000
    send.buffer.bytes: 131072
    # ...
  resources: ③
    requests:
      cpu: 1
      memory: 512Mi
    limits:
      cpu: 2
      memory: 2Gi
  logging: ④
    type: inline
    loggers:
      cruisecontrol.root.logger: "INFO"
  template: ⑤
    pod:
      metadata:
        labels:
          label1: value1
      securityContext:
        runAsUser: 1000001
        fsGroup: 0
      terminationGracePeriodSeconds: 120
  readinessProbe: ⑥
    initialDelaySeconds: 15
    timeoutSeconds: 5
  livenessProbe: ⑦
    initialDelaySeconds: 15
    timeoutSeconds: 5
# ...

```

① Specifies capacity limits for broker resources. For more information, see [Capacity configuration](#).

- ② Defines the Cruise Control configuration, including the default optimization goals (in `default.goals`) and any customizations to the master optimization goals (in `goals`) or the hard goals (in `hard.goals`). You can provide any [standard Cruise Control configuration option](#) apart from those managed directly by Strimzi. For more information on configuring optimization goals, see [Optimization goals overview](#).
- ③ CPU and memory resources reserved for Cruise Control. For more information, see [resources](#).
- ④ Defined loggers and log levels added directly (inline) or indirectly (external) through a ConfigMap. A custom ConfigMap must be placed under the `log4j.properties` key. Cruise Control has a single logger named `cruisecontrol.root.logger`. You can set the log level to INFO, ERROR, WARN, TRACE, DEBUG, FATAL or OFF. For more information, see [Logging configuration](#).
- ⑤ [Customization of deployment templates and pods](#).
- ⑥ [Healthcheck readiness probes](#).
- ⑦ [Healthcheck liveness probes](#).

2. Create or update the resource:

```
kubectl apply -f kafka.yaml
```

3. Verify that Cruise Control was successfully deployed:

```
kubectl get deployments -l app.kubernetes.io/name=strimzi
```

## Auto-created topics

The following table shows the three topics that are automatically created when Cruise Control is deployed. These topics are required for Cruise Control to work properly and must not be deleted or changed.

*Table 2. Auto-created topics*

Auto-created topic	Created by	Function
<code>strimzi.cruisecontrol.metrics</code>	Strimzi Metrics Reporter	Stores the raw metrics from the Metrics Reporter in each Kafka broker.
<code>strimzi.cruisecontrol.partitionmetricsamples</code>	Cruise Control	Stores the derived metrics for each partition. These are created by the <a href="#">Metric Sample Aggregator</a> .
<code>strimzi.cruisecontrol.modeltrainingsamples</code>	Cruise Control	Stores the metrics samples used to create the <a href="#">Cluster Workload Model</a> .

To prevent the removal of records that are needed by Cruise Control, log compaction is disabled in

the auto-created topics.

#### What to do next

After configuring and deploying Cruise Control, you can [generate optimization proposals](#).

#### Additional resources

[CruiseControlTemplate schema reference](#).

## 7.7. Generating optimization proposals

When you create or update a [KafkaRebalance](#) resource, Cruise Control generates an [optimization proposal](#) for the Kafka cluster based on the configured [optimization goals](#).

Analyze the information in the optimization proposal and decide whether to approve it.

#### Prerequisites

- You have [deployed Cruise Control](#) to your Strimzi cluster.
- You have configured [optimization goals](#) and, optionally, [capacity limits on broker resources](#).

#### Procedure

1. Create a [KafkaRebalance](#) resource:

- To use the *default optimization goals* defined in the [Kafka](#) resource, leave the [spec](#) property empty:

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaRebalance
metadata:
  name: my-rebalance
  labels:
    strimzi.io/cluster: my-cluster
spec: {}
```

- To configure *user-provided optimization goals* instead of using the default goals, add the [goals](#) property and enter one or more goals.

In the following example, rack awareness and replica capacity are configured as user-provided optimization goals:

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaRebalance
metadata:
  name: my-rebalance
  labels:
    strimzi.io/cluster: my-cluster
spec:
  goals:
    - RackAwareGoal
    - ReplicaCapacityGoal
```

- c. To ignore the configured hard goals, add the `skipHardGoalCheck: true` property:

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaRebalance
metadata:
  name: my-rebalance
  labels:
    strimzi.io/cluster: my-cluster
spec:
  goals:
    - RackAwareGoal
    - ReplicaCapacityGoal
  skipHardGoalCheck: true
```

2. Create or update the resource:

```
kubectl apply -f your-file
```

The Cluster Operator requests the optimization proposal from Cruise Control. This might take a few minutes depending on the size of the Kafka cluster.

3. Check the status of the `KafkaRebalance` resource:

```
kubectl describe kafkarebalance rebalance-cr-name
```

Cruise Control returns one of two statuses:

- `PendingProposal`: The rebalance operator is polling the Cruise Control API to check if the optimization proposal is ready.
- `ProposalReady`: The optimization proposal is ready for review and, if desired, approval. The optimization proposal is contained in the `Status.Optimization Result` property of the `KafkaRebalance` resource.

4. Review the optimization proposal.

```
kubectl describe kafka-rebalance rebalance-cr-name
```

Here is an example proposal:

```
Status:  
Conditions:  
  Last Transition Time: 2020-05-19T13:50:12.533Z  
  Status: ProposalReady  
  Type: State  
Observed Generation: 1  
Optimization Result:  
  Data To Move MB: 0  
  Excluded Brokers For Leadership:  
  Excluded Brokers For Replica Move:  
  Excluded Topics:  
    Intra Broker Data To Move MB: 0  
    Monitored Partitions Percentage: 100  
    Num Intra Broker Replica Movements: 0  
    Num Leader Movements: 0  
    Num Replica Movements: 26  
    On Demand Balancedness Score After: 81.8666802863978  
    On Demand Balancedness Score Before: 78.01176356230222  
    Recent Windows: 1  
Session Id: 05539377-ca7b-45ef-b359-e13564f1458c
```

The properties in the [Optimization Result](#) section describe the pending cluster rebalance operation. For descriptions of each property, see [Contents of optimization proposals](#).

*What to do next*

[Approving an optimization proposal](#)

*Additional resources*

- [Optimization proposals overview](#)

## 7.8. Approving an optimization proposal

You can approve an [optimization proposal](#) generated by Cruise Control, if its status is [ProposalReady](#). Cruise Control will then apply the optimization proposal to the Kafka cluster, reassigning partitions to brokers and changing partition leadership.

**This is not a dry run.** Before you approve an optimization proposal, you must:

**CAUTION**

- Refresh the proposal in case it has become out of date.
- Carefully review the [contents of the proposal](#).

*Prerequisites*

- You have [generated an optimization proposal](#) from Cruise Control.
- The [KafkaRebalance](#) custom resource status is [ProposalReady](#).

#### *Procedure*

Perform these steps for the optimization proposal that you want to approve:

1. Unless the optimization proposal is newly generated, check that it is based on current information about the state of the Kafka cluster. To do so, refresh the optimization proposal to make sure it uses the latest cluster metrics:

- a. Annotate the [KafkaRebalance](#) resource in Kubernetes with [refresh](#):

```
kubectl annotate kafka-rebalance rebalance-cr-name strimzi.io/rebalance=refresh
```

- b. Check the status of the [KafkaRebalance](#) resource:

```
kubectl describe kafka-rebalance rebalance-cr-name
```

- c. Wait until the status changes to [ProposalReady](#).

2. Approve the optimization proposal that you want Cruise Control to apply.

Annotate the [KafkaRebalance](#) resource in Kubernetes:

```
kubectl annotate kafka-rebalance rebalance-cr-name strimzi.io/rebalance=approve
```

3. The Cluster Operator detects the annotated resource and instructs Cruise Control to rebalance the Kafka cluster.

4. Check the status of the [KafkaRebalance](#) resource:

```
kubectl describe kafka-rebalance rebalance-cr-name
```

5. Cruise Control returns one of three statuses:

- Rebalancing: The cluster rebalance operation is in progress.
- Ready: The cluster rebalancing operation completed successfully. The [KafkaRebalance](#) custom resource cannot be reused.
- NotReady: An error occurred—see [Fixing problems with a KafkaRebalance resource](#).

#### *Additional resources*

- [Optimization proposals overview](#)
- [Stopping a cluster rebalance](#)

## 7.9. Stopping a cluster rebalance

Once started, a cluster rebalance operation might take some time to complete and affect the overall performance of the Kafka cluster.

If you want to stop a cluster rebalance operation that is in progress, apply the `stop` annotation to the `KafkaRebalance` custom resource. This instructs Cruise Control to finish the current batch of partition reassessments and then stop the rebalance. When the rebalance has stopped, completed partition reassessments have already been applied; therefore, the state of the Kafka cluster is different when compared to prior to the start of the rebalance operation. If further rebalancing is required, you should generate a new optimization proposal.

**NOTE**

The performance of the Kafka cluster in the intermediate (stopped) state might be worse than in the initial state.

*Prerequisites*

- You have [approved the optimization proposal](#) by annotating the `KafkaRebalance` custom resource with `approve`.
- The status of the `KafkaRebalance` custom resource is `Rebalancing`.

*Procedure*

1. Annotate the `KafkaRebalance` resource in Kubernetes:

```
kubectl annotate kafka-rebalance rebalance-cr-name strimzi.io/rebalance=stop
```

2. Check the status of the `KafkaRebalance` resource:

```
kubectl describe kafka-rebalance rebalance-cr-name
```

3. Wait until the status changes to `Stopped`.

*Additional resources*

- [Optimization proposals overview](#)

## 7.10. Fixing problems with a `KafkaRebalance` resource

If an issue occurs when creating a `KafkaRebalance` resource or interacting with Cruise Control, the error is reported in the resource status, along with details of how to fix it. The resource also moves to the `NotReady` state.

To continue with the cluster rebalance operation, you must fix the problem in the `KafkaRebalance` resource itself or with the overall Cruise Control deployment. Problems might include the following:

- A misconfigured parameter in the `KafkaRebalance` resource.

- The `strimzi.io/cluster` label for specifying the Kafka cluster in the `KafkaRebalance` resource is missing.
- The Cruise Control server is not deployed as the `cruiseControl` property in the `Kafka` resource is missing.
- The Cruise Control server is not reachable.

After fixing the issue, you need to add the `refresh` annotation to the `KafkaRebalance` resource. During a “refresh”, a new optimization proposal is requested from the Cruise Control server.

#### *Prerequisites*

- You have [approved an optimization proposal](#).
- The status of the `KafkaRebalance` custom resource for the rebalance operation is `NotReady`.

#### *Procedure*

1. Get information about the error from the `KafkaRebalance` status:

```
kubectl describe kafka-rebalance-rebalance-cr-name
```

2. Attempt to resolve the issue in the `KafkaRebalance` resource.
3. Annotate the `KafkaRebalance` resource in Kubernetes:

```
kubectl annotate kafka-rebalance-rebalance-cr-name strimzi.io/rebalance=refresh
```

4. Check the status of the `KafkaRebalance` resource:

```
kubectl describe kafka-rebalance-rebalance-cr-name
```

5. Wait until the status changes to `PendingProposal`, or directly to `ProposalReady`.

#### *Additional resources*

- [Optimization proposals overview](#)

# Chapter 8. Distributed tracing

Distributed tracing allows you to track the progress of transactions between applications in a distributed system. In a microservices architecture, tracing tracks the progress of transactions between services. Trace data is useful for monitoring application performance and investigating issues with target systems and end-user applications.

In Strimzi, tracing facilitates the end-to-end tracking of messages: from source systems to Kafka, and then from Kafka to target systems and applications. It complements the metrics that are available to view in [Grafana dashboards](#), as well as the component loggers.

## How Strimzi supports tracing

Support for tracing is built in to the following components:

- Kafka Connect (including Kafka Connect with Source2Image support)
- MirrorMaker
- MirrorMaker 2.0
- Strimzi Kafka Bridge

You enable and configure tracing for these components using template configuration properties in their custom resources.

To enable tracing in Kafka producers, consumers, and Kafka Streams API applications, you *instrument* application code using the [OpenTracing Apache Kafka Client Instrumentation](#) library (included with Strimzi). When instrumented, clients generate trace data; for example, when producing messages or writing offsets to the log.

Traces are sampled according to a sampling strategy and then visualized in the Jaeger user interface.

Tracing is not supported for Kafka brokers.

**NOTE** Setting up tracing for applications and systems beyond Strimzi is outside the scope of this chapter. To learn more about this subject, search for "inject and extract" in the [OpenTracing documentation](#).

## Outline of procedures

To set up tracing for Strimzi, follow these procedures in order:

- Set up tracing for clients:
  - [Initialize a Jaeger tracer for Kafka clients](#)
- Instrument clients with tracers:
  - [Instrument producers and consumers for tracing](#)

- Instrument Kafka Streams applications for tracing
- Set up tracing for MirrorMaker, Kafka Connect, and the Kafka Bridge

#### Prerequisites

- The Jaeger backend components are deployed to your Kubernetes cluster. For deployment instructions, see the [Jaeger deployment documentation](#).

## 8.1. Overview of OpenTracing and Jaeger


Strimzi uses the OpenTracing and Jaeger projects.

OpenTracing is an API specification that is independent from the tracing or monitoring system.

- The OpenTracing APIs are used to *instrument* application code
- Instrumented applications generate *traces* for individual transactions across the distributed system
- Traces are composed of *spans* that define specific units of work over time

Jaeger is a tracing system for microservices-based distributed systems.

- Jaeger implements the OpenTracing APIs and provides client libraries for instrumentation
- The Jaeger user interface allows you to query, filter, and analyze trace data



#### Additional resources

- [OpenTracing](#)
- [Jaeger](#)

## 8.2. Setting up tracing for Kafka clients

Initialize a Jaeger tracer to instrument your client applications for distributed tracing.

## 8.2.1. Initializing a Jaeger tracer for Kafka clients

Configure and initialize a Jaeger tracer using a set of [tracing environment variables](#).

### Procedure

In each client application:

1. Add Maven dependencies for Jaeger to the `pom.xml` file for the client application:

```
<dependency>
  <groupId>io.jaegertracing</groupId>
  <artifactId>jaeger-client</artifactId>
  <version>1.3.2</version>
</dependency>
```

2. Define the configuration of the Jaeger tracer using the [tracing environment variables](#).

3. Create the Jaeger tracer from the environment variables that you defined in step two:

```
Tracer tracer = Configuration.fromEnv().getTracer();
```

**NOTE**

For alternative ways to initialize a Jaeger tracer, see the [Java OpenTracing library](#) documentation.

4. Register the Jaeger tracer as a global tracer:

```
GlobalTracer.register(tracer);
```

A Jaeger tracer is now initialized for the client application to use.

## 8.2.2. Environment variables for tracing

Use these environment variables when configuring a Jaeger tracer for Kafka clients.

**NOTE**

The tracing environment variables are part of the Jaeger project and are subject to change. For the latest environment variables, see the [Jaeger documentation](#).

Property	Required	Description
<code>JAEGER_SERVICE_NAME</code>	Yes	The name of the Jaeger tracer service.
<code>JAEGER_AGENT_HOST</code>	No	The hostname for communicating with the <code>jaeger-agent</code> through the User Datagram Protocol (UDP).
<code>JAEGER_AGENT_PORT</code>	No	The port used for communicating with the <code>jaeger-agent</code> through UDP.

Property	Required	Description
JAEGER_ENDPOINT	No	The <code>traces</code> endpoint. Only define this variable if the client application will bypass the <code>jaeger-agent</code> and connect directly to the <code>jaeger-collector</code> .
JAEGER_AUTH_TOKEN	No	The authentication token to send to the endpoint as a bearer token.
JAEGER_USER	No	The username to send to the endpoint if using basic authentication.
JAEGER_PASSWORD	No	The password to send to the endpoint if using basic authentication.
JAEGER_PROPAGATION	No	A comma-separated list of formats to use for propagating the trace context. Defaults to the standard Jaeger format. Valid values are <code>jaeger</code> , <code>b3</code> , and <code>w3c</code> .
JAEGER_REPORTER_LOG_SPANS	No	Indicates whether the reporter should also log the spans.
JAEGER_REPORTER_MAX_QUEUE_SIZE	No	The reporter's maximum queue size.
JAEGER_REPORTER_FLUSH_INTERVAL	No	The reporter's flush interval, in ms. Defines how frequently the Jaeger reporter flushes span batches.
JAEGER_SAMPLER_TYPE	No	<p>The sampling strategy to use for client traces:</p> <ul style="list-style-type: none"> <li>• Constant</li> <li>• Probabilistic</li> <li>• Rate Limiting</li> <li>• Remote (the default)</li> </ul> <p>To sample all traces, use the Constant sampling strategy with a parameter of 1.</p> <p>For more information, see the <a href="#">Jaeger documentation</a>.</p>
JAEGER_SAMPLER_PARAM	No	The sampler parameter (number).
JAEGER_SAMPLER_MANAGER_HOST_PORT	No	The hostname and port to use if a Remote sampling strategy is selected.
JAEGER_TAGS	No	<p>A comma-separated list of tracer-level tags that are added to all reported spans.</p> <p>The value can also refer to an environment variable using the format <code> \${envVarName:default}</code>. <code>:default</code> is optional and identifies a value to use if the environment variable cannot be found.</p>

*Additional resources*

- Initializing a Jaeger tracer for Kafka clients

## 8.3. Instrumenting Kafka clients with tracers

Instrument Kafka producer and consumer clients, and Kafka Streams API applications for distributed tracing.

### 8.3.1. Instrumenting producers and consumers for tracing

Use a Decorator pattern or Interceptors to instrument your Java producer and consumer application code for tracing.

#### *Procedure*

In the application code of each producer and consumer application:

1. Add the Maven dependency for OpenTracing to the producer or consumer's `pom.xml` file.

```
<dependency>
    <groupId>io.opentracing.contrib</groupId>
    <artifactId>opentracing-kafka-client</artifactId>
    <version>0.1.15</version>
</dependency>
```

2. Instrument your client application code using either a Decorator pattern or Interceptors.

- To use a Decorator pattern:

```

// Create an instance of the KafkaProducer:
KafkaProducer<Integer, String> producer = new KafkaProducer<>(senderProps);

// Create an instance of the TracingKafkaProducer:
TracingKafkaProducer<Integer, String> tracingProducer = new
TracingKafkaProducer<>(producer,
    tracer);

// Send:
tracingProducer.send(...);

// Create an instance of the KafkaConsumer:
KafkaConsumer<Integer, String> consumer = new KafkaConsumer<>(consumerProps);

// Create an instance of the TracingKafkaConsumer:
TracingKafkaConsumer<Integer, String> tracingConsumer = new
TracingKafkaConsumer<>(consumer,
    tracer);

// Subscribe:
tracingConsumer.subscribe(Collections.singletonList("messages"));

// Get messages:
ConsumerRecords<Integer, String> records = tracingConsumer.poll(1000);

// Retrieve SpanContext from polled record (consumer side):
ConsumerRecord<Integer, String> record = ...
SpanContext spanContext = TracingKafkaUtils.extractSpanContext(record.headers(),
    tracer);

```

- To use Interceptors:

```

// Register the tracer with GlobalTracer:
GlobalTracer.register(tracer);

// Add the TracingProducerInterceptor to the sender properties:
senderProps.put(ProducerConfig.INTERCEPTOR_CLASSES_CONFIG,
    TracingProducerInterceptor.class.getName());

// Create an instance of the KafkaProducer:
KafkaProducer<Integer, String> producer = new KafkaProducer<>(senderProps);

// Send:
producer.send(...);

// Add the TracingConsumerInterceptor to the consumer properties:
consumerProps.put(ConsumerConfig.INTERCEPTOR_CLASSES_CONFIG,
    TracingConsumerInterceptor.class.getName());

// Create an instance of the KafkaConsumer:
KafkaConsumer<Integer, String> consumer = new KafkaConsumer<>(consumerProps);

// Subscribe:
consumer.subscribe(Collections.singletonList("messages"));

// Get messages:
ConsumerRecords<Integer, String> records = consumer.poll(1000);

// Retrieve the SpanContext from a polled message (consumer side):
ConsumerRecord<Integer, String> record = ...
SpanContext spanContext = TracingKafkaUtils.extractSpanContext(record.headers(),
    tracer);

```

## Custom span names in a Decorator pattern

A *span* is a logical unit of work in Jaeger, with an operation name, start time, and duration.

To use a Decorator pattern to instrument your producer and consumer applications, define custom span names by passing a `BiFunction` object as an additional argument when creating the `TracingKafkaProducer` and `TracingKafkaConsumer` objects. The OpenTracing Apache Kafka Client Instrumentation library includes several built-in span names.

*Example: Using custom span names to instrument client application code in a Decorator pattern*

```
// Create a BiFunction for the KafkaProducer that operates on (String operationName,
ProducerRecord consumerRecord) and returns a String to be used as the name:
BiFunction<String, ProducerRecord, String> producerSpanNameProvider =
(operationName, producerRecord) -> "CUSTOM_PRODUCER_NAME";

// Create an instance of the KafkaProducer:
KafkaProducer<Integer, String> producer = new KafkaProducer<>(senderProps);

// Create an instance of the TracingKafkaProducer
TracingKafkaProducer<Integer, String> tracingProducer = new
TracingKafkaProducer<>(producer,
    tracer,
    producerSpanNameProvider);

// Spans created by the tracingProducer will now have "CUSTOM_PRODUCER_NAME" as the
// span name.

// Create a BiFunction for the KafkaConsumer that operates on (String operationName,
ConsumerRecord consumerRecord) and returns a String to be used as the name:
BiFunction<String, ConsumerRecord, String> consumerSpanNameProvider =
(operationName, consumerRecord) -> operationName.toUpperCase();

// Create an instance of the KafkaConsumer:
KafkaConsumer<Integer, String> consumer = new KafkaConsumer<>(consumerProps);

// Create an instance of the TracingKafkaConsumer, passing in the
// consumerSpanNameProvider BiFunction:
TracingKafkaConsumer<Integer, String> tracingConsumer = new
TracingKafkaConsumer<>(consumer,
    tracer,
    consumerSpanNameProvider);

// Spans created by the tracingConsumer will have the operation name as the span name,
// in upper-case.
// "receive" -> "RECEIVE"
```

## Built-in span names

When defining custom span names, you can use the following [BiFunctions](#) in the [ClientSpanNameProvider](#) class. If no [spanNameProvider](#) is specified, [CONSUMER\\_OPERATION\\_NAME](#) and [PRODUCER\\_OPERATION\\_NAME](#) are used.

BiFunction	Description
<code>CONSUMER_OPERATION_NAME,</code> <code>PRODUCER_OPERATION_NAME</code>	Returns the <code>operationName</code> as the span name: "receive" for consumers and "send" for producers.
<code>CONSUMER_PREFIXED_OPERATION_NAME(String prefix),</code> <code>PRODUCER_PREFIXED_OPERATION_NAME(String prefix)</code>	Returns a String concatenation of <code>prefix</code> and <code>operationName</code> .
<code>CONSUMER_TOPIC, PRODUCER_TOPIC</code>	Returns the name of the topic that the message was sent to or retrieved from in the format ( <code>record.topic()</code> ).
<code>PREFIXED_CONSUMER_TOPIC(String prefix),</code> <code>PREFIXED_PRODUCER_TOPIC(String prefix)</code>	Returns a String concatenation of <code>prefix</code> and the topic name in the format ( <code>record.topic()</code> ).
<code>CONSUMER_OPERATION_NAME_TOPIC,</code> <code>PRODUCER_OPERATION_NAME_TOPIC</code>	Returns the operation name and the topic name: " <code>operationName - record.topic()</code> ".
<code>CONSUMER_PREFIXED_OPERATION_NAME_TOPIC(String prefix),</code> <code>PRODUCER_PREFIXED_OPERATION_NAME_TOPIC(String prefix)</code>	Returns a String concatenation of <code>prefix</code> and " <code>operationName - record.topic()</code> ".

### 8.3.2. Instrumenting Kafka Streams applications for tracing

This section describes how to instrument Kafka Streams API applications for distributed tracing.

#### Procedure

In each Kafka Streams API application:

1. Add the `opentracing-kafka-streams` dependency to the pom.xml file for your Kafka Streams API application:

```
<dependency>
  <groupId>io.opentracing.contrib</groupId>
  <artifactId>opentracing-kafka-streams</artifactId>
  <version>0.1.15</version>
</dependency>
```

2. Create an instance of the `TracingKafkaClientSupplier` supplier interface:

```
KafkaClientSupplier supplier = new TracingKafkaClientSupplier(tracer);
```

3. Provide the supplier interface to `KafkaStreams`:

```
KafkaStreams streams = new KafkaStreams(builder.build(), new StreamsConfig(config),  
supplier);  
streams.start();
```

## 8.4. Setting up tracing for MirrorMaker, Kafka Connect, and the Kafka Bridge

Distributed tracing is supported for MirrorMaker, MirrorMaker 2.0, Kafka Connect (including Kafka Connect with Source2Image support), and the Strimzi Kafka Bridge.

### *Tracing in MirrorMaker and MirrorMaker 2.0*

For MirrorMaker and MirrorMaker 2.0, messages are traced from the source cluster to the target cluster. The trace data records messages entering and leaving the MirrorMaker or MirrorMaker 2.0 component.

### *Tracing in Kafka Connect*

Only messages produced and consumed by Kafka Connect itself are traced. To trace messages sent between Kafka Connect and external systems, you must configure tracing in the connectors for those systems. For more information, see [Configuring Kafka Connect](#).

### *Tracing in the Kafka Bridge*

Messages produced and consumed by the Kafka Bridge are traced. Incoming HTTP requests from client applications to send and receive messages through the Kafka Bridge are also traced. To have end-to-end tracing, you must configure tracing in your HTTP clients.

### 8.4.1. Enabling tracing in MirrorMaker, Kafka Connect, and Kafka Bridge resources

Update the configuration of [KafkaMirrorMaker](#), [KafkaMirrorMaker2](#), [KafkaConnect](#), [KafkaConnectS2I](#), and [KafkaBridge](#) custom resources to specify and configure a Jaeger tracer service for each resource. Updating a tracing-enabled resource in your Kubernetes cluster triggers two events:

- Interceptor classes are updated in the integrated consumers and producers in MirrorMaker, MirrorMaker 2.0, Kafka Connect, or the Strimzi Kafka Bridge.
- For MirrorMaker, MirrorMaker 2.0, and Kafka Connect, the tracing agent initializes a Jaeger tracer based on the tracing configuration defined in the resource.
- For the Kafka Bridge, a Jaeger tracer based on the tracing configuration defined in the resource is initialized by the Kafka Bridge itself.

### *Procedure*

Perform these steps for each [KafkaMirrorMaker](#), [KafkaMirrorMaker2](#), [KafkaConnect](#), [KafkaConnectS2I](#), and [KafkaBridge](#) resource.

1. In the `spec.template` property, configure the Jaeger tracer service. For example:

### *Jaeger tracer configuration for Kafka Connect*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect-cluster
spec:
  #...
  template:
    connectContainer: ①
      env:
        - name: JAEGER_SERVICE_NAME
          value: my-jaeger-service
        - name: JAEGER_AGENT_HOST
          value: jaeger-agent-name
        - name: JAEGER_AGENT_PORT
          value: "6831"
    tracing: ②
      type: jaeger
    #...
```

### *Jaeger tracer configuration for MirrorMaker*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaMirrorMaker
metadata:
  name: my-mirror-maker
spec:
  #...
  template:
    mirrorMakerContainer:
      env:
        - name: JAEGER_SERVICE_NAME
          value: my-jaeger-service
        - name: JAEGER_AGENT_HOST
          value: jaeger-agent-name
        - name: JAEGER_AGENT_PORT
          value: "6831"
    tracing:
      type: jaeger
    #...
```

### *Jaeger tracer configuration for MirrorMaker 2.0*

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaMirrorMaker2
metadata:
  name: my-mm2-cluster
spec:
  #...
  template:
    connectContainer:
      env:
        - name: JAEGER_SERVICE_NAME
          value: my-jaeger-service
        - name: JAEGER_AGENT_HOST
          value: jaeger-agent-name
        - name: JAEGER_AGENT_PORT
          value: "6831"
    tracing:
      type: jaeger
#...
```

### *Jaeger tracer configuration for the Kafka Bridge*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaBridge
metadata:
  name: my-bridge
spec:
  #...
  template:
    bridgeContainer:
      env:
        - name: JAEGER_SERVICE_NAME
          value: my-jaeger-service
        - name: JAEGER_AGENT_HOST
          value: jaeger-agent-name
        - name: JAEGER_AGENT_PORT
          value: "6831"
    tracing:
      type: jaeger
#...
```

- ① Use the [tracing environment variables](#) as template configuration properties.
- ② Set the `spec.tracing.type` property to `jaeger`.

2. Create or update the resource:

```
kubectl apply -f your-file
```

*Additional resources*

- [ContainerTemplate schema reference](#)
- [Customizing Kubernetes resources](#)

# Chapter 9. Managing TLS certificates

Strimzi supports encrypted communication between the Kafka and Strimzi components using the TLS protocol. Communication between Kafka brokers (interbroker communication), between ZooKeeper nodes (internodal communication), and between these and the Strimzi operators is always encrypted. Communication between Kafka clients and Kafka brokers is encrypted according to how the cluster is configured. For the Kafka and Strimzi components, TLS certificates are also used for authentication.

The Cluster Operator automatically sets up and renews TLS certificates to enable encryption and authentication within your cluster. It also sets up other TLS certificates if you want to enable encryption or TLS authentication between Kafka brokers and clients. Certificates provided by users are not renewed.

You can provide your own server certificates, called *Kafka listener certificates*, for TLS listeners or external listeners which have TLS encryption enabled. For more information, see [Kafka listener certificates](#).




Figure 5. Example architecture of the communication secured by TLS

## 9.1. Certificate Authorities

To support encryption, each Strimzi component needs its own private keys and public key certificates. All component certificates are signed by an internal Certificate Authority (CA) called the *cluster CA*.

Similarly, each Kafka client application connecting to Strimzi using TLS client authentication needs to provide private keys and certificates. A second internal CA, named the *clients CA*, is used to sign

certificates for the Kafka clients.

### 9.1.1. CA certificates

Both the cluster CA and clients CA have a self-signed public key certificate.

Kafka brokers are configured to trust certificates signed by either the cluster CA or clients CA. Components that clients do not need to connect to, such as ZooKeeper, only trust certificates signed by the cluster CA. Unless TLS encryption for external listeners is disabled, client applications must trust certificates signed by the cluster CA. This is also true for client applications that perform [mutual TLS authentication](#).

By default, Strimzi automatically generates and renews CA certificates issued by the cluster CA or clients CA. You can configure the management of these CA certificates in the `Kafka.spec.clusterCa` and `Kafka.spec.clientsCa` objects. Certificates provided by users are not renewed.

You can provide your own CA certificates for the cluster CA or clients CA. For more information, see [Installing your own CA certificates](#). If you provide your own certificates, you must manually renew them when needed.

### 9.1.2. Installing your own CA certificates

This procedure describes how to install your own CA certificates and keys instead of using the CA certificates and private keys generated by the Cluster Operator.

You can use this procedure to install your own cluster or client CA certificates.

The procedure describes renewal of CA certificates in PEM format. You can also use certificates in PKCS #12 format.

#### *Prerequisites*

- The Cluster Operator is running.
- A Kafka cluster is not yet deployed.
- Your own X.509 certificates and keys in PEM format for the cluster CA or clients CA.
  - If you want to use a cluster or clients CA which is not a Root CA, you have to include the whole chain in the certificate file. The chain should be in the following order:
    1. The cluster or clients CA
    2. One or more intermediate CAs
    3. The root CA
  - All CAs in the chain should be configured as a CA in the X509v3 Basic Constraints.

#### *Procedure*

1. Put your CA certificate in the corresponding `Secret`.
  - a. Delete the existing secret:

```
kubectl delete secret CA-CERTIFICATE-SECRET
```

*CA-CERTIFICATE-SECRET* is the name of the *Secret*, which is *CLUSTER-NAME-cluster-ca-cert* for the cluster CA certificate and *CLUSTER-NAME-clients-ca-cert* for the clients CA certificate.

Ignore any "Not Exists" errors.

- b. Create and label the new secret

```
kubectl create secret generic CA-CERTIFICATE-SECRET --from-file=ca.crt=CA-CERTIFICATE-FILENAME
```

2. Put your CA key in the corresponding *Secret*.

- a. Delete the existing secret:

```
kubectl delete secret CA-KEY-SECRET
```

*CA-KEY-SECRET* is the name of CA key, which is *CLUSTER-NAME-cluster-ca* for the cluster CA key and *CLUSTER-NAME-clients-ca* for the clients CA key.

- b. Create the new secret:

```
kubectl create secret generic CA-KEY-SECRET --from-file=ca.key=CA-KEY-SECRET-FILENAME
```

3. Label the secrets with the labels *strimzi.io/kind=Kafka* and *strimzi.io/cluster=CLUSTER-NAME*:

```
kubectl label secret CA-CERTIFICATE-SECRET strimzi.io/kind=Kafka  
strimzi.io/cluster=CLUSTER-NAME  
kubectl label secret CA-KEY-SECRET strimzi.io/kind=Kafka  
strimzi.io/cluster=CLUSTER-NAME
```

4. Create the *Kafka* resource for your cluster, configuring either the *Kafka.spec.clusterCa* or the *Kafka.spec.clientsCa* object to *not* use generated CAs:

*Example fragment Kafka resource configuring the cluster CA to use certificates you supply for yourself*

```
kind: Kafka  
version: kafka.strimzi.io/v1beta1  
spec:  
  # ...  
  clusterCa:  
    generateCertificateAuthority: false
```

## Additional resources

- To renew CA certificates you have previously installed, see [Renewing your own CA certificates](#).
- [Providing your own Kafka listener certificates](#).

## 9.2. Secrets

Strimzi uses *Secrets* to store private keys and certificates for Kafka cluster components and clients. Secrets are used for establishing TLS encrypted connections between Kafka brokers, and between brokers and clients. They are also used for mutual TLS authentication.

- A *Cluster Secret* contains a cluster CA certificate to sign Kafka broker certificates, and is used by a connecting client to establish a TLS encrypted connection with the Kafka cluster to validate broker identity.
- A *Client Secret* contains a client CA certificate for a user to sign its own client certificate to allow mutual authentication against the Kafka cluster. The broker validates the client identity through the client CA certificate itself.
- A *User Secret* contains a private key and certificate, which are generated and signed by the client CA certificate when a new user is created. The key and certificate are used for authentication and authorization when accessing the cluster.

Secrets provide private keys and certificates in PEM and PKCS #12 formats. Using private keys and certificates in PEM format means that users have to get them from the Secrets, and generate a corresponding truststore (or keystore) to use in their Java applications. PKCS #12 storage provides a truststore (or keystore) that can be used directly.

All keys are 2048 bits in size.

### 9.2.1. PKCS #12 storage

PKCS #12 defines an archive file format ([.p12](#)) for storing cryptography objects into a single file with password protection. You can use PKCS #12 to manage certificates and keys in one place.

Each Secret contains fields specific to PKCS #12.

- The `.p12` field contains the certificates and keys.
- The `.password` field is the password that protects the archive.

### 9.2.2. Cluster CA Secrets

The following tables describe the Cluster Secrets that are managed by the Cluster Operator in a Kafka cluster.

Only the `<cluster>-cluster-ca-cert` Secret needs to be used by clients. All other *Secrets* described only need to be accessed by the Strimzi components. You can enforce this using Kubernetes role-based access controls, if necessary.

*Table 3. Fields in the `<cluster>-cluster-ca` Secret*

Field	Description
ca.key	The current private key for the cluster CA.

Table 4. Fields in the <cluster>-cluster-ca-cert Secret

Field	Description
ca.p12	PKCS #12 archive file for storing certificates and keys.
ca.password	Password for protecting the PKCS #12 archive file.
ca.crt	The current certificate for the cluster CA.

**NOTE** The CA certificates in <cluster>-cluster-ca-cert must be trusted by Kafka client applications so that they validate the Kafka broker certificates when connecting to Kafka brokers over TLS.

Table 5. Fields in the <cluster>-kafka-brokers Secret

Field	Description
<cluster>-kafka-<num>.p12	PKCS #12 archive file for storing certificates and keys.
<cluster>-kafka-<num>.password	Password for protecting the PKCS #12 archive file.
<cluster>-kafka-<num>.crt	Certificate for Kafka broker pod <num>. Signed by a current or former cluster CA private key in <cluster>-cluster-ca.
<cluster>-kafka-<num>.key	Private key for Kafka broker pod <num>.

Table 6. Fields in the <cluster>-zookeeper-nodes Secret

Field	Description
<cluster>-zookeeper-<num>.p12	PKCS #12 archive file for storing certificates and keys.
<cluster>-zookeeper-<num>.password	Password for protecting the PKCS #12 archive file.
<cluster>-zookeeper-<num>.crt	Certificate for ZooKeeper node <num>. Signed by a current or former cluster CA private key in <cluster>-cluster-ca.
<cluster>-zookeeper-<num>.key	Private key for ZooKeeper pod <num>.

Table 7. Fields in the <cluster>-entity-operator-certs Secret

Field	Description
entity-operator_.p12	PKCS #12 archive file for storing certificates and keys.
entity-operator_.password	Password for protecting the PKCS #12 archive file.
entity-operator_.crt	Certificate for TLS communication between the Entity Operator and Kafka or ZooKeeper. Signed by a current or former cluster CA private key in <cluster>-cluster-ca.
entity-operator.key	Private key for TLS communication between the Entity Operator and Kafka or ZooKeeper.

### 9.2.3. Client CA Secrets

Table 8. Clients CA [Secrets](#) managed by the Cluster Operator in <cluster>

Secret name	Field within Secret	Description
<cluster>-clients-ca	ca.key	The current private key for the clients CA.
<cluster>-clients-ca-cert	ca.p12	PKCS #12 archive file for storing certificates and keys.
	ca.password	Password for protecting the PKCS #12 archive file.
	ca.crt	The current certificate for the clients CA.

The certificates in <cluster>-clients-ca-cert are those which the Kafka brokers trust.

**NOTE**

<cluster>-clients-ca is used to sign certificates of client applications. It needs to be accessible to the Strimzi components and for administrative access if you are intending to issue application certificates without using the User Operator. You can enforce this using Kubernetes role-based access controls if necessary.

### 9.2.4. Disabling ownerReference in the CA Secrets

By default, the CA secrets created by Strimzi have an `ownerReference` property set to the [Kafka](#) custom resource. This means that when the [Kafka](#) custom resource is deleted, the CA secrets are garbage collected (deleted) by Kubernetes as well.

If you want to reuse the CA for a new cluster, you can disable the `ownerReference` by setting the `generateSecretOwnerReference` property for the cluster and client CAs to `false` in the [Kafka](#) configuration. When disabled, CA secrets are not deleted by Kubernetes when the corresponding [Kafka](#) custom resource is deleted.

*Example Kafka configuration with disabled ownerReference for cluster and client CAs*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
# ...
spec:
# ...
  clusterCa:
    generateSecretOwnerReference: false
  clientsCa:
    generateSecretOwnerReference: false
# ...
```

## 9.2.5. User Secrets

Table 9. Secrets managed by the User Operator

Secret name	Field within Secret	Description
<user>	<code>user.p12</code>	PKCS #12 archive file for storing certificates and keys.
	<code>user.password</code>	Password for protecting the PKCS #12 archive file.
	<code>user.crt</code>	Certificate for the user, signed by the clients CA
	<code>user.key</code>	Private key for the user

## 9.3. Certificate renewal and validity periods

Cluster CA and clients CA certificates are only valid for a limited time period, known as the validity period. This is usually defined as a number of days since the certificate was generated.

For auto-generated CA certificates, you can configure the validity period of:

- Cluster CA certificates in `Kafka.spec.clusterCa.validityDays`
- Client CA certificates in `Kafka.spec.clientsCa.validityDays`

The default validity period for both certificates is 365 days. Manually-installed CA certificates should have their own validity periods defined.

When a CA certificate expires, components and clients that still trust that certificate will not accept TLS connections from peers whose certificates were signed by the CA private key. The components and clients need to trust the *new* CA certificate instead.

To allow the renewal of CA certificates without a loss of service, the Cluster Operator will initiate certificate renewal before the old CA certificates expire.


You can configure the renewal period of:

- Cluster CA certificates in `Kafka.spec.clusterCa.renewalDays`
- Client CA certificates in `Kafka.spec.clientsCa.renewalDays`

The default renewal period for both certificates is 30 days.

The renewal period is measured backwards, from the expiry date of the current certificate.

*Validity period against renewal period*



To make a change to the validity and renewal periods after creating the Kafka cluster, you configure and apply the [Kafka](#) custom resource, and [manually renew the CA certificates](#). If you do not manually renew the certificates, the new periods will be used the next time the certificate is renewed automatically.

*Example Kafka configuration for certificate validity and renewal periods*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
# ...
spec:
# ...
  clusterCa:
    renewalDays: 30
    validityDays: 365
    generateCertificateAuthority: true
  clientsCa:
    renewalDays: 30
    validityDays: 365
    generateCertificateAuthority: true
# ...
```

The behavior of the Cluster Operator during the renewal period depends on whether the relevant setting is enabled in either `Kafka.spec.clusterCa.generateCertificateAuthority` or `Kafka.spec.clientsCa.generateCertificateAuthority`.

### 9.3.1. Renewal process with generated CAs

The Cluster Operator performs the following process to renew CA certificates:

1. Generate a new CA certificate, but retain the existing key. The new certificate replaces the old one with the name `ca.crt` within the corresponding `Secret`.
2. Generate new client certificates (for ZooKeeper nodes, Kafka brokers, and the Entity Operator). This is not strictly necessary because the signing key has not changed, but it keeps the validity period of the client certificate in sync with the CA certificate.
3. Restart ZooKeeper nodes so that they will trust the new CA certificate and use the new client certificates.
4. Restart Kafka brokers so that they will trust the new CA certificate and use the new client certificates.
5. Restart the Topic and User Operators so that they will trust the new CA certificate and use the new client certificates.

### 9.3.2. Client applications

The Cluster Operator is not aware of the client applications using the Kafka cluster.

When connecting to the cluster, and to ensure they operate correctly, client applications must:

- Trust the cluster CA certificate published in the <cluster>-cluster-ca-cert Secret.
- Use the credentials published in their <user-name> Secret to connect to the cluster.

The User Secret provides credentials in PEM and PKCS #12 format, or it can provide a password when using SCRAM-SHA authentication. The User Operator creates the user credentials when a user is created.

For workloads running inside the same Kubernetes cluster and namespace, Secrets can be mounted as a volume so the client Pods construct their keystores and truststores from the current state of the Secrets. For more details on this procedure, see [Configuring internal clients to trust the cluster CA](#).

### Client certificate renewal

You must ensure clients continue to work after certificate renewal. The renewal process depends on how the clients are configured.

If you are provisioning client certificates and keys manually, you must generate new client certificates and ensure the new certificates are used by clients within the renewal period. Failure to do this by the end of the renewal period could result in client applications being unable to connect to the cluster.

#### 9.3.3. Renewing CA certificates manually

Cluster and clients CA certificates auto-renew at the start of their respective certificate renewal periods. If `Kafka.spec.clusterCa.generateCertificateAuthority` and `Kafka.spec.clientsCa.generateCertificateAuthority` are set to `false`, the CA certificates do not auto-renew.

You can manually renew one or both of these certificates before the certificate renewal period starts. You might do this for security reasons, or if you have [changed the renewal or validity periods for the certificates](#).

A renewed certificate uses the same private key as the old certificate.

##### Prerequisites

- The Cluster Operator is running.
- A Kafka cluster in which CA certificates and private keys are installed.

##### Procedure

1. Apply the `strimzi.io/force-renew` annotation to the `Secret` that contains the CA certificate that you want to renew.

*Table 10. Annotation for the Secret that forces renewal of certificates*

Certificate	Secret	Annotate command
Cluster CA	<code>KAFKA-CLUSTER-NAME</code> -cluster-ca-cert	<code>kubectl annotate secret KAFKA-CLUSTER-NAME-cluster-ca-cert strimzi.io/force-renew=true</code>

Certificate	Secret	Annotate command
Clients CA	<code>KAFKA-CLUSTER-NAME-clients-ca-cert</code>	<code>kubectl annotate secret KAFKA-CLUSTER-NAME-clients-ca-cert strimzi.io/force-renew=true</code>

At the next reconciliation the Cluster Operator will generate a new CA certificate for the `Secret` that you annotated. If maintenance time windows are configured, the Cluster Operator will generate the new CA certificate at the first reconciliation within the next maintenance time window.

Client applications must reload the cluster and clients CA certificates that were renewed by the Cluster Operator.

## 2. Check the period the CA certificate is valid:

For example, using an `openssl` command:

```
kubectl get secret CA-CERTIFICATE-SECRET -o 'jsonpath={.data.CA-CERTIFICATE}' | base64 -d | openssl x509 -subject -issuer -startdate -enddate -noout
```

`CA-CERTIFICATE-SECRET` is the name of the `Secret`, which is `KAFKA-CLUSTER-NAME-cluster-ca-cert` for the cluster CA certificate and `KAFKA-CLUSTER-NAME-clients-ca-cert` for the clients CA certificate.

`CA-CERTIFICATE` is the name of the CA certificate, such as `jsonpath={.data.ca\.crt}`.

The command returns a `notBefore` and `notAfter` date, which is the validity period for the CA certificate.

For example, for a cluster CA certificate:

```
subject=O = io.strimzi, CN = cluster-ca v0
issuer=O = io.strimzi, CN = cluster-ca v0
notBefore=Jun 30 09:43:54 2020 GMT
notAfter=Jun 30 09:43:54 2021 GMT
```

## 3. Delete old certificates from the Secret.

When components are using the new certificates, older certificates might still be active. Delete the old certificates to remove any potential security risk.

### Additional resources

- [Secrets](#)
- [Maintenance time windows for rolling updates](#)
- [CertificateAuthority schema reference](#)

### 9.3.4. Renewing your own CA certificates

This procedure describes how to renew CA certificates and keys that you previously installed. You need to follow this procedure during the renewal period in order to replace CA certificates which will soon expire.

You can use this procedure to install your own cluster or client CA certificates.

The procedure describes renewal of CA certificates in PEM format. You can also use certificates in PKCS #12 format.

#### *Prerequisites*

- The Cluster Operator is running.
- A Kafka cluster in which you previously installed your own CA certificates and private keys.
- New cluster and clients X.509 certificates and keys in PEM format.

These could be generated using an `openssl` command, such as:

```
openssl req -x509 -new -days NUMBER-OF-DAYS-VALID --nodes -out ca.crt -keyout ca.key
```

#### *Procedure*

1. Check the details of the current CA certificates in the `Secret`:

```
kubectl describe secret CA-CERTIFICATE-SECRET
```

`CA-CERTIFICATE-SECRET` is the name of the `Secret`, which is `KAFKA-CLUSTER-NAME-cluster-ca-cert` for the cluster CA certificate and `KAFKA-CLUSTER-NAME-clients-ca-cert` for the clients CA certificate.

2. Create a directory to contain the existing CA certificates in the secret.

```
mkdir new-ca-cert-secret  
cd new-ca-cert-secret
```

3. Fetch the secret for each CA certificate you wish to renew:

```
kubectl get secret CA-CERTIFICATE-SECRET -o 'jsonpath={.data.CA-CERTIFICATE}' |  
base64 -d > CA-CERTIFICATE
```

Replace `CA-CERTIFICATE` with the name of each CA certificate.

4. Rename the old `ca.crt` file as `ca-DATE.crt`, where `DATE` is the certificate expiry date in the format `YEAR-MONTH-DAYTHOUR-MINUTE-SECONDZ`.

For example `ca-2018-09-27T17-32-00Z.crt`.

```
mv ca.crt ca-$(date -u -d$(openssl x509 -enddate -noout -in ca.crt | sed 's/.*=//') +'%Y-%m-%dT%H-%M-%SZ').crt
```

5. Copy your new CA certificate into the directory, naming it `ca.crt`:

```
cp PATH-TO-NEW-CERTIFICATE ca.crt
```

6. Put your CA certificate in the corresponding `Secret`.

- Delete the existing secret:

```
kubectl delete secret CA-CERTIFICATE-SECRET
```

`CA-CERTIFICATE-SECRET` is the name of the `Secret`, as returned in the first step.

Ignore any "Not Exists" errors.

- Recreate the secret:

```
kubectl create secret generic CA-CERTIFICATE-SECRET --from-file=.
```

7. Delete the directory you created:

```
cd ..  
rm -r new-ca-cert-secret
```

8. Put your CA key in the corresponding `Secret`.

- Delete the existing secret:

```
kubectl delete secret CA-KEY-SECRET
```

`CA-KEY-SECRET` is the name of CA key, which is `KAFKA-CLUSTER-NAME-cluster-ca` for the cluster CA key and `KAFKA-CLUSTER-NAME-clients-ca` for the clients CA key.

- Recreate the secret with the new CA key:

```
kubectl create secret generic CA-KEY-SECRET --from-file=ca.key=CA-KEY-SECRET-  
FILENAME
```

9. Label the secrets with the labels `strimzi.io/kind=Kafka` and `strimzi.io/cluster=KAFKA-CLUSTER-NAME`:

```

kubectl label secret CA-CERTIFICATE-SECRET strimzi.io/kind=Kafka
strimzi.io/cluster=KAFKA-CLUSTER-NAME
kubectl label secret CA-KEY-SECRET strimzi.io/kind=Kafka strimzi.io/cluster=KAFKA-
CLUSTER-NAME

```

## 9.4. Replacing private keys

You can replace the private keys used by the cluster CA and clients CA certificates. When a private key is replaced, the Cluster Operator generates a new CA certificate for the new private key.

### *Prerequisites*

- The Cluster Operator is running.
- A Kafka cluster in which CA certificates and private keys are installed.

### *Procedure*

- Apply the `strimzi.io/force-replace` annotation to the `Secret` that contains the private key that you want to renew.

*Table 11. Commands for replacing private keys*

Private key for	Secret	Annotate command
Cluster CA	<code>&lt;cluster-name&gt;-cluster-ca</code>	<code>kubectl annotate secret &lt;cluster-name&gt;-cluster-ca strimzi.io/force-replace=true</code>
Clients CA	<code>&lt;cluster-name&gt;-clients-ca</code>	<code>kubectl annotate secret &lt;cluster-name&gt;-clients-ca strimzi.io/force-replace=true</code>

At the next reconciliation the Cluster Operator will:

- Generate a new private key for the `Secret` that you annotated
- Generate a new CA certificate

If maintenance time windows are configured, the Cluster Operator will generate the new private key and CA certificate at the first reconciliation within the next maintenance time window.

Client applications must reload the cluster and clients CA certificates that were renewed by the Cluster Operator.

### *Additional resources*

- [Secrets](#)
- [Maintenance time windows for rolling updates](#)

## 9.5. TLS connections

### 9.5.1. ZooKeeper communication

Communication between the ZooKeeper nodes on all ports as well as between clients and ZooKeeper is encrypted.

### 9.5.2. Kafka interbroker communication

Communication between Kafka brokers is done through an internal listener on port 9091, which is encrypted by default and not accessible to Kafka clients.

Communication between Kafka brokers and ZooKeeper nodes is also encrypted.

### 9.5.3. Topic and User Operators

All Operators use encryption for communication with both Kafka and ZooKeeper. In Topic and User Operators, a TLS sidecar is used when communicating with ZooKeeper.

### 9.5.4. Cruise Control

Cruise Control uses encryption for communication with both Kafka and ZooKeeper. A TLS sidecar is used when communicating with ZooKeeper.

### 9.5.5. Kafka Client connections

Encrypted or unencrypted communication between Kafka brokers and clients is configured using the `tls` property for `spec.kafka.listeners`.

## 9.6. Configuring internal clients to trust the cluster CA

This procedure describes how to configure a Kafka client that resides inside the Kubernetes cluster — connecting to a TLS listener — to trust the cluster CA certificate.

The easiest way to achieve this for an internal client is to use a volume mount to access the `Secrets` containing the necessary certificates and keys.

Follow the steps to configure trust certificates that are signed by the cluster CA for Java-based Kafka Producer, Consumer, and Streams APIs.

Choose the steps to follow according to the certificate format of the cluster CA: PKCS #12 (`.p12`) or PEM (`.crt`).

The steps describe how to mount the Cluster Secret that verifies the identity of the Kafka cluster to the client pod.

#### *Prerequisites*

- The Cluster Operator must be running.
- There needs to be a `Kafka` resource within the Kubernetes cluster.
- You need a Kafka client application inside the Kubernetes cluster that will connect using TLS,

and needs to trust the cluster CA certificate.

- The client application must be running in the same namespace as the **Kafka** resource.

#### *Using PKCS #12 format (.p12)*

1. Mount the cluster Secret as a volume when defining the client pod.

For example:

```
kind: Pod
apiVersion: v1
metadata:
  name: client-pod
spec:
  containers:
    - name: client-name
      image: client-name
      volumeMounts:
        - name: secret-volume
          mountPath: /data/p12
      env:
        - name: SECRET_PASSWORD
          valueFrom:
            secretKeyRef:
              name: my-secret
              key: my-password
      volumes:
        - name: secret-volume
          secret:
            secretName: my-cluster-cluster-ca-cert
```

Here we're mounting:

- The PKCS #12 file into an exact path, which can be configured
  - The password into an environment variable, where it can be used for Java configuration
2. Configure the Kafka client with the following properties:
    - A security protocol option:
      - **security.protocol: SSL** when using TLS for encryption (with or without TLS authentication).
      - **security.protocol: SASL\_SSL** when using SCRAM-SHA authentication over TLS.
    - **ssl.truststore.location** with the truststore location where the certificates were imported.
    - **ssl.truststore.password** with the password for accessing the truststore.
    - **ssl.truststore.type=PKCS12** to identify the truststore type.

#### *Using PEM format (.crt)*

1. Mount the cluster Secret as a volume when defining the client pod.

For example:

```
kind: Pod
apiVersion: v1
metadata:
  name: client-pod
spec:
  containers:
    - name: client-name
      image: client-name
      volumeMounts:
        - name: secret-volume
          mountPath: /data/crt
  volumes:
    - name: secret-volume
      secret:
        secretName: my-cluster-cluster-ca-cert
```

2. Use the certificate with clients that use certificates in X.509 format.

## 9.7. Configuring external clients to trust the cluster CA

This procedure describes how to configure a Kafka client that resides outside the Kubernetes cluster – connecting to an **external** listener – to trust the cluster CA certificate. Follow this procedure when setting up the client and during the renewal period, when the old clients CA certificate is replaced.

Follow the steps to configure trust certificates that are signed by the cluster CA for Java-based Kafka Producer, Consumer, and Streams APIs.

Choose the steps to follow according to the certificate format of the cluster CA: PKCS #12 (**.p12**) or PEM (**.crt**).

The steps describe how to obtain the certificate from the Cluster Secret that verifies the identity of the Kafka cluster.

**IMPORTANT**

The **<cluster-name>-cluster-ca-cert Secret** will contain more than one CA certificate during the CA certificate renewal period. Clients must add *all* of them to their truststores.

*Prerequisites*

- The Cluster Operator must be running.
- There needs to be a **Kafka** resource within the Kubernetes cluster.
- You need a Kafka client application outside the Kubernetes cluster that will connect using TLS, and needs to trust the cluster CA certificate.

*Using PKCS #12 format (.p12)*

1. Extract the cluster CA certificate and password from the generated <cluster-name>-cluster-ca-cert Secret.

```
kubectl get secret <cluster-name>-cluster-ca-cert -o jsonpath='{.data.ca\.p12}' | base64 -d > ca.p12
```

```
kubectl get secret <cluster-name>-cluster-ca-cert -o jsonpath='{.data.ca\.password}' | base64 -d > ca.password
```

2. Configure the Kafka client with the following properties:

- A security protocol option:
  - **security.protocol**: `SSL` when using TLS for encryption (with or without TLS authentication).
  - **security.protocol**: `SASL_SSL` when using SCRAM-SHA authentication over TLS.
- **ssl.truststore.location** with the truststore location where the certificates were imported.
- **ssl.truststore.password** with the password for accessing the truststore. This property can be omitted if it is not needed by the truststore.
- **ssl.truststore.type=PKCS12** to identify the truststore type.

*Using PEM format (.crt)*

1. Extract the cluster CA certificate from the generated <cluster-name>-cluster-ca-cert Secret.

```
kubectl get secret <cluster-name>-cluster-ca-cert -o jsonpath='{.data.ca\.crt}' | base64 -d > ca.crt
```

2. Use the certificate with clients that use certificates in X.509 format.

## 9.8. Kafka listener certificates

You can provide your own server certificates and private keys for the following types of listeners:

- Internal TLS listeners for communication within the Kubernetes cluster
- External listeners (`route`, `loadbalancer`, `ingress`, and `nodeport` types), which have TLS encryption enabled, for communication between Kafka clients and Kafka brokers

These user-provided certificates are called *Kafka listener certificates*.

Providing Kafka listener certificates for external listeners allows you to leverage existing security infrastructure, such as your organization's private CA or a public CA. Kafka clients will connect to Kafka brokers using Kafka listener certificates rather than certificates signed by the cluster CA or clients CA.

You must manually renew Kafka listener certificates when needed.

## 9.8.1. Providing your own Kafka listener certificates

This procedure shows how to configure a listener to use your own private key and server certificate, called a [Kafka listener certificate](#).

Your client applications should use the CA public key as a trusted certificate in order to verify the identity of the Kafka broker.

### *Prerequisites*

- A Kubernetes cluster.
- The Cluster Operator is running.
- For each listener, a compatible server certificate signed by an external CA.
  - Provide an X.509 certificate in PEM format.
  - Specify the correct Subject Alternative Names (SANs) for each listener. For more information, see [Alternative subjects in server certificates for Kafka listeners](#).
  - You can provide a certificate that includes the whole CA chain in the certificate file.

### *Procedure*

1. Create a **Secret** containing your private key and server certificate:

```
kubectl create secret generic my-secret --from-file=my-listener-key.key --from-file=my-listener-certificate.crt
```

2. Edit the **Kafka** resource for your cluster. Configure the listener to use your **Secret**, certificate file, and private key file in the **configuration.brokerCertChainAndKey** property.

*Example configuration for a **loadbalancer** external listener with TLS encryption enabled*

```
# ...
listeners:
  - name: plain
    port: 9092
    type: internal
    tls: false
  - name: external
    port: 9094
    type: loadbalancer
    tls: true
    authentication:
      type: tls
    configuration:
      brokerCertChainAndKey:
        secretName: my-secret
        certificate: my-listener-certificate.crt
        key: my-listener-key.key
# ...
```

### *Example configuration for a TLS listener*

```
# ...
listeners:
  - name: plain
    port: 9092
    type: internal
    tls: false
  - name: tls
    port: 9093
    type: internal
    tls: true
    authentication:
      type: tls
    configuration:
      brokerCertChainAndKey:
        secretName: my-secret
        certificate: my-listener-certificate.crt
        key: my-listener-key.key
# ...
```

3. Apply the new configuration to create or update the resource:

```
kubectl apply -f kafka.yaml
```

The Cluster Operator starts a rolling update of the Kafka cluster, which updates the configuration of the listeners.

**NOTE**

A rolling update is also started if you update a Kafka listener certificate in a [Secret](#) that is already used by a TLS or external listener.

#### *Additional resources*

- [Alternative subjects in server certificates for Kafka listeners](#)
- [GenericKafkaListener schema reference](#)
- [Kafka listener certificates](#)

## **9.8.2. Alternative subjects in server certificates for Kafka listeners**

In order to use TLS hostname verification with your own [Kafka listener certificates](#), you must use the correct Subject Alternative Names (SANs) for each listener. The certificate SANs must specify hostnames for:

- All of the Kafka brokers in your cluster
- The Kafka cluster bootstrap service

You can use wildcard certificates if they are supported by your CA.

## TLS listener SAN examples

Use the following examples to help you specify hostnames of the SANs in your certificates for TLS listeners.

### *Wildcards example*

```
//Kafka brokers
*.<cluster-name>-kafka-brokers
*.<cluster-name>-kafka-brokers.<namespace>.svc

// Bootstrap service
<cluster-name>-kafka-bootstrap
<cluster-name>-kafka-bootstrap.<namespace>.svc
```

### *Non-wildcards example*

```
// Kafka brokers
<cluster-name>-kafka-0.<cluster-name>-kafka-brokers
<cluster-name>-kafka-0.<cluster-name>-kafka-brokers.<namespace>.svc
<cluster-name>-kafka-1.<cluster-name>-kafka-brokers
<cluster-name>-kafka-1.<cluster-name>-kafka-brokers.<namespace>.svc
# ...

// Bootstrap service
<cluster-name>-kafka-bootstrap
<cluster-name>-kafka-bootstrap.<namespace>.svc
```

## External listener SAN examples

For external listeners which have TLS encryption enabled, the hostnames you need to specify in certificates depends on the external listener [type](#).

*Table 12. SANs for each type of external listener*

External listener type	In the SANs, specify...
Route	Addresses of all Kafka broker <a href="#">Routes</a> and the address of the bootstrap <a href="#">Route</a> .  You can use a matching wildcard name.
loadbalancer	Addresses of all Kafka broker <a href="#">loadbalancers</a> and the bootstrap <a href="#">loadbalancer</a> address.  You can use a matching wildcard name.
NodePort	Addresses of all Kubernetes worker nodes that the Kafka broker pods might be scheduled to.  You can use a matching wildcard name.

*Additional resources*

- [Providing your own Kafka listener certificates](#)

# Chapter 10. Managing Strimzi

This chapter covers tasks to maintain a deployment of Strimzi.

## 10.1. Working with custom resources

You can use `kubectl` commands to retrieve information and perform other operations on Strimzi custom resources.

Using `kubectl` with the `status` subresource of a custom resource allows you to get the information about the resource.

### 10.1.1. Performing `kubectl` operations on custom resources

Use `kubectl` commands, such as `get`, `describe`, `edit`, or `delete`, to perform operations on resource types. For example, `kubectl get kafkatopics` retrieves a list of all Kafka topics and `kubectl get kafkas` retrieves all deployed Kafka clusters.

When referencing resource types, you can use both singular and plural names: `kubectl get kafkas` gets the same results as `kubectl get kafka`.

You can also use the *short name* of the resource. Learning short names can save you time when managing Strimzi. The short name for `Kafka` is `k`, so you can also run `kubectl get k` to list all Kafka clusters.

```
kubectl get k

NAME      DESIRED KAFKA REPLICAS  DESIRED ZK REPLICAS
my-cluster  3                      3
```

Table 13. Long and short names for each Strimzi resource

Strimzi resource	Long name	Short name
Kafka	kafka	k
Kafka Topic	kafkatopic	kt
Kafka User	kafkauser	ku
Kafka Connect	kafkaconnect	kc
Kafka Connect S2I	kafkaconnects2i	kcs2i
Kafka Connector	kafkaconnector	kctr
Kafka Mirror Maker	kafkamirrormaker	kmm
Kafka Mirror Maker 2	kafkamirrormaker2	kmm2
Kafka Bridge	kafkabridge	kb
Kafka Rebalance	kafkarebalance	kr

## Resource categories

Categories of custom resources can also be used in `kubectl` commands.

All Strimzi custom resources belong to the category `strimzi`, so you can use `strimzi` to get all the Strimzi resources with one command.

For example, running `kubectl get strimzi` lists all Strimzi custom resources in a given namespace.

```
kubectl get strimzi

NAME                                     DESIRED KAFKA REPLICAS DESIRED ZK REPLICAS
kafka.kafka.strimzi.io/my-cluster         3                      3

NAME                                     PARTITIONS REPLICATION FACTOR
kafkatopic.kafka.strimzi.io/kafka-apps   3                      3

NAME                                     AUTHENTICATION AUTHORIZATION
kafkauser.kafka.strimzi.io/my-user        tls                   simple
```

The `kubectl get strimzi -o name` command returns all resource types and resource names. The `-o name` option fetches the output in the *type/name* format

```
kubectl get strimzi -o name

kafka.kafka.strimzi.io/my-cluster
kafkatopic.kafka.strimzi.io/kafka-apps
kafkauser.kafka.strimzi.io/my-user
```

You can combine this `strimzi` command with other commands. For example, you can pass it into a `kubectl delete` command to delete all resources in a single command.

```
kubectl delete $(kubectl get strimzi -o name)

kafka.kafka.strimzi.io "my-cluster" deleted
kafkatopic.kafka.strimzi.io "kafka-apps" deleted
kafkauser.kafka.strimzi.io "my-user" deleted
```

Deleting all resources in a single operation might be useful, for example, when you are testing new Strimzi features.

## Querying the status of sub-resources

There are other values you can pass to the `-o` option. For example, by using `-o yaml` you get the output in YAML format. Using `-o json` will return it as JSON.

You can see all the options in `kubectl get --help`.

One of the most useful options is the [JSONPath support](#), which allows you to pass JSONPath expressions to query the Kubernetes API. A JSONPath expression can extract or navigate specific parts of any resource.

For example, you can use the JSONPath expression `{.status.listeners[?(@.type=="tls")].bootstrapServers}` to get the bootstrap address from the status of the Kafka custom resource and use it in your Kafka clients.

Here, the command finds the `bootstrapServers` value of the `tls` listeners.

```
kubectl get kafka my-cluster  
-o=jsonpath='{.status.listeners[?(@.type=="tls")].bootstrapServers}{ "\n"}'  
  
my-cluster-kafka-bootstrap.myproject.svc:9093
```

By changing the type condition to `@.type=="external"` or `@.type=="plain"` you can also get the address of the other Kafka listeners.

```
kubectl get kafka my-cluster  
-o=jsonpath='{.status.listeners[?(@.type=="external")].bootstrapServers}{ "\n"}'  
  
192.168.1.247:9094
```

You can use `jsonpath` to extract any other property or group of properties from any custom resource.

### 10.1.2. Strimzi custom resource status information

Several resources have a `status` property, as described in the following table.

*Table 14. Custom resource status properties*

Strimzi resource	Schema reference	Publishes status information on...
Kafka	<a href="#">KafkaStatus schema reference</a>	The Kafka cluster.
KafkaConnect	<a href="#">KafkaConnectStatus schema reference</a>	The Kafka Connect cluster, if deployed.
KafkaConnectS2I	<a href="#">KafkaConnectS2IStatus schema reference</a>	The Kafka Connect cluster with Source-to-Image support, if deployed.
KafkaConnector	<a href="#">KafkaConnectorStatus schema reference</a>	<a href="#">KafkaConnector</a> resources, if deployed.
KafkaMirrorMaker	<a href="#">KafkaMirrorMakerStatus schema reference</a>	The Kafka MirrorMaker tool, if deployed.

Strimzi resource	Schema reference	Publishes status information on...
KafkaTopic	<a href="#">KafkaTopicStatus schema reference</a>	Kafka topics in your Kafka cluster.
KafkaUser	<a href="#">KafkaUserStatus schema reference</a>	Kafka users in your Kafka cluster.
KafkaBridge	<a href="#">KafkaBridgeStatus schema reference</a>	The Strimzi Kafka Bridge, if deployed.

The `status` property of a resource provides information on the resource's:

- *Current state*, in the `status.conditions` property
- *Last observed generation*, in the `status.observedGeneration` property

The `status` property also provides resource-specific information. For example:

- `KafkaConnectStatus` provides the REST API endpoint for Kafka Connect connectors.
- `KafkaUserStatus` provides the user name of the Kafka user and the `Secret` in which their credentials are stored.
- `KafkaBridgeStatus` provides the HTTP address at which external client applications can access the Bridge service.

A resource's *current state* is useful for tracking progress related to the resource achieving its *desired state*, as defined by the `spec` property. The status conditions provide the time and reason the state of the resource changed and details of events preventing or delaying the operator from realizing the resource's desired state.

The *last observed generation* is the generation of the resource that was last reconciled by the Cluster Operator. If the value of `observedGeneration` is different from the value of `metadata.generation`, the operator has not yet processed the latest update to the resource. If these values are the same, the status information reflects the most recent changes to the resource.

Strimzi creates and maintains the status of custom resources, periodically evaluating the current state of the custom resource and updating its status accordingly. When performing an update on a custom resource using `kubectl edit`, for example, its `status` is not editable. Moreover, changing the `status` would not affect the configuration of the Kafka cluster.

Here we see the `status` property specified for a Kafka custom resource.

## Kafka custom resource with status

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
spec:
# ...
status:
  conditions: ①
  - lastTransitionTime: 2019-07-23T23:46:57+0000
    status: "True"
    type: Ready ②
  observedGeneration: 4 ③
  listeners: ④
  - addresses:
    - host: my-cluster-kafka-bootstrap.myproject.svc
      port: 9092
      type: plain
  - addresses:
    - host: my-cluster-kafka-bootstrap.myproject.svc
      port: 9093
  certificates:
  - |
    -----BEGIN CERTIFICATE-----
    ...
    -----END CERTIFICATE-----
  type: tls
  - addresses:
    - host: 172.29.49.180
      port: 9094
  certificates:
  - |
    -----BEGIN CERTIFICATE-----
    ...
    -----END CERTIFICATE-----
  type: external
# ...
```

- ① Status **conditions** describe criteria related to the status that cannot be deduced from the existing resource information, or are specific to the instance of a resource.
- ② The **Ready** condition indicates whether the Cluster Operator currently considers the Kafka cluster able to handle traffic.
- ③ The **observedGeneration** indicates the generation of the **Kafka** custom resource that was last reconciled by the Cluster Operator.
- ④ The **listeners** describe the current Kafka bootstrap addresses by type.

**IMPORTANT**

The address in the custom resource status for external listeners with type **nodeport** is currently not supported.

**NOTE**

The Kafka bootstrap addresses listed in the status do not signify that those endpoints or the Kafka cluster is in a ready state.

#### *Accessing status information*

You can access status information for a resource from the command line. For more information, see [Finding the status of a custom resource](#).

### **10.1.3. Finding the status of a custom resource**

This procedure describes how to find the status of a custom resource.

#### *Prerequisites*

- A Kubernetes cluster.
- The Cluster Operator is running.

#### *Procedure*

- Specify the custom resource and use the `-o jsonpath` option to apply a standard JSONPath expression to select the `status` property:

```
kubectl get kafka <kafka_resource_name> -o jsonpath='{.status}'
```

This expression returns all the status information for the specified custom resource. You can use dot notation, such as `status.listeners` or `status.observedGeneration`, to fine-tune the status information you wish to see.

#### *Additional resources*

- [Strimzi custom resource status information](#)
- For more information about using JSONPath, see [JSONPath support](#).

## **10.2. Manually starting rolling updates of Kafka and ZooKeeper clusters**

Strimzi supports the use of annotations on `StatefulSet` and `Pod` resources to manually trigger a rolling update of Kafka and ZooKeeper clusters through the Cluster Operator. Rolling updates restart the pods of the resource with new ones.

Manually performing a rolling update on a specific pod or set of pods from the same `StatefulSet` is usually only required in exceptional circumstances. However, rather than deleting the pods directly, if you perform the rolling update through the Cluster Operator you ensure that:

- The manual deletion of the pod does not conflict with simultaneous Cluster Operator operations, such as deleting other pods in parallel.
- The Cluster Operator logic handles the Kafka configuration specifications, such as the number of in-sync replicas.

## 10.2.1. Prerequisites

To perform a manual rolling update, you need a running Cluster Operator and Kafka cluster.

See the *Deploying and Upgrading Strimzi* guide for instructions on running a:

- [Cluster Operator](#)
- [Kafka cluster](#)

## 10.2.2. Performing a rolling update using a StatefulSet annotation

This procedure describes how to manually trigger a rolling update of an existing Kafka cluster or ZooKeeper cluster using a Kubernetes **StatefulSet** annotation.

*Procedure*

1. Find the name of the **StatefulSet** that controls the Kafka or ZooKeeper pods you want to manually update.

For example, if your Kafka cluster is named *my-cluster*, the corresponding **StatefulSet** names are *my-cluster-kafka* and *my-cluster-zookeeper*.

2. Annotate the **StatefulSet** resource in Kubernetes.

Use `kubectl annotate`:

```
kubectl annotate statefulset cluster-name-kafka strimzi.io/manual-rolling-
update=true

kubectl annotate statefulset cluster-name-zookeeper strimzi.io/manual-rolling-
update=true
```

3. Wait for the next reconciliation to occur (every two minutes by default). A rolling update of all pods within the annotated **StatefulSet** is triggered, as long as the annotation was detected by the reconciliation process. When the rolling update of all the pods is complete, the annotation is removed from the **StatefulSet**.

## 10.2.3. Performing a rolling update using a Pod annotation

This procedure describes how to manually trigger a rolling update of an existing Kafka cluster or ZooKeeper cluster using a Kubernetes **Pod** annotation. When multiple pods from the same **StatefulSet** are annotated, consecutive rolling updates are performed within the same reconciliation run.

*Procedure*

1. Find the name of the Kafka or ZooKeeper **Pod** you want to manually update.

For example, if your Kafka cluster is named *my-cluster*, the corresponding **Pod** names are *my-cluster-kafka-index* and *my-cluster-zookeeper-index*. The *index* starts at zero and ends at the total number of replicas.

## 2. Annotate the **Pod** resource in Kubernetes.

Use `kubectl annotate`:

```
kubectl annotate pod cluster-name-kafka-index strimzi.io/manual-rolling-update=true
```

```
kubectl annotate pod cluster-name-zookeeper-index strimzi.io/manual-rolling-update=true
```

## 3. Wait for the next reconciliation to occur (every two minutes by default). A rolling update of the annotated **Pod** is triggered, as long as the annotation was detected by the reconciliation process. When the rolling update of a pod is complete, the annotation is removed from the **Pod**.

## 10.3. Discovering services using labels and annotations

Service discovery makes it easier for client applications running in the same Kubernetes cluster as Strimzi to interact with a Kafka cluster.

A *service discovery* label and annotation is generated for services used to access the Kafka cluster:

- Internal Kafka bootstrap service
- HTTP Bridge service

The label helps to make the service discoverable, and the annotation provides connection details that a client application can use to make the connection.

The service discovery label, `strimzi.io/discovery`, is set as `true` for the **Service** resources. The service discovery annotation has the same key, providing connection details in JSON format for each service.

### Example internal Kafka bootstrap service

```

apiVersion: v1
kind: Service
metadata:
  annotations:
    strimzi.io/discovery: |-
      [ {
        "port" : 9092,
        "tls" : false,
        "protocol" : "kafka",
        "auth" : "scram-sha-512"
      }, {
        "port" : 9093,
        "tls" : true,
        "protocol" : "kafka",
        "auth" : "tls"
      } ]
  labels:
    strimzi.io/cluster: my-cluster
    strimzi.io/discovery: "true"
    strimzi.io/kind: Kafka
    strimzi.io/name: my-cluster-kafka-bootstrap
  name: my-cluster-kafka-bootstrap
spec:
  #...

```

## Example HTTP Bridge service

```

apiVersion: v1
kind: Service
metadata:
  annotations:
    strimzi.io/discovery: |-
      [ {
        "port" : 8080,
        "tls" : false,
        "auth" : "none",
        "protocol" : "http"
      } ]
  labels:
    strimzi.io/cluster: my-bridge
    strimzi.io/discovery: "true"
    strimzi.io/kind: KafkaBridge
    strimzi.io/name: my-bridge-bridge-service

```

### 10.3.1. Returning connection details on services

You can find the services by specifying the discovery label when fetching services from the command line or a corresponding API call.

```
kubectl get service -l strimzi.io/discovery=true
```

The connection details are returned when retrieving the service discovery label.

## 10.4. Recovering a cluster from persistent volumes

You can recover a Kafka cluster from persistent volumes (PVs) if they are still present.

You might want to do this, for example, after:

- A namespace was deleted unintentionally
- A whole Kubernetes cluster is lost, but the PVs remain in the infrastructure

### 10.4.1. Recovery from namespace deletion

Recovery from namespace deletion is possible because of the relationship between persistent volumes and namespaces. A **PersistentVolume** (PV) is a storage resource that lives outside of a namespace. A PV is mounted into a Kafka pod using a **PersistentVolumeClaim** (PVC), which lives inside a namespace.

The reclaim policy for a PV tells a cluster how to act when a namespace is deleted. If the reclaim policy is set as:

- *Delete* (default), PVs are deleted when PVCs are deleted within a namespace
- *Retain*, PVs are not deleted when a namespace is deleted

To ensure that you can recover from a PV if a namespace is deleted unintentionally, the policy must be reset from *Delete* to *Retain* in the PV specification using the **persistentVolumeReclaimPolicy** property:

```
apiVersion: v1
kind: PersistentVolume
# ...
spec:
  # ...
  persistentVolumeReclaimPolicy: Retain
```

Alternatively, PVs can inherit the reclaim policy of an associated storage class. Storage classes are used for dynamic volume allocation.

By configuring the **reclaimPolicy** property for the storage class, PVs that use the storage class are created with the appropriate reclaim policy. The storage class is configured for the PV using the **storageClassName** property.

```
apiVersion: v1
kind: StorageClass
metadata:
  name: gp2-retain
parameters:
  # ...
# ...
reclaimPolicy: Retain
```

```
apiVersion: v1
kind: PersistentVolume
# ...
spec:
  # ...
storageClassName: gp2-retain
```

**NOTE** If you are using *Retain* as the reclaim policy, but you want to delete an entire cluster, you need to delete the PVs manually. Otherwise they will not be deleted, and may cause unnecessary expenditure on resources.

## 10.4.2. Recovery from loss of a Kubernetes cluster

When a cluster is lost, you can use the data from disks/volumes to recover the cluster if they were preserved within the infrastructure. The recovery procedure is the same as with namespace deletion, assuming PVs can be recovered and they were created manually.

## 10.4.3. Recovering a deleted cluster from persistent volumes

This procedure describes how to recover a deleted cluster from persistent volumes (PVs).

In this situation, the Topic Operator identifies that topics exist in Kafka, but the [KafkaTopic](#) resources do not exist.

When you get to the step to recreate your cluster, you have two options:

1. Use *Option 1* when you can recover all [KafkaTopic](#) resources.

The [KafkaTopic](#) resources must therefore be recovered before the cluster is started so that the corresponding topics are not deleted by the Topic Operator.

2. Use *Option 2* when you are unable to recover all [KafkaTopic](#) resources.

This time you deploy your cluster without the Topic Operator, delete the Topic Operator data in ZooKeeper, and then redeploy it so that the Topic Operator can recreate the [KafkaTopic](#) resources from the corresponding topics.

**NOTE** If the Topic Operator is not deployed, you only need to recover the **PersistentVolumeClaim** (PVC) resources.

### Before you begin

In this procedure, it is essential that PVs are mounted into the correct PVC to avoid data corruption. A **volumeName** is specified for the PVC and this must match the name of the PV.

For more information, see:

- [Persistent Volume Claim naming](#)
- [JBOD and Persistent Volume Claims](#)

**NOTE** The procedure does not include recovery of **KafkaUser** resources, which must be recreated manually. If passwords and certificates need to be retained, secrets must be recreated before creating the **KafkaUser** resources.

### Procedure

1. Check information on the PVs in the cluster:

```
kubectl get pv
```

Information is presented for PVs with data.

Example output showing columns important to this procedure:

NAME	RECLAIMPOLICY	CLAIM
pvc-5e9c5c7f-3317-11ea-a650-06e1eadd9a4c	... Retain ...	myproject/data-my-cluster-zookeeper-1
pvc-5e9cc72d-3317-11ea-97b0-0aef8816c7ea	... Retain ...	myproject/data-my-cluster-zookeeper-0
pvc-5ead43d1-3317-11ea-97b0-0aef8816c7ea	... Retain ...	myproject/data-my-cluster-zookeeper-2
pvc-7e1f67f9-3317-11ea-a650-06e1eadd9a4c	... Retain ...	myproject/data-0-my-cluster-kafka-0
pvc-7e21042e-3317-11ea-9786-02deaf9aa87e	... Retain ...	myproject/data-0-my-cluster-kafka-1
pvc-7e226978-3317-11ea-97b0-0aef8816c7ea	... Retain ...	myproject/data-0-my-cluster-kafka-2

- *NAME* shows the name of each PV.
- *RECLAIM POLICY* shows that PVs are *retained*.
- *CLAIM* shows the link to the original PVCs.

2. Recreate the original namespace:

```
kubectl create namespace myproject
```

3. Recreate the original PVC resource specifications, linking the PVCs to the appropriate PV:

For example:

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: data-0-my-cluster-kafka-0
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 100Gi
  storageClassName: gp2-retain
  volumeMode: Filesystem
  volumeName: pvc-7e1f67f9-3317-11ea-a650-06e1eadd9a4c
```

4. Edit the PV specifications to delete the **claimRef** properties that bound the original PVC.

For example:

```

apiVersion: v1
kind: PersistentVolume
metadata:
  annotations:
    kubernetes.io/createdby: aws-ebs-dynamic-provisioner
    pv.kubernetes.io/bound-by-controller: "yes"
    pv.kubernetes.io/provisioned-by: kubernetes.io/aws-ebs
  creationTimestamp: "<date>"
  finalizers:
  - kubernetes.io/pv-protection
  labels:
    failure-domain.beta.kubernetes.io/region: eu-west-1
    failure-domain.beta.kubernetes.io/zone: eu-west-1c
  name: pvc-7e226978-3317-11ea-97b0-0aef8816c7ea
  resourceVersion: "39431"
  selfLink: /api/v1/persistentvolumes/pvc-7e226978-3317-11ea-97b0-0aef8816c7ea
  uid: 7efe6b0d-3317-11ea-a650-06e1eadd9a4c
spec:
  accessModes:
  - ReadWriteOnce
  awsElasticBlockStore:
    fsType: xfs
    volumeID: aws://eu-west-1c/vol-09db3141656d1c258
  capacity:
    storage: 100Gi
  claimRef:
    apiVersion: v1
    kind: PersistentVolumeClaim
    name: data-0-my-cluster-kafka-2
    namespace: myproject
    resourceVersion: "39113"
    uid: 54be1c60-3319-11ea-97b0-0aef8816c7ea
  nodeAffinity:
    required:
      nodeSelectorTerms:
      - matchExpressions:
        - key: failure-domain.beta.kubernetes.io/zone
          operator: In
          values:
          - eu-west-1c
        - key: failure-domain.beta.kubernetes.io/region
          operator: In
          values:
          - eu-west-1
  persistentVolumeReclaimPolicy: Retain
  storageClassName: gp2-retain
  volumeMode: Filesystem

```

In the example, the following properties are deleted:

```
claimRef:  
  apiVersion: v1  
  kind: PersistentVolumeClaim  
  name: data-0-my-cluster-kafka-2  
  namespace: myproject  
  resourceVersion: "39113"  
  uid: 54be1c60-3319-11ea-97b0-0aef8816c7ea
```

## 5. Deploy the Cluster Operator.

```
kubectl apply -f install/cluster-operator -n my-project
```

## 6. Recreate your cluster.

Follow the steps depending on whether or not you have all the **KafkaTopic** resources needed to recreate your cluster.

**Option 1:** If you have **all** the **KafkaTopic** resources that existed before you lost your cluster, including internal topics such as committed offsets from **\_consumer\_offsets**:

### 1. Recreate all **KafkaTopic** resources.

It is essential that you recreate the resources before deploying the cluster, or the Topic Operator will delete the topics.

### 2. Deploy the Kafka cluster.

For example:

```
kubectl apply -f kafka.yaml
```

**Option 2:** If you do not have all the **KafkaTopic** resources that existed before you lost your cluster:

### 1. Deploy the Kafka cluster, as with the first option, but without the Topic Operator by removing the **topicOperator** property from the Kafka resource before deploying.

If you include the Topic Operator in the deployment, the Topic Operator will delete all the topics.

### 2. Run an **exec** command to one of the Kafka broker pods to open the ZooKeeper shell script.

For example, where *my-cluster-kafka-0* is the name of the broker pod:

```
kubectl exec -ti my-cluster-zookeeper-0 -- bin/zookeeper-shell.sh  
localhost:12181
```

3. Delete the whole `/strimzi` path to remove the Topic Operator storage:

```
deleteall /strimzi
```

4. Enable the Topic Operator by redeploying the Kafka cluster with the `topicOperator` property to recreate the `KafkaTopic` resources.

For example:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  #...
  entityOperator:
    topicOperator: {} ①
  #...
```

① Here we show the default configuration, which has no additional properties. You specify the required configuration using the properties described in [EntityTopicOperatorSpec schema reference](#).

7. Verify the recovery by listing the `KafkaTopic` resources:

```
kubectl get KafkaTopic
```

## 10.5. Tuning client configuration

Use configuration properties to optimize the performance of Kafka producers and consumers.

A minimum set of configuration properties is required, but you can add or adjust properties to change how producers and consumers interact with Kafka. For example, for producers you can tune latency and throughput of messages so that clients can respond to data in real time. Or you can change the configuration to provide stronger message durability guarantees.

You might start by analyzing client metrics to gauge where to make your initial configurations, then make incremental changes and further comparisons until you have the configuration you need.

### 10.5.1. Kafka producer configuration tuning

Use a basic producer configuration with optional properties that are tailored to specific use cases.

Adjusting your configuration to maximize throughput might increase latency or vice versa. You will need to experiment and tune your producer configuration to get the balance you need.

## Basic producer configuration

Connection and serializer properties are required for every producer. Generally, it is good practice to add a client id for tracking, and use compression on the producer to reduce batch sizes in requests.

In a basic producer configuration:

- The order of messages in a partition is not guaranteed.
- The acknowledgment of messages reaching the broker does not guarantee durability.

```
# ...
bootstrap.servers=localhost:9092 ①
key.serializer=org.apache.kafka.common.serialization.StringSerializer ②
value.serializer=org.apache.kafka.common.serialization.StringSerializer ③
client.id=my-client ④
compression.type=gzip ⑤
# ...
```

- ① (Required) Tells the producer to connect to a Kafka cluster using a *host:port* bootstrap server address for a Kafka broker. The producer uses the address to discover and connect to all brokers in the cluster. Use a comma-separated list to specify two or three addresses in case a server is down, but it's not necessary to provide a list of all the brokers in the cluster.
- ② (Required) Serializer to transform the key of each message to bytes prior to them being sent to a broker.
- ③ (Required) Serializer to transform the value of each message to bytes prior to them being sent to a broker.
- ④ (Optional) The logical name for the client, which is used in logs and metrics to identify the source of a request.
- ⑤ (Optional) The codec for compressing messages, which are sent and might be stored in compressed format and then decompressed when reaching a consumer. Compression is useful for improving throughput and reducing the load on storage, but might not be suitable for low latency applications where the cost of compression or decompression could be prohibitive.

## Data durability

You can apply greater data durability, to minimize the likelihood that messages are lost, using message delivery acknowledgments.

```
# ...
acks=all ①
# ...
```

- ① Specifying `acks=all` forces a partition leader to replicate messages to a certain number of followers before acknowledging that the message request was successfully received. Because of the additional checks, `acks=all` increases the latency between the producer sending a message

and receiving acknowledgment.

The number of brokers which need to have appended the messages to their logs before the acknowledgment is sent to the producer is determined by the topic's `min.insync.replicas` configuration. A typical starting point is to have a topic replication factor of 3, with two in-sync replicas on other brokers. In this configuration, the producer can continue unaffected if a single broker is unavailable. If a second broker becomes unavailable, the producer won't receive acknowledgments and won't be able to produce more messages.

*Topic configuration to support `acks=all`*

```
# ...
min.insync.replicas=2 ①
# ...
```

① Use 2 in-sync replicas. The default is 1.

**NOTE** If the system fails, there is a risk of unsent data in the buffer being lost.

## Ordered delivery

Idempotent producers avoid duplicates as messages are delivered exactly once. IDs and sequence numbers are assigned to messages to ensure the order of delivery, even in the event of failure. If you are using `acks=all` for data consistency, enabling idempotency makes sense for ordered delivery.

*Ordered delivery with idempotency*

```
# ...
enable.idempotence=true ①
max.in.flight.requests.per.connection=5 ②
acks=all ③
retries=2147483647 ④
# ...
```

① Set to `true` to enable the idempotent producer.

② With idempotent delivery the number of in-flight requests may be greater than 1 while still providing the message ordering guarantee. The default is 5 in-flight requests.

③ Set `acks` to `all`.

④ Set the number of attempts to resend a failed message request.

If you are not using `acks=all` and idempotency because of the performance cost, set the number of in-flight (unacknowledged) requests to 1 to preserve ordering. Otherwise, a situation is possible where *Message-A* fails only to succeed after *Message-B* was already written to the broker.

## Ordered delivery without idempotency

```
# ...
enable.idempotence=false ①
max.in.flight.requests.per.connection=1 ②
retries=2147483647
# ...
```

① Set to `false` to disable the idempotent producer.

② Set the number of in-flight requests to exactly `1`.

## Reliability guarantees

Idempotence is useful for exactly once writes to a single partition. Transactions, when used with idempotence, allow exactly once writes across multiple partitions.

Transactions guarantee that messages using the same transactional ID are produced once, and either *all* are successfully written to the respective logs or *none* of them are.

```
# ...
enable.idempotence=true
max.in.flight.requests.per.connection=5
acks=all
retries=2147483647
transactional.id=UNIQUE-ID ①
transaction.timeout.ms=900000 ②
# ...
```

① Specify a unique transactional ID.

② Set the maximum allowed time for transactions in milliseconds before a timeout error is returned. The default is `900000` or 15 minutes.

The choice of `transactional.id` is important in order that the transactional guarantee is maintained. Each transactional id should be used for a unique set of topic partitions. For example, this can be achieved using an external mapping of topic partition names to transactional ids, or by computing the transactional id from the topic partition names using a function that avoids collisions.

## Optimizing throughput and latency

Usually, the requirement of a system is to satisfy a particular throughput target for a proportion of messages within a given latency. For example, targeting 500,000 messages per second with 95% of messages being acknowledged within 2 seconds.

It's likely that the messaging semantics (message ordering and durability) of your producer are defined by the requirements for your application. For instance, it's possible that you don't have the option of using `acks=0` or `acks=1` without breaking some important property or guarantee provided by your application.

Broker restarts have a significant impact on high percentile statistics. For example, over a long

period the 99th percentile latency is dominated by behavior around broker restarts. This is worth considering when designing benchmarks or comparing performance numbers from benchmarking with performance numbers seen in production.

Depending on your objective, Kafka offers a number of configuration parameters and techniques for tuning producer performance for throughput and latency.

### Message batching (`linger.ms` and `batch.size`)

Message batching delays sending messages in the hope that more messages destined for the same broker will be sent, allowing them to be batched into a single produce request. Batching is a compromise between higher latency in return for higher throughput. Time-based batching is configured using `linger.ms`, and size-based batching is configured using `batch.size`.

### Compression (`compression.type`)

Message compression adds latency in the producer (CPU time spent compressing the messages), but makes requests (and potentially disk writes) smaller, which can increase throughput. Whether compression is worthwhile, and the best compression to use, will depend on the messages being sent. Compression happens on the thread which calls `KafkaProducer.send()`, so if the latency of this method matters for your application you should consider using more threads.

### Pipelining (`max.in.flight.requests.per.connection`)

Pipelining means sending more requests before the response to a previous request has been received. In general more pipelining means better throughput, up to a threshold at which other effects, such as worse batching, start to counteract the effect on throughput.

#### *Lowering latency*

When your application calls `KafkaProducer.send()` the messages are:

- Processed by any interceptors
- Serialized
- Assigned to a partition
- Compressed
- Added to a batch of messages in a per-partition queue

At which point the `send()` method returns. So the time `send()` is blocked is determined by:

- The time spent in the interceptors, serializers and partitioner
- The compression algorithm used
- The time spent waiting for a buffer to use for compression

Batches will remain in the queue until one of the following occurs:

- The batch is full (according to `batch.size`)
- The delay introduced by `linger.ms` has passed
- The sender is about to send message batches for other partitions to the same broker, and it is possible to add this batch too

- The producer is being flushed or closed

Look at the configuration for batching and buffering to mitigate the impact of `send()` blocking on latency.

```
# ...
linger.ms=100 ①
batch.size=16384 ②
buffer.memory=33554432 ③
# ...
```

- ① The `linger` property adds a delay in milliseconds so that larger batches of messages are accumulated and sent in a request. The default is `0`.
- ② If a maximum `batch.size` in bytes is used, a request is sent when the maximum is reached, or messages have been queued for longer than `linger.ms` (whichever comes sooner). Adding the delay allows batches to accumulate messages up to the batch size.
- ③ The buffer size must be at least as big as the batch size, and be able to accommodate buffering, compression and in-flight requests.

#### *Increasing throughput*

Improve throughput of your message requests by adjusting the maximum time to wait before a message is delivered and completes a send request.

You can also direct messages to a specified partition by writing a custom partitioner to replace the default.

```
# ...
delivery.timeout.ms=120000 ①
partitioner.class=my-custom-partitioner ②
# ...
```

- ① The maximum time in milliseconds to wait for a complete send request. You can set the value to `MAX_LONG` to delegate to Kafka an indefinite number of retries. The default is `120000` or 2 minutes.
- ② Specify the class name of the custom partitioner.

### 10.5.2. Kafka consumer configuration tuning

Use a basic consumer configuration with optional properties that are tailored to specific use cases.

When tuning your consumers your primary concern will be ensuring that they cope efficiently with the amount of data ingested. As with the producer tuning, be prepared to make incremental changes until the consumers operate as expected.

#### Basic consumer configuration

Connection and deserializer properties are required for every consumer. Generally, it is good

practice to add a client id for tracking.

In a consumer configuration, irrespective of any subsequent configuration:

- The consumer fetches from a given offset and consumes the messages in order, unless the offset is changed to skip or re-read messages.
- The broker does not know if the consumer processed the responses, even when committing offsets to Kafka, because the offsets might be sent to a different broker in the cluster.

```
# ...
bootstrap.servers=localhost:9092 ①
key.deserializer=org.apache.kafka.common.serialization.StringDeserializer ②
value.deserializer=org.apache.kafka.common.serialization.StringDeserializer ③
client.id=my-client ④
group.id=my-group-id ⑤
# ...
```

① (Required) Tells the consumer to connect to a Kafka cluster using a *host:port* bootstrap server address for a Kafka broker. The consumer uses the address to discover and connect to all brokers in the cluster. Use a comma-separated list to specify two or three addresses in case a server is down, but it is not necessary to provide a list of all the brokers in the cluster. If you are using a loadbalancer service to expose the Kafka cluster, you only need the address for the service because the availability is handled by the loadbalancer.

② (Required) Deserializer to transform the bytes fetched from the Kafka broker into message keys.

③ (Required) Deserializer to transform the bytes fetched from the Kafka broker into message values.

④ (Optional) The logical name for the client, which is used in logs and metrics to identify the source of a request. The id can also be used to throttle consumers based on processing time quotas.

⑤ (Conditional) A group id is *required* for a consumer to be able to join a consumer group.

Consumer groups are used to share a typically large data stream generated by multiple producers from a given topic. Consumers are grouped using a `group.id`, allowing messages to be spread across the members.

## Scaling data consumption using consumer groups

Consumer groups share a typically large data stream generated by one or multiple producers from a given topic. Consumers with the same `group.id` property are in the same group. One of the consumers in the group is elected leader and decides how the partitions are assigned to the consumers in the group. Each partition can only be assigned to a single consumer.

If you do not already have as many consumers as partitions, you can scale data consumption by adding more consumer instances with the same `group.id`. Adding more consumers to a group than there are partitions will not help throughput, but it does mean that there are consumers on standby should one stop functioning. If you can meet throughput goals with fewer consumers, you save on resources.

Consumers within the same consumer group send offset commits and heartbeats to the same broker. So the greater the number of consumers in the group, the higher the request load on the broker.

```
# ...
group.id=my-group-id ①
# ...
```

① Add a consumer to a consumer group using a group id.

## Message ordering guarantees

Kafka brokers receive fetch requests from consumers that ask the broker to send messages from a list of topics, partitions and offset positions.

A consumer observes messages in a single partition in the same order that they were committed to the broker, which means that Kafka **only** provides ordering guarantees for messages in a single partition. Conversely, if a consumer is consuming messages from multiple partitions, the order of messages in different partitions as observed by the consumer does not necessarily reflect the order in which they were sent.

If you want a strict ordering of messages from one topic, use one partition per consumer.

## Optimizing throughput and latency

Control the number of messages returned when your client application calls `KafkaConsumer.poll()`.

Use the `fetch.max.wait.ms` and `fetch.min.bytes` properties to increase the minimum amount of data fetched by the consumer from the Kafka broker. Time-based batching is configured using `fetch.max.wait.ms`, and size-based batching is configured using `fetch.min.bytes`.

If CPU utilization in the consumer or broker is high, it might be because there are too many requests from the consumer. You can adjust `fetch.max.wait.ms` and `fetch.min.bytes` properties higher so that there are fewer requests and messages are delivered in bigger batches. By adjusting higher, throughput is improved with some cost to latency. You can also adjust higher if the amount of data being produced is low.

For example, if you set `fetch.max.wait.ms` to 500ms and `fetch.min.bytes` to 16384 bytes, when Kafka receives a fetch request from the consumer it will respond when the first of either threshold is reached.

Conversely, you can adjust the `fetch.max.wait.ms` and `fetch.min.bytes` properties lower to improve end-to-end latency.

```
# ...
fetch.max.wait.ms=500 ①
fetch.min.bytes=16384 ②
# ...
```

- ① The maximum time in milliseconds the broker will wait before completing fetch requests. The default is **500** milliseconds.
- ② If a minimum batch size in bytes is used, a request is sent when the minimum is reached, or messages have been queued for longer than **fetch.max.wait.ms** (whichever comes sooner). Adding the delay allows batches to accumulate messages up to the batch size.

#### *Lowering latency by increasing the fetch request size*

Use the **fetch.max.bytes** and **max.partition.fetch.bytes** properties to increase the maximum amount of data fetched by the consumer from the Kafka broker.

The **fetch.max.bytes** property sets a maximum limit in bytes on the amount of data fetched from the broker at one time.

The **max.partition.fetch.bytes** sets a maximum limit in bytes on how much data is returned for each partition, which must always be larger than the number of bytes set in the broker or topic configuration for **max.message.bytes**.

The maximum amount of memory a client can consume is calculated approximately as:

```
NUMBER-OF-BROKERS * fetch.max.bytes and NUMBER-OF-PARTITIONS *
max.partition.fetch.bytes
```

If memory usage can accommodate it, you can increase the values of these two properties. By allowing more data in each request, latency is improved as there are fewer fetch requests.

```
# ...
fetch.max.bytes=52428800 ①
max.partition.fetch.bytes=1048576 ②
# ...
```

① The maximum amount of data in bytes returned for a fetch request.

② The maximum amount of data in bytes returned for each partition.

#### **Avoiding data loss or duplication when committing offsets**

The Kafka *auto-commit mechanism* allows a consumer to commit the offsets of messages automatically. If enabled, the consumer will commit offsets received from polling the broker at 5000ms intervals.

The auto-commit mechanism is convenient, but it introduces a risk of data loss and duplication. If a consumer has fetched and transformed a number of messages, but the system crashes with processed messages in the consumer buffer when performing an auto-commit, that data is lost. If the system crashes after processing the messages, but before performing the auto-commit, the data is duplicated on another consumer instance after rebalancing.

Auto-committing can avoid data loss only when all messages are processed before the next poll to the broker, or the consumer closes.

To minimize the likelihood of data loss or duplication, you can set `enable.auto.commit` to `false` and develop your client application to have more control over committing offsets. Or you can use `auto.commit.interval.ms` to decrease the intervals between commits.

```
# ...
enable.auto.commit=false ①
# ...
```

① Auto commit is set to false to provide more control over committing offsets.

By setting to `enable.auto.commit` to `false`, you can commit offsets after **all** processing has been performed and the message has been consumed. For example, you can set up your application to call the Kafka `commitSync` and `commitAsync` commit APIs.

The `commitSync` API commits the offsets in a message batch returned from polling. You call the API when you are finished processing all the messages in the batch. If you use the `commitSync` API, the application will not poll for new messages until the last offset in the batch is committed. If this negatively affects throughput, you can commit less frequently, or you can use the `commitAsync` API. The `commitAsync` API does not wait for the broker to respond to a commit request, but risks creating more duplicates when rebalancing. A common approach is to combine both commit APIs in an application, with the `commitSync` API used just before shutting the consumer down or rebalancing to make sure the final commit is successful.

### Controlling transactional messages

Consider using transactional ids and enabling idempotence (`enable.idempotence=true`) on the producer side to guarantee exactly-once delivery. On the consumer side, you can then use the `isolation.level` property to control how transactional messages are read by the consumer.

The `isolation.level` property has two valid values:

- `read_committed`
- `read_uncommitted` (default)

Use `read_committed` to ensure that only transactional messages that have been committed are read by the consumer. However, this will cause an increase in end-to-end latency, because the consumer will not be able to return a message until the brokers have written the transaction markers that record the result of the transaction (*committed* or *aborted*).

```
# ...
enable.auto.commit=false
isolation.level=read_committed ①
# ...
```

① Set to `read_committed` so that only committed messages are read by the consumer.

## Recovering from failure to avoid data loss

Use the `session.timeout.ms` and `heartbeat.interval.ms` properties to configure the time taken to check and recover from consumer failure within a consumer group.

The `session.timeout.ms` property specifies the maximum amount of time in milliseconds a consumer within a consumer group can be out of contact with a broker before being considered inactive and a *rebalancing* is triggered between the active consumers in the group. When the group rebalances, the partitions are reassigned to the members of the group.

The `heartbeat.interval.ms` property specifies the interval in milliseconds between *heartbeat* checks to the consumer group coordinator to indicate that the consumer is active and connected. The heartbeat interval must be lower, usually by a third, than the session timeout interval.

If you set the `session.timeout.ms` property lower, failing consumers are detected earlier, and rebalancing can take place quicker. However, take care not to set the timeout so low that the broker fails to receive a heartbeat in time and triggers an unnecessary rebalance.

Decreasing the heartbeat interval reduces the chance of accidental rebalancing, but more frequent heartbeats increases the overhead on broker resources.

## Managing offset policy

Use the `auto.offset.reset` property to control how a consumer behaves when no offsets have been committed, or a committed offset is no longer valid or deleted.

Suppose you deploy a consumer application for the first time, and it reads messages from an existing topic. Because this is the first time the `group.id` is used, the `__consumer_offsets` topic does not contain any offset information for this application. The new application can start processing all existing messages from the start of the log or only new messages. The default reset value is `latest`, which starts at the end of the partition, and consequently means some messages are missed. To avoid data loss, but increase the amount of processing, set `auto.offset.reset` to `earliest` to start at the beginning of the partition.

Also consider using the `earliest` option to avoid messages being lost when the offsets retention period (`offsets.retention.minutes`) configured for a broker has ended. If a consumer group or standalone consumer is inactive and commits no offsets during the retention period, previously committed offsets are deleted from `__consumer_offsets`.

```
# ...
heartbeat.interval.ms=3000 ①
session.timeout.ms=10000 ②
auto.offset.reset=earliest ③
# ...
```

① Adjust the heartbeat interval lower according to anticipated rebalances.

② If no heartbeats are received by the Kafka broker before the timeout duration expires, the consumer is removed from the consumer group and a rebalance is initiated. If the broker configuration has a `group.min.session.timeout.ms` and `group.max.session.timeout.ms`, the session

timeout value must be within that range.

- ③ Set to `earliest` to return to the start of a partition and avoid data loss if offsets were not committed.

If the amount of data returned in a single fetch request is large, a timeout might occur before the consumer has processed it. In this case, you can lower `max.partition.fetch.bytes` or increase `session.timeout.ms`.

## Minimizing the impact of rebalances

The rebalancing of a partition between active consumers in a group is the time it takes for:

- Consumers to commit their offsets
- The new consumer group to be formed
- The group leader to assign partitions to group members
- The consumers in the group to receive their assignments and start fetching

Clearly, the process increases the downtime of a service, particularly when it happens repeatedly during a rolling restart of a consumer group cluster.

In this situation, you can use the concept of *static membership* to reduce the number of rebalances. Rebalancing assigns topic partitions evenly among consumer group members. Static membership uses persistence so that a consumer instance is recognized during a restart after a session timeout.

The consumer group coordinator can identify a new consumer instance using a unique id that is specified using the `group.instance.id` property. During a restart, the consumer is assigned a new member id, but as a static member it continues with the same instance id, and the same assignment of topic partitions is made.

If the consumer application does not make a call to poll at least every `max.poll.interval.ms` milliseconds, the consumer is considered to be failed, causing a rebalance. If the application cannot process all the records returned from poll in time, you can avoid a rebalance by using the `max.poll.interval.ms` property to specify the interval in milliseconds between polls for new messages from a consumer. Or you can use the `max.poll.records` property to set a maximum limit on the number of records returned from the consumer buffer, allowing your application to process fewer records within the `max.poll.interval.ms` limit.

```
# ...
group.instance.id=UNIQUE-ID ①
max.poll.interval.ms=300000 ②
max.poll.records=500 ③
# ...
```

- ① The unique instance id ensures that a new consumer instance receives the same assignment of topic partitions.
- ② Set the interval to check the consumer is continuing to process messages.
- ③ Sets the number of processed records returned from the consumer.

## 10.6. Uninstalling Strimzi

This procedure describes how to uninstall Strimzi and remove resources related to the deployment.

### *Prerequisites*

In order to perform this procedure, identify resources created specifically for a deployment and referenced from the Strimzi resource.

Such resources include:

- Secrets (Custom CAs and certificates, Kafka Connect secrets, and other Kafka secrets)
- Logging [ConfigMaps](#) (of type `external`)

These are resources referenced by [Kafka](#), [KafkaConnect](#), [KafkaConnectS2I](#), [KafkaMirrorMaker](#), or [KafkaBridge](#) configuration.

### *Procedure*

1. Delete the Cluster Operator [Deployment](#), related [CustomResourceDefinitions](#), and [RBAC](#) resources:

```
kubectl delete -f install/cluster-operator
```

#### **WARNING**

Deleting [CustomResourceDefinitions](#) results in the garbage collection of the corresponding custom resources ([Kafka](#), [KafkaConnect](#), [KafkaConnectS2I](#), [KafkaMirrorMaker](#), or [KafkaBridge](#)) and the resources dependent on them (Deployments, StatefulSets, and other dependent resources).

2. Delete the resources you identified in the prerequisites.

# Appendix A: Frequently asked questions

## A.1. Questions related to the Cluster Operator

### A.1.1. Why do I need cluster administrator privileges to install Strimzi?

To install Strimzi, you need to be able to create the following cluster-scoped resources:

- Custom Resource Definitions (CRDs) to instruct Kubernetes about resources that are specific to Strimzi, such as [Kafka](#) and [KafkaConnect](#)
- [ClusterRoles](#) and [ClusterRoleBindings](#)

Cluster-scoped resources, which are not scoped to a particular Kubernetes namespace, typically require *cluster administrator* privileges to install.

As a cluster administrator, you can inspect all the resources being installed (in the `/install/` directory) to ensure that the [ClusterRoles](#) do not grant unnecessary privileges.

After installation, the Cluster Operator runs as a regular [Deployment](#), so any standard (non-admin) Kubernetes user with privileges to access the [Deployment](#) can configure it. The cluster administrator can grant standard users the privileges necessary to manage [Kafka](#) custom resources.

See also:

- [Why does the Cluster Operator need to create ClusterRoleBindings?](#)
- [Can standard Kubernetes users create Kafka custom resources?](#)

### A.1.2. Why does the Cluster Operator need to create ClusterRoleBindings?

Kubernetes has built-in [privilege escalation prevention](#), which means that the Cluster Operator cannot grant privileges it does not have itself, specifically, it cannot grant such privileges in a namespace it cannot access. Therefore, the Cluster Operator must have the privileges necessary for *all* the components it orchestrates.

The Cluster Operator needs to be able to grant access so that:

- The Topic Operator can manage [KafkaTopics](#), by creating [Roles](#) and [RoleBindings](#) in the namespace that the operator runs in
- The User Operator can manage [KafkaUsers](#), by creating [Roles](#) and [RoleBindings](#) in the namespace that the operator runs in
- The failure domain of a [Node](#) is discovered by Strimzi, by creating a [ClusterRoleBinding](#)

When using rack-aware partition assignment, the broker pod needs to be able to get information about the [Node](#) it is running on, for example, the Availability Zone in Amazon AWS. A [Node](#) is a cluster-scoped resource, so access to it can only be granted through a [ClusterRoleBinding](#), not a namespace-scoped [RoleBinding](#).

### A.1.3. Can standard Kubernetes users create Kafka custom resources?

By default, standard Kubernetes users will not have the privileges necessary to manage the custom resources handled by the Cluster Operator. The cluster administrator can grant a user the necessary privileges using Kubernetes RBAC resources.

For more information, see [Designating Strimzi administrators](#) in the *Deploying and Upgrading Strimzi* guide.

### A.1.4. What do the *failed to acquire lock* warnings in the log mean?

For each cluster, the Cluster Operator executes only one operation at a time. The Cluster Operator uses locks to make sure that there are never two parallel operations running for the same cluster. Other operations must wait until the current operation completes before the lock is released.

#### INFO

Examples of cluster operations include *cluster creation*, *rolling update*, *scale down*, and *scale up*.

If the waiting time for the lock takes too long, the operation times out and the following warning message is printed to the log:

```
2018-03-04 17:09:24 WARNING AbstractClusterOperations:290 - Failed to acquire lock for
kafka cluster lock::kafka::myproject::my-cluster
```

Depending on the exact configuration of `STRIMZI_FULL_RECONCILIATION_INTERVAL_MS` and `STRIMZI_OPERATION_TIMEOUT_MS`, this warning message might appear occasionally without indicating any underlying issues. Operations that time out are picked up in the next periodic reconciliation, so that the operation can acquire the lock and execute again.

Should this message appear periodically, even in situations when there should be no other operations running for a given cluster, it might indicate that the lock was not properly released due to an error. If this is the case, try restarting the Cluster Operator.

### A.1.5. Why is hostname verification failing when connecting to NodePorts using TLS?

Currently, off-cluster access using NodePorts with TLS encryption enabled does not support TLS hostname verification. As a result, the clients that verify the hostname will fail to connect. For example, the Java client will fail with the following exception:

```
Caused by: java.security.cert.CertificateException: No subject alternative names  
matching IP address 168.72.15.231 found  
at sun.security.util.HostnameChecker.matchIP(HostnameChecker.java:168)  
at sun.security.util.HostnameChecker.match(HostnameChecker.java:94)  
at sun.security.ssl.X509TrustManagerImpl.checkIdentity(X509TrustManagerImpl.java:455)  
at sun.security.ssl.X509TrustManagerImpl.checkIdentity(X509TrustManagerImpl.java:436)  
at sun.security.ssl.X509TrustManagerImpl.checkTrusted(X509TrustManagerImpl.java:252)  
at  
sun.security.ssl.X509TrustManagerImpl.checkServerTrusted(X509TrustManagerImpl.java:136)  
)  
at sun.security.ssl.ClientHandshaker.serverCertificate(ClientHandshaker.java:1501)  
... 17 more
```

To connect, you must disable hostname verification. In the Java client, you can do this by setting the configuration option `ssl.endpoint.identification.algorithm` to an empty string.

When configuring the client using a properties file, you can do it this way:

```
ssl.endpoint.identification.algorithm=
```

When configuring the client directly in Java, set the configuration option to an empty string:

```
props.put("ssl.endpoint.identification.algorithm", "");
```

# Appendix B: Custom Resource API Reference

## B.1. Common configuration properties

Common configuration properties apply to more than one resource.

### B.1.1. `replicas`

Use the `replicas` property to configure replicas.

The type of replication depends on the resource.

- `KafkaTopic` uses a replication factor to configure the number of replicas of each partition within a Kafka cluster.
- Kafka components use replicas to configure the number of pods in a deployment to provide better availability and scalability.

**NOTE** When running a Kafka component on Kubernetes it may not be necessary to run multiple replicas for high availability. When the node where the component is deployed crashes, Kubernetes will automatically reschedule the Kafka component pod to a different node. However, running Kafka components with multiple replicas can provide faster failover times as the other nodes will be up and running.

### B.1.2. `bootstrapServers`

Use the `bootstrapServers` property to configure a list of bootstrap servers.

The bootstrap server lists can refer to Kafka clusters that are not deployed in the same Kubernetes cluster. They can also refer to a Kafka cluster not deployed by Strimzi.

If on the same Kubernetes cluster, each list must ideally contain the Kafka cluster bootstrap service which is named `CLUSTER-NAME-kafka-bootstrap` and a port number. If deployed by Strimzi but on different Kubernetes clusters, the list content depends on the approach used for exposing the clusters (routes, ingress, nodeports or loadbalancers).

When using Kafka with a Kafka cluster not managed by Strimzi, you can specify the bootstrap servers list according to the configuration of the given cluster.

### B.1.3. `ssl`

Use the three allowed `ssl` configuration options for client connection using a specific *cipher suite* for a TLS version. A cipher suite combines algorithms for secure connection and data transfer.

You can also configure the `ssl.endpoint.identification.algorithm` property to enable or disable hostname verification.

### *Example SSL configuration*

```
# ...
spec:
  config:
    ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384" ①
    ssl.enabled.protocols: "TLSv1.2" ②
    ssl.protocol: "TLSv1.2" ③
    ssl.endpoint.identification.algorithm: HTTPS ④
# ...
```

- ① The cipher suite for TLS using a combination of `ECDHE` key exchange mechanism, `RSA` authentication algorithm, `AES` bulk encryption algorithm and `SHA384` MAC algorithm.
- ② The SSL protocol `TLSv1.2` is enabled.
- ③ Specifies the `TLSv1.2` protocol to generate the SSL context. Allowed values are `TLSv1.1` and `TLSv1.2`.
- ④ Hostname verification is enabled by setting to `HTTPS`. An empty string disables the verification.

### **B.1.4. `trustedCertificates`**

Having set `tls` to configure TLS encryption, use the `trustedCertificates` property to provide a list of secrets with key names under which the certificates are stored in X.509 format.

You can use the secrets created by the Cluster Operator for the Kafka cluster, or you can create your own TLS certificate file, then create a `Secret` from the file:

```
kubectl create secret generic MY-SECRET \
--from-file=MY-TLS-CERTIFICATE-FILE.crt
```

### *Example TLS encryption configuration*

```
tls:
  trustedCertificates:
    - secretName: my-cluster-cluster-cert
      certificate: ca.crt
    - secretName: my-cluster-cluster-cert
      certificate: ca2.crt
```

If certificates are stored in the same secret, it can be listed multiple times.

If you want to enable TLS, but use the default set of public certification authorities shipped with Java, you can specify `trustedCertificates` as an empty array:

### *Example of enabling TLS with the default Java certificates*

```
tls:
  trustedCertificates: []
```

For information on configuring TLS client authentication, see [KafkaClientAuthenticationTls schema reference](#).

### B.1.5. resources

You request CPU and memory resources for components. Limits specify the maximum resources that can be consumed by a given container.

Resource requests and limits for the Topic Operator and User Operator are set in the [Kafka](#) resource.

Use the `resources.requests` and `resources.limits` properties to configure resource requests and limits.

For every deployed container, Strimzi allows you to request specific resources and define the maximum consumption of those resources.

Strimzi supports requests and limits for the following types of resources:

- `cpu`
- `memory`

Strimzi uses the Kubernetes syntax for specifying these resources.

For more information about managing computing resources on Kubernetes, see [Managing Compute Resources for Containers](#).

#### Resource requests

Requests specify the resources to reserve for a given container. Reserving the resources ensures that they are always available.

**IMPORTANT**

If the resource request is for more than the available free resources in the Kubernetes cluster, the pod is not scheduled.

A request may be configured for one or more supported resources.

*Example resource requests configuration*

```
# ...
resources:
  requests:
    cpu: 12
    memory: 64Gi
# ...
```

#### Resource limits

Limits specify the maximum resources that can be consumed by a given container. The limit is not reserved and might not always be available. A container can use the resources up to the limit only when they are available. Resource limits should be always higher than the resource requests.

A resource may be configured for one or more supported limits.

#### Example resource limits configuration

```
# ...
resources:
  limits:
    cpu: 12
    memory: 64Gi
# ...
```

### Supported CPU formats

CPU requests and limits are supported in the following formats:

- Number of CPU cores as integer (`5` CPU core) or decimal (`2.5` CPU core).
- Number or *millicpus / millicores* (`100m`) where 1000 *millicores* is the same `1` CPU core.

#### Example CPU units

```
# ...
resources:
  requests:
    cpu: 500m
  limits:
    cpu: 2.5
# ...
```

**NOTE**

The computing power of 1 CPU core may differ depending on the platform where Kubernetes is deployed.

For more information on CPU specification, see the [Meaning of CPU](#).

### Supported memory formats

Memory requests and limits are specified in megabytes, gigabytes, mebibytes, and gibibytes.

- To specify memory in megabytes, use the `M` suffix. For example `1000M`.
- To specify memory in gigabytes, use the `G` suffix. For example `1G`.
- To specify memory in mebibytes, use the `Mi` suffix. For example `1000Mi`.
- To specify memory in gibibytes, use the `Gi` suffix. For example `1Gi`.

```
# ...
resources:
  requests:
    memory: 512Mi
  limits:
    memory: 2Gi
# ...
```

For more details about memory specification and additional supported units, see [Meaning of memory](#).

### B.1.6. `image`

Use the `image` property to configure the container image used by the component.

Overriding container images is recommended only in special situations where you need to use a different container registry or a customized image.

For example, if your network does not allow access to the container repository used by Strimzi, you can copy the Strimzi images or build them from the source. However, if the configured image is not compatible with Strimzi images, it might not work properly.

A copy of the container image might also be customized and used for debugging.

You can specify which container image to use for a component using the `image` property in the following resources:

- `Kafka.spec.kafka`
- `Kafka.spec.zookeeper`
- `Kafka.spec.entityOperator.topicOperator`
- `Kafka.spec.entityOperator.userOperator`
- `Kafka.spec.entityOperator.tlsSidecar`
- `Kafka.spec.jmxTrans`
- `KafkaConnect.spec`
- `KafkaConnectS2I.spec`
- `KafkaMirrorMaker.spec`
- `KafkaMirrorMaker2.spec`
- `KafkaBridge.spec`

### Configuring the `image` property for Kafka, Kafka Connect, and Kafka MirrorMaker

Kafka, Kafka Connect (including Kafka Connect with S2I support), and Kafka MirrorMaker support multiple versions of Kafka. Each component requires its own image. The default images for the different Kafka versions are configured in the following environment variables:

- `STRIMZI_KAFKA_IMAGES`
- `STRIMZI_KAFKA_CONNECT_IMAGES`
- `STRIMZI_KAFKA_CONNECT_S2I_IMAGES`
- `STRIMZI_KAFKA_MIRROR MAKER_IMAGES`

These environment variables contain mappings between the Kafka versions and their corresponding images. The mappings are used together with the `image` and `version` properties:

- If neither `image` nor `version` are given in the custom resource then the `version` will default to the Cluster Operator's default Kafka version, and the image will be the one corresponding to this version in the environment variable.
- If `image` is given but `version` is not, then the given image is used and the `version` is assumed to be the Cluster Operator's default Kafka version.
- If `version` is given but `image` is not, then the image that corresponds to the given version in the environment variable is used.
- If both `version` and `image` are given, then the given image is used. The image is assumed to contain a Kafka image with the given version.

The `image` and `version` for the different components can be configured in the following properties:

- For Kafka in `spec.kafka.image` and `spec.kafka.version`.
- For Kafka Connect, Kafka Connect S2I, and Kafka MirrorMaker in `spec.image` and `spec.version`.

#### **WARNING**

It is recommended to provide only the `version` and leave the `image` property unspecified. This reduces the chance of making a mistake when configuring the custom resource. If you need to change the images used for different versions of Kafka, it is preferable to configure the Cluster Operator's environment variables.

### Configuring the `image` property in other resources

For the `image` property in the other custom resources, the given value will be used during deployment. If the `image` property is missing, the `image` specified in the Cluster Operator configuration will be used. If the `image` name is not defined in the Cluster Operator configuration, then the default value will be used.

- For Topic Operator:
  1. Container image specified in the `STRIMZI_DEFAULT_TOPIC_OPERATOR_IMAGE` environment variable from the Cluster Operator configuration.
  2. `quay.io/strimzi/operator:0.21.1` container image.
- For User Operator:
  1. Container image specified in the `STRIMZI_DEFAULT_USER_OPERATOR_IMAGE` environment variable from the Cluster Operator configuration.
  2. `quay.io/strimzi/operator:0.21.1` container image.

- For Entity Operator TLS sidecar:
  1. Container image specified in the `STRIMZI_DEFAULT_TLS_SIDECAR_ENTITY_OPERATOR_IMAGE` environment variable from the Cluster Operator configuration.
  2. `quay.io/strimzi/kafka:0.21.1-kafka-2.7.0` container image.
- For Kafka Exporter:
  1. Container image specified in the `STRIMZI_DEFAULT_KAFKA_EXPORTER_IMAGE` environment variable from the Cluster Operator configuration.
  2. `quay.io/strimzi/kafka:0.21.1-kafka-2.7.0` container image.
- For Kafka Bridge:
  1. Container image specified in the `STRIMZI_DEFAULT_KAFKA_BRIDGE_IMAGE` environment variable from the Cluster Operator configuration.
  2. `quay.io/strimzi/kafka-bridge:0.19.0` container image.
- For Kafka broker initializer:
  1. Container image specified in the `STRIMZI_DEFAULT_KAFKA_INIT_IMAGE` environment variable from the Cluster Operator configuration.
  2. `quay.io/strimzi/operator:0.21.1` container image.
- For Kafka jmxTrans:
  1. Container image specified in the `STRIMZI_DEFAULT_JMXTRANS_IMAGE` environment variable from the Cluster Operator configuration.
  2. `quay.io/strimzi/jmxtrans:0.21.1` container image.

*Example of container image configuration*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    image: my-org/my-image:latest
    # ...
  zookeeper:
    # ...
```

## B.1.7. livenessProbe and readinessProbe healthchecks

Use the `livenessProbe` and `readinessProbe` properties to configure healthcheck probes supported in Strimzi.

Healthchecks are periodical tests which verify the health of an application. When a Healthcheck probe fails, Kubernetes assumes that the application is not healthy and attempts to fix it.

For more details about the probes, see [Configure Liveness and Readiness Probes](#).

Both `livenessProbe` and `readinessProbe` support the following options:

- `initialDelaySeconds`
- `timeoutSeconds`
- `periodSeconds`
- `successThreshold`
- `failureThreshold`

*Example of liveness and readiness probe configuration*

```
# ...
readinessProbe:
  initialDelaySeconds: 15
  timeoutSeconds: 5
livenessProbe:
  initialDelaySeconds: 15
  timeoutSeconds: 5
# ...
```

For more information about the `livenessProbe` and `readinessProbe` options, see [Probe schema reference](#).

### B.1.8. `metricsConfig`

Use the `metricsConfig` property to enable and configure Prometheus metrics.

The `metricsConfig` property contains a reference to a ConfigMap containing additional configuration for the [Prometheus JMX exporter](#). Strimzi supports Prometheus metrics using Prometheus JMX exporter to convert the JMX metrics supported by Apache Kafka and ZooKeeper to Prometheus metrics.

To enable Prometheus metrics export without further configuration, you can reference a ConfigMap containing an empty file under `metricsConfig.valueFrom.configMapKeyRef.key`. When referencing an empty file, all metrics are exposed as long as they have not been renamed.

### Example ConfigMap with metrics configuration for Kafka

```
kind: ConfigMap
apiVersion: kafka.strimzi.io/v1beta1
metadata:
  name: my-configmap
data:
  my-key: |
    lowercaseOutputName: true
    rules:
      # Special cases and very specific rules
      - pattern: kafka.server<type=(.+), name=(.+), clientId=(.+), topic=(.+), partition=(.+)><>Value
        name: kafka_server_$1_$2
        type: GAUGE
        labels:
          clientId: "$3"
          topic: "$4"
          partition: "$5"
    # further configuration
```

### Example metrics configuration for Kafka

```
[source,shell,subs="+attributes"]
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    metricsConfig:
      type: jmxPrometheusExporter
      valueFrom:
        configMapKeyRef:
          name: my-config-map
          key: my-key
    # ...
  zookeeper:
    # ...
```

When metrics are enabled, they are exposed on port 9404.

When the `metricsConfig` (or deprecated `metrics`) property is not defined in the resource, the Prometheus metrics are disabled.

For more information about setting up and deploying Prometheus and Grafana, see [Introducing Metrics to Kafka](#) in the *Deploying and Upgrading Strimzi* guide.

## B.1.9. `jvmOptions`

The following components of Strimzi run inside a Virtual Machine (VM):

- Apache Kafka
- Apache ZooKeeper
- Apache Kafka Connect
- Apache Kafka MirrorMaker
- Strimzi Kafka Bridge

JVM configuration options optimize the performance for different platforms and architectures, and can be configured using the `jvmOptions` property in following resources:

- `Kafka.spec.kafka`
- `Kafka.spec.zookeeper`
- `KafkaConnect.spec`
- `KafkaConnectS2I.spec`
- `KafkaMirrorMaker.spec`
- `KafkaMirrorMaker2.spec`
- `KafkaBridge.spec`

Only the following JVM options are supported:

### `-Xms`

Configures the minimum initial allocation heap size when the JVM starts.

### `-Xmx`

Configures the maximum heap size.

**NOTE** The units accepted by JVM settings such as `-Xmx` and `-Xms` are those accepted by the JDK `java` binary in the corresponding image. Accordingly, `1g` or `1G` means 1,073,741,824 bytes, and `Gi` is not a valid unit suffix. This is in contrast to the units used for [memory requests and limits](#), which follow the Kubernetes convention where `1G` means 1,000,000,000 bytes, and `1Gi` means 1,073,741,824 bytes

The default values used for `-Xms` and `-Xmx` depends on whether there is a [memory request](#) limit configured for the container.

- If there is a memory limit, the JVM's minimum and maximum memory is set to a value corresponding to the limit.
- If there is no memory limit, the JVM's minimum memory is set to `128M`. The JVM's maximum memory is not defined to allow the memory to grow as needed, which is ideal for single node environments in test and development.

**IMPORTANT**

Setting `-Xmx` explicitly requires some care:

- The JVM's overall memory usage will be approximately  $4 \times$  the maximum heap, as configured by `-Xmx`.
- If `-Xmx` is set without also setting an appropriate Kubernetes memory limit, it is possible that the container will be killed should the Kubernetes node experience memory pressure (from other Pods running on it).
- If `-Xmx` is set without also setting an appropriate Kubernetes memory request, it is possible that the container will be scheduled to a node with insufficient memory. In this case, the container will not start but crash (immediately if `-Xms` is set to `-Xmx`, or some later time if not).

When setting `-Xmx` explicitly, it is recommended to:

- Set the memory request and the memory limit to the same value
- Use a memory request that is at least  $4.5 \times$  the `-Xmx`
- Consider setting `-Xms` to the same value as `-Xmx`

**IMPORTANT**

Containers doing lots of disk I/O (such as Kafka broker containers) will need to leave some memory available for use as an operating system page cache. On such containers, the requested memory should be significantly higher than the memory used by the JVM.

*Example fragment configuring `-Xmx` and `-Xms`*

```
# ...
jvmOptions:
  "-Xmx": "2g"
  "-Xms": "2g"
# ...
```

In the above example, the JVM will use 2 GiB (=2,147,483,648 bytes) for its heap. Its total memory usage will be approximately 8GiB.

Setting the same value for initial (`-Xms`) and maximum (`-Xmx`) heap sizes avoids the JVM having to allocate memory after startup, at the cost of possibly allocating more heap than is really needed. For Kafka and ZooKeeper pods such allocation could cause unwanted latency. For Kafka Connect avoiding over allocation may be the most important concern, especially in distributed mode where the effects of over-allocation is multiplied by the number of consumers.

**-XX**

`-XX` object can be used for configuring advanced runtime options of a JVM. The `-XX` options are used to configure the `KAFKA_JVM_PERFORMANCE_OPTS` option of Apache Kafka.

*Example showing the use of the `-XX` object*

```
jvmOptions:  
  "-XX":  
    "UseG1GC": true  
    "MaxGCPauseMillis": 20  
    "InitiatingHeapOccupancyPercent": 35  
    "ExplicitGCInvokesConcurrent": true
```

The example configuration above will result in the following JVM options:

```
-XX:+UseG1GC -XX:MaxGCPauseMillis=20 -XX:InitiatingHeapOccupancyPercent=35  
-XX:+ExplicitGCInvokesConcurrent -XX:-UseParNewGC
```

**NOTE**

When no `-XX` options are specified, the default Apache Kafka configuration of `KAFKA_JVM_PERFORMANCE_OPTS` is used.

### B.1.10. Garbage collector logging

The `jvmOptions` property also allows you to enable and disable garbage collector (GC) logging. GC logging is disabled by default. To enable it, set the `gcLoggingEnabled` property as follows:

*Example of enabling GC logging*

```
# ...  
jvmOptions:  
  gcLoggingEnabled: true  
# ...
```

## B.2. Kafka schema reference

Property	Description
spec	The specification of the Kafka and ZooKeeper clusters, and Topic Operator.
<code>KafkaSpec</code>	
status	The status of the Kafka and ZooKeeper clusters, and Topic Operator.
<code>KafkaStatus</code>	

## B.3. KafkaSpec schema reference

Used in: [Kafka](#)

Property	Description
kafka	Configuration of the Kafka cluster.
<a href="#">KafkaClusterSpec</a>	
zookeeper	Configuration of the ZooKeeper cluster.
<a href="#">ZookeeperClusterSpec</a>	
topicOperator	The property <code>topicOperator</code> has been deprecated. This feature should now be configured at path <code>spec.entityOperator.topicOperator</code> . Configuration of the Topic Operator.
<a href="#">TopicOperatorSpec</a>	Configuration of the Topic Operator.
entityOperator	Configuration of the Entity Operator.
<a href="#">EntityOperatorSpec</a>	
clusterCa	Configuration of the cluster certificate authority.
<a href="#">CertificateAuthority</a>	
clientsCa	Configuration of the clients certificate authority.
<a href="#">CertificateAuthority</a>	
cruiseControl	Configuration for Cruise Control deployment. Deploys a Cruise Control instance when specified.
<a href="#">CruiseControlSpec</a>	
jmxTrans	Configuration for JmxTrans. When the property is present a JmxTrans deployment is created for gathering JMX metrics from each Kafka broker. For more information see <a href="#">JmxTrans GitHub</a> .
<a href="#">JmxTransSpec</a>	
kafkaExporter	Configuration of the Kafka Exporter. Kafka Exporter can provide additional metrics, for example lag of consumer group at topic/partition.
<a href="#">KafkaExporterSpec</a>	
maintenanceTimeWindows	A list of time windows for maintenance tasks (that is, certificates renewal). Each time window is defined by a cron expression.
string array	

## B.4. KafkaClusterSpec schema reference

Used in: [KafkaSpec](#)

Configures a Kafka cluster.

### B.4.1. listeners

Use the `listeners` property to configure listeners to provide access to Kafka brokers.

## *Example configuration of a plain (unencrypted) listener without authentication*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
    listeners:
      - name: plain
        port: 9092
        type: internal
        tls: false
    # ...
  zookeeper:
    # ...
```

### B.4.2. config

Use the `config` properties to configure Kafka broker options as keys.

Standard Apache Kafka configuration may be provided, restricted to those properties not managed directly by Strimzi.

Configuration options that cannot be configured relate to:

- Security (Encryption, Authentication, and Authorization)
- Listener configuration
- Broker ID configuration
- Configuration of log data directories
- Inter-broker communication
- ZooKeeper connectivity

The values can be one of the following JSON types:

- String
- Number
- Boolean

You can specify and configure the options listed in the [Apache Kafka documentation](#) with the exception of those options that are managed directly by Strimzi. Specifically, all configuration options with keys equal to or starting with one of the following strings are forbidden:

- `listeners`
- `advertised.`
- `broker.`
- `listener.`

- `host.name`
- `port`
- `inter.broker.listener.name`
- `sasl.`
- `ssl.`
- `security.`
- `password.`
- `principal.builder.class`
- `log.dir`
- `zookeeper.connect`
- `zookeeper.set.acl`
- `authorizer.`
- `super.user`

When a forbidden option is present in the `config` property, it is ignored and a warning message is printed to the Cluster Operator log file. All other supported options are passed to Kafka.

There are exceptions to the forbidden options. For client connection using a specific *cipher suite* for a TLS version, you can [configure allowed ssl properties](#). You can also configure the `zookeeper.connection.timeout.ms` property to set the maximum time allowed for establishing a ZooKeeper connection.

## *Example Kafka broker configuration*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    config:
      num.partitions: 1
      num.recovery.threads.per.data.dir: 1
      default.replication.factor: 3
      offsets.topic.replication.factor: 3
      transaction.state.log.replication.factor: 3
      transaction.state.log.min.isr: 1
      log.retention.hours: 168
      log.segment.bytes: 1073741824
      log.retention.check.interval.ms: 300000
      num.network.threads: 3
      num.io.threads: 8
      socket.send.buffer.bytes: 102400
      socket.receive.buffer.bytes: 102400
      socket.request.max.bytes: 104857600
      group.initial.rebalance.delay.ms: 0
      ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384"
      ssl.enabled.protocols: "TLSv1.2"
      ssl.protocol: "TLSv1.2"
      zookeeper.connection.timeout.ms: 6000
    # ...
```

### **B.4.3. `brokerRackInitImage`**

When rack awareness is enabled, Kafka broker pods use init container to collect the labels from the Kubernetes cluster nodes. The container image used for this container can be configured using the `brokerRackInitImage` property. When the `brokerRackInitImage` field is missing, the following images are used in order of priority:

1. Container image specified in `STRIMZI_DEFAULT_KAFKA_INIT_IMAGE` environment variable in the Cluster Operator configuration.
2. `quay.io/strimzi/operator:0.21.1` container image.

*Example brokerRackInitImage configuration*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    rack:
      topologyKey: topology.kubernetes.io/zone
  brokerRackInitImage: my-org/my-image:latest
  # ...
```

**NOTE** Overriding container images is recommended only in special situations, where you need to use a different container registry. For example, because your network does not allow access to the container registry used by Strimzi. In this case, you should either copy the Strimzi images or build them from the source. If the configured image is not compatible with Strimzi images, it might not work properly.

#### B.4.4. jmxOptions

JMX metrics are obtained from Kafka brokers by opening a JMX port on 9999. Use the `jmxOptions` property to configure a password-protected JMX port, to prevent unauthorized pods from accessing the port, or an unprotected JMX port.

You can then obtain metrics about each Kafka broker, such as usage data from the `BytesPerSecond` value, or the request rate of the network of the broker.

To enable security for the JMX port, set the `type` parameter in the `authentication` field to `password`.

*Example password-protected JMX configuration*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    jmxOptions:
      authentication:
        type: "password"
    # ...
  zookeeper:
    # ...
```

You can then deploy a pod into a cluster and obtain JMX metrics using the headless service by specifying which broker you want to address.

For example, to get JMX metrics from broker 0 you specify:

```
"CLUSTER-NAME-kafka-0-CLUSTER-NAME-HEADLESS-SERVICE-NAME"
```

If the JMX port is secured, you can get the username and password by referencing them from the JMX secret in the deployment of your pod.

For an unprotected JMX port, use an empty object `{}` to open the JMX Port on the headless service. You deploy a pod and obtain metrics in the same way as for the protected port, but in this case any pod can read from the JMX port.

*Example open port JMX configuration*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    jmxOptions: {}
    # ...
  zookeeper:
    # ...
```

#### B.4.5. logging

Kafka has its own configurable loggers:

- `log4j.logger.org.I0Itec.zkclient.ZkClient`
- `log4j.logger.org.apache.zookeeper`
- `log4j.logger.kafka`
- `log4j.logger.org.apache.kafka`
- `log4j.logger.kafka.request.logger`
- `log4j.logger.kafka.network.Processor`
- `log4j.logger.kafka.server.KafkaApis`
- `log4j.logger.kafka.network.RequestChannel$`
- `log4j.logger.kafka.controller`
- `log4j.logger.kafka.log.LogCleaner`
- `log4j.logger.state.change.logger`
- `log4j.logger.kafka.authorizer.logger`

Kafka uses the Apache `log4j` logger implementation.

Use the `logging` property to configure loggers and logger levels.

You can set the log levels by specifying the logger and level directly (inline) or use a custom (external) ConfigMap. If a ConfigMap is used, you set `logging.name` property to the name of the ConfigMap containing the external logging configuration. Inside the ConfigMap, the logging configuration is described using `log4j.properties`. For more information about log levels, see [Apache logging services](#).

Here we see examples of `inline` and `external` logging.

#### *Inline logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  # ...
  kafka:
    # ...
    logging:
      type: inline
      loggers:
        kafka.root.logger.level: "INFO"
  # ...
```

#### *External logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  # ...
  logging:
    type: external
    name: customConfigMap
  # ...
```

Any available loggers that are not configured have their level set to `OFF`.

If Kafka was deployed using the Cluster Operator, changes to Kafka logging levels are applied dynamically.

If you use external logging, a rolling update is triggered when logging appenders are changed.

#### *Garbage collector (GC)*

Garbage collector logging can also be enabled (or disabled) using the `jvmOptions` property.

### B.4.6. Schema reference

Property	Description
replicas	The number of pods in the cluster.
integer	
image	The docker image for the pods. The default value depends on the configured <code>Kafka.spec.kafka.version</code> .
string	
storage	Storage configuration (disk). Cannot be updated. The type depends on the value of the <code>storage.type</code> property within the given object, which must be one of [ephemeral, persistent-claim, jbod].
EphemeralStorage, PersistentClaimStorage, JbodStorage	
listeners	Configures listeners of Kafka brokers.
GenericKafkaListener array or KafkaListeners	
authorization	Authorization configuration for Kafka brokers. The type depends on the value of the <code>authorization.type</code> property within the given object, which must be one of [simple, opa, keycloak].
KafkaAuthorizationSimple, KafkaAuthorizationOpa, KafkaAuthorizationKeycloak	
config	Kafka broker config properties with the following prefixes cannot be set: listeners, advertised., broker., listener., host.name, port, inter.broker.listener.name, sasl., ssl., security., password., principal.builder.class, log.dir, zookeeper.connect, zookeeper.set.acl, zookeeper.ssl, zookeeper.clientCnxnSocket, authorizer., super.user, cruise.control.metrics.topic, cruise.control.metrics.reporter.bootstrap.servers (with the exception of: zookeeper.connection.timeout.ms, ssl.cipher.suites, ssl.protocol, ssl.enabled.protocols, cruise.control.metrics.topic.num.partitions, cruise.control.metrics.topic.replication.factor, cruise.control.metrics.topic.retention.ms, cruise.control.metrics.topic.auto.create.retries, cruise.control.metrics.topic.auto.create.timeout.ms, cruise.control.metrics.topic.min.insync.replicas).
map	
rack	Configuration of the <code>broker.rack</code> broker config.
Rack	
brokerRackInitImage	The image of the init container used for initializing the <code>broker.rack</code> .
string	

Property	Description
affinity	The property <b>affinity</b> has been deprecated. This feature should now be configured at path <code>spec.kafka.template.pod.affinity</code> . The pod's affinity rules. For more information, see the <a href="#">external documentation for core/v1 affinity</a> .
Affinity	
tolerations	The property <b>tolerations</b> has been deprecated. This feature should now be configured at path <code>spec.kafka.template.pod.tolerations</code> . The pod's tolerations. For more information, see the <a href="#">external documentation for core/v1 toleration</a> .
Toleration array	
livenessProbe	Pod liveness checking.
Probe	
readinessProbe	Pod readiness checking.
Probe	
jvmOptions	JVM Options for pods.
JvmOptions	
jmxOptions	JMX Options for Kafka brokers.
KafkaJmxOptions	
resources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
ResourceRequirements	
metrics	The property <b>metrics</b> has been deprecated. This feature should now be configured at path <code>spec.kafka.metricsConfig</code> . The Prometheus JMX Exporter configuration. See <a href="https://github.com/prometheus/jmx_exporter">https://github.com/prometheus/jmx_exporter</a> for details of the structure of this configuration.
map	
metricsConfig	Metrics configuration. The type depends on the value of the <code>metricsConfig.type</code> property within the given object, which must be one of [jmxPrometheusExporter].
JmxPrometheusExporterMetrics	
logging	Logging configuration for Kafka. The type depends on the value of the <code>logging.type</code> property within the given object, which must be one of [inline, external].
InlineLogging, ExternalLogging	
tlsSidecar	The property <b>tlsSidecar</b> has been deprecated. TLS sidecar configuration.
TlsSidecar	
template	Template for Kafka cluster resources. The template allows users to specify how are the <code>StatefulSet</code> , <code>Pods</code> and <code>Services</code> generated.
KafkaClusterTemplate	

Property	Description
version	The kafka broker version. Defaults to 2.7.0. Consult the user documentation to understand the process required to upgrade or downgrade the version.
string	

## B.5. EphemeralStorage schema reference

Used in: [JbodStorage](#), [KafkaClusterSpec](#), [ZookeeperClusterSpec](#)

The `type` property is a discriminator that distinguishes use of the [EphemeralStorage](#) type from [PersistentClaimStorage](#). It must have the value `ephemeral` for the type [EphemeralStorage](#).

Property	Description
id	Storage identification number. It is mandatory only for storage volumes defined in a storage of type 'jbod'.
integer	
sizeLimit	When type=ephemeral, defines the total amount of local storage required for this EmptyDir volume (for example 1Gi).
string	
type	Must be <code>ephemeral</code> .
string	

## B.6. PersistentClaimStorage schema reference

Used in: [JbodStorage](#), [KafkaClusterSpec](#), [ZookeeperClusterSpec](#)

The `type` property is a discriminator that distinguishes use of the [PersistentClaimStorage](#) type from [EphemeralStorage](#). It must have the value `persistent-claim` for the type [PersistentClaimStorage](#).

Property	Description
type	Must be <code>persistent-claim</code> .
string	
size	When type=persistent-claim, defines the size of the persistent volume claim (i.e 1Gi). Mandatory when type=persistent-claim.
string	
selector	Specifies a specific persistent volume to use. It contains key:value pairs representing labels for selecting such a volume.
map	
deleteClaim	Specifies if the persistent volume claim has to be deleted when the cluster is un-deployed.
boolean	

Property	Description
class	The storage class to use for dynamic volume allocation.
string	
id	Storage identification number. It is mandatory only for storage volumes defined in a storage of type 'jbod'.
integer	
overrides	Overrides for individual brokers. The <code>overrides</code> field allows to specify a different configuration for different brokers.
<code>PersistentClaimStorageOverride</code> array	

## B.7. PersistentClaimStorageOverride schema reference

Used in: [PersistentClaimStorage](#)

Property	Description
class	The storage class to use for dynamic volume allocation for this broker.
string	
broker	Id of the kafka broker (broker identifier).
integer	

## B.8. JbodStorage schema reference

Used in: [KafkaClusterSpec](#)

The `type` property is a discriminator that distinguishes use of the `JbodStorage` type from `EphemeralStorage`, `PersistentClaimStorage`. It must have the value `jbod` for the type `JbodStorage`.

Property	Description
type	Must be <code>jbod</code> .
string	
volumes	List of volumes as Storage objects representing the JBOD disks array.
<code>EphemeralStorage</code> , <code>PersistentClaimStorage</code> array	

## B.9. GenericKafkaListener schema reference

Used in: [KafkaClusterSpec](#)

Configures listeners to connect to Kafka brokers within and outside Kubernetes.

You configure the listeners in the `Kafka` resource.

## *Example Kafka resource showing listener configuration*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    #...
    listeners:
      - name: plain
        port: 9092
        type: internal
        tls: false
      - name: tls
        port: 9093
        type: internal
        tls: true
        authentication:
          type: tls
      - name: external1
        port: 9094
        type: route
        tls: true
      - name: external2
        port: 9095
        type: ingress
        tls: false
        authentication:
          type: tls
    configuration:
      bootstrap:
        host: bootstrap.myingress.com
      brokers:
        - broker: 0
          host: broker-0.myingress.com
        - broker: 1
          host: broker-1.myingress.com
        - broker: 2
          host: broker-2.myingress.com
    #...
```

### **B.9.1. `listeners`**

You configure Kafka broker listeners using the `listeners` property in the `Kafka` resource. Listeners are defined as an array.

### *Example listener configuration*

```
listeners:
- name: plain
  port: 9092
  type: internal
  tls: false
```

The name and port must be unique within the Kafka cluster. The name can be up to 25 characters long, comprising lower-case letters and numbers. Allowed port numbers are 9092 and higher with the exception of ports 9404 and 9999, which are already used for Prometheus and JMX.

By specifying a unique name and port for each listener, you can configure multiple listeners.

#### **B.9.2. type**

The type is set as `internal`, or for external listeners, as `route`, `loadbalancer`, `nodeport` or `ingress`.

##### **internal**

You can configure internal listeners with or without encryption using the `tls` property.

##### *Example `internal` listener configuration*

```
#...
spec:
kafka:
#...
listeners:
#...
- name: plain
  port: 9092
  type: internal
  tls: false
- name: tls
  port: 9093
  type: internal
  tls: true
  authentication:
    type: tls
#...
```

##### **route**

Configures an external listener to expose Kafka using OpenShift `Routes` and the HAProxy router.

A dedicated `Route` is created for every Kafka broker pod. An additional `Route` is created to serve as a Kafka bootstrap address. Kafka clients can use these `Routes` to connect to Kafka on port 443. The client connects on port 443, the default router port, but traffic is then routed to the port you configure, which is `9094` in this example.

*Example `route` listener configuration*

```
#...
spec:
  kafka:
    #...
    listeners:
      #...
      - name: external1
        port: 9094
        type: route
        tls: true
      #...
```

## ingress

Configures an external listener to expose Kafka using Kubernetes [Ingress](#) and the [NGINX Ingress Controller for Kubernetes](#).

A dedicated [Ingress](#) resource is created for every Kafka broker pod. An additional [Ingress](#) resource is created to serve as a Kafka bootstrap address. Kafka clients can use these [Ingress](#) resources to connect to Kafka on port 443. The client connects on port 443, the default controller port, but traffic is then routed to the port you configure, which is [9095](#) in the following example.

You must specify the hostnames used by the bootstrap and per-broker services using [GenericKafkaListenerConfigurationBootstrap](#) and [GenericKafkaListenerConfigurationBroker](#) properties.

### Example `ingress` listener configuration

```
#...
spec:
  kafka:
    #...
    listeners:
      #...
      - name: external2
        port: 9095
        type: ingress
        tls: false
        authentication:
          type: tls
        configuration:
          bootstrap:
            host: bootstrap.myingress.com
        brokers:
          - broker: 0
            host: broker-0.myingress.com
          - broker: 1
            host: broker-1.myingress.com
          - broker: 2
            host: broker-2.myingress.com
    #...
```

#### NOTE

External listeners using `Ingress` are currently only tested with the [NGINX Ingress Controller for Kubernetes](#).

## loadbalancer

Configures an external listener to expose Kafka `Loadbalancer` type `Services`.

A new loadbalancer service is created for every Kafka broker pod. An additional loadbalancer is created to serve as a Kafka `bootstrap` address. Loadbalancers listen to the specified port number, which is port `9094` in the following example.

You can use the `loadBalancerSourceRanges` property to configure `source ranges` to restrict access to the specified IP addresses.

*Example `loadbalancer` listener configuration*

```
#...
spec:
  kafka:
    #...
    listeners:
      - name: external3
        port: 9094
        type: loadbalancer
        tls: true
        configuration:
          loadBalancerSourceRanges:
            - 10.0.0.0/8
            - 88.208.76.87/32
    #...
```

## nodeport

Configures an external listener to expose Kafka using `NodePort` type [Services](#).

Kafka clients connect directly to the nodes of Kubernetes. An additional `NodePort` type of service is created to serve as a Kafka bootstrap address.

When configuring the advertised addresses for the Kafka broker pods, Strimzi uses the address of the node on which the given pod is running. You can use `preferredNodePortAddressType` property to configure the [first address type checked as the node address](#).

*Example `nodeport` listener configuration*

```
#...
spec:
  kafka:
    #...
    listeners:
      #...
      - name: external4
        port: 9095
        type: nodeport
        tls: false
        configuration:
          preferredNodePortAddressType: InternalDNS
    #...
```

### NOTE

TLS hostname verification is not currently supported when exposing Kafka clusters using node ports.

### B.9.3. port

The port number is the port used in the Kafka cluster, which might not be the same port used for

access by a client.

- `loadbalancer` listeners use the specified port number, as do `internal` listeners
- `ingress` and `route` listeners use port 443 for access
- `nodeport` listeners use the port number assigned by Kubernetes

For client connection, use the address and port for the bootstrap service of the listener. You can retrieve this from the status of the `Kafka` resource.

*Example command to retrieve the address and port for client connection*

```
kubectl get kafka KAFKA-CLUSTER-NAME  
-o=jsonpath='{.status.listeners[?(@.type=="external")].bootstrapServers}{"\n"}'
```

**NOTE**

Listeners cannot be configured to use the ports set aside for interbroker communication (9091) and metrics (9404).

#### B.9.4. `tls`

The TLS property is required.

By default, TLS encryption is not enabled. To enable it, set the `tls` property to `true`.

TLS encryption is always used with `route` listeners.

#### B.9.5. `authentication`

Authentication for the listener can be specified as:

- Mutual TLS (`tls`)
- SCRAM-SHA-512 (`scram-sha-512`)
- Token-based OAuth 2.0 (`oauth`).

#### B.9.6. `networkPolicyPeers`

Use `networkPolicyPeers` to configure network policies that restrict access to a listener at the network level. The following example shows a `networkPolicyPeers` configuration for a `plain` and a `tls` listener.

```

listeners:
#...
- name: plain
  port: 9092
  type: internal
  tls: true
  authentication:
    type: scram-sha-512
networkPolicyPeers:
- podSelector:
  matchLabels:
    app: kafka-sasl-consumer
- podSelector:
  matchLabels:
    app: kafka-sasl-producer
- name: tls
  port: 9093
  type: internal
  tls: true
  authentication:
    type: tls
networkPolicyPeers:
- namespaceSelector:
  matchLabels:
    project: myproject
- namespaceSelector:
  matchLabels:
    project: myproject2
# ...

```

In the example:

- Only application pods matching the labels `app: kafka-sasl-consumer` and `app: kafka-sasl-producer` can connect to the `plain` listener. The application pods must be running in the same namespace as the Kafka broker.
- Only application pods running in namespaces matching the labels `project: myproject` and `project: myproject2` can connect to the `tls` listener.

The syntax of the `networkPolicyPeers` field is the same as the `from` field in `NetworkPolicy` resources.

*Backwards compatibility with KafkaListeners*

`GenericKafkaListener` replaces the `KafkaListeners` schema, which is now deprecated.

To convert the listeners configured using the `KafkaListeners` schema into the format of the `GenericKafkaListener` schema, with backwards compatibility, use the following names, ports and types:

```

listeners:
#...
- name: plain
  port: 9092
  type: internal
  tls: false
- name: tls
  port: 9093
  type: internal
  tls: true
- name: external
  port: 9094
  type: EXTERNAL-LISTENER-TYPE ①
  tls: true
# ...

```

① Options: `ingress`, `loadbalancer`, `nodeport`, `route`

### B.9.7. Schema reference

Property	Description
name	Name of the listener. The name will be used to identify the listener and the related Kubernetes objects. The name has to be unique within given a Kafka cluster. The name can consist of lowercase characters and numbers and be up to 11 characters long.
string	
port	Port number used by the listener inside Kafka. The port number has to be unique within a given Kafka cluster. Allowed port numbers are 9092 and higher with the exception of ports 9404 and 9999, which are already used for Prometheus and JMX. Depending on the listener type, the port number might not be the same as the port number that connects Kafka clients.
integer	
type	Type of the listener. Currently the supported types are <code>internal</code> , <code>route</code> , <code>loadbalancer</code> , <code>nodeport</code> and <code>ingress</code> . <ul style="list-style-type: none"> <li>* <code>internal</code> type exposes Kafka internally only within the Kubernetes cluster.</li> <li>* <code>route</code> type uses OpenShift Routes to expose Kafka.</li> <li>* <code>loadbalancer</code> type uses LoadBalancer type services to expose Kafka.</li> <li>* <code>nodeport</code> type uses NodePort type services to expose Kafka.</li> <li>* <code>ingress</code> type uses Kubernetes Nginx Ingress to expose Kafka..</li> </ul>
string (one of [ingress, internal, route, loadbalancer, nodeport])	

Property	Description
tls	Enables TLS encryption on the listener. This is a required property.
boolean	
authentication	Authentication configuration for this listener. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, oauth].
<code>KafkaListenerAuthenticationTls</code> , <code>KafkaListenerAuthenticationScramSha512</code> , <code>KafkaListenerAuthenticationOAuth</code>	
configuration	Additional listener configuration.
<code>GenericKafkaListenerConfiguration</code>	
networkPolicyPeers	List of peers which should be able to connect to this listener. Peers in this list are combined using a logical OR operation. If this field is empty or missing, all connections will be allowed for this listener. If this field is present and contains at least one item, the listener only allows the traffic which matches at least one item in this list. For more information, see the <a href="#">external documentation for networking.k8s.io/v1 networkpolicypeer</a> .
<code>NetworkPolicyPeer</code> array	

## B.10. KafkaListenerAuthenticationTls schema reference

Used in: `GenericKafkaListener`, `KafkaListenerExternalIngress`, `KafkaListenerExternalLoadBalancer`, `KafkaListenerExternalNodePort`, `KafkaListenerExternalRoute`, `KafkaListenerPlain`, `KafkaListenerTls`

The `type` property is a discriminator that distinguishes use of the `KafkaListenerAuthenticationTls` type from `KafkaListenerAuthenticationScramSha512`, `KafkaListenerAuthenticationOAuth`. It must have the value `tls` for the type `KafkaListenerAuthenticationTls`.

Property	Description
<code>type</code>	Must be <code>tls</code> .
<code>string</code>	

## B.11. KafkaListenerAuthenticationScramSha512 schema reference

Used in: `GenericKafkaListener`, `KafkaListenerExternalIngress`, `KafkaListenerExternalLoadBalancer`, `KafkaListenerExternalNodePort`, `KafkaListenerExternalRoute`, `KafkaListenerPlain`, `KafkaListenerTls`

The `type` property is a discriminator that distinguishes use of the `KafkaListenerAuthenticationScramSha512` type from `KafkaListenerAuthenticationTls`, `KafkaListenerAuthenticationOAuth`. It must have the value `scram-sha-512` for the type

## KafkaListenerAuthenticationScramSha512.

Property	Description
type	Must be <code>scram-sha-512</code> .
string	

## B.12. KafkaListenerAuthenticationOAuth schema reference

Used in: `GenericKafkaListener`, `KafkaListenerExternalIngress`, `KafkaListenerExternalLoadBalancer`, `KafkaListenerExternalNodePort`, `KafkaListenerExternalRoute`, `KafkaListenerPlain`, `KafkaListenerTls`

The `type` property is a discriminator that distinguishes use of the `KafkaListenerAuthenticationOAuth` type from `KafkaListenerAuthenticationTls`, `KafkaListenerAuthenticationScramSha512`. It must have the value `oauth` for the type `KafkaListenerAuthenticationOAuth`.

Property	Description
accessTokenIsJwt	Configure whether the access token is treated as JWT. This must be set to <code>false</code> if the authorization server returns opaque tokens. Defaults to <code>true</code> .
boolean	
checkAccessTokenType	Configure whether the access token type check is performed or not. This should be set to <code>false</code> if the authorization server does not include 'typ' claim in JWT token. Defaults to <code>true</code> .
boolean	
checkIssuer	Enable or disable issuer checking. By default issuer is checked using the value configured by <code>validIssuerUri</code> . Default value is <code>true</code> .
boolean	
clientId	OAuth Client ID which the Kafka broker can use to authenticate against the authorization server and use the introspect endpoint URI.
string	
clientSecret	Link to Kubernetes Secret containing the OAuth client secret which the Kafka broker can use to authenticate against the authorization server and use the introspect endpoint URI.
GenericSecretSource	
disableTlsHostnameVerification	Enable or disable TLS hostname verification. Default value is <code>false</code> .
boolean	
enableECDSA	Enable or disable ECDSA support by installing BouncyCastle crypto provider. Default value is <code>false</code> .
boolean	

Property	Description
fallbackUserNameClaim	The fallback username claim to be used for the user id if the claim specified by <code>userNameClaim</code> is not present. This is useful when <code>client_credentials</code> authentication only results in the client id being provided in another claim. It only takes effect if <code>userNameClaim</code> is set.
string	
fallbackUserNamePrefix	The prefix to use with the value of <code>fallbackUserNameClaim</code> to construct the user id. This only takes effect if <code>fallbackUserNameClaim</code> is true, and the value is present for the claim. Mapping usernames and client ids into the same user id space is useful in preventing name collisions.
string	
introspectionEndpointUri	URI of the token introspection endpoint which can be used to validate opaque non-JWT tokens.
string	
jwksEndpointUri	URI of the JWKS certificate endpoint, which can be used for local JWT validation.
string	
jwksExpirySeconds	Configures how often are the JWKS certificates considered valid. The expiry interval has to be at least 60 seconds longer than the refresh interval specified in <code>jwksRefreshSeconds</code> . Defaults to 360 seconds.
integer	
jwksMinRefreshPauseSeconds	The minimum pause between two consecutive refreshes. When an unknown signing key is encountered the refresh is scheduled immediately, but will always wait for this minimum pause. Defaults to 1 second.
integer	
jwksRefreshSeconds	Configures how often are the JWKS certificates refreshed. The refresh interval has to be at least 60 seconds shorter than the expiry interval specified in <code>jwksExpirySeconds</code> . Defaults to 300 seconds.
integer	
maxSecondsWithoutReauthentication	Maximum number of seconds the authenticated session remains valid without re-authentication.
integer	
This enables Apache Kafka re-authentication feature, and causes sessions to expire when the access token expires. If the access token expires before max time or if max time is reached, the client has to re-authenticate, otherwise the server will drop the connection. Not set by default - the authenticated session does not expire when the access token expires.	

Property	Description
tlsTrustedCertificates <a href="#">CertSecretSource</a> array	Trusted certificates for TLS connection to the OAuth server.
type	Must be <a href="#">oauth</a> .
string	
userInfoEndpointUri	URI of the User Info Endpoint to use as a fallback to obtaining the user id when the Introspection Endpoint does not return information that can be used for the user id.
string	
userNameClaim	Name of the claim from the JWT authentication token, Introspection Endpoint response or User Info Endpoint response which will be used to extract the user id. Defaults to <a href="#">sub</a> .
string	
validIssuerUri	URI of the token issuer used for authentication.
string	
validTokenType	Valid value for the <a href="#">token_type</a> attribute returned by the Introspection Endpoint. No default value, and not checked by default.
string	

## B.13. [GenericSecretSource](#) schema reference

Used in: [KafkaClientAuthenticationOAuth](#), [KafkaListenerAuthenticationOAuth](#)

Property	Description
key	The key under which the secret value is stored in the Kubernetes Secret.
string	
secretName	The name of the Kubernetes Secret containing the secret value.
string	

## B.14. [CertSecretSource](#) schema reference

Used in: [KafkaAuthorizationKeycloak](#), [KafkaBridgeTls](#), [KafkaClientAuthenticationOAuth](#), [KafkaConnectTls](#), [KafkaListenerAuthenticationOAuth](#), [KafkaMirrorMaker2Tls](#), [KafkaMirrorMakerTls](#)

Property	Description
certificate	The name of the file certificate in the Secret.
string	
secretName	The name of the Secret containing the certificate.
string	

## B.15. GenericKafkaListenerConfiguration schema reference

Used in: [GenericKafkaListener](#)

Configuration for Kafka listeners.

### B.15.1. brokerCertChainAndKey

The `brokerCertChainAndKey` property is only used with listeners that have TLS encryption enabled. You can use the property to providing your own Kafka listener certificates.

*Example configuration for a `loadbalancer` external listener with TLS encryption enabled*

```
listeners:  
#...  
- name: external  
  port: 9094  
  type: loadbalancer  
  tls: true  
  authentication:  
    type: tls  
  configuration:  
    brokerCertChainAndKey:  
      secretName: my-secret  
      certificate: my-listener-certificate.crt  
      key: my-listener-key.key  
# ...
```

### B.15.2. externalTrafficPolicy

The `externalTrafficPolicy` property is used with `loadbalancer` and `nodeport` listeners. When exposing Kafka outside of Kubernetes you can choose `Local` or `Cluster`. `Local` avoids hops to other nodes and preserves the client IP, whereas `Cluster` does neither. The default is `Cluster`.

### B.15.3. loadBalancerSourceRanges

The `loadBalancerSourceRanges` property is only used with `loadbalancer` listeners. When exposing Kafka outside of Kubernetes use source ranges, in addition to labels and annotations, to customize how a service is created.

*Example source ranges configured for a loadbalancer listener*

```
listeners:  
#...  
- name: external  
  port: 9094  
  type: loadbalancer  
  tls: false  
  configuration:  
    externalTrafficPolicy: Local  
    loadBalancerSourceRanges:  
      - 10.0.0.0/8  
      - 88.208.76.87/32  
    # ...  
# ...
```

#### B.15.4. class

The `class` property is only used with `ingress` listeners.

By default, the `Ingress` class is set to `nginx`. You can change the `Ingress` class using the `class` property.

*Example of an external listener of type ingress using Ingress class nginx-internal*

```
listeners:  
#...  
- name: external  
  port: 9094  
  type: ingress  
  tls: false  
  configuration:  
    class: nginx-internal  
  # ...  
# ...
```

#### B.15.5. preferredNodePortAddressType

The `preferredNodePortAddressType` property is only used with `nodeport` listeners.

Use the `preferredNodePortAddressType` property in your listener configuration to specify the first address type checked as the node address. This property is useful, for example, if your deployment does not have DNS support, or you only want to expose a broker internally through an internal DNS or IP address. If an address of this type is found, it is used. If the preferred address type is not found, Strimzi proceeds through the types in the standard order of priority:

1. ExternalDNS
2. ExternalIP

3. Hostname
4. InternalDNS
5. InternalIP

*Example of an external listener configured with a preferred node port address type*

```
listeners:
#...
- name: external
  port: 9094
  type: nodeport
  tls: false
  configuration:
    preferredNodePortAddressType: InternalDNS
    # ...
# ...
```

### B.15.6. `useServiceDnsDomain`

The `useServiceDnsDomain` property is only used with `internal` listeners. It defines whether the fully-qualified DNS names that include the cluster service suffix (usually `.cluster.local`) are used. With `useServiceDnsDomain` set as `false`, the advertised addresses are generated without the service suffix; for example, `my-cluster-kafka-0.my-cluster-kafka-brokers.myproject.svc`. With `useServiceDnsDomain` set as `true`, the advertised addresses are generated with the service suffix; for example, `my-cluster-kafka-0.my-cluster-kafka-brokers.myproject.svc.cluster.local`. Default is `false`.

*Example of an internal listener configured to use the Service DNS domain*

```
listeners:
#...
- name: plain
  port: 9092
  type: internal
  tls: false
  configuration:
    useServiceDnsDomain: true
    # ...
# ...
```

If your Kubernetes cluster uses a different service suffix than `.cluster.local`, you can configure the suffix using the `KUBERNETES_SERVICE_DNS_DOMAIN` environment variable in the Cluster Operator configuration. See [Cluster Operator configuration](#) for more details.

### B.15.7. Schema reference

Property	Description
brokerCertChainAndKey  <b>CertAndKeySecretSource</b>	Reference to the <b>Secret</b> which holds the certificate and private key pair which will be used for this listener. The certificate can optionally contain the whole chain. This field can be used only with listeners with enabled TLS encryption.
externalTrafficPolicy  string (one of [Local, Cluster])	Specifies whether the service routes external traffic to node-local or cluster-wide endpoints. <b>Cluster</b> may cause a second hop to another node and obscures the client source IP. <b>Local</b> avoids a second hop for LoadBalancer and Nodeport type services and preserves the client source IP (when supported by the infrastructure). If unspecified, Kubernetes will use <b>Cluster</b> as the default. This field can be used only with <b>loadbalancer</b> or <b>nodeport</b> type listener.
loadBalancerSourceRanges  string array	A list of CIDR ranges (for example <b>10.0.0.0/8</b> or <b>130.211.204.1/32</b> ) from which clients can connect to load balancer type listeners. If supported by the platform, traffic through the loadbalancer is restricted to the specified CIDR ranges. This field is applicable only for loadbalancer type services and is ignored if the cloud provider does not support the feature. For more information, see <a href="https://v1-17.docs.kubernetes.io/docs/tasks/access-application-cluster/configure-cloud-provider-firewall/">https://v1-17.docs.kubernetes.io/docs/tasks/access-application-cluster/configure-cloud-provider-firewall/</a> . This field can be used only with <b>loadbalancer</b> type listener.
bootstrap  <b>GenericKafkaListenerConfigurationBootstrap</b>	Bootstrap configuration.
brokers  <b>GenericKafkaListenerConfigurationBroker</b> array	Per-broker configurations.
class	Configures the <b>Ingress</b> class that defines which <b>Ingress</b> controller will be used. If not set, the <b>Ingress</b> class is set to <b>nginx</b> . This field can be used only with <b>ingress</b> type listener.
string	

Property	Description
preferredNodePortAddressType	<p>Defines which address type should be used as the node address. Available types are: <code>ExternalDNS</code>, <code>ExternalIP</code>, <code>InternalDNS</code>, <code>InternalIP</code> and <code>Hostname</code>. By default, the addresses will be used in the following order (the first one found will be used): * <code>ExternalDNS</code> * <code>ExternalIP</code> * <code>InternalDNS</code> * <code>InternalIP</code> * <code>Hostname</code></p>
string (one of [ExternalDNS, ExternalIP, Hostname, InternalIP, InternalDNS])	<p>This field can be used to select the address type which will be used as the preferred type and checked first. In case no address will be found for this address type, the other types will be used in the default order. This field can be used only with <code>nodeport</code> type listener..</p>
useServiceDnsDomain	<p>Configures whether the Kubernetes service DNS domain should be used or not. If set to <code>true</code>, the generated addresses will contain the service DNS domain suffix (by default <code>.cluster.local</code>, can be configured using environment variable <code>KUBERNETES_SERVICE_DNS_DOMAIN</code>). Defaults to <code>false</code>. This field can be used only with <code>internal</code> type listener.</p>
boolean	

## B.16. CertAndKeySecretSource schema reference

Used in: `GenericKafkaListenerConfiguration`, `IngressListenerConfiguration`, `KafkaClientAuthenticationTls`, `KafkaListenerExternalConfiguration`, `NodePortListenerConfiguration`, `TlsListenerConfiguration`

Property	Description
certificate	The name of the file certificate in the Secret.
string	
key	The name of the private key in the Secret.
string	
secretName	The name of the Secret containing the certificate.
string	

## B.17. GenericKafkaListenerConfigurationBootstrap schema reference

Used in: [GenericKafkaListenerConfiguration](#)

Configures bootstrap service overrides for external listeners.

Broker service equivalents of `nodePort`, `host`, `loadBalancerIP` and `annotations` properties are configured in the [GenericKafkaListenerConfigurationBroker schema](#).

### B.17.1. alternativeNames

You can specify alternative names for the bootstrap service. The names are added to the broker certificates and can be used for TLS hostname verification. The `alternativeNames` property is applicable to all types of external listeners.

*Example of an external `route` listener configured with an additional bootstrap address*

```
listeners:  
#...  
- name: external  
  port: 9094  
  type: route  
  tls: true  
  authentication:  
    type: tls  
  configuration:  
    bootstrap:  
      alternativeNames:  
        - example.hostname1  
        - example.hostname2  
# ...
```

### B.17.2. host

The `host` property is used with `route` and `ingress` listeners to specify the hostnames used by the bootstrap and per-broker services.

A `host` property value is mandatory for `ingress` listener configuration, as the Ingress controller does not assign any hostnames automatically. Make sure that the hostnames resolve to the Ingress endpoints. Strimzi will not perform any validation that the requested hosts are available and properly routed to the Ingress endpoints.

*Example of host configuration for an ingress listener*

```
listeners:  
#...  
- name: external  
  port: 9094  
  type: ingress  
  tls: true  
  authentication:  
    type: tls  
  configuration:  
    bootstrap:  
      host: bootstrap.myingress.com  
    brokers:  
      - broker: 0  
        host: broker-0.myingress.com  
      - broker: 1  
        host: broker-1.myingress.com  
      - broker: 2  
        host: broker-2.myingress.com  
# ...
```

By default, **route** listener hosts are automatically assigned by OpenShift. However, you can override the assigned route hosts by specifying hosts.

Strimzi does not perform any validation that the requested hosts are available. You must ensure that they are free and can be used.

*Example of host configuration for a route listener*

```
# ...  
listeners:  
#...  
- name: external  
  port: 9094  
  type: route  
  tls: true  
  authentication:  
    type: tls  
  configuration:  
    bootstrap:  
      host: bootstrap.myrouter.com  
    brokers:  
      - broker: 0  
        host: broker-0.myrouter.com  
      - broker: 1  
        host: broker-1.myrouter.com  
      - broker: 2  
        host: broker-2.myrouter.com  
# ...
```

### B.17.3. nodePort

By default, the port numbers used for the bootstrap and broker services are automatically assigned by Kubernetes. You can override the assigned node ports for `nodeport` listeners by specifying the requested port numbers.

Strimzi does not perform any validation on the requested ports. You must ensure that they are free and available for use.

*Example of an external listener configured with overrides for node ports*

```
# ...
listeners:
#...
- name: external
  port: 9094
  type: nodeport
  tls: true
  authentication:
    type: tls
  configuration:
    bootstrap:
      nodePort: 32100
    brokers:
      - broker: 0
        nodePort: 32000
      - broker: 1
        nodePort: 32001
      - broker: 2
        nodePort: 32002
# ...
```

### B.17.4. loadBalancerIP

Use the `loadBalancerIP` property to request a specific IP address when creating a loadbalancer. Use this property when you need to use a loadbalancer with a specific IP address. The `loadBalancerIP` field is ignored if the cloud provider does not support the feature.

*Example of an external listener of type `loadbalancer` with specific loadbalancer IP address requests*

```
# ...
listeners:
#...
- name: external
  port: 9094
  type: loadbalancer
  tls: true
  authentication:
    type: tls
  configuration:
    bootstrap:
      loadBalancerIP: 172.29.3.10
    brokers:
      - broker: 0
        loadBalancerIP: 172.29.3.1
      - broker: 1
        loadBalancerIP: 172.29.3.2
      - broker: 2
        loadBalancerIP: 172.29.3.3
# ...
```

## B.17.5. [annotations](#)

Use the `annotations` property to add annotations to `loadbalancer`, `nodeport` or `ingress` listeners. You can use these annotations to instrument DNS tooling such as [External DNS](#), which automatically assigns DNS names to the loadbalancer services.

*Example of an external listener of type `loadbalancer` using annotations*

```
# ...
listeners:
#...
- name: external
  port: 9094
  type: loadbalancer
  tls: true
  authentication:
    type: tls
  configuration:
    bootstrap:
      annotations:
        external-dns.alpha.kubernetes.io/hostname: kafka-bootstrap.mydomain.com.
        external-dns.alpha.kubernetes.io/ttl: "60"
    brokers:
      - broker: 0
        annotations:
          external-dns.alpha.kubernetes.io/hostname: kafka-broker-0.mydomain.com.
          external-dns.alpha.kubernetes.io/ttl: "60"
      - broker: 1
        annotations:
          external-dns.alpha.kubernetes.io/hostname: kafka-broker-1.mydomain.com.
          external-dns.alpha.kubernetes.io/ttl: "60"
      - broker: 2
        annotations:
          external-dns.alpha.kubernetes.io/hostname: kafka-broker-2.mydomain.com.
          external-dns.alpha.kubernetes.io/ttl: "60"
    # ...
```

## B.17.6. Schema reference

Property	Description
alternativeNames string array	Additional alternative names for the bootstrap service. The alternative names will be added to the list of subject alternative names of the TLS certificates.
host string	The bootstrap host. This field will be used in the Ingress resource or in the Route resource to specify the desired hostname. This field can be used only with <code>route</code> (optional) or <code>ingress</code> (required) type listeners.
nodePort integer	Node port for the bootstrap service. This field can be used only with <code>nodeport</code> type listener.

Property	Description
loadBalancerIP	The loadbalancer is requested with the IP address specified in this field. This feature depends on whether the underlying cloud provider supports specifying the <code>loadBalancerIP</code> when a load balancer is created. This field is ignored if the cloud provider does not support the feature. This field can be used only with <code>loadbalancer</code> type listener.
string	
annotations	Annotations that will be added to the <code>Ingress</code> or <code>Service</code> resource. You can use this field to configure DNS providers such as External DNS. This field can be used only with <code>loadbalancer</code> , <code>nodeport</code> , or <code>ingress</code> type listeners.
map	

## B.18. GenericKafkaListenerConfigurationBroker schema reference

Used in: [GenericKafkaListenerConfiguration](#)

Configures broker service overrides for external listeners.

You can see example configuration for the `nodePort`, `host`, `loadBalancerIP` and `annotations` properties in the [GenericKafkaListenerConfigurationBootstrap schema](#), which configures bootstrap service overrides for external listeners.

### *Advertised addresses for brokers*

By default, Strimzi tries to automatically determine the hostnames and ports that your Kafka cluster advertises to its clients. This is not sufficient in all situations, because the infrastructure on which Strimzi is running might not provide the right hostname or port through which Kafka can be accessed.

You can specify a broker ID and customize the advertised hostname and port in the `configuration` property of the external listener. Strimzi will then automatically configure the advertised address in the Kafka brokers and add it to the broker certificates so it can be used for TLS hostname verification. Overriding the advertised host and ports is available for all types of external listeners.

*Example of an external `route` listener configured with overrides for advertised addresses*

```
listeners:  
#...  
- name: external  
  port: 9094  
  type: route  
  tls: true  
  authentication:  
    type: tls  
  configuration:  
    brokers:  
      - broker: 0  
        advertisedHost: example.hostname.0  
        advertisedPort: 12340  
      - broker: 1  
        advertisedHost: example.hostname.1  
        advertisedPort: 12341  
      - broker: 2  
        advertisedHost: example.hostname.2  
        advertisedPort: 12342  
# ...
```

## B.18.1. Schema reference

Property	Description
broker	ID of the kafka broker (broker identifier). Broker IDs start from 0 and correspond to the number of broker replicas.
integer	
advertisedHost	The host name which will be used in the brokers' <code>advertised.brokers</code> .
string	
advertisedPort	The port number which will be used in the brokers' <code>advertised.brokers</code> .
integer	
host	The broker host. This field will be used in the Ingress resource or in the Route resource to specify the desired hostname. This field can be used only with <code>route</code> (optional) or <code>ingress</code> (required) type listeners.
string	
nodePort	Node port for the per-broker service. This field can be used only with <code>nodeport</code> type listener.
integer	

Property	Description
loadBalancerIP	The loadbalancer is requested with the IP address specified in this field. This feature depends on whether the underlying cloud provider supports specifying the <code>loadBalancerIP</code> when a load balancer is created. This field is ignored if the cloud provider does not support the feature. This field can be used only with <code>loadbalancer</code> type listener.
string	
annotations	Annotations that will be added to the <code>Ingress</code> or <code>Service</code> resource. You can use this field to configure DNS providers such as External DNS. This field can be used only with <code>loadbalancer</code> , <code>nodeport</code> , or <code>ingress</code> type listeners.
map	

## B.19. KafkaListeners schema reference

The type `KafkaListeners` has been deprecated. Please use `GenericKafkaListener` instead.

Used in: `KafkaClusterSpec`

Refer to [previous documentation](#) for example configuration.

Property	Description
plain	Configures plain listener on port 9092.
<code>KafkaListenerPlain</code>	
tls	Configures TLS listener on port 9093.
<code>KafkaListenerTls</code>	
external	Configures external listener on port 9094. The type depends on the value of the <code>external.type</code> property within the given object, which must be one of [route, loadbalancer, nodeport, ingress].
<code>KafkaListenerExternalRoute</code> , <code>KafkaListenerExternalLoadBalancer</code> , <code>KafkaListenerExternalNodePort</code> , <code>KafkaListenerExternalIngress</code>	

## B.20. KafkaListenerPlain schema reference

Used in: `KafkaListeners`

Property	Description
authentication	Authentication configuration for this listener. Since this listener does not use TLS transport you cannot configure an authentication with <code>type: tls</code> . The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, oauth].
<code>KafkaListenerAuthenticationTls</code> , <code>KafkaListenerAuthenticationScramSha512</code> , <code>KafkaListenerAuthenticationOAuth</code>	
networkPolicyPeers	List of peers which should be able to connect to this listener. Peers in this list are combined using a logical OR operation. If this field is empty or missing, all connections will be allowed for this listener. If this field is present and contains at least one item, the listener only allows the traffic which matches at least one item in this list. For more information, see the <a href="#">external documentation for networking.k8s.io/v1 networkpolicypeer</a> .
<code>NetworkPolicyPeer</code> array	

## B.21. KafkaListenerTls schema reference

Used in: [KafkaListeners](#)

Property	Description
authentication	Authentication configuration for this listener. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, oauth].
<code>KafkaListenerAuthenticationTls</code> , <code>KafkaListenerAuthenticationScramSha512</code> , <code>KafkaListenerAuthenticationOAuth</code>	
configuration	Configuration of TLS listener.
<code>TlsListenerConfiguration</code>	
networkPolicyPeers	List of peers which should be able to connect to this listener. Peers in this list are combined using a logical OR operation. If this field is empty or missing, all connections will be allowed for this listener. If this field is present and contains at least one item, the listener only allows the traffic which matches at least one item in this list. For more information, see the <a href="#">external documentation for networking.k8s.io/v1 networkpolicypeer</a> .
<code>NetworkPolicyPeer</code> array	

## B.22. TlsListenerConfiguration schema reference

Used in: [KafkaListenerTls](#)

Property	Description
brokerCertChainAndKey	Reference to the <a href="#">Secret</a> which holds the certificate and private key pair. The certificate can optionally contain the whole chain.
CertAndKeySecretSource	

## B.23. KafkaListenerExternalRoute schema reference

Used in: [KafkaListeners](#)

The `type` property is a discriminator that distinguishes use of the [KafkaListenerExternalRoute](#) type from [KafkaListenerExternalLoadBalancer](#), [KafkaListenerExternalNodePort](#), [KafkaListenerExternalIngress](#). It must have the value `route` for the type [KafkaListenerExternalRoute](#).

Property	Description
type	Must be <code>route</code> .
string	
authentication	Authentication configuration for Kafka brokers. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, oauth].
KafkaListenerAuthenticationTls, KafkaListenerAuthenticationScramSha512, KafkaListenerAuthenticationOAuth	
overrides	Overrides for external bootstrap and broker services and externally advertised addresses.
RouteListenerOverride	
configuration	External listener configuration.
KafkaListenerExternalConfiguration	
networkPolicyPeers	List of peers which should be able to connect to this listener. Peers in this list are combined using a logical OR operation. If this field is empty or missing, all connections will be allowed for this listener. If this field is present and contains at least one item, the listener only allows the traffic which matches at least one item in this list. For more information, see the <a href="#">external documentation for networking.k8s.io/v1 networkpolicypeer</a> .
NetworkPolicyPeer array	

## B.24. RouteListenerOverride schema reference

Used in: [KafkaListenerExternalRoute](#)

Property	Description
bootstrap	External bootstrap service configuration.
<a href="#">RouteListenerBootstrapOverride</a>	
brokers	External broker services configuration.
<a href="#">RouteListenerBrokerOverride</a> array	

## B.25. RouteListenerBootstrapOverride schema reference

Used in: [RouteListenerOverride](#)

Property	Description
address	
string	Additional address name for the bootstrap service. The address will be added to the list of subject alternative names of the TLS certificates.
host	
string	Host for the bootstrap route. This field will be used in the <code>spec.host</code> field of the OpenShift Route.

## B.26. RouteListenerBrokerOverride schema reference

Used in: [RouteListenerOverride](#)

Property	Description
broker	Id of the kafka broker (broker identifier).
integer	
advertisedHost	The host name which will be used in the brokers' <code>advertised.brokers</code> .
string	
advertisedPort	The port number which will be used in the brokers' <code>advertised.brokers</code> .
integer	
host	Host for the broker route. This field will be used in the <code>spec.host</code> field of the OpenShift Route.
string	

## B.27. KafkaListenerExternalConfiguration schema reference

Used in: [KafkaListenerExternalLoadBalancer](#), [KafkaListenerExternalRoute](#)

Property	Description
brokerCertChainAndKey	Reference to the <a href="#">Secret</a> which holds the certificate and private key pair. The certificate can optionally contain the whole chain.
CertAndKeySecretSource	

## B.28. KafkaListenerExternalLoadBalancer schema reference

Used in: [KafkaListeners](#)

The `type` property is a discriminator that distinguishes use of the [KafkaListenerExternalLoadBalancer](#) type from [KafkaListenerExternalRoute](#), [KafkaListenerExternalNodePort](#), [KafkaListenerExternalIngress](#). It must have the value `loadbalancer` for the type [KafkaListenerExternalLoadBalancer](#).

Property	Description
type	Must be <code>loadbalancer</code> .
string	
authentication	Authentication configuration for Kafka brokers. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, oauth].
KafkaListenerAuthenticationTls, KafkaListenerAuthenticationScramSha512, KafkaListenerAuthenticationOAuth	
overrides	Overrides for external bootstrap and broker services and externally advertised addresses.
LoadBalancerListenerOverride	
configuration	External listener configuration.
KafkaListenerExternalConfiguration	
networkPolicyPeers	List of peers which should be able to connect to this listener. Peers in this list are combined using a logical OR operation. If this field is empty or missing, all connections will be allowed for this listener. If this field is present and contains at least one item, the listener only allows the traffic which matches at least one item in this list. For more information, see the <a href="#">external documentation for networking.k8s.io/v1 networkpolicypeer</a> .
NetworkPolicyPeer array	

Property	Description
tls	Enables TLS encryption on the listener. By default set to <code>true</code> for enabled TLS encryption.
boolean	

## B.29. LoadBalancerListenerOverride schema reference

Used in: [KafkaListenerExternalLoadBalancer](#)

Property	Description
bootstrap	External bootstrap service configuration.
<a href="#">LoadBalancerListenerBootstrapOverride</a>	
brokers	External broker services configuration.
<a href="#">LoadBalancerListenerBrokerOverride</a> array	

## B.30. LoadBalancerListenerBootstrapOverride schema reference

Used in: [LoadBalancerListenerOverride](#)

Property	Description
address	Additional address name for the bootstrap service. The address will be added to the list of subject alternative names of the TLS certificates.
string	
dnsAnnotations	Annotations that will be added to the <a href="#">Service</a> resource. You can use this field to configure DNS providers such as External DNS.
map	
loadBalancerIP	The loadbalancer is requested with the IP address specified in this field. This feature depends on whether the underlying cloud provider supports specifying the <a href="#">loadBalancerIP</a> when a load balancer is created. This field is ignored if the cloud provider does not support the feature.
string	

## B.31. LoadBalancerListenerBrokerOverride schema reference

Used in: [LoadBalancerListenerOverride](#)

Property	Description
broker	Id of the kafka broker (broker identifier).
integer	
advertisedHost	The host name which will be used in the brokers' <code>advertised.brokers</code> .
string	
advertisedPort	The port number which will be used in the brokers' <code>advertised.brokers</code> .
integer	
dnsAnnotations	Annotations that will be added to the <code>Service</code> resources for individual brokers. You can use this field to configure DNS providers such as External DNS.
map	
loadBalancerIP	The loadbalancer is requested with the IP address specified in this field. This feature depends on whether the underlying cloud provider supports specifying the <code>loadBalancerIP</code> when a load balancer is created. This field is ignored if the cloud provider does not support the feature.
string	

## B.32. KafkaListenerExternalNodePort schema reference

Used in: [KafkaListeners](#)

The `type` property is a discriminator that distinguishes use of the `KafkaListenerExternalNodePort` type from `KafkaListenerExternalRoute`, `KafkaListenerExternalLoadBalancer`, `KafkaListenerExternalIngress`. It must have the value `nodeport` for the type `KafkaListenerExternalNodePort`.

Property	Description
type	Must be <code>nodeport</code> .
string	
authentication	Authentication configuration for Kafka brokers. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, oauth].
KafkaListenerAuthenticationTls, KafkaListenerAuthenticationScramSha512, KafkaListenerAuthenticationOAuth	
overrides	Overrides for external bootstrap and broker services and externally advertised addresses.
NodePortListenerOverride	
configuration	External listener configuration.
NodePortListenerConfiguration	

Property	Description
networkPolicyPeers	List of peers which should be able to connect to this listener. Peers in this list are combined using a logical OR operation. If this field is empty or missing, all connections will be allowed for this listener. If this field is present and contains at least one item, the listener only allows the traffic which matches at least one item in this list. For more information, see the <a href="#">external documentation for networking.k8s.io/v1 networkpolicypeer</a> .
NetworkPolicyPeer array	
tls	Enables TLS encryption on the listener. By default set to <code>true</code> for enabled TLS encryption.
boolean	

## B.33. NodePortListenerOverride schema reference

Used in: [KafkaListenerExternalNodePort](#)

Property	Description
bootstrap	External bootstrap service configuration.
<a href="#">NodePortListenerBootstrapOverride</a>	
brokers	External broker services configuration.
<a href="#">NodePortListenerBrokerOverride</a> array	

## B.34. NodePortListenerBootstrapOverride schema reference

Used in: [NodePortListenerOverride](#)

Property	Description
address	Additional address name for the bootstrap service. The address will be added to the list of subject alternative names of the TLS certificates.
string	
dnsAnnotations	Annotations that will be added to the <a href="#">Service</a> resource. You can use this field to configure DNS providers such as External DNS.
map	
nodePort	Node port for the bootstrap service.
integer	

## B.35. NodePortListenerBrokerOverride schema reference

Used in: [NodePortListenerOverride](#)

Property	Description
broker	Id of the kafka broker (broker identifier).
integer	
advertisedHost	The host name which will be used in the brokers' <a href="#">advertised.brokers</a> .
string	
advertisedPort	The port number which will be used in the brokers' <a href="#">advertised.brokers</a> .
integer	
nodePort	Node port for the broker service.
integer	
dnsAnnotations	Annotations that will be added to the <a href="#">Service</a> resources for individual brokers. You can use this field to configure DNS providers such as External DNS.
map	

## B.36. NodePortListenerConfiguration schema reference

Used in: [KafkaListenerExternalNodePort](#)

Property	Description
brokerCertChainAndKey	Reference to the <a href="#">Secret</a> which holds the certificate and private key pair. The certificate can optionally contain the whole chain.
<a href="#">CertAndKeySecretSource</a>	
preferredAddressType	Defines which address type should be used as the node address. Available types are: <a href="#">ExternalDNS</a> , <a href="#">ExternalIP</a> , <a href="#">InternalDNS</a> , <a href="#">InternalIP</a> and <a href="#">Hostname</a> . By default, the addresses will be used in the following order (the first one found will be used): * <a href="#">ExternalDNS</a> * <a href="#">ExternalIP</a> * <a href="#">InternalDNS</a> * <a href="#">InternalIP</a> * <a href="#">Hostname</a>
string (one of [ExternalDNS, ExternalIP, Hostname, InternalIP, InternalDNS])	This field can be used to select the address type which will be used as the preferred type and checked first. In case no address will be found for this address type, the other types will be used in the default order..

## B.37. KafkaListenerExternalIngress schema reference

Used in: [KafkaListeners](#)

The `type` property is a discriminator that distinguishes use of the `KafkaListenerExternalIngress` type from `KafkaListenerExternalRoute`, `KafkaListenerExternalLoadBalancer`, `KafkaListenerExternalNodePort`. It must have the value `ingress` for the type `KafkaListenerExternalIngress`.

Property	Description
<code>type</code>	Must be <code>ingress</code> .
<code>string</code>	
<code>authentication</code>	Authentication configuration for Kafka brokers. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, oauth].
<code>class</code>	Configures the <code>Ingress</code> class that defines which <code>Ingress</code> controller will be used. If not set, the <code>Ingress</code> class is set to <code>nginx</code> .
<code>string</code>	
<code>configuration</code>	External listener configuration.
<a href="#">IngressListenerConfiguration</a>	
<code>networkPolicyPeers</code>	List of peers which should be able to connect to this listener. Peers in this list are combined using a logical OR operation. If this field is empty or missing, all connections will be allowed for this listener. If this field is present and contains at least one item, the listener only allows the traffic which matches at least one item in this list. For more information, see the <a href="#">external documentation for networking.k8s.io/v1 networkpolicypeer</a> .
<code>NetworkPolicyPeer</code> array	

## B.38. IngressListenerConfiguration schema reference

Used in: [KafkaListenerExternalIngress](#)

Property	Description
<code>bootstrap</code>	External bootstrap ingress configuration.
<a href="#">IngressListenerBootstrapConfiguration</a>	
<code>brokers</code>	External broker ingress configuration.
<a href="#">IngressListenerBrokerConfiguration</a> array	

Property	Description
brokerCertChainAndKey	Reference to the <a href="#">Secret</a> which holds the certificate and private key pair. The certificate can optionally contain the whole chain.
CertAndKeySecretSource	

## B.39. IngressListenerBootstrapConfiguration schema reference

Used in: [IngressListenerConfiguration](#)

Property	Description
address	Additional address name for the bootstrap service. The address will be added to the list of subject alternative names of the TLS certificates.
string	
dnsAnnotations	Annotations that will be added to the <a href="#">Ingress</a> resource. You can use this field to configure DNS providers such as External DNS.
map	
host	Host for the bootstrap route. This field will be used in the Ingress resource.
string	

## B.40. IngressListenerBrokerConfiguration schema reference

Used in: [IngressListenerConfiguration](#)

Property	Description
broker	Id of the kafka broker (broker identifier).
integer	
advertisedHost	The host name which will be used in the brokers' <a href="#">advertised.brokers</a> .
string	
advertisedPort	The port number which will be used in the brokers' <a href="#">advertised.brokers</a> .
integer	
host	Host for the broker ingress. This field will be used in the Ingress resource.
string	
dnsAnnotations	Annotations that will be added to the <a href="#">Ingress</a> resources for individual brokers. You can use this field to configure DNS providers such as External DNS.
map	

## B.41. KafkaAuthorizationSimple schema reference

Used in: [KafkaClusterSpec](#)

Simple authorization in Strimzi uses the [AclAuthorizer](#) plugin, the default Access Control Lists (ACLs) authorization plugin provided with Apache Kafka. ACLs allow you to define which users have access to which resources at a granular level.

Configure the [Kafka](#) custom resource to use simple authorization. Set the `type` property in the `authorization` section to the value `simple`, and configure a list of super users.

Access rules are configured for the [KafkaUser](#), as described in the [ACLRule schema reference](#).

### B.41.1. superUsers

A list of user principals treated as super users, so that they are always allowed without querying ACL rules. For more information see [Kafka authorization](#).

*An example of simple authorization configuration*

```
authorization:  
  type: simple  
  superUsers:  
    - CN=client_1  
    - user_2  
    - CN=client_3
```

**NOTE** The `superUser` configuration option in the `config` property in `Kafka.spec.kafka` is ignored. Designate super users in the `authorization` property instead. For more information, see [Kafka broker configuration](#).

### B.41.2. Schema reference

The `type` property is a discriminator that distinguishes use of the [KafkaAuthorizationSimple](#) type from [KafkaAuthorizationOpa](#), [KafkaAuthorizationKeycloak](#). It must have the value `simple` for the type [KafkaAuthorizationSimple](#).

Property	Description
<code>type</code>	Must be <code>simple</code> .
<code>string</code>	
<code>superUsers</code>	List of super users. Should contain list of user principals which should get unlimited access rights.
<code>string array</code>	

## B.42. KafkaAuthorizationOpa schema reference

Used in: [KafkaClusterSpec](#)

To use [Open Policy Agent](#) authorization, set the `type` property in the `authorization` section to the value `opa`, and configure OPA properties as required.

### B.42.1. url

The URL used to connect to the Open Policy Agent server. The URL has to include the policy which will be queried by the authorizer. **Required**.

### B.42.2. allowOnError

Defines whether a Kafka client should be allowed or denied by default when the authorizer fails to query the Open Policy Agent, for example, when it is temporarily unavailable. Defaults to `false` - all actions will be denied.

### B.42.3. initialCacheCapacity

Initial capacity of the local cache used by the authorizer to avoid querying the Open Policy Agent for every request. Defaults to `5000`.

### B.42.4. maximumCacheSize

Maximum capacity of the local cache used by the authorizer to avoid querying the Open Policy Agent for every request. Defaults to `50000`.

### B.42.5. expireAfterMs

The expiration of the records kept in the local cache to avoid querying the Open Policy Agent for every request. Defines how often the cached authorization decisions are reloaded from the Open Policy Agent server. In milliseconds. Defaults to `3600000` milliseconds (1 hour).

### B.42.6. superUsers

A list of user principals treated as super users, so that they are always allowed without querying the open Policy Agent policy. For more information see [Kafka authorization](#).

## An example of Open Policy Agent authorizer configuration

```
authorization:  
  type: opa  
  url: http://opa:8181/v1/data/kafka/allow  
  allowOnError: false  
  initialCacheCapacity: 1000  
  maximumCacheSize: 10000  
  expireAfterMs: 60000  
  superUsers:  
    - CN=fred  
    - sam  
    - CN=edward
```

### B.42.7. Schema reference

The `type` property is a discriminator that distinguishes use of the `KafkaAuthorizationOpa` type from `KafkaAuthorizationSimple`, `KafkaAuthorizationKeycloak`. It must have the value `opa` for the type `KafkaAuthorizationOpa`.

Property	Description
type	Must be <code>opa</code> .
string	
url	The URL used to connect to the Open Policy Agent server. The URL has to include the policy which will be queried by the authorizer. This option is required.
string	
allowOnError	Defines whether a Kafka client should be allowed or denied by default when the authorizer fails to query the Open Policy Agent, for example, when it is temporarily unavailable). Defaults to <code>false</code> - all actions will be denied.
boolean	
initialCacheCapacity	Initial capacity of the local cache used by the authorizer to avoid querying the Open Policy Agent for every request. Defaults to <code>5000</code> .
integer	
maximumCacheSize	Maximum capacity of the local cache used by the authorizer to avoid querying the Open Policy Agent for every request. Defaults to <code>50000</code> .
integer	
expireAfterMs	The expiration of the records kept in the local cache to avoid querying the Open Policy Agent for every request. Defines how often the cached authorization decisions are reloaded from the Open Policy Agent server. In milliseconds. Defaults to <code>3600000</code> .
integer	

Property	Description
superUsers	
string array	List of super users, which is specifically a list of user principals that have unlimited access rights.

## B.43. KafkaAuthorizationKeycloak schema reference

Used in: [KafkaClusterSpec](#)

The `type` property is a discriminator that distinguishes use of the [KafkaAuthorizationKeycloak](#) type from [KafkaAuthorizationSimple](#), [KafkaAuthorizationOpa](#). It must have the value `keycloak` for the type [KafkaAuthorizationKeycloak](#).

Property	Description
type	Must be <code>keycloak</code> .
string	
clientId	OAuth Client ID which the Kafka client can use to authenticate against the OAuth server and use the token endpoint URI.
string	
tokenEndpointUri	Authorization server token endpoint URI.
string	
tlsTrustedCertificates	Trusted certificates for TLS connection to the OAuth server.
<a href="#">CertSecretSource</a> array	
disableTlsHostnameVerification	Enable or disable TLS hostname verification. Default value is <code>false</code> .
boolean	
delegateToKafkaAcls	Whether authorization decision should be delegated to the 'Simple' authorizer if DENIED by Keycloak Authorization Services policies. Default value is <code>false</code> .
boolean	
grantsRefreshPeriodSeconds	The time between two consecutive grants refresh runs in seconds. The default value is 60.
integer	
grantsRefreshPoolSize	The number of threads to use to refresh grants for active sessions. The more threads, the more parallelism, so the sooner the job completes. However, using more threads places a heavier load on the authorization server. The default value is 5.
integer	
superUsers	List of super users. Should contain list of user principals which should get unlimited access rights.
string array	

## B.44. Rack schema reference

Used in: [KafkaClusterSpec](#), [KafkaConnectS2ISpec](#), [KafkaConnectSpec](#)

Configures rack awareness to spread partition replicas across different racks.

A *rack* can represent an availability zone, data center, or an actual rack in your data center. By configuring a `rack` for a Kafka cluster, consumers can fetch data from the closest replica. This is useful for reducing the load on your network when a Kafka cluster spans multiple datacenters.

To configure Kafka brokers for rack awareness, you specify a `topologyKey` value to match the label of the cluster node used by Kubernetes when scheduling Kafka broker pods to nodes.

If the Kubernetes cluster is running on a cloud provider platform, the label must represent the availability zone where the node is running. Usually, nodes are labeled with the `topology.kubernetes.io/zone` label (or `failure-domain.beta.kubernetes.io/zone` on older Kubernetes versions), which can be used as the `topologyKey` value.

The rack awareness configuration spreads the broker pods and partition replicas across zones, improving resiliency, and also sets a `broker.rack` configuration for each Kafka broker. The `broker.rack` configuration assigns a rack ID to each broker.

Consult your Kubernetes administrator regarding the node label that represents the zone or rack into which the node is deployed.

*Example rack configuration for Kafka*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
    rack:
      topologyKey: topology.kubernetes.io/zone
    config:
      # ...
      replica.selector.class: org.apache.kafka.common.replica.RackAwareReplicaSelector
    # ...
```

Use the `RackAwareReplicaSelector` implementation for the Kafka `ReplicaSelector` plugin if you want clients to consume from the closest replica. The `ReplicaSelector` plugin provides the logic that enables clients to consume from the nearest replica. Specify `RackAwareReplicaSelector` for the `replica.selector.class` to switch from the default implementation. The default implementation uses `LeaderSelector` to always select the leader replica for the client. By switching from the leader replica to the replica follower, there is some cost to latency. If required, you can also customize your own implementation.

For clients, including Kafka Connect, you specify the same topology key as the broker that the client

will use to consume messages.

#### Example `rack` configuration for Kafka Connect

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
# ...
spec:
  kafka:
    # ...
    rack:
      topologyKey: topology.kubernetes.io/zone
    # ...
```

The client is assigned a `client.rack` ID.

`RackAwareReplicaSelector` associates matching `broker.rack` and `client.rack` IDs, so the client can consume from the nearest replica.




Figure 6. Example showing client consuming from replicas in the same availability zone

If there are multiple replicas in the same rack, `RackAwareReplicaSelector` always selects the most up-to-date replica. If the rack ID is not specified, or if it cannot find a replica with the same rack ID, it will fall back to the leader replica.

For more information about Kubernetes node labels, see [Well-Known Labels, Annotations and Taints](#).

## B.44.1. Schema reference

Property	Description
topologyKey	A key that matches labels assigned to the Kubernetes cluster nodes. The value of the label is used to set the broker's <code>broker.rack</code> config and <code>client.rack</code> in Kafka Connect.
string	

## B.45. Probe schema reference

Used in: `CruiseControlSpec`, `EntityTopicOperatorSpec`, `EntityUserOperatorSpec`, `KafkaBridgeSpec`, `KafkaClusterSpec`, `KafkaConnectS2ISpec`, `KafkaConnectSpec`, `KafkaExporterSpec`, `KafkaMirrorMaker2Spec`, `KafkaMirrorMakerSpec`, `TlsSidecar`, `TopicOperatorSpec`, `ZookeeperClusterSpec`

Property	Description
failureThreshold	Minimum consecutive failures for the probe to be considered failed after having succeeded. Defaults to 3. Minimum value is 1.
integer	
initialDelaySeconds	The initial delay before first the health is first checked. Default to 15 seconds. Minimum value is 0.
integer	
periodSeconds	How often (in seconds) to perform the probe. Default to 10 seconds. Minimum value is 1.
integer	
successThreshold	Minimum consecutive successes for the probe to be considered successful after having failed. Defaults to 1. Must be 1 for liveness. Minimum value is 1.
integer	
timeoutSeconds	The timeout for each attempted health check. Default to 5 seconds. Minimum value is 1.
integer	

## B.46. JvmOptions schema reference

Used in: `CruiseControlSpec`, `EntityTopicOperatorSpec`, `EntityUserOperatorSpec`, `KafkaBridgeSpec`, `KafkaClusterSpec`, `KafkaConnectS2ISpec`, `KafkaConnectSpec`, `KafkaMirrorMaker2Spec`, `KafkaMirrorMakerSpec`, `TopicOperatorSpec`, `ZookeeperClusterSpec`

Property	Description
-XX	A map of -XX options to the JVM.
map	
-Xms	-Xms option to the JVM.
string	

Property	Description
-Xmx	-Xmx option to to the JVM.
string	
gcLoggingEnabled	Specifies whether the Garbage Collection logging is enabled. The default is false.
boolean	
javaSystemProperties	A map of additional system properties which will be passed using the <code>-D</code> option to the JVM.
SystemProperty array	

## B.47. SystemProperty schema reference

Used in: [JvmOptions](#)

Property	Description
name	The system property name.
string	
value	The system property value.
string	

## B.48. KafkaJmxOptions schema reference

Used in: [KafkaClusterSpec](#)

Property	Description
authentication	Authentication configuration for connecting to the Kafka JMX port. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [password].
KafkaJmxAuthenticationPassword	

## B.49. KafkaJmxAuthenticationPassword schema reference

Used in: [KafkaJmxOptions](#)

The `type` property is a discriminator that distinguishes use of the `KafkaJmxAuthenticationPassword` type from other subtypes which may be added in the future. It must have the value `password` for the type `KafkaJmxAuthenticationPassword`.

Property	Description
type	Must be <code>password</code> .
string	

## B.50. JmxPrometheusExporterMetrics schema reference

Used in: [CruiseControlSpec](#), [KafkaClusterSpec](#), [KafkaConnectS2ISpec](#), [KafkaConnectSpec](#), [KafkaMirrorMaker2Spec](#), [KafkaMirrorMakerSpec](#), [ZookeeperClusterSpec](#)

The `type` property is a discriminator that distinguishes use of the `JmxPrometheusExporterMetrics` type from other subtypes which may be added in the future. It must have the value `jmxPrometheusExporter` for the type `JmxPrometheusExporterMetrics`.

Property	Description
type	Must be <code>jmxPrometheusExporter</code> .
string	
valueFrom	ConfigMap where the Prometheus JMX Exporter configuration is stored. For details of the structure of this configuration, see the <a href="#">JMX Exporter documentation</a> .
<code>ExternalConfigurationMetrics</code>	

## B.51. ExternalConfigurationMetrics schema reference

Used in: [JmxPrometheusExporterMetrics](#)

Property	Description
configMapKeyRef	Reference to the key in the ConfigMap containing the metrics configuration. For more information, see the <a href="#">external documentation for core/v1 configmapkeyselector</a> .
ConfigMapKeySelector	

## B.52. InlineLogging schema reference

Used in: [CruiseControlSpec](#), [EntityTopicOperatorSpec](#), [EntityUserOperatorSpec](#), [KafkaBridgeSpec](#), [KafkaClusterSpec](#), [KafkaConnectS2ISpec](#), [KafkaConnectSpec](#), [KafkaMirrorMaker2Spec](#), [KafkaMirrorMakerSpec](#), [TopicOperatorSpec](#), [ZookeeperClusterSpec](#)

The `type` property is a discriminator that distinguishes use of the `InlineLogging` type from `ExternalLogging`. It must have the value `inline` for the type `InlineLogging`.

Property	Description
type	Must be <code>inline</code> .
string	
loggers	A Map from logger name to logger level.
map	

## B.53. ExternalLogging schema reference

Used in: `CruiseControlSpec`, `EntityTopicOperatorSpec`, `EntityUserOperatorSpec`, `KafkaBridgeSpec`, `KafkaClusterSpec`, `KafkaConnectS2ISpec`, `KafkaConnectSpec`, `KafkaMirrorMaker2Spec`, `KafkaMirrorMakerSpec`, `TopicOperatorSpec`, `ZookeeperClusterSpec`

The `type` property is a discriminator that distinguishes use of the `ExternalLogging` type from `InlineLogging`. It must have the value `external` for the type `ExternalLogging`.

Property	Description
<code>type</code>	Must be <code>external</code> .
<code>string</code>	
<code>name</code>	The name of the <code>ConfigMap</code> from which to get the logging configuration.
<code>string</code>	

## B.54. TlsSidecar schema reference

Used in: `CruiseControlSpec`, `EntityOperatorSpec`, `KafkaClusterSpec`, `TopicOperatorSpec`, `ZookeeperClusterSpec`

Configures a TLS sidecar, which is a container that runs in a pod, but serves a supporting purpose. In Strimzi, the TLS sidecar uses TLS to encrypt and decrypt communication between components and ZooKeeper.

The TLS sidecar is used in:

- Entity Operator
- Cruise Control

The TLS sidecar is configured using the `tlsSidecar` property in:

- `Kafka.spec.entityOperator`
- `Kafka.spec.cruiseControl`

The TLS sidecar supports the following additional options:

- `image`
- `resources`
- `logLevel`
- `readinessProbe`
- `livenessProbe`

The `resources` property specifies the `memory` and `CPU resources` allocated for the TLS sidecar.

The `image` property configures the `container image` which will be used.

The `readinessProbe` and `livenessProbe` properties configure `healthcheck probes` for the TLS sidecar.

The `logLevel` property specifies the logging level. The following logging levels are supported:

- emerg
- alert
- crit
- err
- warning
- notice
- info
- debug

The default value is *notice*.

### Example TLS sidecar configuration

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  # ...
  entityOperator:
    # ...
    tlsSidecar:
      resources:
        requests:
          cpu: 200m
          memory: 64Mi
        limits:
          cpu: 500m
          memory: 128Mi
    # ...
  cruiseControl:
    # ...
    tlsSidecar:
      image: my-org/my-image:latest
      resources:
        requests:
          cpu: 200m
          memory: 64Mi
        limits:
          cpu: 500m
          memory: 128Mi
      logLevel: debug
      readinessProbe:
        initialDelaySeconds: 15
        timeoutSeconds: 5
      livenessProbe:
        initialDelaySeconds: 15
        timeoutSeconds: 5
    # ...
```

### B.54.1. Schema reference

Property	Description
image	The docker image for the container.
string	
livenessProbe	Pod liveness checking.
Probe	

Property	Description
logLevel string (one of [emerg, debug, crit, err, alert, warning, notice, info])	The log level for the TLS sidecar. Default value is <a href="#">notice</a> .
readinessProbe	Pod readiness checking.
<b>Probe</b>	
resources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
<b>ResourceRequirements</b>	

## B.55. KafkaClusterTemplate schema reference

Used in: [KafkaClusterSpec](#)

Property	Description
statefulset	Template for Kafka <a href="#">StatefulSet</a> .
<b>StatefulSetTemplate</b>	
pod	Template for Kafka <a href="#">Pods</a> .
<b>PodTemplate</b>	
bootstrapService	Template for Kafka bootstrap <a href="#">Service</a> .
<b>ResourceTemplate</b>	
brokersService	Template for Kafka broker <a href="#">Service</a> .
<b>ResourceTemplate</b>	
externalBootstrapService	Template for Kafka external bootstrap <a href="#">Service</a> .
<b>ExternalServiceTemplate</b>	
perPodService	Template for Kafka per-pod <a href="#">Services</a> used for access from outside of Kubernetes.
<b>ExternalServiceTemplate</b>	
externalBootstrapRoute	Template for Kafka external bootstrap <a href="#">Route</a> .
<b>ResourceTemplate</b>	
perPodRoute	Template for Kafka per-pod <a href="#">Routes</a> used for access from outside of OpenShift.
<b>ResourceTemplate</b>	
externalBootstrapIngress	Template for Kafka external bootstrap <a href="#">Ingress</a> .
<b>ResourceTemplate</b>	
perPodIngress	Template for Kafka per-pod <a href="#">Ingress</a> used for access from outside of Kubernetes.
<b>ResourceTemplate</b>	
persistentVolumeClaim	Template for all Kafka <a href="#">PersistentVolumeClaims</a> .
<b>ResourceTemplate</b>	

Property	Description
podDisruptionBudget	Template for Kafka <a href="#">PodDisruptionBudget</a> .
<a href="#">PodDisruptionBudgetTemplate</a>	
kafkaContainer	Template for the Kafka broker container.
<a href="#">ContainerTemplate</a>	
tlsSidecarContainer	The property <a href="#">tlsSidecarContainer</a> has been deprecated. Template for the Kafka broker TLS sidecar container.
<a href="#">ContainerTemplate</a>	
initContainer	Template for the Kafka init container.
<a href="#">ContainerTemplate</a>	
clusterCaCert	Template for Secret with Kafka Cluster certificate public key.
<a href="#">ResourceTemplate</a>	

## B.56. StatefulSetTemplate schema reference

Used in: [KafkaClusterTemplate](#), [ZookeeperClusterTemplate](#)

Property	Description
metadata	Metadata applied to the resource.
<a href="#">MetadataTemplate</a>	
podManagementPolicy	PodManagementPolicy which will be used for this StatefulSet. Valid values are <a href="#">Parallel</a> and <a href="#">OrderedReady</a> . Defaults to <a href="#">Parallel</a> .
string (one of [OrderedReady, Parallel])	

## B.57. MetadataTemplate schema reference

Used in: [DeploymentTemplate](#), [ExternalServiceTemplate](#), [PodDisruptionBudgetTemplate](#), [PodTemplate](#), [ResourceTemplate](#), [StatefulSetTemplate](#)

Labels and Annotations are used to identify and organize resources, and are configured in the metadata property.

For example:

```

# ...
template:
  statefulset:
    metadata:
      labels:
        label1: value1
        label2: value2
      annotations:
        annotation1: value1
        annotation2: value2
# ...

```

The `labels` and `annotations` fields can contain any labels or annotations that do not contain the reserved string `strimzi.io`. Labels and annotations containing `strimzi.io` are used internally by Strimzi and cannot be configured.

### B.57.1. Schema reference

Property	Description
labels	Labels added to the resource template. Can be applied to different resources such as <a href="#">StatefulSets</a> , <a href="#">Deployments</a> , <a href="#">Pods</a> , and <a href="#">Services</a> .
map	
annotations	Annotations added to the resource template. Can be applied to different resources such as <a href="#">StatefulSets</a> , <a href="#">Deployments</a> , <a href="#">Pods</a> , and <a href="#">Services</a> .
map	

## B.58. PodTemplate schema reference

Used in: [CruiseControlTemplate](#), [EntityOperatorTemplate](#), [JmxTransTemplate](#), [KafkaBridgeTemplate](#), [KafkaClusterTemplate](#), [KafkaConnectTemplate](#), [KafkaExporterTemplate](#), [KafkaMirrorMakerTemplate](#), [ZookeeperClusterTemplate](#)

*Example PodTemplate configuration*

```
# ...
template:
  pod:
    metadata:
      labels:
        label1: value1
      annotations:
        anno1: value1
    imagePullSecrets:
      - name: my-docker-credentials
    securityContext:
      runAsUser: 1000001
      fsGroup: 0
    terminationGracePeriodSeconds: 120
# ...
```

### B.58.1. hostAliases

Use the `hostAliases` property to specify a list of hosts and IP addresses, which are injected into the `/etc/hosts` file of the pod.

This configuration is especially useful for Kafka Connect or MirrorMaker when a connection outside of the cluster is also requested by users.

*Example hostAliases configuration*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
#...
spec:
  # ...
  template:
    pod:
      hostAliases:
        - ip: "192.168.1.86"
          hostnames:
            - "my-host-1"
            - "my-host-2"
#...
```

### B.58.2. Schema reference

Property	Description
metadata	Metadata applied to the resource.
MetadataTemplate	

Property	Description
imagePullSecrets	List of references to secrets in the same namespace to use for pulling any of the images used by this Pod. When the <code>STRIMZI_IMAGE_PULL_SECRETS</code> environment variable in Cluster Operator and the <code>imagePullSecrets</code> option are specified, only the <code>imagePullSecrets</code> variable is used and the <code>STRIMZI_IMAGE_PULL_SECRETS</code> variable is ignored.
<code>LocalObjectReference</code> array	For more information, see the <a href="#">external documentation for core/v1 localobjectreference</a> .
securityContext	Configures pod-level security attributes and common container settings. For more information, see the <a href="#">external documentation for core/v1 podsecuritycontext</a> .
<code>PodSecurityContext</code>	
terminationGracePeriodSeconds	The grace period is the duration in seconds after the processes running in the pod are sent a termination signal, and the time when the processes are forcibly halted with a kill signal. Set this value to longer than the expected cleanup time for your process. Value must be a non-negative integer. A zero value indicates delete immediately. You might need to increase the grace period for very large Kafka clusters, so that the Kafka brokers have enough time to transfer their work to another broker before they are terminated. Defaults to 30 seconds.
integer	
affinity	The pod's affinity rules. For more information, see the <a href="#">external documentation for core/v1 affinity</a> .
<code>Affinity</code>	
tolerations	The pod's tolerations. For more information, see the <a href="#">external documentation for core/v1 toleration</a> .
<code>Toleration</code> array	
priorityClassName	The name of the priority class used to assign priority to the pods. For more information about priority classes, see <a href="#">Pod Priority and Preemption</a> .
string	
schedulerName	The name of the scheduler used to dispatch this <code>Pod</code> . If not specified, the default scheduler will be used.
string	

Property	Description
hostAliases	The pod's HostAliases. HostAliases is an optional list of hosts and IPs that will be injected into the pod's hosts file if specified. For more information, see the <a href="#">external documentation for core/v1 HostAlias</a> .
HostAlias array	
topologySpreadConstraints	The pod's topology spread constraints. For more information, see the <a href="#">external documentation for core/v1 TopologySpreadConstraint</a> .
TopologySpreadConstraint array	

## B.59. ResourceTemplate schema reference

Used in: [CruiseControlTemplate](#), [EntityOperatorTemplate](#), [JmxTransTemplate](#), [KafkaBridgeTemplate](#), [KafkaClusterTemplate](#), [KafkaConnectTemplate](#), [KafkaExporterTemplate](#), [KafkaUserTemplate](#), [ZookeeperClusterTemplate](#)

Property	Description
metadata	Metadata applied to the resource.
MetadataTemplate	

## B.60. ExternalServiceTemplate schema reference

Used in: [KafkaClusterTemplate](#)

When exposing Kafka outside of Kubernetes using loadbalancers or node ports, you can use properties, in addition to labels and annotations, to customize how a Service is created.

*An example showing customized external services*

```
# ...
template:
  externalBootstrapService:
    externalTrafficPolicy: Local
    loadBalancerSourceRanges:
      - 10.0.0.0/8
      - 88.208.76.87/32
  perPodService:
    externalTrafficPolicy: Local
    loadBalancerSourceRanges:
      - 10.0.0.0/8
      - 88.208.76.87/32
# ...
```

### B.60.1. Schema reference

Property	Description
metadata <small>MetadataTemplate</small>	Metadata applied to the resource.
externalTrafficPolicy <small>string (one of [Local, Cluster])</small>	<b>The property <code>externalTrafficPolicy</code> has been deprecated.</b> Specifies whether the service routes external traffic to node-local or cluster-wide endpoints. <code>Cluster</code> may cause a second hop to another node and obscures the client source IP. <code>Local</code> avoids a second hop for LoadBalancer and Nodeport type services and preserves the client source IP (when supported by the infrastructure). If unspecified, Kubernetes will use <code>Cluster</code> as the default.
loadBalancerSourceRanges <small>string array</small>	<b>The property <code>loadBalancerSourceRanges</code> has been deprecated.</b> A list of CIDR ranges (for example <code>10.0.0.0/8</code> or <code>130.211.204.1/32</code> ) from which clients can connect to load balancer type listeners. If supported by the platform, traffic through the loadbalancer is restricted to the specified CIDR ranges. This field is applicable only for loadbalancer type services and is ignored if the cloud provider does not support the feature. For more information, see <a href="https://v1-17.docs.kubernetes.io/docs/tasks/access-application-cluster/configure-cloud-provider-firewall/">https://v1-17.docs.kubernetes.io/docs/tasks/access-application-cluster/configure-cloud-provider-firewall/</a> .

## B.61. PodDisruptionBudgetTemplate schema reference

Used in: `CruiseControlTemplate`, `KafkaBridgeTemplate`, `KafkaClusterTemplate`, `KafkaConnectTemplate`, `KafkaMirrorMakerTemplate`, `ZookeeperClusterTemplate`

Strimzi creates a `PodDisruptionBudget` for every new `StatefulSet` or `Deployment`. By default, pod disruption budgets only allow a single pod to be unavailable at a given time. You can increase the amount of unavailable pods allowed by changing the default value of the `maxUnavailable` property in the `PodDisruptionBudget.spec` resource.

An example of [PodDisruptionBudget template](#)

```
# ...
template:
  podDisruptionBudget:
    metadata:
      labels:
        key1: label1
        key2: label2
      annotations:
        key1: label1
        key2: label2
    maxUnavailable: 1
# ...
```

### B.61.1. Schema reference

Property	Description
metadata	Metadata to apply to the <a href="#">PodDisruptionBugetTemplate</a> resource.
MetadataTemplate	
maxUnavailable	Maximum number of unavailable pods to allow automatic Pod eviction. A Pod eviction is allowed when the <code>maxUnavailable</code> number of pods or fewer are unavailable after the eviction. Setting this value to 0 prevents all voluntary evictions, so the pods must be evicted manually. Defaults to 1.
integer	

## B.62. ContainerTemplate schema reference

Used in: [CruiseControlTemplate](#), [EntityOperatorTemplate](#), [JmxTransTemplate](#), [KafkaBridgeTemplate](#), [KafkaClusterTemplate](#), [KafkaConnectTemplate](#), [KafkaExporterTemplate](#), [KafkaMirrorMakerTemplate](#), [ZookeeperClusterTemplate](#)

You can set custom security context and environment variables for a container.

The environment variables are defined under the `env` property as a list of objects with `name` and `value` fields. The following example shows two custom environment variables and a custom security context set for the Kafka broker containers:

```

# ...
template:
  kafkaContainer:
    env:
      - name: EXAMPLE_ENV_1
        value: example.env.one
      - name: EXAMPLE_ENV_2
        value: example.env.two
    securityContext:
      runAsUser: 2000
# ...

```

Environment variables prefixed with `KAFKA_` are internal to Strimzi and should be avoided. If you set a custom environment variable that is already in use by Strimzi, it is ignored and a warning is recorded in the log.

### B.62.1. Schema reference

Property	Description
env	Environment variables which should be applied to the container.
<code>ContainerEnvVar</code> array	
securityContext	Security context for the container. For more information, see the <a href="#">external documentation for core/v1 securitycontext</a> .
<code>SecurityContext</code>	

## B.63. `ContainerEnvVar` schema reference

Used in: [ContainerTemplate](#)

Property	Description
name	The environment variable key.
string	
value	The environment variable value.
string	

## B.64. `ZookeeperClusterSpec` schema reference

Used in: [KafkaSpec](#)

Configures a ZooKeeper cluster.

### B.64.1. `config`

Use the `config` properties to configure ZooKeeper options as keys.

Standard Apache ZooKeeper configuration may be provided, restricted to those properties not managed directly by Strimzi.

Configuration options that cannot be configured relate to:

- Security (Encryption, Authentication, and Authorization)
- Listener configuration
- Configuration of data directories
- ZooKeeper cluster composition

The values can be one of the following JSON types:

- String
- Number
- Boolean

You can specify and configure the options listed in the [ZooKeeper documentation](#) with the exception of those managed directly by Strimzi. Specifically, all configuration options with keys equal to or starting with one of the following strings are forbidden:

- `server.`
- `dataDir`
- `dataLogDir`
- `clientPort`
- `authProvider`
- `quorum.auth`
- `requireClientAuthScheme`

When a forbidden option is present in the `config` property, it is ignored and a warning message is printed to the Cluster Operator log file. All other supported options are passed to ZooKeeper.

There are exceptions to the forbidden options. For client connection using a specific *cipher suite* for a TLS version, you can [configure allowed `ssl` properties](#).

## Example ZooKeeper configuration

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  config:
    autopurge.snapRetainCount: 3
    autopurge.purgeInterval: 1
    ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384"
    ssl.enabled.protocols: "TLSv1.2"
    ssl.protocol: "TLSv1.2"
    # ...
```

### B.64.2. logging

ZooKeeper has a configurable logger:

- `zookeeper.root.logger`

ZooKeeper uses the Apache `log4j` logger implementation.

Use the `logging` property to configure loggers and logger levels.

You can set the log levels by specifying the logger and level directly (inline) or use a custom (external) ConfigMap. If a ConfigMap is used, you set `logging.name` property to the name of the ConfigMap containing the external logging configuration. Inside the ConfigMap, the logging configuration is described using `log4j.properties`. For more information about log levels, see [Apache logging services](#).

Here we see examples of `inline` and `external` logging.

#### Inline logging

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  # ...
  zookeeper:
    # ...
  logging:
    type: inline
    loggers:
      zookeeper.root.logger: "INFO"
    # ...
```

## External logging

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  # ...
  zookeeper:
    # ...
    logging:
      type: external
      name: customConfigMap
  # ...
```

## Garbage collector (GC)

Garbage collector logging can also be enabled (or disabled) using the [jvmOptions](#) property.

### B.64.3. Schema reference

Property	Description
replicas	The number of pods in the cluster.
integer	
image	The docker image for the pods.
string	
storage	Storage configuration (disk). Cannot be updated. The type depends on the value of the <a href="#">storage.type</a> property within the given object, which must be one of [ephemeral, persistent-claim].
EphemeralStorage, PersistentClaimStorage	
config	The ZooKeeper broker config. Properties with the following prefixes cannot be set: server., dataDir, dataLogDir, clientPort, authProvider, quorum.auth, requireClientAuthScheme, snapshot.trust.empty, standaloneEnabled, reconfigEnabled, 4lw.commands.whitelist, secureClientPort, ssl., serverCnxnFactory, sslQuorum (with the exception of: ssl.protocol, ssl.quorum.protocol, ssl.enabledProtocols, ssl.quorum.enabledProtocols, ssl.ciphersuites, ssl.quorum.ciphersuites, ssl.hostnameVerification, ssl.quorum.hostnameVerification).
map	

Property	Description
affinity	The property <b>affinity</b> has been deprecated. This feature should now be configured at path <code>spec.zookeeper.template.pod.affinity</code> . The pod's affinity rules. For more information, see the <a href="#">external documentation for core/v1 affinity</a> .
Affinity	
tolerations	The property <b>tolerations</b> has been deprecated. This feature should now be configured at path <code>spec.zookeeper.template.pod.tolerations</code> . The pod's tolerations. For more information, see the <a href="#">external documentation for core/v1 toleration</a> .
Toleration array	
livenessProbe	Pod liveness checking.
Probe	
readinessProbe	Pod readiness checking.
Probe	
jvmOptions	JVM Options for pods.
JvmOptions	
resources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
ResourceRequirements	
metrics	The property <b>metrics</b> has been deprecated. This feature should now be configured at path <code>spec.zookeeper.metricsConfig</code> . The Prometheus JMX Exporter configuration. See <a href="https://github.com/prometheus/jmx_exporter">https://github.com/prometheus/jmx_exporter</a> for details of the structure of this configuration.
map	
metricsConfig	Metrics configuration. The type depends on the value of the <code>metricsConfig.type</code> property within the given object, which must be one of [jmxPrometheusExporter].
JmxPrometheusExporterMetrics	
logging	Logging configuration for ZooKeeper. The type depends on the value of the <code>logging.type</code> property within the given object, which must be one of [inline, external].
InlineLogging, ExternalLogging	
template	Template for ZooKeeper cluster resources. The template allows users to specify how are the <code>StatefulSet</code> , <code>Pods</code> and <code>Services</code> generated.
ZookeeperClusterTemplate	
tlsSidecar	The property <b>tlsSidecar</b> has been deprecated. TLS sidecar configuration. The TLS sidecar is not used anymore and this option will be ignored.
TlsSidecar	

## B.65. ZookeeperClusterTemplate schema reference

Used in: [ZookeeperClusterSpec](#)

Property	Description
statefulset	Template for ZooKeeper <a href="#">StatefulSet</a> .
<a href="#">StatefulSetTemplate</a>	
pod	Template for ZooKeeper <a href="#">Pods</a> .
<a href="#">PodTemplate</a>	
clientService	Template for ZooKeeper client <a href="#">Service</a> .
<a href="#">ResourceTemplate</a>	
nodesService	Template for ZooKeeper nodes <a href="#">Service</a> .
<a href="#">ResourceTemplate</a>	
persistentVolumeClaim	Template for all ZooKeeper <a href="#">PersistentVolumeClaims</a> .
<a href="#">ResourceTemplate</a>	
podDisruptionBudget	Template for ZooKeeper <a href="#">PodDisruptionBudget</a> .
<a href="#">PodDisruptionBudgetTemplate</a>	
zookeeperContainer	Template for the ZooKeeper container.
<a href="#">ContainerTemplate</a>	
tlsSidecarContainer	The property <a href="#">tlsSidecarContainer</a> has been deprecated. Template for the Zookeeper server TLS sidecar container. The TLS sidecar is not used anymore and this option will be ignored.
<a href="#">ContainerTemplate</a>	

## B.66. TopicOperatorSpec schema reference

The type [TopicOperatorSpec](#) has been deprecated. Please use [EntityTopicOperatorSpec](#) instead.

Used in: [KafkaSpec](#)

Property	Description
watchedNamespace	The namespace the Topic Operator should watch.
string	
image	The image to use for the Topic Operator.
string	
reconciliationIntervalSeconds	Interval between periodic reconciliations.
integer	
zookeeperSessionTimeoutSeconds	Timeout for the ZooKeeper session.
integer	

Property	Description
affinity	Pod affinity rules. For more information, see the <a href="#">external documentation for core/v1 affinity</a> .
Affinity	
resources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resource requirements</a> .
ResourceRequirements	
topicMetadataMaxAttempts	The number of attempts at getting topic metadata.
integer	
tlsSidecar	TLS sidecar configuration.
TlsSidecar	
logging	Logging configuration. The type depends on the value of the <code>logging.type</code> property within the given object, which must be one of [inline, external].
InlineLogging, ExternalLogging	
jvmOptions	JVM Options for pods.
JvmOptions	
livenessProbe	Pod liveness checking.
Probe	
readinessProbe	Pod readiness checking.
Probe	

## B.67. EntityOperatorSpec schema reference

Used in: [KafkaSpec](#)

Property	Description
topicOperator	Configuration of the Topic Operator.
EntityTopicOperatorSpec	
userOperator	Configuration of the User Operator.
EntityUserOperatorSpec	
affinity	<b>The property <code>affinity</code> has been deprecated. This feature should now be configured at path <code>spec.template.pod.affinity</code>.</b> The pod's affinity rules. For more information, see the <a href="#">external documentation for core/v1 affinity</a> .
Affinity	

Property	Description
tolerations	The property <b>tolerations</b> has been deprecated. This feature should now be configured at path <a href="#">spec.template.pod.tolerations</a> . The pod's tolerations. For more information, see the <a href="#">external documentation for core/v1 toleration</a> .
Toleration array	
tlsSidecar	TLS sidecar configuration.
TlsSidecar	
template	Template for Entity Operator resources. The template allows users to specify how is the <a href="#">Deployment</a> and <a href="#">Pods</a> generated.
EntityOperatorTemplate	

## B.68. EntityTopicOperatorSpec schema reference

Used in: [EntityOperatorSpec](#)

Configures the Topic Operator.

### B.68.1. logging

The Topic Operator has a configurable logger:

- `rootLogger.level`

The Topic Operator uses the Apache [log4j2](#) logger implementation.

Use the `logging` property in the `entityOperator.topicOperator` field of the Kafka resource [Kafka](#) resource to configure loggers and logger levels.

You can set the log levels by specifying the logger and level directly (inline) or use a custom (external) ConfigMap. If a ConfigMap is used, you set `logging.name` property to the name of the ConfigMap containing the external logging configuration. Inside the ConfigMap, the logging configuration is described using [log4j2.properties](#). For more information about log levels, see [Apache logging services](#).

Here we see examples of [inline](#) and [external](#) logging.

### *Inline logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  entityOperator:
    # ...
  topicOperator:
    watchedNamespace: my-topic-namespace
    reconciliationIntervalSeconds: 60
    logging:
      type: inline
      loggers:
        rootLogger.level: INFO
# ...
```

### *External logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  entityOperator:
    # ...
  topicOperator:
    watchedNamespace: my-topic-namespace
    reconciliationIntervalSeconds: 60
    logging:
      type: external
      name: customConfigMap
# ...
```

### *Garbage collector (GC)*

Garbage collector logging can also be enabled (or disabled) using the [jvmOptions](#) property.

## B.68.2. Schema reference

Property	Description
watchedNamespace string	The namespace the Topic Operator should watch.
image string	The image to use for the Topic Operator.
reconciliationIntervalSeconds integer	Interval between periodic reconciliations.
zookeeperSessionTimeoutSeconds integer	Timeout for the ZooKeeper session.
livenessProbe <b>Probe</b>	Pod liveness checking.
readinessProbe <b>Probe</b>	Pod readiness checking.
resources <b>ResourceRequirements</b>	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
topicMetadataMaxAttempts integer	The number of attempts at getting topic metadata.
logging <b>InlineLogging, ExternalLogging</b>	Logging configuration. The type depends on the value of the <code>logging.type</code> property within the given object, which must be one of [inline, external].
jvmOptions <b>JvmOptions</b>	JVM Options for pods.

## B.69. EntityUserOperatorSpec schema reference

Used in: [EntityOperatorSpec](#)

Configures the User Operator.

### B.69.1. logging

The User Operator has a configurable logger:

- `rootLogger.level`

The User Operator uses the Apache `log4j2` logger implementation.

Use the `logging` property in the `entityOperator.userOperator` field of the `Kafka` resource to configure loggers and logger levels.

You can set the log levels by specifying the logger and level directly (inline) or use a custom (external) ConfigMap. If a ConfigMap is used, you set `logging.name` property to the name of the ConfigMap containing the external logging configuration. Inside the ConfigMap, the logging configuration is described using `log4j2.properties`. For more information about log levels, see [Apache logging services](#).

Here we see examples of `inline` and `external` logging.

#### *Inline logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  entityOperator:
    # ...
  userOperator:
    watchedNamespace: my-topic-namespace
    reconciliationIntervalSeconds: 60
    logging:
      type: inline
      loggers:
        rootLogger.level: INFO
# ...
```

#### *External logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    # ...
  zookeeper:
    # ...
  entityOperator:
    # ...
  userOperator:
    watchedNamespace: my-topic-namespace
    reconciliationIntervalSeconds: 60
    logging:
      type: external
      name: customConfigMap
# ...
```

## *Garbage collector (GC)*

Garbage collector logging can also be enabled (or disabled) using the [jvmOptions](#) property.

### B.69.2. Schema reference

Property	Description
watchedNamespace	The namespace the User Operator should watch.
string	
image	The image to use for the User Operator.
string	
reconciliationIntervalSeconds	Interval between periodic reconciliations.
integer	
zookeeperSessionTimeoutSeconds	Timeout for the ZooKeeper session.
integer	
secretPrefix	The prefix that will be added to the KafkaUser name to be used as the Secret name.
string	
livenessProbe	Pod liveness checking.
<a href="#">Probe</a>	
readinessProbe	Pod readiness checking.
<a href="#">Probe</a>	
resources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
<a href="#">ResourceRequirements</a>	
logging	Logging configuration. The type depends on the value of the <a href="#">logging.type</a> property within the given object, which must be one of [inline, external].
<a href="#">InlineLogging</a> , <a href="#">ExternalLogging</a>	
jvmOptions	JVM Options for pods.
<a href="#">JvmOptions</a>	

## B.70. EntityOperatorTemplate schema reference

Used in: [EntityOperatorSpec](#)

Property	Description
deployment	Template for Entity Operator <a href="#">Deployment</a> .
<a href="#">ResourceTemplate</a>	
pod	Template for Entity Operator <a href="#">Pods</a> .
<a href="#">PodTemplate</a>	

Property	Description
tlsSidecarContainer	Template for the Entity Operator TLS sidecar container.
<a href="#">ContainerTemplate</a>	
topicOperatorContainer	Template for the Entity Topic Operator container.
<a href="#">ContainerTemplate</a>	
userOperatorContainer	Template for the Entity User Operator container.
<a href="#">ContainerTemplate</a>	

## B.71. `CertificateAuthority` schema reference

Used in: [KafkaSpec](#)

Configuration of how TLS certificates are used within the cluster. This applies to certificates used for both internal communication within the cluster and to certificates used for client access via [Kafka.spec.kafka.listeners.tls](#).

Property	Description
generateCertificateAuthority	If true then Certificate Authority certificates will be generated automatically. Otherwise the user will need to provide a Secret with the CA certificate. Default is true.
boolean	
generateSecretOwnerReference	If true then the Certificate Authority certificates secrets owner reference is set to the Kafka object. Otherwise no owner reference is set so deleting the Kafka object won't delete the secrets. Default is true.
boolean	
validityDays	The number of days generated certificates should be valid for. The default is 365.
integer	
renewalDays	The number of days in the certificate renewal period. This is the number of days before the a certificate expires during which renewal actions may be performed. When <a href="#">generateCertificateAuthority</a> is true, this will cause the generation of a new certificate. When <a href="#">generateCertificateAuthority</a> is true, this will cause extra logging at WARN level about the pending certificate expiry. Default is 30.
integer	
certificateExpirationPolicy	How should CA certificate expiration be handled when <a href="#">generateCertificateAuthority=true</a> . The default is for a new CA certificate to be generated reusing the existing private key.
string (one of [replace-key, renew-certificate])	

## B.72. CruiseControlSpec schema reference

Used in: [KafkaSpec](#)

Property	Description
image	The docker image for the pods.
string	
tlsSidecar	TLS sidecar configuration.
<a href="#">TlsSidecar</a>	
resources	CPU and memory resources to reserve for the Cruise Control container. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
<a href="#">ResourceRequirements</a>	
livenessProbe	Pod liveness checking for the Cruise Control container.
<a href="#">Probe</a>	
readinessProbe	Pod readiness checking for the Cruise Control container.
<a href="#">Probe</a>	
jvmOptions	JVM Options for the Cruise Control container.
<a href="#">JvmOptions</a>	
logging	Logging configuration (log4j1) for Cruise Control. The type depends on the value of the <a href="#">logging.type</a> property within the given object, which must be one of [inline, external].
<a href="#">InlineLogging</a> , <a href="#">ExternalLogging</a>	
template	Template to specify how Cruise Control resources, <a href="#">Deployments</a> and <a href="#">Pods</a> , are generated.
<a href="#">CruiseControlTemplate</a>	
brokerCapacity	The Cruise Control <a href="#">brokerCapacity</a> configuration.
<a href="#">BrokerCapacity</a>	

Property	Description
config	The Cruise Control configuration. For a full list of configuration options refer to <a href="https://github.com/linkedin/cruise-control/wiki/Configurations">https://github.com/linkedin/cruise-control/wiki/Configurations</a> . Note that properties with the following prefixes cannot be set: bootstrap.servers, client.id, zookeeper., network., security., failed.brokers.zk.path, webserver.http., webserver.api.urlprefix, webserver.session.path, webserver.accesslog., two.step., request.reason.required, metric.reporter.sample.r.bootstrap.servers, metric.reporter.topic, partition.metric.sample.store.topic, broker.metric.sample.store.topic, capacity.config.file, self.healing., anomaly.detection., ssl. (with the exception of: ssl.cipher.suites, ssl.protocol, ssl.enabled.protocols, webserver.http.cors.enabled, webserver.http(cors.origin, webserver.http.cors.exposeheaders).
map	
metrics	<b>The property metrics has been deprecated.</b> This feature should now be configured at path <code>spec.cruiseControl.metricsConfig</code> . The Prometheus JMX Exporter configuration. See <a href="https://github.com/prometheus/jmx_exporter">https://github.com/prometheus/jmx_exporter</a> for details of the structure of this configuration.
map	
metricsConfig	Metrics configuration. The type depends on the value of the <code>metricsConfig.type</code> property within the given object, which must be one of [jmxPrometheusExporter].
JmxPrometheusExporterMetrics	

## B.73. `CruiseControlTemplate` schema reference

Used in: `CruiseControlSpec`

Property	Description
deployment	Template for Cruise Control <a href="#">Deployment</a> .
ResourceTemplate	
pod	Template for Cruise Control <a href="#">Pods</a> .
PodTemplate	
apiService	Template for Cruise Control <a href="#">API Service</a> .
ResourceTemplate	

Property	Description
podDisruptionBudget	Template for Cruise Control <a href="#">PodDisruptionBudget</a> .
<a href="#">PodDisruptionBudgetTemplate</a>	
cruiseControlContainer	Template for the Cruise Control container.
<a href="#">ContainerTemplate</a>	
tlsSidecarContainer	Template for the Cruise Control TLS sidecar container.
<a href="#">ContainerTemplate</a>	

## B.74. BrokerCapacity schema reference

Used in: [CruiseControlSpec](#)

Property	Description
disk	Broker capacity for disk in bytes, for example, 100Gi.
string	
cpuUtilization	Broker capacity for CPU resource utilization as a percentage (0 - 100).
integer	
inboundNetwork	Broker capacity for inbound network throughput in bytes per second, for example, 10000KB/s.
string	
outboundNetwork	Broker capacity for outbound network throughput in bytes per second, for example 10000KB/s.
string	

## B.75. JmxTransSpec schema reference

Used in: [KafkaSpec](#)

Property	Description
image	The image to use for the JmxTrans.
string	
outputDefinitions	Defines the output hosts that will be referenced later on. For more information on these properties see, <a href="#">JmxTransOutputDefinitionTemplate schema reference</a> .
<a href="#">JmxTransOutputDefinitionTemplate</a> array	
logLevel	Sets the logging level of the JmxTrans deployment. For more information see, <a href="#">JmxTrans Logging Level</a> .
string	

Property	Description
kafkaQueries	Queries to send to the Kafka brokers to define what data should be read from each broker. For more information on these properties see, <a href="#">JmxTransQueryTemplate schema reference</a> .
JmxTransQueryTemplate array	
resources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
ResourceRequirements	
template	Template for JmxTrans resources.
JmxTransTemplate	

## B.76. JmxTransOutputDefinitionTemplate schema reference

Used in: [JmxTransSpec](#)

Property	Description
outputType	Template for setting the format of the data that will be pushed. For more information see <a href="#">JmxTrans OutputWriters</a> .
string	
host	The DNS/hostname of the remote host that the data is pushed to.
string	
port	The port of the remote host that the data is pushed to.
integer	
flushDelayInSeconds	How many seconds the JmxTrans waits before pushing a new set of data out.
integer	
typeNames	Template for filtering data to be included in response to a wildcard query. For more information see <a href="#">JmxTrans queries</a> .
string array	
name	Template for setting the name of the output definition. This is used to identify where to send the results of queries should be sent.
string	

## B.77. JmxTransQueryTemplate schema reference

Used in: [JmxTransSpec](#)

Property	Description
targetMBean	If using wildcards instead of a specific MBean then the data is gathered from multiple MBeans. Otherwise if specifying an MBean then data is gathered from that specified MBean.
string	
attributes	Determine which attributes of the targeted MBean should be included.
string array	
outputs	List of the names of output definitions specified in the spec.kafka.jmxTrans.outputDefinitions that have defined where JMX metrics are pushed to, and in which data format.
string array	

## B.78. JmxTransTemplate schema reference

Used in: [JmxTransSpec](#)

Property	Description
deployment	Template for JmxTrans <a href="#">Deployment</a> .
<a href="#">ResourceTemplate</a>	
pod	Template for JmxTrans <a href="#">Pods</a> .
<a href="#">PodTemplate</a>	
container	Template for JmxTrans container.
<a href="#">ContainerTemplate</a>	

## B.79. KafkaExporterSpec schema reference

Used in: [KafkaSpec](#)

Property	Description
image	The docker image for the pods.
string	
groupRegex	Regular expression to specify which consumer groups to collect. Default value is <code>.*</code> .
string	
topicRegex	Regular expression to specify which topics to collect. Default value is <code>.*</code> .
string	
resources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
<a href="#">ResourceRequirements</a>	

Property	Description
logging	Only log messages with the given severity or above. Valid levels: [debug, info, warn, error, fatal]. Default log level is info.
string	
enableSaramaLogging	Enable Sarama logging, a Go client library used by the Kafka Exporter.
boolean	
template	Customization of deployment templates and pods.
<a href="#">KafkaExporterTemplate</a>	
livenessProbe	Pod liveness check.
<a href="#">Probe</a>	
readinessProbe	Pod readiness check.
<a href="#">Probe</a>	

## B.80. KafkaExporterTemplate schema reference

Used in: [KafkaExporterSpec](#)

Property	Description
deployment	Template for Kafka Exporter <a href="#">Deployment</a> .
<a href="#">ResourceTemplate</a>	
pod	Template for Kafka Exporter <a href="#">Pods</a> .
<a href="#">PodTemplate</a>	
service	Template for Kafka Exporter <a href="#">Service</a> .
<a href="#">ResourceTemplate</a>	
container	Template for the Kafka Exporter container.
<a href="#">ContainerTemplate</a>	

## B.81. KafkaStatus schema reference

Used in: [Kafka](#)

Property	Description
conditions	List of status conditions.
<a href="#">Condition</a> array	
observedGeneration	The generation of the CRD that was last reconciled by the operator.
integer	
listeners	Addresses of the internal and external listeners.
<a href="#">ListenerStatus</a> array	

Property	Description
clusterId	Kafka cluster Id.
string	

## B.82. Condition schema reference

Used in: [KafkaBridgeStatus](#), [KafkaConnectorStatus](#), [KafkaConnectS2IStatus](#), [KafkaConnectStatus](#), [KafkaMirrorMaker2Status](#), [KafkaMirrorMakerStatus](#), [KafkaRebalanceStatus](#), [KafkaStatus](#), [KafkaTopicStatus](#), [KafkaUserStatus](#)

Property	Description
type	The unique identifier of a condition, used to distinguish between other conditions in the resource.
string	
status	The status of the condition, either True, False or Unknown.
string	
lastTransitionTime	Last time the condition of a type changed from one status to another. The required format is 'yyyy-MM-ddTHH:mm:ssZ', in the UTC time zone.
string	
reason	The reason for the condition's last transition (a single word in CamelCase).
string	
message	Human-readable message indicating details about the condition's last transition.
string	

## B.83. ListenerStatus schema reference

Used in: [KafkaStatus](#)

Property	Description
type	The type of the listener. Can be one of the following three types: <a href="#">plain</a> , <a href="#">tls</a> , and <a href="#">external</a> .
string	
addresses	A list of the addresses for this listener.
<a href="#">ListenerAddress</a> array	
bootstrapServers	A comma-separated list of <code>host:port</code> pairs for connecting to the Kafka cluster using this listener.
string	
certificates	A list of TLS certificates which can be used to verify the identity of the server when connecting to the given listener. Set only for <a href="#">tls</a> and <a href="#">external</a> listeners.
string array	

## B.84. ListenerAddress schema reference

Used in: [ListenerStatus](#)

Property	Description
host	The DNS name or IP address of the Kafka bootstrap service.
string	
port	The port of the Kafka bootstrap service.
integer	

## B.85. KafkaConnect schema reference

Property	Description
spec	The specification of the Kafka Connect cluster.
<a href="#">KafkaConnectSpec</a>	
status	The status of the Kafka Connect cluster.
<a href="#">KafkaConnectStatus</a>	

## B.86. KafkaConnectSpec schema reference

Used in: [KafkaConnect](#)

Configures a Kafka Connect cluster.

### B.86.1. config

Use the `config` properties to configure Kafka options as keys.

Standard Apache Kafka Connect configuration may be provided, restricted to those properties not managed directly by Strimzi.

Configuration options that cannot be configured relate to:

- Kafka cluster bootstrap address
- Security (Encryption, Authentication, and Authorization)
- Listener / REST interface configuration
- Plugin path configuration

The values can be one of the following JSON types:

- String
- Number
- Boolean

You can specify and configure the options listed in the [Apache Kafka documentation](#) with the exception of those options that are managed directly by Strimzi. Specifically, configuration options with keys equal to or starting with one of the following strings are forbidden:

- `ssl.`
- `sasl.`
- `security.`
- `listeners`
- `plugin.path`
- `rest.`
- `bootstrap.servers`

When a forbidden option is present in the `config` property, it is ignored and a warning message is printed to the Cluster Operator log file. All other options are passed to Kafka Connect.

**IMPORTANT**

The Cluster Operator does not validate keys or values in the `config` object provided. When an invalid configuration is provided, the Kafka Connect cluster might not start or might become unstable. In this circumstance, fix the configuration in the `KafkaConnect.spec.config` or `KafkaConnectS2I.spec.config` object, then the Cluster Operator can roll out the new configuration to all Kafka Connect nodes.

Certain options have default values:

- `group.id` with default value `connect-cluster`
- `offset.storage.topic` with default value `connect-cluster-offsets`
- `config.storage.topic` with default value `connect-cluster-configs`
- `status.storage.topic` with default value `connect-cluster-status`
- `key.converter` with default value `org.apache.kafka.connect.json.JsonConverter`
- `value.converter` with default value `org.apache.kafka.connect.json.JsonConverter`

These options are automatically configured in case they are not present in the `KafkaConnect.spec.config` or `KafkaConnectS2I.spec.config` properties.

There are exceptions to the forbidden options. You can use three allowed `ssl` configuration options for client connection using a specific *cipher suite* for a TLS version. A cipher suite combines algorithms for secure connection and data transfer. You can also configure the `ssl.endpoint.identification.algorithm` property to enable or disable hostname verification.

## Example Kafka Connect configuration

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect
spec:
  # ...
  config:
    group.id: my-connect-cluster
    offset.storage.topic: my-connect-cluster-offsets
    config.storage.topic: my-connect-cluster-configs
    status.storage.topic: my-connect-cluster-status
    key.converter: org.apache.kafka.connect.json.JsonConverter
    value.converter: org.apache.kafka.connect.json.JsonConverter
    key.converter.schemas.enable: true
    value.converter.schemas.enable: true
    config.storage.replication.factor: 3
    offset.storage.replication.factor: 3
    status.storage.replication.factor: 3
    ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384"
    ssl.enabled.protocols: "TLSv1.2"
    ssl.protocol: "TLSv1.2"
    ssl.endpoint.identification.algorithm: HTTPS
  # ...
```

For client connection using a specific *cipher suite* for a TLS version, you can [configure allowed ssl properties](#). You can also [configure the ssl.endpoint.identification.algorithm property](#) to enable or disable hostname verification.

### B.86.2. logging

Kafka Connect (and Kafka Connect with Source2Image support) has its own configurable loggers:

- `connect.root.logger.level`
- `log4j.logger.org.reflections`

Further loggers are added depending on the Kafka Connect plugins running.

Use a curl request to get a complete list of Kafka Connect loggers running from any Kafka broker pod:

```
curl -s http://<connect-cluster-name>-connect-api:8083/admin/loggers/
```

Kafka Connect uses the Apache `log4j` logger implementation.

Use the `logging` property to configure loggers and logger levels.

You can set the log levels by specifying the logger and level directly (inline) or use a custom

(external) ConfigMap. If a ConfigMap is used, you set `logging.name` property to the name of the ConfigMap containing the external logging configuration. Inside the ConfigMap, the logging configuration is described using `log4j.properties`. For more information about log levels, see [Apache logging services](#).

Here we see examples of `inline` and `external` logging.

#### *Inline logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
spec:
  # ...
  logging:
    type: inline
    loggers:
      connect.root.logger.level: "INFO"
  # ...
```

#### *External logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
spec:
  # ...
  logging:
    type: external
    name: customConfigMap
  # ...
```

Any available loggers that are not configured have their level set to `OFF`.

If Kafka Connect was deployed using the Cluster Operator, changes to Kafka Connect logging levels are applied dynamically.

If you use external logging, a rolling update is triggered when logging appenders are changed.

#### *Garbage collector (GC)*

Garbage collector logging can also be enabled (or disabled) using the `jvmOptions` property.

### B.86.3. Schema reference

Property	Description
<code>replicas</code>	The number of pods in the Kafka Connect group.
<code>integer</code>	

Property	Description
version string	The Kafka Connect version. Defaults to 2.7.0. Consult the user documentation to understand the process required to upgrade or downgrade the version.
image string	The docker image for the pods.
bootstrapServers string	Bootstrap servers to connect to. This should be given as a comma separated list of <hostname>:<port> pairs.
tls <b>KafkaConnectTls</b>	TLS configuration.
authentication  <b>KafkaClientAuthenticationTls</b> , <b>KafkaClientAuthenticationScramSha512</b> , <b>KafkaClientAuthenticationPlain</b> , <b>KafkaClientAuthenticationOAuth</b>	Authentication configuration for Kafka Connect. The type depends on the value of the <b>authentication.type</b> property within the given object, which must be one of [tls, scram-sha-512, plain, oauth].
config	The Kafka Connect configuration. Properties with the following prefixes cannot be set: ssl., sasl., security., listeners, plugin.path, rest., bootstrap.servers, consumer.interceptor.classes, producer.interceptor.classes (with the exception of: ssl.endpoint.identification.algorithm, ssl.cipher.suites, ssl.protocol, ssl.enabled.protocols).
map	
resources  <b>ResourceRequirements</b>	The maximum limits for CPU and memory resources and the requested initial resources. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
livenessProbe  <b>Probe</b>	Pod liveness checking.
readinessProbe  <b>Probe</b>	Pod readiness checking.
jvmOptions  <b>JvmOptions</b>	JVM Options for pods.

Property	Description
affinity	The property <b>affinity</b> has been deprecated. This feature should now be configured at path <code>spec.template.pod.affinity</code> . The pod's affinity rules. For more information, see the <a href="#">external documentation for core/v1 affinity</a> .
Affinity	
tolerations	The property <b>tolerations</b> has been deprecated. This feature should now be configured at path <code>spec.template.pod.tolerations</code> . The pod's tolerations. For more information, see the <a href="#">external documentation for core/v1 toleration</a> .
Toleration array	
logging	Logging configuration for Kafka Connect. The type depends on the value of the <code>logging.type</code> property within the given object, which must be one of [inline, external].
InlineLogging, ExternalLogging	
metrics	The property <b>metrics</b> has been deprecated. This feature should now be configured at path <code>spec.metricsConfig</code> . The Prometheus JMX Exporter configuration. See <a href="https://github.com/prometheus/jmx_exporter">https://github.com/prometheus/jmx_exporter</a> for details of the structure of this configuration.
map	
tracing	The configuration of tracing in Kafka Connect. The type depends on the value of the <code>tracing.type</code> property within the given object, which must be one of [jaeger].
JaegerTracing	
template	Template for Kafka Connect and Kafka Connect S2I resources. The template allows users to specify how the <code>Deployment</code> , <code>Pods</code> and <code>Service</code> are generated.
KafkaConnectTemplate	
externalConfiguration	Pass data from Secrets or ConfigMaps to the Kafka Connect pods and use them to configure connectors.
ExternalConfiguration	
build	Configures how the Connect container image should be built. Optional.
Build	
clientRackInitImage	The image of the init container used for initializing the <code>client.rack</code> .
string	
metricsConfig	Metrics configuration. The type depends on the value of the <code>metricsConfig.type</code> property within the given object, which must be one of [jmxPrometheusExporter].
JmxPrometheusExporterMetrics	

Property	Description
rack	Configuration of the node label which will be used as the client.rack consumer configuration.
Rack	

## B.87. KafkaConnectTls schema reference

Used in: [KafkaConnectS2ISpec](#), [KafkaConnectSpec](#)

Configures TLS trusted certificates for connecting Kafka Connect to the cluster.

### B.87.1. trustedCertificates

Provide a list of secrets using the [trustedCertificates](#) property.

### B.87.2. Schema reference

Property	Description
trustedCertificates	Trusted certificates for TLS connection.
<a href="#">CertSecretSource</a> array	

## B.88. KafkaClientAuthenticationTls schema reference

Used in: [KafkaBridgeSpec](#), [KafkaConnectS2ISpec](#), [KafkaConnectSpec](#), [KafkaMirrorMaker2ClusterSpec](#), [KafkaMirrorMakerConsumerSpec](#), [KafkaMirrorMakerProducerSpec](#)

To use TLS client authentication, set the [type](#) property to the value [tls](#). TLS client authentication uses a TLS certificate to authenticate.

### B.88.1. certificateAndKey

The certificate is specified in the [certificateAndKey](#) property and is always loaded from a Kubernetes secret. In the secret, the certificate must be stored in X509 format under two different keys: public and private.

You can use the secrets created by the User Operator, or you can create your own TLS certificate file, with the keys used for authentication, then create a [Secret](#) from the file:

```
kubectl create secret generic MY-SECRET \
--from-file=MY-PUBLIC-TLS-CERTIFICATE-FILE.crt \
--from-file=MY-PRIVATE.key
```

**NOTE** TLS client authentication can only be used with TLS connections.

## *Example TLS client authentication configuration*

```
authentication:  
  type: tls  
  certificateAndKey:  
    secretName: my-secret  
    certificate: my-public-tls-certificate-file.crt  
    key: private.key
```

### B.88.2. Schema reference

The `type` property is a discriminator that distinguishes use of the `KafkaClientAuthenticationTls` type from `KafkaClientAuthenticationScramSha512`, `KafkaClientAuthenticationPlain`, `KafkaClientAuthenticationOAuth`. It must have the value `tls` for the type `KafkaClientAuthenticationTls`.

Property	Description
certificateAndKey	Reference to the <code>Secret</code> which holds the certificate and private key pair.
<code>CertAndKeySecretSource</code>	
type	Must be <code>tls</code> .
string	

## B.89. KafkaClientAuthenticationScramSha512 schema reference

Used in: `KafkaBridgeSpec`, `KafkaConnectS2ISpec`, `KafkaConnectSpec`, `KafkaMirrorMaker2ClusterSpec`, `KafkaMirrorMakerConsumerSpec`, `KafkaMirrorMakerProducerSpec`

To configure SASL-based SCRAM-SHA-512 authentication, set the `type` property to `scram-sha-512`. The SCRAM-SHA-512 authentication mechanism requires a username and password.

### B.89.1. `username`

Specify the username in the `username` property.

### B.89.2. `passwordSecret`

In the `passwordSecret` property, specify a link to a `Secret` containing the password.

You can use the secrets created by the User Operator.

If required, you can create a text file that contains the password, in cleartext, to use for authentication:

```
echo -n PASSWORD > MY-PASSWORD.txt
```

You can then create a **Secret** from the text file, setting your own field name (key) for the password:

```
kubectl create secret generic MY-CONNECT-SECRET-NAME --from-file=MY-PASSWORD-FIELD-  
NAME=./MY-PASSWORD.txt
```

*Example Secret for SCRAM-SHA-512 client authentication for Kafka Connect*

```
apiVersion: v1  
kind: Secret  
metadata:  
  name: my-connect-secret-name  
type: Opaque  
data:  
  my-connect-password-field: LFTIyFRF1MmU2N2Tm
```

The **secretName** property contains the name of the **Secret**, and the **password** property contains the name of the key under which the password is stored inside the **Secret**.

**IMPORTANT** Do not specify the actual password in the **password** property.

*Example SASL-based SCRAM-SHA-512 client authentication configuration for Kafka Connect*

```
authentication:  
  type: scram-sha-512  
  username: my-connect-username  
  passwordSecret:  
    secretName: my-connect-secret-name  
    password: my-connect-password-field
```

### B.89.3. Schema reference

The **type** property is a discriminator that distinguishes use of the **KafkaClientAuthenticationScramSha512** type from **KafkaClientAuthenticationTls**, **KafkaClientAuthenticationPlain**, **KafkaClientAuthenticationOAuth**. It must have the value **scram-sha-512** for the type **KafkaClientAuthenticationScramSha512**.

Property	Description
passwordSecret	Reference to the <b>Secret</b> which holds the password.
PasswordSecretSource	
type	Must be <b>scram-sha-512</b> .
string	
username	Username used for the authentication.
string	

## B.90. PasswordSecretSource schema reference

Used in: [KafkaClientAuthenticationPlain](#), [KafkaClientAuthenticationScramSha512](#)

Property	Description
password	The name of the key in the Secret under which the password is stored.
secretName	The name of the Secret containing the password.
string	

## B.91. KafkaClientAuthenticationPlain schema reference

Used in: [KafkaBridgeSpec](#), [KafkaConnectS2ISpec](#), [KafkaConnectSpec](#), [KafkaMirrorMaker2ClusterSpec](#), [KafkaMirrorMakerConsumerSpec](#), [KafkaMirrorMakerProducerSpec](#)

To configure SASL-based PLAIN authentication, set the `type` property to `plain`. SASL PLAIN authentication mechanism requires a username and password.

**WARNING**

The SASL PLAIN mechanism will transfer the username and password across the network in cleartext. Only use SASL PLAIN authentication if TLS encryption is enabled.

### B.91.1. `username`

Specify the username in the `username` property.

### B.91.2. `passwordSecret`

In the `passwordSecret` property, specify a link to a `Secret` containing the password.

You can use the secrets created by the User Operator.

If required, create a text file that contains the password, in cleartext, to use for authentication:

```
echo -n PASSWORD > MY-PASSWORD.txt
```

You can then create a `Secret` from the text file, setting your own field name (key) for the password:

```
kubectl create secret generic MY-CONNECT-SECRET-NAME --from-file=MY-PASSWORD-FIELD-NAME=./MY-PASSWORD.txt
```

*Example Secret for PLAIN client authentication for Kafka Connect*

```
apiVersion: v1
kind: Secret
metadata:
  name: my-connect-secret-name
type: Opaque
data:
  my-password-field-name: LFTIyFRF1MmU2N2Tm
```

The `secretName` property contains the name of the `Secret` and the `password` property contains the name of the key under which the password is stored inside the `Secret`.

**IMPORTANT** Do not specify the actual password in the `password` property.

*An example SASL based PLAIN client authentication configuration*

```
authentication:
  type: plain
  username: my-connect-username
  passwordSecret:
    secretName: my-connect-secret-name
    password: my-password-field-name
```

### B.91.3. Schema reference

The `type` property is a discriminator that distinguishes use of the `KafkaClientAuthenticationPlain` type from `KafkaClientAuthenticationTls`, `KafkaClientAuthenticationScramSha512`, `KafkaClientAuthenticationOAuth`. It must have the value `plain` for the type `KafkaClientAuthenticationPlain`.

Property	Description
passwordSecret	Reference to the <code>Secret</code> which holds the password.
PasswordSecretSource	
type	Must be <code>plain</code> .
string	
username	Username used for the authentication.
string	

## B.92. KafkaClientAuthenticationOAuth schema reference

Used in: `KafkaBridgeSpec`, `KafkaConnectS2ISpec`, `KafkaConnectSpec`, `KafkaMirrorMaker2ClusterSpec`, `KafkaMirrorMakerConsumerSpec`, `KafkaMirrorMakerProducerSpec`

To use OAuth client authentication, set the `type` property to the value `oauth`.

OAuth authentication can be configured using one of the following options:

- Client ID and secret
- Client ID and refresh token
- Access token
- TLS

#### *Client ID and secret*

You can configure the address of your authorization server in the `tokenEndpointUri` property together with the client ID and client secret used in authentication. The OAuth client will connect to the OAuth server, authenticate using the client ID and secret and get an access token which it will use to authenticate with the Kafka broker. In the `clientSecret` property, specify a link to a `Secret` containing the client secret.

*An example of OAuth client authentication using client ID and client secret*

```
authentication:  
  type: oauth  
  tokenEndpointUri:  
    https://sso.myproject.svc:8443/auth/realms/internal/protocol/openid-connect/token  
    clientId: my-client-id  
    clientSecret:  
      secretName: my-client-oauth-secret  
      key: client-secret
```

#### *Client ID and refresh token*

You can configure the address of your OAuth server in the `tokenEndpointUri` property together with the OAuth client ID and refresh token. The OAuth client will connect to the OAuth server, authenticate using the client ID and refresh token and get an access token which it will use to authenticate with the Kafka broker. In the `refreshToken` property, specify a link to a `Secret` containing the refresh token.

+ .An example of OAuth client authentication using client ID and refresh token

```
authentication:  
  type: oauth  
  tokenEndpointUri:  
    https://sso.myproject.svc:8443/auth/realms/internal/protocol/openid-connect/token  
    clientId: my-client-id  
    refreshToken:  
      secretName: my-refresh-token-secret  
      key: refresh-token
```

#### *Access token*

You can configure the access token used for authentication with the Kafka broker directly. In this case, you do not specify the `tokenEndpointUri`. In the `accessToken` property, specify a link to a `Secret` containing the access token.

*An example of OAuth client authentication using only an access token*

```
authentication:  
  type: oauth  
  accessToken:  
    secretName: my-access-token-secret  
    key: access-token
```

## TLS

Accessing the OAuth server using the HTTPS protocol does not require any additional configuration as long as the TLS certificates used by it are signed by a trusted certification authority and its hostname is listed in the certificate.

If your OAuth server is using certificates which are self-signed or are signed by a certification authority which is not trusted, you can configure a list of trusted certificates in the custom resource. The `tlsTrustedCertificates` property contains a list of secrets with key names under which the certificates are stored. The certificates must be stored in X509 format.

*An example of TLS certificates provided*

```
authentication:  
  type: oauth  
  tokenEndpointUri:  
    https://sso.myproject.svc:8443/auth/realm/internal/protocol/openid-connect/token  
    clientId: my-client-id  
  refreshToken:  
    secretName: my-refresh-token-secret  
    key: refresh-token  
  tlsTrustedCertificates:  
    - secretName: oauth-server-ca  
      certificate: tls.crt
```

The OAuth client will by default verify that the hostname of your OAuth server matches either the certificate subject or one of the alternative DNS names. If it is not required, you can disable the hostname verification.

*An example of disabled TLS hostname verification*

```
authentication:  
  type: oauth  
  tokenEndpointUri:  
    https://sso.myproject.svc:8443/auth/realm/internal/protocol/openid-connect/token  
    clientId: my-client-id  
  refreshToken:  
    secretName: my-refresh-token-secret  
    key: refresh-token  
  disableTlsHostnameVerification: true
```

## B.92.1. Schema reference

The `type` property is a discriminator that distinguishes use of the `KafkaClientAuthenticationOAuth` type from `KafkaClientAuthenticationTls`, `KafkaClientAuthenticationSramSha512`, `KafkaClientAuthenticationPlain`. It must have the value `oauth` for the type `KafkaClientAuthenticationOAuth`.

Property	Description
accessToken	Link to Kubernetes Secret containing the access token which was obtained from the authorization server.
<code>GenericSecretSource</code>	
accessTokenIsJwt	Configure whether access token should be treated as JWT. This should be set to <code>false</code> if the authorization server returns opaque tokens. Defaults to <code>true</code> .
boolean	
clientId	OAuth Client ID which the Kafka client can use to authenticate against the OAuth server and use the token endpoint URI.
string	
clientSecret	Link to Kubernetes Secret containing the OAuth client secret which the Kafka client can use to authenticate against the OAuth server and use the token endpoint URI.
<code>GenericSecretSource</code>	
disableTlsHostnameVerification	Enable or disable TLS hostname verification. Default value is <code>false</code> .
boolean	
maxTokenExpirySeconds	Set or limit time-to-live of the access tokens to the specified number of seconds. This should be set if the authorization server returns opaque tokens.
integer	
refreshToken	Link to Kubernetes Secret containing the refresh token which can be used to obtain access token from the authorization server.
<code>GenericSecretSource</code>	
scope	OAuth scope to use when authenticating against the authorization server. Some authorization servers require this to be set. The possible values depend on how authorization server is configured. By default <code>scope</code> is not specified when doing the token endpoint request.
string	
tlsTrustedCertificates	Trusted certificates for TLS connection to the OAuth server.
<code>CertSecretSource</code> array	
tokenEndpointUri	Authorization server token endpoint URI.
string	

Property	Description
type	Must be <code>oauth</code> .
string	

## B.93. JaegerTracing schema reference

Used in: `KafkaBridgeSpec`, `KafkaConnectS2ISpec`, `KafkaConnectSpec`, `KafkaMirrorMaker2Spec`, `KafkaMirrorMakerSpec`

The `type` property is a discriminator that distinguishes use of the `JaegerTracing` type from other subtypes which may be added in the future. It must have the value `jaeger` for the type `JaegerTracing`.

Property	Description
type	Must be <code>jaeger</code> .
string	

## B.94. KafkaConnectTemplate schema reference

Used in: `KafkaConnectS2ISpec`, `KafkaConnectSpec`, `KafkaMirrorMaker2Spec`

Property	Description
deployment	Template for Kafka Connect <code>Deployment</code> .
<code>DeploymentTemplate</code>	
pod	Template for Kafka Connect <code>Pods</code> .
<code>PodTemplate</code>	
apiService	Template for Kafka Connect API <code>Service</code> .
<code>ResourceTemplate</code>	
buildConfig	Template for the Kafka Connect <code>BuildConfig</code> used to build new container images. The <code>BuildConfig</code> is used only on OpenShift.
<code>ResourceTemplate</code>	
buildContainer	Template for the Kafka Connect Build container. The build container is used only on Kubernetes.
<code>ContainerTemplate</code>	
buildPod	Template for Kafka Connect Build <code>Pods</code> . The build pod is used only on Kubernetes.
<code>PodTemplate</code>	
connectContainer	Template for the Kafka Connect container.
<code>ContainerTemplate</code>	
initContainer	Template for the Kafka init container.
<code>ContainerTemplate</code>	

Property	Description
podDisruptionBudget	Template for Kafka Connect <a href="#">PodDisruptionBudget</a> .
<a href="#">PodDisruptionBudgetTemplate</a>	

## B.95. DeploymentTemplate schema reference

Used in: [KafkaBridgeTemplate](#), [KafkaConnectTemplate](#), [KafkaMirrorMakerTemplate](#)

Property	Description
metadata	Metadata applied to the resource.
<a href="#">MetadataTemplate</a>	
deploymentStrategy	DeploymentStrategy which will be used for this Deployment. Valid values are <a href="#">RollingUpdate</a> and <a href="#">Recreate</a> . Defaults to <a href="#">RollingUpdate</a> .
string (one of [RollingUpdate, Recreate])	

## B.96. ExternalConfiguration schema reference

Used in: [KafkaConnectS2ISpec](#), [KafkaConnectSpec](#), [KafkaMirrorMaker2Spec](#)

Configures external storage properties that define configuration options for Kafka Connect connectors.

You can mount ConfigMaps or Secrets into a Kafka Connect pod as environment variables or volumes. Volumes and environment variables are configured in the `externalConfiguration` property in [KafkaConnect.spec](#) and [KafkaConnectS2I.spec](#).

When applied, the environment variables and volumes are available for use when developing your connectors.

### B.96.1. env

The `env` property is used to specify one or more environment variables. These variables can contain a value from either a ConfigMap or a Secret.

*Example Secret containing values for environment variables*

```
apiVersion: v1
kind: Secret
metadata:
  name: aws-creds
type: Opaque
data:
  awsAccessKey: QUtJQVhYWFhYWFhYWFhYWFg=
  awsSecretAccessKey: Ylhsd1lYTnpkMj15WkE=
```

**NOTE**

The names of user-defined environment variables cannot start with `KAFKA_` or `STRIMZI_`.

To mount a value from a Secret to an environment variable, use the `valueFrom` property and the `secretKeyRef`.

*Example environment variables set to values from a Secret*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect
spec:
  # ...
  externalConfiguration:
    env:
      - name: AWS_ACCESS_KEY_ID
        valueFrom:
          secretKeyRef:
            name: aws-creds
            key: awsAccessKey
      - name: AWS_SECRET_ACCESS_KEY
        valueFrom:
          secretKeyRef:
            name: aws-creds
            key: awsSecretAccessKey
```

A common use case for mounting Secrets to environment variables is when your connector needs to communicate with Amazon AWS and needs to read the `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` environment variables with credentials.

To mount a value from a ConfigMap to an environment variable, use `configMapKeyRef` in the `valueFrom` property as shown in the following example.

*Example environment variables set to values from a ConfigMap*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect
spec:
  # ...
  externalConfiguration:
    env:
      - name: MY_ENVIRONMENT_VARIABLE
        valueFrom:
          configMapKeyRef:
            name: my-config-map
            key: my-key
```

## B.96.2. volumes

You can also mount ConfigMaps or Secrets to a Kafka Connect pod as volumes.

Using volumes instead of environment variables is useful in the following scenarios:

- Mounting truststores or keystores with TLS certificates
- Mounting a properties file that is used to configure Kafka Connect connectors

*Example Secret with properties*

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
stringData:
  connector.properties: |- ①
    dbUsername: my-user ②
    dbPassword: my-password
```

① The connector configuration in properties file format.

② Database username and password properties used in the configuration.

In this example, a Secret named `mysecret` is mounted to a volume named `connector-config`. In the `config` property, a configuration provider (`FileConfigProvider`) is specified, which will load configuration values from external sources. The Kafka `FileConfigProvider` is given the alias `file`, and will read and extract database `username` and `password` property values from the file to use in the connector configuration.

*Example external volumes set to values from a Secret*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect
spec:
  # ...
  config:
    config.providers: file ①
    config.providers.file.class:
      org.apache.kafka.common.config.provider.FileConfigProvider ②
    #...
    externalConfiguration:
      volumes:
        - name: connector-config ③
          secret:
            secretName: mysecret ④
```

① The alias for the configuration provider, which is used to define other configuration parameters.

Use a comma-separated list if you want to add more than one provider.

- ② The `FileConfigProvider` is the configuration provider that provides values from properties files. The parameter uses the alias from `config.providers`, taking the form `config.providers.${alias}.class`.
- ③ The name of the volume containing the Secret. Each volume must specify a name in the `name` property and a reference to ConfigMap or Secret.
- ④ The name of the Secret.

The volumes are mounted inside the Kafka Connect containers in the path `/opt/kafka/external-configuration/<volume-name>`. For example, the files from a volume named `connector-config` would appear in the directory `/opt/kafka/external-configuration/connector-config`.

The `FileConfigProvider` is used to read the values from the mounted properties files in connector configurations.

### B.96.3. Schema reference

Property	Description
<code>env</code>	Allows to pass data from Secret or ConfigMap to the Kafka Connect pods as environment variables.
<code>ExternalConfigurationEnv</code> array	
<code>volumes</code>	Allows to pass data from Secret or ConfigMap to the Kafka Connect pods as volumes.
<code>ExternalConfigurationVolumeSource</code> array	

## B.97. `ExternalConfigurationEnv` schema reference

Used in: `ExternalConfiguration`

Property	Description
<code>name</code>	Name of the environment variable which will be passed to the Kafka Connect pods. The name of the environment variable cannot start with <code>KAFKA_</code> or <code>STRIMZI_</code> .
<code>string</code>	
<code>valueFrom</code>	Value of the environment variable which will be passed to the Kafka Connect pods. It can be passed either as a reference to Secret or ConfigMap field. The field has to specify exactly one Secret or ConfigMap.
<code>ExternalConfigurationEnvVarSource</code>	

## B.98. ExternalConfigurationEnvVarSource schema reference

Used in: [ExternalConfigurationEnv](#)

Property	Description
configMapKeyRef	Reference to a key in a ConfigMap. For more information, see the <a href="#">external documentation for core/v1 configmapkeyselector</a> .
ConfigMapKeySelector	
secretKeyRef	Reference to a key in a Secret. For more information, see the <a href="#">external documentation for core/v1 secretkeyselector</a> .
SecretKeySelector	

## B.99. ExternalConfigurationVolumeSource schema reference

Used in: [ExternalConfiguration](#)

Property	Description
configMap	Reference to a key in a ConfigMap. Exactly one Secret or ConfigMap has to be specified. For more information, see the <a href="#">external documentation for core/v1 configmapvolumesource</a> .
ConfigMapVolumeSource	
name	Name of the volume which will be added to the Kafka Connect pods.
string	
secret	Reference to a key in a Secret. Exactly one Secret or ConfigMap has to be specified. For more information, see the <a href="#">external documentation for core/v1 secretvolumesource</a> .
SecretVolumeSource	

## B.100. Build schema reference

Used in: [KafkaConnectS2ISpec](#), [KafkaConnectSpec](#)

Configures additional connectors for Kafka Connect deployments.

### B.100.1. output

To build new container images with additional connector plugins, Strimzi requires a container registry where the images can be pushed to, stored, and pulled from. Strimzi does not run its own container registry, so a registry must be provided. Strimzi supports private container registries as well as public registries such as [Quay](#) or [Docker Hub](#). The container registry is configured in the

.spec.build.output section of the `KafkaConnect` custom resource. The `output` configuration, which is required, supports two types: `docker` and `imagestream`.

#### Using Docker registry

To use a Docker registry, you have to specify the `type` as `docker`, and the `image` field with the full name of the new container image. The full name must include:

- The address of the registry
- Port number (if listening on a non-standard port)
- The tag of the new container image

Example valid container image names:

- `docker.io/my-org/my-image/my-tag`
- `quay.io/my-org/my-image/my-tag`
- `image-registry.image-registry.svc:5000/myproject/kafka-connect-build:latest`

Each Kafka Connect deployment must use a separate image, which can mean different tags at the most basic level.

If the registry requires authentication, use the `pushSecret` to set a name of the Secret with the registry credentials. For the Secret, use the `kubernetes.io/dockerconfigjson` type and a `.dockerconfigjson` file to contain the Docker credentials. For more information on pulling an image from a private registry, see [Create a Secret based on existing Docker credentials](#).

#### Example output configuration

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect-cluster
spec:
  #...
  build:
    output:
      type: docker ①
      image: my-registry.io/my-org/my-connect-cluster:latest ②
      pushSecret: my-registry-credentials ③
  #...
```

① (Required) Type of output used by Strimzi.

② (Required) Full name of the image used, including the repository and tag.

③ (Optional) Name of the secret with the container registry credentials.

#### Using OpenShift ImageStream

Instead of Docker, you can use OpenShift ImageStream to store a new container image. The ImageStream has to be created manually before deploying Kafka Connect. To use ImageStream, set the `type` to `imagestream`, and use the `image` property to specify the name of the ImageStream and the

tag used. For example, `my-connect-image-stream:latest`.

#### Example output configuration

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect-cluster
spec:
  #...
  build:
    output:
      type: imagestream ①
      image: my-connect-build:latest ②
  #...
```

① (Required) Type of output used by Strimzi.

② (Required) Name of the ImageStream and tag.

## B.100.2. `plugins`

Connector plugins are a set of files that define the implementation required to connect to certain types of external system. The connector plugins required for a container image must be configured using the `.spec.build.plugins` property of the `KafkaConnect` custom resource. Each connector plugin must have a name which is unique within the Kafka Connect deployment. Additionally, the plugin artifacts must be listed. These artifacts are downloaded by Strimzi, added to the new container image, and used in the Kafka Connect deployment. The connector plugin artifacts can also include additional components, such as (de)serializers. Each connector plugin is downloaded into a separate directory so that the different connectors and their dependencies are properly *sandboxed*. Each plugin must be configured with at least one `artifact`.

## Example `plugins` configuration with two connector plugins

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect-cluster
spec:
  #...
  build:
    output:
      #...
    plugins: ①
      - name: debezium-postgres-connector
        artifacts:
          - type: tgz
            url: https://repo1.maven.org/maven2/io/debezium/debezium-connector-
postgres/1.3.1.Final/debezium-connector-postgres-1.3.1.Final-plugin.tar.gz
            sha512sum:
962a12151bdf9a5a30627eebac739955a4fd95a08d373b86bdcea2b4d0c27dd6e1edd5cb548045e115e33a
9e69b1b2a352bee24df035a0447cb820077af00c03
          - name: camel-telegram
            artifacts:
              - type: tgz
                url:
https://repo.maven.apache.org/maven2/org/apache/camel/kafkaconnector/camel-telegram-
kafka-connector/0.7.0/camel-telegram-kafka-connector-0.7.0-package.tar.gz
                sha512sum:
a9b1ac63e3284bea7836d7d24d84208c49cdf5600070e6bd1535de654f6920b74ad950d51733e8020bf418
7870699819f54ef5859c7846ee4081507f48873479
  #...
```

① (Required) List of connector plugins and their artifacts.

Strimzi supports two types of artifacts: \* JAR files, which are downloaded and used directly \* TGZ archives, which are downloaded and unpacked

### IMPORTANT

Strimzi does not perform any security scanning of the downloaded artifacts. For security reasons, you should first verify the artifacts manually, and configure the checksum verification to make sure the same artifact is used in the automated build and in the Kafka Connect deployment.

### Using JAR artifacts

JAR artifacts represent a resource which is downloaded and added to a container image. JAR artifacts are mainly used for downloading JAR files, but they can also be used to download other file types. To use a JAR artifact, set the `type` property to `jar`, and specify the download location using the `url` property.

Additionally, you can specify a SHA-512 checksum of the artifact. If specified, Strimzi will verify the checksum of the artifact while building the new container image.

## Example JAR artifact

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect-cluster
spec:
  #...
  build:
    output:
      #...
    plugins:
      - name: my-plugin
        artifacts:
          - type: jar ①
            url: https://my-domain.tld/my-jar.jar ②
            sha512sum: 589...ab4 ③
          - type: jar
            url: https://my-domain.tld/my-jar2.jar
  #...
```

① (Required) Type of artifact.

② (Required) URL from which the artifact is downloaded.

③ (Optional) SHA-512 checksum to verify the artifact.

## Using TGZ artifacts

TGZ artifacts are used to download TAR archives that have been compressed using Gzip compression. The TGZ artifact can contain the whole Kafka Connect connector, even when comprising multiple different files. The TGZ artifact is automatically downloaded and unpacked by Strimzi while building the new container image. To use TGZ artifacts, set the `type` property to `tgz`, and specify the download location using the `url` property.

Additionally, you can specify a SHA-512 checksum of the artifact. If specified, Strimzi will verify the checksum before unpacking it and building the new container image.

### Example TGZ artifact

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
metadata:
  name: my-connect-cluster
spec:
  #...
  build:
    output:
      #...
    plugins:
      - name: my-plugin
        artifacts:
          - type: tgz ①
            url: https://my-domain.tld/my-connector-archive.jar ②
            sha512sum: 158...jg10 ③
  #...
```

① (Required) Type of artifact.

② (Required) URL from which the archive is downloaded.

③ (Optional) SHA-512 checksum to verify the artifact.

### B.100.3. Schema reference

Property	Description
output	Configures where should the newly built image be stored. Required. The type depends on the value of the <code>output.type</code> property within the given object, which must be one of [docker, imagestream].
DockerOutput, ImageStreamOutput	
resources	CPU and memory resources to reserve for the build. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
ResourceRequirements	
plugins	List of connector plugins which should be added to the Kafka Connect. Required.
Plugin array	

## B.101. DockerOutput schema reference

Used in: [Build](#)

The `type` property is a discriminator that distinguishes use of the `DockerOutput` type from `ImageStreamOutput`. It must have the value `docker` for the type `DockerOutput`.

Property	Description
image	The full name which should be used for tagging and pushing the newly built image. For example <code>quay.io/my-organization/my-custom-connect:latest</code> . Required.
string	
pushSecret	Container Registry Secret with the credentials for pushing the newly built image.
string	
type	Must be <code>docker</code> .
string	

## B.102. `ImageStreamOutput` schema reference

Used in: [Build](#)

The `type` property is a discriminator that distinguishes use of the `ImageStreamOutput` type from `DockerOutput`. It must have the value `imagestream` for the type `ImageStreamOutput`.

Property	Description
image	The name and tag of the ImageStream where the newly built image will be pushed. For example <code>my-custom-connect:latest</code> . Required.
string	
type	Must be <code>imagestream</code> .
string	

## B.103. `Plugin` schema reference

Used in: [Build](#)

Property	Description
name	The unique name of the connector plugin. Will be used to generate the path where the connector artifacts will be stored. The name has to be unique within the KafkaConnect resource. The name has to follow the following pattern: <code>^[a-z][-_a-z0-9]*[a-z]\$</code> . Required.
string	
artifacts	List of artifacts which belong to this connector plugin. Required.
<code>JarArtifact</code> , <code>TgzArtifact</code> array	

## B.104. `JarArtifact` schema reference

Used in: [Plugin](#)

Property	Description
url	URL of the artifact which will be downloaded. Strimzi does not do any security scanning of the downloaded artifacts. For security reasons, you should first verify the artifacts manually and configure the checksum verification to make sure the same artifact is used in the automated build. Required.
string	
sha512sum	SHA512 checksum of the artifact. Optional. If specified, the checksum will be verified while building the new container. If not specified, the downloaded artifact will not be verified.
string	
type	Must be <a href="#">jar</a> .
string	

## B.105. [TgzArtifact](#) schema reference

Used in: [Plugin](#)

Property	Description
url	URL of the artifact which will be downloaded. Strimzi does not do any security scanning of the downloaded artifacts. For security reasons, you should first verify the artifacts manually and configure the checksum verification to make sure the same artifact is used in the automated build. Required.
string	
sha512sum	SHA512 checksum of the artifact. Optional. If specified, the checksum will be verified while building the new container. If not specified, the downloaded artifact will not be verified.
string	
type	Must be <a href="#">tgz</a> .
string	

## B.106. [KafkaConnectStatus](#) schema reference

Used in: [KafkaConnect](#)

Property	Description
conditions	List of status conditions.
<a href="#">Condition</a> array	

Property	Description
observedGeneration integer	The generation of the CRD that was last reconciled by the operator.
url string	The URL of the REST API endpoint for managing and monitoring Kafka Connect connectors.
connectorPlugins <a href="#">ConnectorPlugin</a> array	The list of connector plugins available in this Kafka Connect deployment.
labelSelector string	Label selector for pods providing this resource.
replicas integer	The current number of pods being used to provide this resource.

## B.107. [ConnectorPlugin](#) schema reference

Used in: [KafkaConnectS2IStatus](#), [KafkaConnectStatus](#), [KafkaMirrorMaker2Status](#)

Property	Description
type string	The type of the connector plugin. The available types are <a href="#">sink</a> and <a href="#">source</a> .
version string	The version of the connector plugin.
class string	The class of the connector plugin.

## B.108. [KafkaConnectS2I](#) schema reference

Property	Description
spec <a href="#">KafkaConnectS2ISpec</a>	The specification of the Kafka Connect Source-to-Image (S2I) cluster.
status <a href="#">KafkaConnectS2IStatus</a>	The status of the Kafka Connect Source-to-Image (S2I) cluster.

## B.109. [KafkaConnectS2ISpec](#) schema reference

Used in: [KafkaConnectS2I](#)

Configures a Kafka Connect cluster with Source-to-Image (S2I) support.

When extending Kafka Connect with connector plugins on OpenShift (only), you can use OpenShift

builds and S2I to create a container image that is used by the Kafka Connect deployment.

The configuration options are similar to Kafka Connect configuration using the [KafkaConnectSpec schema](#).

### B.109.1. Schema reference

Property	Description
replicas	The number of pods in the Kafka Connect group.
integer	
image	The docker image for the pods.
string	
buildResources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
<a href="#">ResourceRequirements</a>	
livenessProbe	Pod liveness checking.
<a href="#">Probe</a>	
readinessProbe	Pod readiness checking.
<a href="#">Probe</a>	
jvmOptions	JVM Options for pods.
<a href="#">JvmOptions</a>	
affinity	<b>The property affinity has been deprecated. This feature should now be configured at path spec.template.pod.affinity.</b> The pod's affinity rules. For more information, see the <a href="#">external documentation for core/v1 affinity</a> .
<a href="#">Affinity</a>	
logging	Logging configuration for Kafka Connect. The type depends on the value of the <a href="#">logging.type</a> property within the given object, which must be one of [inline, external].
<a href="#">InlineLogging, ExternalLogging</a>	
metrics	<b>The property metrics has been deprecated. This feature should now be configured at path spec.metricsConfig.</b> The Prometheus JMX Exporter configuration. See <a href="https://github.com/prometheus/jmx_exporter">https://github.com/prometheus/jmx_exporter</a> for details of the structure of this configuration.
map	
template	Template for Kafka Connect and Kafka Connect S2I resources. The template allows users to specify how the <a href="#">Deployment, Pods and Service</a> are generated.
<a href="#">KafkaConnectTemplate</a>	

Property	Description
authentication	Authentication configuration for Kafka Connect. The type depends on the value of the <a href="#">authentication.type</a> property within the given object, which must be one of [tls, scram-sha-512, plain, oauth].
<a href="#">KafkaClientAuthenticationTls</a> , <a href="#">KafkaClientAuthenticationScramSha512</a> , <a href="#">KafkaClientAuthenticationPlain</a> , <a href="#">KafkaClientAuthenticationOAuth</a>	
bootstrapServers	Bootstrap servers to connect to. This should be given as a comma separated list of <hostname>:<port> pairs.
string	
build	Configures how the Connect container image should be built. Optional.
<a href="#">Build</a>	
clientRackInitImage	The image of the init container used for initializing the <a href="#">client.rack</a> .
string	
config	The Kafka Connect configuration. Properties with the following prefixes cannot be set: ssl., sasl., security., listeners, plugin.path, rest., bootstrap.servers, consumer.interceptor.classes, producer.interceptor.classes (with the exception of: ssl.endpoint.identification.algorithm, ssl.cipher.suites, ssl.protocol, ssl.enabled.protocols).
map	
externalConfiguration	Pass data from Secrets or ConfigMaps to the Kafka Connect pods and use them to configure connectors.
<a href="#">ExternalConfiguration</a>	
insecureSourceRepository	When true this configures the source repository with the 'Local' reference policy and an import policy that accepts insecure source tags.
boolean	
metricsConfig	Metrics configuration. The type depends on the value of the <a href="#">metricsConfig.type</a> property within the given object, which must be one of [jmxPrometheusExporter].
<a href="#">JmxPrometheusExporterMetrics</a>	
rack	Configuration of the node label which will be used as the client.rack consumer configuration.
<a href="#">Rack</a>	
resources	The maximum limits for CPU and memory resources and the requested initial resources. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
<a href="#">ResourceRequirements</a>	
tls	TLS configuration.
<a href="#">KafkaConnectTls</a>	

Property	Description
tolerations	The property <b>tolerations</b> has been deprecated. This feature should now be configured at path <a href="#">spec.template.pod.tolerations</a> . The pod's tolerations. For more information, see the <a href="#">external documentation for core/v1 toleration</a> .
Toleration array	
tracing	The configuration of tracing in Kafka Connect. The type depends on the value of the <b>tracing.type</b> property within the given object, which must be one of [jaeger].
JaegerTracing	
version	The Kafka Connect version. Defaults to 2.7.0. Consult the user documentation to understand the process required to upgrade or downgrade the version.
string	

## B.110. KafkaConnectS2IStatus schema reference

Used in: [KafkaConnectS2I](#)

Property	Description
conditions	List of status conditions.
Condition array	
observedGeneration	The generation of the CRD that was last reconciled by the operator.
integer	
url	The URL of the REST API endpoint for managing and monitoring Kafka Connect connectors.
string	
connectorPlugins	The list of connector plugins available in this Kafka Connect deployment.
ConnectorPlugin array	
buildConfigName	The name of the build configuration.
string	
labelSelector	Label selector for pods providing this resource.
string	
replicas	The current number of pods being used to provide this resource.
integer	

## B.111. KafkaTopic schema reference

Property	Description
spec	The specification of the topic.
<a href="#">KafkaTopicSpec</a>	
status	The status of the topic.
<a href="#">KafkaTopicStatus</a>	

## B.112. KafkaTopicSpec schema reference

Used in: [KafkaTopic](#)

Property	Description
partitions	The number of partitions the topic should have. This cannot be decreased after topic creation. It can be increased after topic creation, but it is important to understand the consequences that has, especially for topics with semantic partitioning.
integer	
replicas	The number of replicas the topic should have.
integer	
config	The topic configuration.
map	
topicName	The name of the topic. When absent this will default to the metadata.name of the topic. It is recommended to not set this unless the topic name is not a valid Kubernetes resource name.
string	

## B.113. KafkaTopicStatus schema reference

Used in: [KafkaTopic](#)

Property	Description
conditions	List of status conditions.
<a href="#">Condition</a> array	
observedGeneration	The generation of the CRD that was last reconciled by the operator.
integer	

## B.114. KafkaUser schema reference

Property	Description
spec	The specification of the user.
KafkaUserSpec	
status	The status of the Kafka User.
KafkaUserStatus	

## B.115. KafkaUserSpec schema reference

Used in: [KafkaUser](#)

Property	Description
authentication	Authentication mechanism enabled for this Kafka user. The type depends on the value of the <a href="#">authentication.type</a> property within the given object, which must be one of [tls, scram-sha-512].
KafkaUserTlsClientAuthentication, KafkaUserScramSha512ClientAuthentication	
authorization	Authorization rules for this Kafka user. The type depends on the value of the <a href="#">authorization.type</a> property within the given object, which must be one of [simple].
KafkaUserAuthorizationSimple	
quotas	Quotas on requests to control the broker resources used by clients. Network bandwidth and request rate quotas can be enforced. Kafka documentation for Kafka User quotas can be found at <a href="http://kafka.apache.org/documentation/#design_quotas">http://kafka.apache.org/documentation/#design_quotas</a> .
KafkaUserQuotas	
template	Template to specify how Kafka User <a href="#">Secrets</a> are generated.
KafkaUserTemplate	

## B.116. KafkaUserTlsClientAuthentication schema reference

Used in: [KafkaUserSpec](#)

The `type` property is a discriminator that distinguishes use of the [KafkaUserTlsClientAuthentication](#) type from [KafkaUserScramSha512ClientAuthentication](#). It must have the value `tls` for the type [KafkaUserTlsClientAuthentication](#).

Property	Description
type	Must be <code>tls</code> .
string	

## B.117. KafkaUserScramSha512ClientAuthentication schema reference

Used in: [KafkaUserSpec](#)

The `type` property is a discriminator that distinguishes use of the `KafkaUserScramSha512ClientAuthentication` type from `KafkaUserTlsClientAuthentication`. It must have the value `scram-sha-512` for the type `KafkaUserScramSha512ClientAuthentication`.

Property	Description
<code>type</code>	Must be <code>scram-sha-512</code> .
<code>string</code>	

## B.118. KafkaUserAuthorizationSimple schema reference

Used in: [KafkaUserSpec](#)

The `type` property is a discriminator that distinguishes use of the `KafkaUserAuthorizationSimple` type from other subtypes which may be added in the future. It must have the value `simple` for the type `KafkaUserAuthorizationSimple`.

Property	Description
<code>type</code>	Must be <code>simple</code> .
<code>string</code>	
<code>acls</code>	List of ACL rules which should be applied to this user.
<code>AclRule</code> array	

## B.119. AclRule schema reference

Used in: [KafkaUserAuthorizationSimple](#)

Configures access control rule for a `KafkaUser` when brokers are using the `AclAuthorizer`.

## Example KafkaUser configuration with authorization

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: my-user
  labels:
    strimzi.io/cluster: my-cluster
spec:
  # ...
  authorization:
    type: simple
    acls:
      - resource:
          type: topic
          name: my-topic
          patternType: literal
          operation: Read
      - resource:
          type: topic
          name: my-topic
          patternType: literal
          operation: Describe
      - resource:
          type: group
          name: my-group
          patternType: prefix
          operation: Read
```

### B.119.1. `resource`

Use the `resource` property to specify the resource that the rule applies to.

Simple authorization supports four resource types, which are specified in the `type` property:

- Topics ([topic](#))
- Consumer Groups ([group](#))
- Clusters ([cluster](#))
- Transactional IDs ([transactionalID](#))

For Topic, Group, and Transactional ID resources you can specify the name of the resource the rule applies to in the `name` property.

Cluster type resources have no name.

A name is specified as a `literal` or a `prefix` using the `patternType` property.

- Literal names are taken exactly as they are specified in the `name` field.
- Prefix names use the value from the `name` as a prefix, and will apply the rule to all resources

with names starting with the value.

### B.119.2. type

The `type` of rule, which is to `allow` or `deny` (not currently supported) an operation.

The `type` field is optional. If `type` is unspecified, the ACL rule is treated as an `allow` rule.

### B.119.3. operation

Specify an `operation` for the rule to allow or deny.

The following operations are supported:

- Read
- Write
- Delete
- Alter
- Describe
- All
- IdempotentWrite
- ClusterAction
- Create
- AlterConfigs
- DescribeConfigs

Only certain operations work with each resource.

For more details about `AclAuthorizer`, ACLs and supported combinations of resources and operations, see [Authorization and ACLs](#).

### B.119.4. host

Use the `host` property to specify a remote host from which the rule is allowed or denied.

Use an asterisk (\*) to allow or deny the operation from all hosts. The `host` field is optional. If `host` is unspecified, the \* value is used by default.

### B.119.5. Schema reference

Property	Description
host	The host from which the action described in the ACL rule is allowed or denied.
string	

Property	Description
operation	Operation which will be allowed or denied. Supported operations are: Read, Write, Create, Delete, Alter, Describe, ClusterAction, AlterConfigs, DescribeConfigs, IdempotentWrite and All.
resource	Indicates the resource for which given ACL rule applies. The type depends on the value of the <code>resource.type</code> property within the given object, which must be one of [topic, group, cluster, transactionalId].
AclRuleTopicResource, AclRuleGroupResource, AclRuleClusterResource, AclRuleTransactionalIdResource	
type	The type of the rule. Currently the only supported type is <code>allow</code> . ACL rules with type <code>allow</code> are used to allow user to execute the specified operations. Default value is <code>allow</code> .
string (one of [allow, deny])	

## B.120. AclRuleTopicResource schema reference

Used in: [AclRule](#)

The `type` property is a discriminator that distinguishes use of the `AclRuleTopicResource` type from `AclRuleGroupResource`, `AclRuleClusterResource`, `AclRuleTransactionalIdResource`. It must have the value `topic` for the type `AclRuleTopicResource`.

Property	Description
type	Must be <code>topic</code> .
string	
name	Name of resource for which given ACL rule applies. Can be combined with <code>patternType</code> field to use prefix pattern.
string	
patternType	Describes the pattern used in the resource field. The supported types are <code>literal</code> and <code>prefix</code> . With <code>literal</code> pattern type, the resource field will be used as a definition of a full topic name. With <code>prefix</code> pattern type, the resource name will be used only as a prefix. Default value is <code>literal</code> .
string (one of [prefix, literal])	

## B.121. AclRuleGroupResource schema reference

Used in: [AclRule](#)

The `type` property is a discriminator that distinguishes use of the `AclRuleGroupResource` type from `AclRuleTopicResource`, `AclRuleClusterResource`, `AclRuleTransactionalIdResource`. It must have the

value `group` for the type `AclRuleGroupResource`.

Property	Description
type	Must be <code>group</code> .
string	
name	Name of resource for which given ACL rule applies. Can be combined with <code>patternType</code> field to use prefix pattern.
string	
patternType	Describes the pattern used in the resource field. The supported types are <code>literal</code> and <code>prefix</code> . With <code>literal</code> pattern type, the resource field will be used as a definition of a full topic name. With <code>prefix</code> pattern type, the resource name will be used only as a prefix. Default value is <code>literal</code> .
string (one of [prefix, literal])	

## B.122. `AclRuleClusterResource` schema reference

Used in: `AclRule`

The `type` property is a discriminator that distinguishes use of the `AclRuleClusterResource` type from `AclRuleTopicResource`, `AclRuleGroupResource`, `AclRuleTransactionalIdResource`. It must have the value `cluster` for the type `AclRuleClusterResource`.

Property	Description
type	Must be <code>cluster</code> .
string	

## B.123. `AclRuleTransactionalIdResource` schema reference

Used in: `AclRule`

The `type` property is a discriminator that distinguishes use of the `AclRuleTransactionalIdResource` type from `AclRuleTopicResource`, `AclRuleGroupResource`, `AclRuleClusterResource`. It must have the value `transactionalId` for the type `AclRuleTransactionalIdResource`.

Property	Description
type	Must be <code>transactionalId</code> .
string	
name	Name of resource for which given ACL rule applies. Can be combined with <code>patternType</code> field to use prefix pattern.
string	

Property	Description
patternType	Describes the pattern used in the resource field. The supported types are <code>literal</code> and <code>prefix</code> . With <code>literal</code> pattern type, the resource field will be used as a definition of a full name. With <code>prefix</code> pattern type, the resource name will be used only as a prefix. Default value is <code>literal</code> .
string (one of [prefix, literal])	

## B.124. KafkaUserQuotas schema reference

Used in: [KafkaUserSpec](#)

Kafka allows a user to set `quotas` to control the use of resources by clients.

### B.124.1. `quotas`

Quotas split into two categories:

- *Network usage* quotas, which are defined as the byte rate threshold for each group of clients sharing a quota
- *CPU utilization* quotas, which are defined as the percentage of time a client can utilize on request handler I/O threads and network threads of each broker within a quota window

Using quotas for Kafka clients might be useful in a number of situations. Consider a wrongly configured Kafka producer which is sending requests at too high a rate. Such misconfiguration can cause a denial of service to other clients, so the problematic client ought to be blocked. By using a network limiting quota, it is possible to prevent this situation from significantly impacting other clients.

Strimzi supports user-level quotas, but not client-level quotas.

An example Kafka user quotas

```
spec:
  quotas:
    producerByteRate: 1048576
    consumerByteRate: 2097152
    requestPercentage: 55
```

For more info about Kafka user quotas, refer to the [Apache Kafka documentation](#).

### B.124.2. Schema reference

Property	Description
consumerByteRate	A quota on the maximum bytes per-second that each client group can fetch from a broker before the clients in the group are throttled. Defined on a per-broker basis.
integer	
producerByteRate	A quota on the maximum bytes per-second that each client group can publish to a broker before the clients in the group are throttled. Defined on a per-broker basis.
integer	
requestPercentage	A quota on the maximum CPU utilization of each client group as a percentage of network and I/O threads.
integer	

## B.125. KafkaUserTemplate schema reference

Used in: [KafkaUserSpec](#)

Specify additional labels and annotations for the secret created by the User Operator.

*An example showing the KafkaUserTemplate*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: my-user
  labels:
    strimzi.io/cluster: my-cluster
spec:
  authentication:
    type: tls
  template:
    secret:
      metadata:
        labels:
          label1: value1
        annotations:
          anno1: value1
    # ...
```

### B.125.1. Schema reference

Property	Description
secret	Template for KafkaUser resources. The template allows users to specify how the <a href="#">Secret</a> with password or TLS certificates is generated.
<a href="#">ResourceTemplate</a>	

## B.126. KafkaUserStatus schema reference

Used in: [KafkaUser](#)

Property	Description
conditions	List of status conditions.
Condition array	
observedGeneration	The generation of the CRD that was last reconciled by the operator.
integer	
username	Username.
string	
secret	The name of <a href="#">Secret</a> where the credentials are stored.
string	

## B.127. KafkaMirrorMaker schema reference

Property	Description
spec	The specification of Kafka MirrorMaker.
<a href="#">KafkaMirrorMakerSpec</a>	
status	The status of Kafka MirrorMaker.
<a href="#">KafkaMirrorMakerStatus</a>	

## B.128. KafkaMirrorMakerSpec schema reference

Used in: [KafkaMirrorMaker](#)

Configures Kafka MirrorMaker.

### B.128.1. whitelist

Use the `whitelist` property to configure a list of topics that Kafka MirrorMaker mirrors from the source to the target Kafka cluster.

The property allows any regular expression from the simplest case with a single topic name to complex patterns. For example, you can mirror topics A and B using "A|B" or all topics using "\*". You can also pass multiple regular expressions separated by commas to the Kafka MirrorMaker.

### B.128.2. KafkaMirrorMakerConsumerSpec and KafkaMirrorMakerProducerSpec

Use the [KafkaMirrorMakerConsumerSpec](#) and [KafkaMirrorMakerProducerSpec](#) to configure source (consumer) and target (producer) clusters.

Kafka MirrorMaker always works together with two Kafka clusters (source and target). To establish

a connection, the bootstrap servers for the source and the target Kafka clusters are specified as comma-separated lists of `HOSTNAME:PORT` pairs. Each comma-separated list contains one or more Kafka brokers or a `Service` pointing to Kafka brokers specified as a `HOSTNAME:PORT` pair.

### B.128.3. logging

Kafka MirrorMaker has its own configurable logger:

- `mirrormaker.root.logger`

MirrorMaker uses the Apache `log4j` logger implementation.

Use the `logging` property to configure loggers and logger levels.

You can set the log levels by specifying the logger and level directly (inline) or use a custom (external) ConfigMap. If a ConfigMap is used, you set `logging.name` property to the name of the ConfigMap containing the external logging configuration. Inside the ConfigMap, the logging configuration is described using `log4j.properties`. For more information about log levels, see [Apache logging services](#).

Here we see examples of `inline` and `external` logging:

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaMirrorMaker
spec:
  # ...
  logging:
    type: inline
    loggers:
      mirrormaker.root.logger: "INFO"
  # ...
```

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaMirrorMaker
spec:
  # ...
  logging:
    type: external
    name: customConfigMap
  # ...
```

*Garbage collector (GC)*

Garbage collector logging can also be enabled (or disabled) using the `jvmOptions` property.

### B.128.4. Schema reference

Property	Description
replicas	The number of pods in the <a href="#">Deployment</a> .
integer	
image	The docker image for the pods.
string	
whitelist	List of topics which are included for mirroring. This option allows any regular expression using Java-style regular expressions. Mirroring two topics named A and B is achieved by using the whitelist ' <code>A B</code> '. Or, as a special case, you can mirror all topics using the whitelist ' <code>*</code> '. You can also specify multiple regular expressions separated by commas.
string	
consumer	Configuration of source cluster.
<a href="#">KafkaMirrorMakerConsumerSpec</a>	
producer	Configuration of target cluster.
<a href="#">KafkaMirrorMakerProducerSpec</a>	
resources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resource requirements</a> .
ResourceRequirements	
affinity	<b>The property <code>affinity</code> has been deprecated.</b>
<a href="#">Affinity</a>	<b>This feature should now be configured at path <code>spec.template.pod.affinity</code>.</b> The pod's affinity rules. For more information, see the <a href="#">external documentation for core/v1 affinity</a> .
tolerations	<b>The property <code>tolerations</code> has been deprecated.</b> This feature should now be configured at path
<a href="#">Toleration array</a>	<b><code>spec.template.pod.tolerations</code>.</b> The pod's tolerations. For more information, see the <a href="#">external documentation for core/v1 toleration</a> .
jvmOptions	JVM Options for pods.
<a href="#">JvmOptions</a>	
logging	Logging configuration for MirrorMaker. The type depends on the value of the <a href="#">logging.type</a> property within the given object, which must be one of [inline, external].
<a href="#">InlineLogging, ExternalLogging</a>	

Property	Description
metrics	The property <code>metrics</code> has been deprecated. This feature should now be configured at path <code>spec.metricsConfig</code> . The Prometheus JMX Exporter configuration. See <a href="#">JMX Exporter documentation</a> for details of the structure of this configuration.
map	
metricsConfig	Metrics configuration. The type depends on the value of the <code>metricsConfig.type</code> property within the given object, which must be one of [ <code>jmxPrometheusExporter</code> ].
<code>JmxPrometheusExporterMetrics</code>	
tracing	The configuration of tracing in Kafka MirrorMaker. The type depends on the value of the <code>tracing.type</code> property within the given object, which must be one of [ <code>jaeger</code> ].
<code>JaegerTracing</code>	
template	Template to specify how Kafka MirrorMaker resources, <code>Deployments</code> and <code>Pods</code> , are generated.
<code>KafkaMirrorMakerTemplate</code>	
livenessProbe	Pod liveness checking.
<code>Probe</code>	
readinessProbe	Pod readiness checking.
<code>Probe</code>	
version	The Kafka MirrorMaker version. Defaults to 2.7.0. Consult the documentation to understand the process required to upgrade or downgrade the version.
string	

## B.129. KafkaMirrorMakerConsumerSpec schema reference

Used in: [KafkaMirrorMakerSpec](#)

Configures a MirrorMaker consumer.

### B.129.1. numStreams

Use the `consumer.numStreams` property to configure the number of streams for the consumer.

You can increase the throughput in mirroring topics by increasing the number of consumer threads. Consumer threads belong to the consumer group specified for Kafka MirrorMaker. Topic partitions are assigned across the consumer threads, which consume messages in parallel.

### B.129.2. offsetCommitInterval

Use the `consumer.offsetCommitInterval` property to configure an offset auto-commit interval for the consumer.

You can specify the regular time interval at which an offset is committed after Kafka MirrorMaker has consumed data from the source Kafka cluster. The time interval is set in milliseconds, with a default value of 60,000.

### B.129.3. `config`

Use the `consumer.config` properties to configure Kafka options for the consumer.

The `config` property contains the Kafka MirrorMaker consumer configuration options as keys, with values set in one of the following JSON types:

- String
- Number
- Boolean

For client connection using a specific *cipher suite* for a TLS version, you can [configure allowed ssl properties](#). You can also [configure the `ssl.endpoint.identification.algorithm` property](#) to enable or disable hostname verification.

### Exceptions

You can specify and configure the options listed in the [Apache Kafka configuration documentation for consumers](#).

However, there are exceptions for options automatically configured and managed directly by Strimzi related to:

- Kafka cluster bootstrap address
- Security (encryption, authentication, and authorization)
- Consumer group identifier
- Interceptors

Specifically, all configuration options with keys equal to or starting with one of the following strings are forbidden:

- `bootstrap.servers`
- `group.id`
- `interceptor.classes`
- `ssl.` (not including specific exceptions)
- `sasl.`
- `security.`

When a forbidden option is present in the `config` property, it is ignored and a warning message is printed to the Cluster Operator log file. All other options are passed to Kafka MirrorMaker.

**IMPORTANT**

The Cluster Operator does not validate keys or values in the provided `config` object. When an invalid configuration is provided, the Kafka MirrorMaker might not start or might become unstable. In such cases, the configuration in the `KafkaMirrorMaker.spec.consumer.config` object should be fixed and the Cluster Operator will roll out the new configuration for Kafka MirrorMaker.

### B.129.4. `groupId`

Use the `consumer.groupId` property to configure a consumer group identifier for the consumer.

Kafka MirrorMaker uses a Kafka consumer to consume messages, behaving like any other Kafka consumer client. Messages consumed from the source Kafka cluster are mirrored to a target Kafka cluster. A group identifier is required, as the consumer needs to be part of a consumer group for the assignment of partitions.

### B.129.5. Schema reference

Property	Description
<code>numStreams</code> integer	Specifies the number of consumer stream threads to create.
<code>offsetCommitInterval</code> integer	Specifies the offset auto-commit interval in ms. Default value is 60000.
<code>groupId</code> string	A unique string that identifies the consumer group this consumer belongs to.
<code>bootstrapServers</code> string	A list of host:port pairs for establishing the initial connection to the Kafka cluster.
<code>authentication</code>  <code>KafkaClientAuthenticationTls</code> , <code>KafkaClientAuthenticationScramSha512</code> , <code>KafkaClientAuthenticationPlain</code> , <code>KafkaClientAuthenticationOAuth</code>	Authentication configuration for connecting to the cluster. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, plain, oauth].
<code>config</code>  map	The MirrorMaker consumer config. Properties with the following prefixes cannot be set: ssl., bootstrap.servers, group.id, sasl., security., interceptor.classes (with the exception of: ssl.endpoint.identification.algorithm, ssl.cipher.suites, ssl.protocol, ssl.enabled.protocols).
<code>tls</code>  <code>KafkaMirrorMakerTls</code>	TLS configuration for connecting MirrorMaker to the cluster.

## B.130. KafkaMirrorMakerTls schema reference

Used in: [KafkaMirrorMakerConsumerSpec](#), [KafkaMirrorMakerProducerSpec](#)

Configures TLS trusted certificates for connecting MirrorMaker to the cluster.

### B.130.1. trustedCertificates

Provide a list of secrets using the [trustedCertificates](#) property.

### B.130.2. Schema reference

Property	Description
trustedCertificates	Trusted certificates for TLS connection.
CertSecretSource array	

## B.131. KafkaMirrorMakerProducerSpec schema reference

Used in: [KafkaMirrorMakerSpec](#)

Configures a MirrorMaker producer.

### B.131.1. abortOnSendFailure

Use the [producer.abortOnSendFailure](#) property to configure how to handle message send failure from the producer.

By default, if an error occurs when sending a message from Kafka MirrorMaker to a Kafka cluster:

- The Kafka MirrorMaker container is terminated in Kubernetes.
- The container is then recreated.

If the [abortOnSendFailure](#) option is set to [false](#), message sending errors are ignored.

### B.131.2. config

Use the [producer.config](#) properties to configure Kafka options for the producer.

The [config](#) property contains the Kafka MirrorMaker producer configuration options as keys, with values set in one of the following JSON types:

- String
- Number
- Boolean

For client connection using a specific *cipher suite* for a TLS version, you can [configure allowed ssl properties](#). You can also [configure the ssl.endpoint.identification.algorithm property](#) to enable or

disable hostname verification.

## Exceptions

You can specify and configure the options listed in the [Apache Kafka configuration documentation for producers](#).

However, there are exceptions for options automatically configured and managed directly by Strimzi related to:

- Kafka cluster bootstrap address
- Security (encryption, authentication, and authorization)
- Interceptors

Specifically, all configuration options with keys equal to or starting with one of the following strings are forbidden:

- `bootstrap.servers`
- `interceptor.classes`
- `ssl.` (not including specific exceptions)
- `sasl.`
- `security.`

When a forbidden option is present in the `config` property, it is ignored and a warning message is printed to the Cluster Operator log file. All other options are passed to Kafka MirrorMaker.

**IMPORTANT**

The Cluster Operator does not validate keys or values in the provided `config` object. When an invalid configuration is provided, the Kafka MirrorMaker might not start or might become unstable. In such cases, the configuration in the `KafkaMirrorMaker.spec.producer.config` object should be fixed and the Cluster Operator will roll out the new configuration for Kafka MirrorMaker.

### B.131.3. Schema reference

Property	Description
<code>bootstrapServers</code> string	A list of host:port pairs for establishing the initial connection to the Kafka cluster.
<code>abortOnSendFailure</code> boolean	Flag to set the MirrorMaker to exit on a failed send. Default value is <code>true</code> .

Property	Description
authentication	Authentication configuration for connecting to the cluster. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, plain, oauth].
<code>KafkaClientAuthenticationTls</code> , <code>KafkaClientAuthenticationScramSha512</code> , <code>KafkaClientAuthenticationPlain</code> , <code>KafkaClientAuthenticationOAuth</code>	
config	The MirrorMaker producer config. Properties with the following prefixes cannot be set: ssl., bootstrap.servers, sasl., security., interceptor.classes (with the exception of: ssl.endpoint.identification.algorithm, ssl.cipher.suites, ssl.protocol, ssl.enabled.protocols).
map	
tls	TLS configuration for connecting MirrorMaker to the cluster.
<code>KafkaMirrorMakerTls</code>	

## B.132. KafkaMirrorMakerTemplate schema reference

Used in: `KafkaMirrorMakerSpec`

Property	Description
deployment	Template for Kafka MirrorMaker <code>Deployment</code> .
<code>DeploymentTemplate</code>	
pod	Template for Kafka MirrorMaker <code>Pods</code> .
<code>PodTemplate</code>	
mirrorMakerContainer	Template for Kafka MirrorMaker container.
<code>ContainerTemplate</code>	
podDisruptionBudget	Template for Kafka MirrorMaker <code>PodDisruptionBudget</code> .
<code>PodDisruptionBudgetTemplate</code>	

## B.133. KafkaMirrorMakerStatus schema reference

Used in: `KafkaMirrorMaker`

Property	Description
conditions	List of status conditions.
<code>Condition</code> array	
observedGeneration	The generation of the CRD that was last reconciled by the operator.
integer	

Property	Description
labelSelector	Label selector for pods providing this resource.
string	
replicas	The current number of pods being used to provide this resource.
integer	

## B.134. KafkaBridge schema reference

Property	Description
spec	The specification of the Kafka Bridge.
<a href="#">KafkaBridgeSpec</a>	
status	The status of the Kafka Bridge.
<a href="#">KafkaBridgeStatus</a>	

## B.135. KafkaBridgeSpec schema reference

Used in: [KafkaBridge](#)

Configures a Kafka Bridge cluster.

Configuration options relate to:

- Kafka cluster bootstrap address
- Security (Encryption, Authentication, and Authorization)
- Consumer configuration
- Producer configuration
- HTTP configuration

### B.135.1. logging

Kafka Bridge has its own configurable loggers:

- `logger.bridge`
- `logger.<operation-id>`

You can replace `<operation-id>` in the `logger.<operation-id>` logger to set log levels for specific operations:

- `createConsumer`
- `deleteConsumer`
- `subscribe`
- `unsubscribe`

- `poll`
- `assign`
- `commit`
- `send`
- `sendToPartition`
- `seekToBeginning`
- `seekToEnd`
- `seek`
- `healthy`
- `ready`
- `openapi`

Each operation is defined according OpenAPI specification, and has a corresponding API endpoint through which the bridge receives requests from HTTP clients. You can change the log level on each endpoint to create fine-grained logging information about the incoming and outgoing HTTP requests.

Each logger has to be configured assigning it a `name` as `http.openapi.operation.<operation-id>`. For example, configuring the logging level for the `send` operation logger means defining the following:

```
logger.send.name = http.openapi.operation.send
logger.send.level = DEBUG
```

Kafka Bridge uses the Apache `log4j2` logger implementation. Loggers are defined in the `log4j2.properties` file, which has the following default configuration for `healthy` and `ready` endpoints:

```
logger.healthy.name = http.openapi.operation.healthy
logger.healthy.level = WARN
logger.ready.name = http.openapi.operation.ready
logger.ready.level = WARN
```

The log level of all other operations is set to `INFO` by default.

Use the `logging` property to configure loggers and logger levels.

You can set the log levels by specifying the logger and level directly (inline) or use a custom (external) ConfigMap. If a ConfigMap is used, you set `logging.name` property to the name of the ConfigMap containing the external logging configuration. Inside the ConfigMap, the logging configuration is described using `log4j.properties`. For more information about log levels, see [Apache logging services](#).

Here we see examples of `inline` and `external` logging.

### *Inline logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaBridge
spec:
  # ...
  logging:
    type: inline
    loggers:
      logger.bridge.level: "INFO"
      # enabling DEBUG just for send operation
      logger.send.name: "http.openapi.operation.send"
      logger.send.level: "DEBUG"
  # ...
```

### *External logging*

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaBridge
spec:
  # ...
  logging:
    type: external
    name: customConfigMap
  # ...
```

Any available loggers that are not configured have their level set to **OFF**.

If the Kafka Bridge was deployed using the Cluster Operator, changes to Kafka Bridge logging levels are applied dynamically.

If you use external logging, a rolling update is triggered when logging appenders are changed.

#### *Garbage collector (GC)*

Garbage collector logging can also be enabled (or disabled) using the `jvmOptions` property.

## B.135.2. Schema reference

Property	Description
replicas	The number of pods in the <a href="#">Deployment</a> .
integer	
image	The docker image for the pods.
string	
bootstrapServers	A list of host:port pairs for establishing the initial connection to the Kafka cluster.
string	

Property	Description
tls	TLS configuration for connecting Kafka Bridge to the cluster.
<a href="#">KafkaBridgeTls</a>	
authentication	Authentication configuration for connecting to the cluster. The type depends on the value of the <a href="#">authentication.type</a> property within the given object, which must be one of [tls, scram-sha-512, plain, oauth].
<a href="#">KafkaClientAuthenticationTls</a> , <a href="#">KafkaClientAuthenticationScramSha512</a> , <a href="#">KafkaClientAuthenticationPlain</a> , <a href="#">KafkaClientAuthenticationOAuth</a>	
http	The HTTP related configuration.
<a href="#">KafkaBridgeHttpConfig</a>	
consumer	Kafka consumer related configuration.
<a href="#">KafkaBridgeConsumerSpec</a>	
producer	Kafka producer related configuration.
<a href="#">KafkaBridgeProducerSpec</a>	
resources	CPU and memory resources to reserve. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
<a href="#">ResourceRequirements</a>	
jvmOptions	<b>Currently not supported</b> JVM Options for pods.
<a href="#">JvmOptions</a>	
logging	Logging configuration for Kafka Bridge. The type depends on the value of the <a href="#">logging.type</a> property within the given object, which must be one of [inline, external].
<a href="#">InlineLogging</a> , <a href="#">ExternalLogging</a>	
enableMetrics	Enable the metrics for the Kafka Bridge. Default is false.
boolean	
livenessProbe	Pod liveness checking.
<a href="#">Probe</a>	
readinessProbe	Pod readiness checking.
<a href="#">Probe</a>	
template	Template for Kafka Bridge resources. The template allows users to specify how is the <a href="#">Deployment</a> and <a href="#">Pods</a> generated.
<a href="#">KafkaBridgeTemplate</a>	
tracing	The configuration of tracing in Kafka Bridge. The type depends on the value of the <a href="#">tracing.type</a> property within the given object, which must be one of [jaeger].
<a href="#">JaegerTracing</a>	

## B.136. KafkaBridgeTls schema reference

Used in: [KafkaBridgeSpec](#)

Property	Description
trustedCertificates	Trusted certificates for TLS connection.
<a href="#">CertSecretSource</a> array	

## B.137. KafkaBridgeHttpConfig schema reference

Used in: [KafkaBridgeSpec](#)

Configures HTTP access to a Kafka cluster for the Kafka Bridge.

The default HTTP configuration is for the Kafka Bridge to listen on port 8080.

### B.137.1. cors

As well as enabling HTTP access to a Kafka cluster, HTTP properties provide the capability to enable and define access control for the Kafka Bridge through Cross-Origin Resource Sharing (CORS). CORS is a HTTP mechanism that allows browser access to selected resources from more than one origin. To configure CORS, you define a list of allowed resource origins and HTTP access methods. For the origins, you can use a URL or a Java regular expression.

*Example Kafka Bridge HTTP configuration*

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaBridge
metadata:
  name: my-bridge
spec:
  # ...
  http:
    port: 8080
    cors:
      allowedOrigins: "https://strimzi.io"
      allowedMethods: "GET,POST,PUT,DELETE,OPTIONS,PATCH"
  # ...
```

### B.137.2. Schema reference

Property	Description
port	The port which is the server listening on.
integer	
cors	CORS configuration for the HTTP Bridge.
<a href="#">KafkaBridgeHttpCors</a>	

## B.138. KafkaBridgeHttpCors schema reference

Used in: [KafkaBridgeHttpConfig](#)

Property	Description
allowedOrigins	List of allowed origins. Java regular expressions can be used.
string array	
allowedMethods	List of allowed HTTP methods.
string array	

## B.139. KafkaBridgeConsumerSpec schema reference

Used in: [KafkaBridgeSpec](#)

Configures consumer options for the Kafka Bridge as keys.

The values can be one of the following JSON types:

- String
- Number
- Boolean

You can specify and configure the options listed in the [Apache Kafka configuration documentation for consumers](#) with the exception of those options which are managed directly by Strimzi. Specifically, all configuration options with keys equal to or starting with one of the following strings are forbidden:

- `ssl.`
- `sasl.`
- `security.`
- `bootstrap.servers`
- `group.id`

When one of the forbidden options is present in the `config` property, it is ignored and a warning message will be printed to the Cluster Operator log file. All other options will be passed to Kafka

**IMPORTANT**

The Cluster Operator does not validate keys or values in the `config` object. If an invalid configuration is provided, the Kafka Bridge cluster might not start or might become unstable. Fix the configuration so that the Cluster Operator can roll out the new configuration to all Kafka Bridge nodes.

There are exceptions to the forbidden options. For client connection using a specific *cipher suite* for a TLS version, you can [configure allowed `ssl` properties](#).

### Example Kafka Bridge consumer configuration

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaBridge
metadata:
  name: my-bridge
spec:
  # ...
  consumer:
    config:
      auto.offset.reset: earliest
      enable.auto.commit: true
      ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384"
      ssl.enabled.protocols: "TLSv1.2"
      ssl.protocol: "TLSv1.2"
      ssl.endpoint.identification.algorithm: HTTPS
  # ...
```

### B.139.1. Schema reference

Property	Description
config	The Kafka consumer configuration used for consumer instances created by the bridge. Properties with the following prefixes cannot be set: ssl., bootstrap.servers, group.id, sasl., security. (with the exception of: ssl.endpoint.identification.algorithm, ssl.cipher.suites, ssl.protocol, ssl.enabled.protocols).
map	

## B.140. KafkaBridgeProducerSpec schema reference

Used in: [KafkaBridgeSpec](#)

Configures producer options for the Kafka Bridge as keys.

The values can be one of the following JSON types:

- String
- Number
- Boolean

You can specify and configure the options listed in the [Apache Kafka configuration documentation for producers](#) with the exception of those options which are managed directly by Strimzi. Specifically, all configuration options with keys equal to or starting with one of the following strings are forbidden:

- `ssl`.
- `sasl`.
- `security`.
- `bootstrap.servers`

When one of the forbidden options is present in the `config` property, it is ignored and a warning message will be printed to the Cluster Operator log file. All other options will be passed to Kafka

#### IMPORTANT

The Cluster Operator does not validate keys or values in the `config` object. If an invalid configuration is provided, the Kafka Bridge cluster might not start or might become unstable. Fix the configuration so that the Cluster Operator can roll out the new configuration to all Kafka Bridge nodes.

There are exceptions to the forbidden options. For client connection using a specific *cipher suite* for a TLS version, you can [configure allowed `ssl` properties](#).

*Example Kafka Bridge producer configuration*

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaBridge
metadata:
  name: my-bridge
spec:
  # ...
  producer:
    config:
      acks: 1
      delivery.timeout.ms: 300000
      ssl.cipher.suites: "TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384"
      ssl.enabled.protocols: "TLSv1.2"
      ssl.protocol: "TLSv1.2"
      ssl.endpoint.identification.algorithm: HTTPS
  # ...
```

### B.140.1. Schema reference

Property	Description
<code>config</code>	The Kafka producer configuration used for producer instances created by the bridge. Properties with the following prefixes cannot be set: <code>ssl</code> ., <code>bootstrap.servers</code> , <code>sasl</code> ., <code>security</code> . (with the exception of: <code>ssl.endpoint.identification.algorithm</code> , <code>ssl.cipher.suites</code> , <code>ssl.protocol</code> , <code>ssl.enabled.protocols</code> ).
<code>map</code>	

## B.141. KafkaBridgeTemplate schema reference

Used in: [KafkaBridgeSpec](#)

Property	Description
deployment	Template for Kafka Bridge <a href="#">Deployment</a> .
<a href="#">DeploymentTemplate</a>	
pod	Template for Kafka Bridge <a href="#">Pods</a> .
<a href="#">PodTemplate</a>	
apiService	Template for Kafka Bridge API <a href="#">Service</a> .
<a href="#">ResourceTemplate</a>	
bridgeContainer	Template for the Kafka Bridge container.
<a href="#">ContainerTemplate</a>	
podDisruptionBudget	Template for Kafka Bridge <a href="#">PodDisruptionBudget</a> .
<a href="#">PodDisruptionBudgetTemplate</a>	

## B.142. KafkaBridgeStatus schema reference

Used in: [KafkaBridge](#)

Property	Description
conditions	List of status conditions.
<a href="#">Condition</a> array	
observedGeneration	The generation of the CRD that was last reconciled by the operator.
integer	
url	The URL at which external client applications can access the Kafka Bridge.
string	
labelSelector	Label selector for pods providing this resource.
string	
replicas	The current number of pods being used to provide this resource.
integer	

## B.143. KafkaConnector schema reference

Property	Description
spec	The specification of the Kafka Connector.
<a href="#">KafkaConnectorSpec</a>	

Property	Description
status	The status of the Kafka Connector.
KafkaConnectorStatus	

## B.144. KafkaConnectorSpec schema reference

Used in: [KafkaConnector](#)

Property	Description
class	The Class for the Kafka Connector.
string	
tasksMax	The maximum number of tasks for the Kafka Connector.
integer	
config	The Kafka Connector configuration. The following properties cannot be set: connector.class, tasks.max.
map	
pause	Whether the connector should be paused. Defaults to false.
boolean	

## B.145. KafkaConnectorStatus schema reference

Used in: [KafkaConnector](#)

Property	Description
conditions	List of status conditions.
Condition array	
observedGeneration	The generation of the CRD that was last reconciled by the operator.
integer	
connectorStatus	The connector status, as reported by the Kafka Connect REST API.
map	
tasksMax	The maximum number of tasks for the Kafka Connector.
integer	

## B.146. KafkaMirrorMaker2 schema reference

Property	Description
spec	The specification of the Kafka MirrorMaker 2.0 cluster.
KafkaMirrorMaker2Spec	

Property	Description
status	The status of the Kafka MirrorMaker 2.0 cluster.
KafkaMirrorMaker2Status	

## B.147. KafkaMirrorMaker2Spec schema reference

Used in: [KafkaMirrorMaker2](#)

Property	Description
replicas	The number of pods in the Kafka Connect group.
integer	
version	The Kafka Connect version. Defaults to 2.7.0. Consult the user documentation to understand the process required to upgrade or downgrade the version.
string	
image	The docker image for the pods.
string	
connectCluster	The cluster alias used for Kafka Connect. The alias must match a cluster in the list at <a href="#">spec.clusters</a> .
string	
clusters	Kafka clusters for mirroring.
<a href="#">KafkaMirrorMaker2ClusterSpec</a> array	
mirrors	Configuration of the MirrorMaker 2.0 connectors.
<a href="#">KafkaMirrorMaker2MirrorSpec</a> array	
resources	The maximum limits for CPU and memory resources and the requested initial resources. For more information, see the <a href="#">external documentation for core/v1 resourcerequirements</a> .
<a href="#">ResourceRequirements</a>	
livenessProbe	Pod liveness checking.
<a href="#">Probe</a>	
readinessProbe	Pod readiness checking.
<a href="#">Probe</a>	
jvmOptions	JVM Options for pods.
<a href="#">JvmOptions</a>	
affinity	<b>The property affinity has been deprecated. This feature should now be configured at path spec.template.pod.affinity.</b> The pod's affinity rules. For more information, see the <a href="#">external documentation for core/v1 affinity</a> .
<a href="#">Affinity</a>	

Property	Description
tolerations	The property <b>tolerations</b> has been deprecated. This feature should now be configured at path <code>spec.template.pod.tolerations</code> . The pod's tolerations. For more information, see the <a href="#">external documentation for core/v1 toleration</a> .
Toleration array	
logging	Logging configuration for Kafka Connect. The type depends on the value of the <code>logging.type</code> property within the given object, which must be one of [inline, external].
InlineLogging, ExternalLogging	
metrics	The property <b>metrics</b> has been deprecated. This feature should now be configured at path <code>spec.metricsConfig</code> . The Prometheus JMX Exporter configuration. See <a href="https://github.com/prometheus/jmx_exporter">https://github.com/prometheus/jmx_exporter</a> for details of the structure of this configuration.
map	
tracing	The configuration of tracing in Kafka Connect. The type depends on the value of the <code>tracing.type</code> property within the given object, which must be one of [jaeger].
JaegerTracing	
template	Template for Kafka Connect and Kafka Connect S2I resources. The template allows users to specify how the <code>Deployment</code> , <code>Pods</code> and <code>Service</code> are generated.
KafkaConnectTemplate	
externalConfiguration	Pass data from Secrets or ConfigMaps to the Kafka Connect pods and use them to configure connectors.
ExternalConfiguration	
metricsConfig	Metrics configuration. The type depends on the value of the <code>metricsConfig.type</code> property within the given object, which must be one of [jmxPrometheusExporter].
JmxPrometheusExporterMetrics	

## B.148. KafkaMirrorMaker2ClusterSpec schema reference

Used in: [KafkaMirrorMaker2Spec](#)

Configures Kafka clusters for mirroring.

### B.148.1. config

Use the `config` properties to configure Kafka options.

Standard Apache Kafka configuration may be provided, restricted to those properties not managed directly by Strimzi.

For client connection using a specific *cipher suite* for a TLS version, you can [configure allowed ssl properties](#). You can also [configure the ssl.endpoint.identification.algorithm property](#) to enable or disable hostname verification.

## B.148.2. Schema reference

Property	Description
alias	Alias used to reference the Kafka cluster.
string	
bootstrapServers	A comma-separated list of <code>host:port</code> pairs for establishing the connection to the Kafka cluster.
string	
config	The MirrorMaker 2.0 cluster config. Properties with the following prefixes cannot be set: <code>ssl.</code> , <code>sasl.</code> , <code>security.</code> , <code>listeners.</code> , <code>plugin.path.</code> , <code>rest.</code> , <code>bootstrap.servers.</code> , <code>consumer.interceptor.classes.</code> , <code>producer.interceptor.classes.</code> (with the exception of: <code>ssl.endpoint.identification.algorithm</code> , <code>ssl.cipher.suites</code> , <code>ssl.protocol</code> , <code>ssl.enabled.protocols</code> ).
map	
tls	TLS configuration for connecting MirrorMaker 2.0 connectors to a cluster.
<a href="#">KafkaMirrorMaker2Tls</a>	
authentication	Authentication configuration for connecting to the cluster. The type depends on the value of the <code>authentication.type</code> property within the given object, which must be one of [tls, scram-sha-512, plain, oauth].
<a href="#">KafkaClientAuthenticationTls</a> , <a href="#">KafkaClientAuthenticationScramSha512</a> , <a href="#">KafkaClientAuthenticationPlain</a> , <a href="#">KafkaClientAuthenticationOAuth</a>	

## B.149. [KafkaMirrorMaker2Tls](#) schema reference

Used in: [KafkaMirrorMaker2ClusterSpec](#)

Property	Description
trustedCertificates	Trusted certificates for TLS connection.
<a href="#">CertSecretSource</a> array	

## B.150. [KafkaMirrorMaker2MirrorSpec](#) schema reference

Used in: [KafkaMirrorMaker2Spec](#)

Property	Description
sourceCluster string	The alias of the source cluster used by the Kafka MirrorMaker 2.0 connectors. The alias must match a cluster in the list at <code>spec.clusters</code> .
targetCluster string	The alias of the target cluster used by the Kafka MirrorMaker 2.0 connectors. The alias must match a cluster in the list at <code>spec.clusters</code> .
sourceConnector <a href="#">KafkaMirrorMaker2ConnectorSpec</a>	The specification of the Kafka MirrorMaker 2.0 source connector.
checkpointConnector <a href="#">KafkaMirrorMaker2ConnectorSpec</a>	The specification of the Kafka MirrorMaker 2.0 checkpoint connector.
heartbeatConnector <a href="#">KafkaMirrorMaker2ConnectorSpec</a>	The specification of the Kafka MirrorMaker 2.0 heartbeat connector.
topicsPattern string	A regular expression matching the topics to be mirrored, for example, "topic1   topic2   topic3". Comma-separated lists are also supported.
topicsBlacklistPattern string	A regular expression matching the topics to exclude from mirroring. Comma-separated lists are also supported.
groupsPattern string	A regular expression matching the consumer groups to be mirrored. Comma-separated lists are also supported.
groupsBlacklistPattern string	A regular expression matching the consumer groups to exclude from mirroring. Comma-separated lists are also supported.

## B.151. [KafkaMirrorMaker2ConnectorSpec](#) schema reference

Used in: [KafkaMirrorMaker2MirrorSpec](#)

Property	Description
tasksMax integer	The maximum number of tasks for the Kafka Connector.
config map	The Kafka Connector configuration. The following properties cannot be set: <code>connector.class</code> , <code>tasks.max</code> .
pause boolean	Whether the connector should be paused. Defaults to false.

## B.152. KafkaMirrorMaker2Status schema reference

Used in: [KafkaMirrorMaker2](#)

Property	Description
conditions	List of status conditions.
<a href="#">Condition</a> array	
observedGeneration	The generation of the CRD that was last reconciled by the operator.
integer	
url	The URL of the REST API endpoint for managing and monitoring Kafka Connect connectors.
string	
connectorPlugins	The list of connector plugins available in this Kafka Connect deployment.
<a href="#">ConnectorPlugin</a> array	
connectors	List of MirrorMaker 2.0 connector statuses, as reported by the Kafka Connect REST API.
map array	
labelSelector	Label selector for pods providing this resource.
string	
replicas	The current number of pods being used to provide this resource.
integer	

## B.153. KafkaRebalance schema reference

Property	Description
spec	The specification of the Kafka rebalance.
<a href="#">KafkaRebalanceSpec</a>	
status	The status of the Kafka rebalance.
<a href="#">KafkaRebalanceStatus</a>	

## B.154. KafkaRebalanceSpec schema reference

Used in: [KafkaRebalance](#)

Property	Description
goals string array	A list of goals, ordered by decreasing priority, to use for generating and executing the rebalance proposal. The supported goals are available at <a href="https://github.com/linkedin/cruise-control#goals">https://github.com/linkedin/cruise-control#goals</a> . If an empty goals list is provided, the goals declared in the default.goals Cruise Control configuration parameter are used.
skipHardGoalCheck boolean	Whether to allow the hard goals specified in the Kafka CR to be skipped in optimization proposal generation. This can be useful when some of those hard goals are preventing a balance solution being found. Default is false.
excludedTopics string	A regular expression where any matching topics will be excluded from the calculation of optimization proposals. This expression will be parsed by the java.util.regex.Pattern class; for more information on the supported formar consult the documentation for that class.
concurrentPartitionMovementsPerBroker integer	The upper bound of ongoing partition replica movements going into/out of each broker. Default is 5.
concurrentIntraBrokerPartitionMovements integer	The upper bound of ongoing partition replica movements between disks within each broker. Default is 2.
concurrentLeaderMovements integer	The upper bound of ongoing partition leadership movements. Default is 1000.
replicationThrottle integer	The upper bound, in bytes per second, on the bandwidth used to move replicas. There is no limit by default.
replicaMovementStrategies string array	A list of strategy class names used to determine the execution order for the replica movements in the generated optimization proposal. By default BaseReplicaMovementStrategy is used, which will execute the replica movements in the order that they were generated.

## B.155. KafkaRebalanceStatus schema reference

Used in: [KafkaRebalance](#)

Property	Description
conditions	List of status conditions.
<b>Condition</b> array	
observedGeneration	The generation of the CRD that was last reconciled by the operator.
integer	
sessionId	The session identifier for requests to Cruise Control pertaining to this KafkaRebalance resource. This is used by the Kafka Rebalance operator to track the status of ongoing rebalancing operations.
string	
optimizationResult	A JSON object describing the optimization result.
map	